
ML - Machine Learning

Time Series

Riccardo Berta

2026.01.23

Contents

1 Time Series	1
1.1 Forecasting	4
1.1.1 Seasonality and Naive Forecasting	4
1.1.2 Performance Evaluation	6
1.1.3 Long-Term Patterns	8
1.2 Moving Average Models	11
1.2.1 Autoregressive Moving Average (ARMA)	11
1.2.2 Autoregressive Integrated Moving Average (ARIMA)	12
1.2.3 Seasonal ARIMA (SARIMA)	13
1.3 Recurrent Neural Network (RNN)	17
1.3.1 Recurrent Neuron	18
1.3.2 Recurrent Layer	22
1.3.3 Recurrent Architectures	27
1.3.4 Backpropagation Through Time	30
1.3.5 Applying RNNs to Time Series Forecasting	38
1.3.6 Deep networks	45
1.4 Extending to Multivariate and Multi-Step Forecasting	51
1.4.1 Multivariate Input	51
1.4.2 Forecasting Several Steps Ahead	55
1.5 Handling Long Sequences	166
1.5.1 Unstable Gradients Problem	166
1.6 Short-Term Memory Problem	169
1.7 Exercise	178

1 Time Series

A **time series** is a sequence of observations collected at successive points in time, typically at regular intervals. Using functional notation, a time series can be written as:

$$y(t), \quad t = 1, 2, \dots, T$$

where $y(t)$ denotes the value of the observed variable at time t . Common examples include the number of daily active users on a website, hourly temperature measurements in a city, a household's daily electricity consumption, or the trajectories of nearby vehicles. What distinguishes time series data is their **temporal ordering**: observations are not independent, and the order in which they occur is essential. In general, the value at time t depends on its past:

$$y(t) = f(y(t-1), y(t-2), \dots) + \varepsilon(t)$$

where

- $f(\cdot)$ captures **temporal patterns** such as trends, cycles, and recurring behaviors
- $\varepsilon(t)$ represents noise

By analyzing how the signal $y(t)$ evolves over time, we can **learn patterns from the past** and use them to **forecast future values**, under the assumption that these patterns **persist**.

As a concrete example, consider the dataset reporting the daily number of passengers using buses and rail services operated by the Chicago Transit Authority. This dataset is publicly available from the Chicago Data Portal and provides a real-world illustration of time series data used for forecasting.

```
import urllib.request
import zipfile
import os

# Define the URL of the dataset
url = "https://www.dropbox.com/s/wiww3a55sgifhb3/ridership.zip?dl=1"

# Create data directory if it does not exist
os.makedirs("./data", exist_ok=True)

# Prepare data if not already existing
if not os.path.exists("./data/ridership.zip"):
    print("Downloading dataset ... ")

    # Download the dataset
    u = urllib.request.urlopen(url)
    data = u.read()
    u.close()

    # Save the dataset to a local file
    with open("./data/ridership.zip", "wb") as f :
        f.write(data)

    # Extract the dataset
    with zipfile.ZipFile("./data/ridership.zip", "r") as zip_ref:
        zip_ref.extractall("./data")

else:
    print("Dataset already exists. Skipping download.")
```

Dataset already exists. Skipping download.

We begin by **loading and cleaning the dataset**. Specifically, we read the CSV file, assign concise column names, sort the observations chronologically, remove the redundant total column, and eliminate any duplicate rows:

```
import pandas as pd

# Load the dataset into a Pandas DataFrame
df = pd.read_csv("./data/ridership/CTA_-_Ridership_-_Daily_Boarding_Totals.csv",
                  parse_dates=[ "service_date" ])

# Assign concise column names
df.columns = [ "date", "day_type", "bus", "rail", "total" ]

# Sort the DataFrame by date and set date as index
df = df.sort_values("date").set_index("date")

# Remove the redundant total column
df = df.drop("total", axis=1)

# Drop duplicate rows
df = df.drop_duplicates()
```

We now inspect the first few rows of the dataset in order to understand its structure and contents:

```
# Inspect the first few rows of the dataset
df.head()
```

	day_type	bus	rail
date			
2001-01-01	U	297192	126455
2001-01-02	W	780827	501952
2001-01-03	W	824923	536432
2001-01-04	W	870021	550011
2001-01-05	W	890426	557917

This dataset represents a **time series**, where observations are recorded at regular time intervals. More specifically, because multiple variables are observed at each time step, it is a **multivariate**

time series. The "day_type" column encodes the type of day: W for weekdays, A for Saturdays, and U for Sundays or public holidays.

1.1 Forecasting

The most common task in time series analysis is **forecasting**, which consists in predicting future values of a signal based solely on its past observations. Formally, given a time series $y(t)$ observed up to time t , the goal of forecasting is to estimate one or more future values:

$$\hat{y}(t+1), \hat{y}(t+2), \dots, \hat{y}(t+H)$$

where H is the forecasting **horizon**.

In other words, forecasting aims to learn a function:

$$\hat{y}(t+h) = f(y(t), y(t-1), \dots)$$

that **maps the historical evolution of the series to its future behavior**, under the assumption that the underlying temporal patterns remain stable.

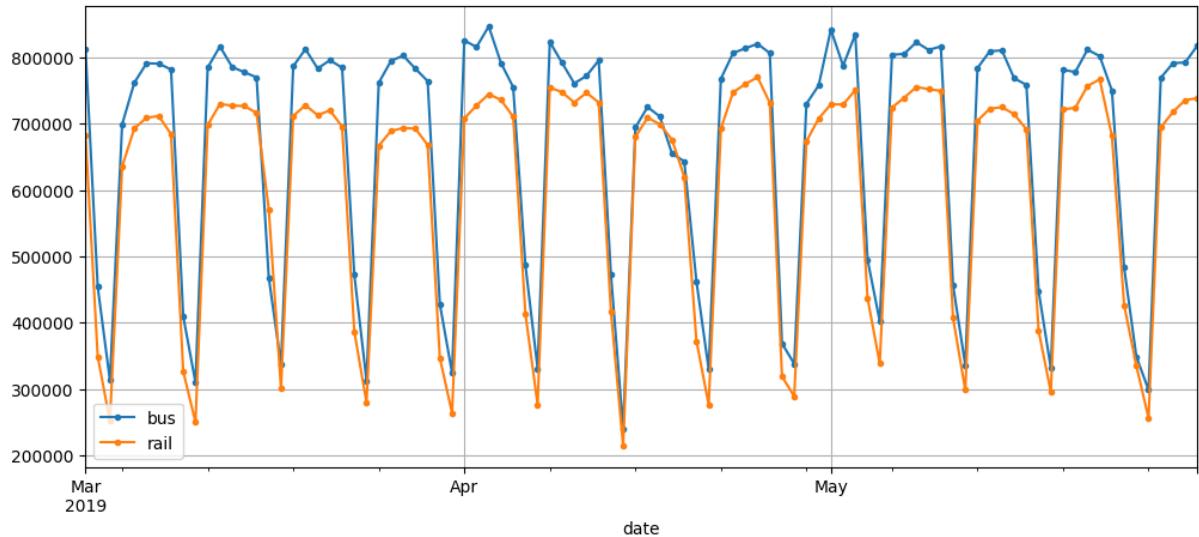
In the concrete example, given the observed ridership up to today, we may want to forecast tomorrow's demand or the ridership over the next week. Such forecasts are essential for operational planning, resource allocation, and service optimization.

1.1.1 Seasonality and Naive Forecasting

To build some intuition, let's plot the bus and rail ridership over a few months in 2019 and observe their behavior over time.

```
import matplotlib.pyplot as plt

# Plot bus and rail ridership from March to May 2019
df["2019-03":"2019-05"].plot(grid=True, marker=".",
    figsize=(12, 5))
plt.show()
```



We can clearly observe a **recurring weekly pattern**, known as **seasonality**. In this case, the seasonal effect is so strong that even a very simple strategy, such as predicting tomorrow's ridership by copying the value observed one week earlier, already yields reasonably good results. This approach is known as **naive forecasting** and predicts future values by reusing past observations from the previous seasonal cycle. For a weekly seasonal pattern, the naive forecast can be written as

$$\hat{y}(t) = y(t - 7)$$

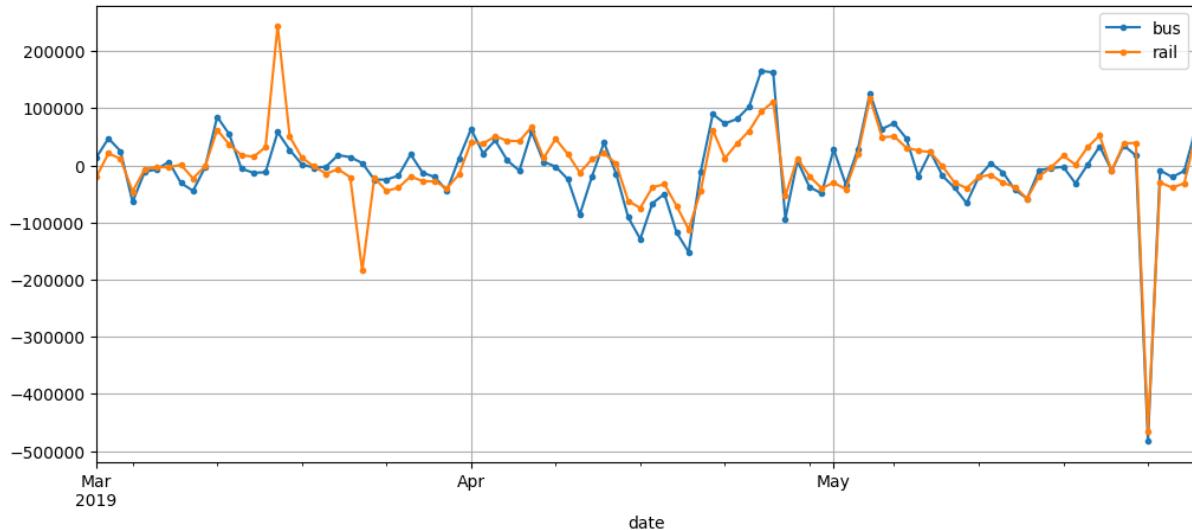
Despite its simplicity, naive forecasting often provides a strong and informative **baseline** against which more sophisticated models can be compared. To make the weekly seasonality explicit, we can compute the **7-day difference**, defined as

$$\Delta_7 y(t) = y(t) - y(t - 7)$$

If the series exhibits a strong weekly pattern, the **differenced signal** will fluctuate around zero, since values one week apart tend to be similar. In practice, this operation can be conveniently performed using the Pandas `diff()` method, which computes the difference between each value and the value observed a specified number of periods earlier:

```
# Compute the difference between each value and the value observed 7 days earlier
diff_7 = df[["bus", "rail"]].diff(7)[ "2019-03": "2019-05"]

# Plot the differenced series to highlight seasonality
diff_7.plot(grid=True, marker=". ", figsize=(12, 5))
plt.show()
```



Notice how closely the lagged time series follows the original one. When a time series shows a **strong correlation with a time-shifted version of itself**, we say that it exhibits **autocorrelation**. In this case, most of the differences remain quite small, indicating a strong weekly autocorrelation. An exception appears toward the end of May, where the differences become larger. This deviation may be due to a disruption of the usual pattern, such as a public holiday. To verify this, let's inspect the day types for that period:

```
# Check the day_type column for the end of May 2019
list(df.loc["2019-05-25":"2019-05-27"]["day_type"])

['A', 'U', 'U']
```

1.1.2 Performance Evaluation

The period with deviations corresponds to a long weekend, with Monday being the Memorial Day holiday. Such **calendar effects** can disrupt the usual seasonal pattern and may be explicitly modeled to improve forecasting accuracy. For now, however, we deliberately keep the model simple. To obtain a first **quantitative assessment** of our forecasting approach, we compute the **Mean Absolute Error (MAE)** over a three-month period. Given a sequence of true values and the corresponding forecasts, the MAE is defined as:

$$\text{MAE} = \frac{1}{T} \sum_{t=1}^T |y(t) - \hat{y}(t)|$$

This metric measures the average magnitude of the forecast errors in the original units of the data, providing an intuitive indication of how far the predictions are, on average, from the true

observations:

```
# Calculate the Mean Absolute Error (MAE) over the three-month period
mae = diff_7.abs().mean();

# Display the MAE results
print("Mean Absolute Error (MAE) from March to May 2019:")
print(f"Bus: {mae['bus']:.2f}")
print(f"Rail: {mae['rail']:.2f}")
```

Mean Absolute Error (MAE) from March to May 2019:

Bus: 43915.61
Rail: 42143.27

At first glance, it is difficult to judge whether these errors are **large or small**. To put them into perspective, we normalize the forecast errors by dividing them by the corresponding target values. This metric is known as the **Mean Absolute Percentage Error (MAPE)** and is defined as:

$$\text{MAPE} = \frac{1}{T} \sum_{t=1}^T \left| \frac{y(t) - \hat{y}(t)}{y(t)} \right| \times 100$$

```
# Get the target values for normalization
targets = df[["bus", "rail"]]["2019-03":"2019-05"]

# Calculate the Normalized Mean Absolute Error (NMAE)
normalized_mae = (diff_7 / targets).abs().mean()

# Display the NMAE results
print("Mean Absolute Percentage Error (MAPE) from March to May 2019:")
print(f"Bus: {normalized_mae['bus']:.2%}")
print(f"Rail: {normalized_mae['rail']:.2%}")
```

Mean Absolute Percentage Error (MAPE) from March to May 2019:

Bus: 8.29%
Rail: 8.99%

An interesting observation is that, while the **MAE** appears slightly better for rail than for bus, the opposite holds when we consider **MAPE**. This happens because bus ridership volumes are generally larger than rail ridership volumes, which naturally leads to larger **absolute errors**:

$$|y(t) - \hat{y}(t)|$$

However, when forecast errors are expressed **relative to the true values**, the bus forecasts turn out to be slightly more accurate than the rail forecasts. For completeness, the **Mean Squared Error (MSE)** is defined as:

$$\text{MSE} = \frac{1}{T} \sum_{t=1}^T (y(t) - \hat{y}(t))^2$$

MSE measures the average **squared** difference between the observed values and the corresponding forecasts over time. MSE **penalizes large forecast errors** more heavily than small ones, making it particularly sensitive to outliers and abrupt deviations. This property is useful when large errors are especially undesirable.

The **MAE**, **MAPE**, and **MSE** are among the most commonly used metrics for evaluating forecasting performance. As always, the choice of metric should be guided by the **scale of the data** and the **specific objectives** of the forecasting task.

1.1.3 Long-Term Patterns

By visually inspecting the time series, it is difficult to identify any clear **long-term trend** or **seasonal pattern** beyond the strong weekly seasonality. To further investigate this aspect, we can analyze the series at different temporal scales by **aggregating** observations over different time intervals.

Let \mathcal{T}_k denote a generic time window (for example, a month, a quarter, or a year). An aggregated version of the time series can then be defined as:

$$y^{(\text{agg})}(k) = \frac{1}{|\mathcal{T}_k|} \sum_{t \in \mathcal{T}_k} y(t)$$

where $|\mathcal{T}_k|$ is the number of observations contained in the time window. By choosing different definitions of time window, we can construct monthly, quarterly, or yearly aggregated series, each emphasizing patterns at a different temporal scale.

In practice, this type of temporal aggregation is conveniently performed using the Pandas **resample()** method, which transforms a time series from one temporal resolution to another by **grouping observations into time-based windows** defined by a chosen resampling rule. The method requires a **DateTimeIndex** and produces a grouped representation of the data, which can then be summarized using aggregation operators such as the mean, sum, or median.

Aggregating a time series over long intervals compresses many observations into a single value. While this helps reveal broad trends, **it also hides information about how the series behaves within each interval**. In other words, aggregation tells us what happens on average, but not

how the process evolves over time. **Rolling statistics** address this limitation by averaging over a moving window: instead of collapsing the data, they slide smoothly along the time axis. This preserves the temporal ordering of the observations while filtering out short-term noise, making it easier to visually track the gradual evolution of long-term trends and seasonal effects as they unfold.

Given a window length m , the rolling average of the series is defined as:

$$\bar{y}^{(m)}(t) = \frac{1}{m} \sum_{i=0}^{m-1} y(t-i)$$

This operation replaces each observation with the average of the previous m values, producing a smoothed version of the series. Smaller values of m highlight medium-term fluctuations, while larger values emphasize long-term trends and low-frequency seasonal components.

In Pandas, rolling statistics are computed using the `rolling()` method, which applies a sliding window over the time series and computes the desired statistic for each position of the window.

By combining **aggregation** and **rolling averages**, we can study the time series at multiple temporal scales and more effectively assess the presence of **long-term trends** and **seasonal patterns** that may not be visible at the original resolution.

In the example, we can analyze data from 2001 to 2019 and visualize both the monthly aggregated series and its 12-month rolling average in order to assess the presence of **yearly seasonality** and **long-term trends** in the data:

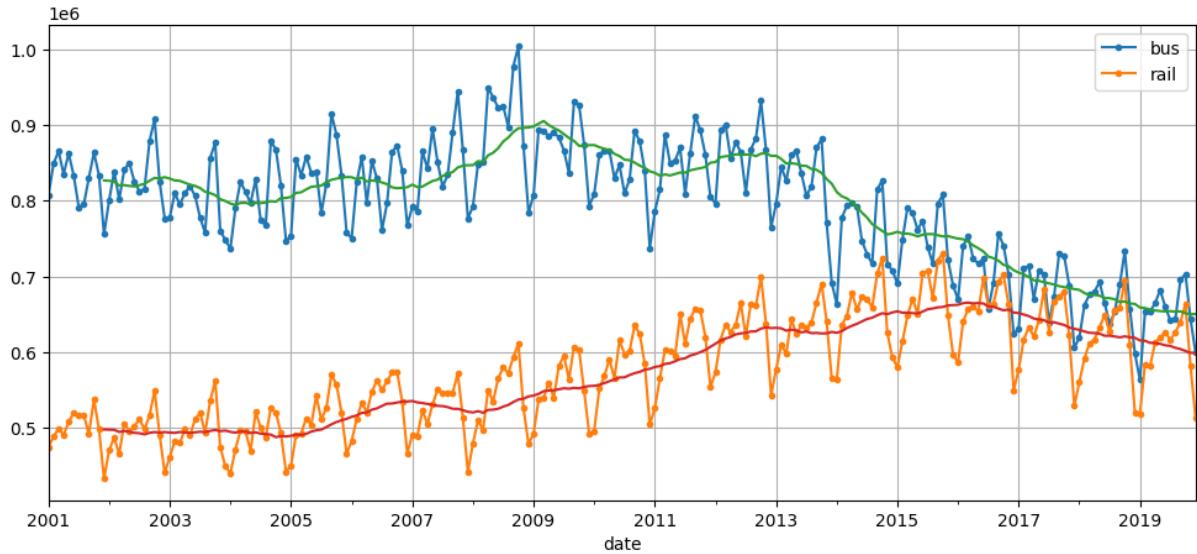
```
# Select a longer period for yearly seasonality analysis
period = slice("2001", "2019")

# Remove the day_type column for resampling
df_monthly = df.loc[ :, df.columns != 'day_type']

# Compute the mean for each month
df_monthly = df_monthly.resample(rule='ME').mean()

# Compute the 12-month rolling average
rolling_average_12_months = df_monthly[period].rolling(window=12).mean()

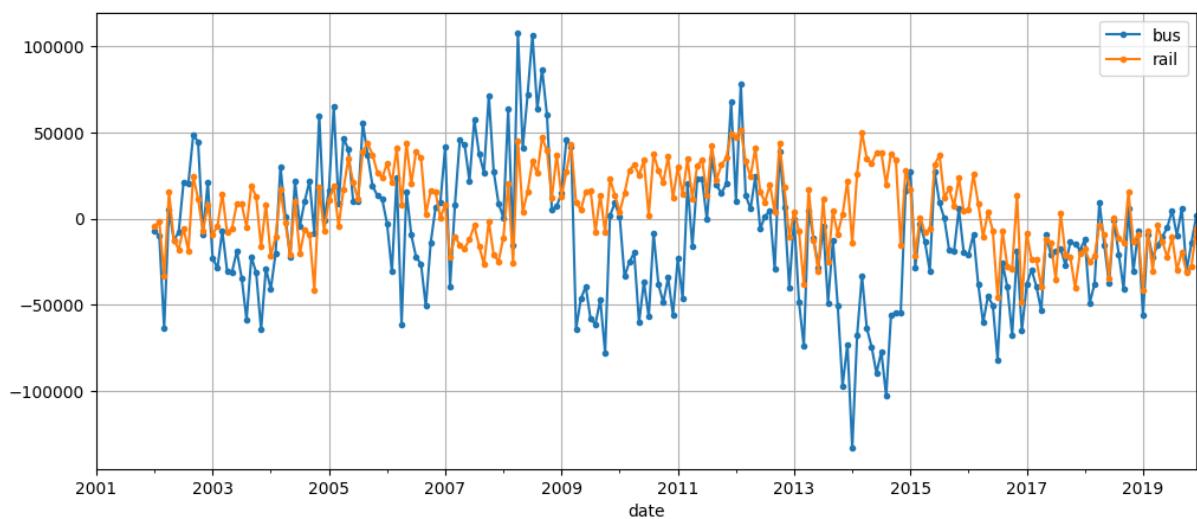
# Plot the monthly data along with rolling average
fig, ax = plt.subplots(figsize=(12, 5))
df_monthly[period].plot(ax=ax, marker=".")
rolling_average_12_months.plot(ax=ax, grid=True, legend=False)
plt.show()
```



The data exhibit a degree of **yearly seasonality**, although it is noisier than the weekly pattern and more pronounced in the rail series than in the bus series. In particular, we observe peaks and troughs occurring at approximately the same times each year. To make this yearly pattern more explicit, let's examine the **12-month difference** of the series:

```
# Compute the 12-month difference
diff_12 = df_monthly.diff(12)[period]

# Plot the differenced series to highlight seasonality
diff_12.plot(grid=True, marker=". ", figsize=(12, 5))
plt.show()
```



After applying the 12-month differencing, the yearly seasonal pattern and long-term trend are largely removed, and the series now fluctuates around zero, making deviations and anomalies much more apparent.

1.2 Moving Average Models

Moving Average models are a class of time series models where the present value is explained not by past values of the series itself, but by how the system has recently deviated from its own predictions, capturing the impact of short-lived random disturbances.

1.2.1 Autoregressive Moving Average (ARMA)

An alternative to naive forecasting is the **Autoregressive Moving Average (ARMA)** model, which generates forecasts by combining two complementary mechanisms. It predicts future values as a **weighted sum of past observations** and refines these predictions by incorporating a **weighted sum of recent forecast errors**:

$$\hat{y}(t) = \sum_{i=1}^p \alpha_i y(t-i) + \sum_{i=1}^q \beta_i \epsilon(t-i)$$

where the forecast error is defined as

$$\epsilon(t) = y(t) - \hat{y}(t)$$

The first term is the **autoregressive (AR)** component, which captures how past values of the series influence the current value. The second term is the **moving average (MA)** component, which corrects the forecast using information contained in past prediction errors. The hyperparameters p and q determine how many past observations and past errors are taken into account, respectively. In practice, there is no single "correct" choice for p and q . A common and effective strategy is to start with **small values**, fit the model to estimate the parameters α_i and β_i (typically using **maximum likelihood estimation**), and then examine the model's **diagnostics**, such as residual autocorrelation and error behavior. If the residuals resemble **white noise**, the model has successfully captured the relevant temporal structure of the data. If systematic patterns remain, the model is too simple and should be refined by increasing p and/or q . Conversely, if increasing model complexity does not lead to a meaningful improvement in the diagnostics, adding more parameters is unlikely to be beneficial and may instead result in **overfitting**, where the model begins to fit noise rather than signal.

This iterative procedure balances **model expressiveness**, the ability to capture temporal dependencies, and **parsimony**, the principle of keeping the model as simple as possible, ensuring that the chosen model explains the data well without introducing unnecessary parameters.

1.2.2 Autoregressive Integrated Moving Average (ARIMA)

The previous model works well when the time series is **stationary**, meaning that its statistical properties, such as mean and variance, do not change over time:

$$\mathbb{E}[y(t)] = \mu, \quad \text{Var}(y(t)) = \sigma^2 \quad \forall, t$$

When a series contains **trends** or evolving patterns, this assumption is violated. In such cases, **differencing** provides a simple yet powerful way to remove non-stationary components. A **one-step difference** is defined as:

$$\Delta y(t) = y(t) - y(t - 1)$$

This operation is analogous to a **first-order derivative** in discrete time: it measures how much the series changes from one time step to the next and removes any **linear trend**. For example, the series:

$$y(t) = [3, 5, 7, 9, 11]$$

grows linearly. Applying first differencing yields:

$$\Delta y(t) = [2, 2, 2, 2]$$

indicating a constant slope and showing that the original linear trend has been removed. If the series follows a **quadratic trend**, a single round of differencing is not sufficient. Consider the series:

$$y(t) = [1, 4, 9, 16, 25, 36]$$

After one differencing step we obtain:

$$\Delta y(t) = [3, 5, 7, 9, 11]$$

which is still trending. Applying differencing a **second time**:

$$\Delta^2 y(t) = \Delta(\Delta y(t)) = y(t) - 2y(t - 1) + y(t - 2)$$

yields:

$$[2, 2, 2, 2]$$

which removes the remaining structure. A convenient notation to express these operations is the **lag operator L**, which represents a shift backward in time:

$$Ly(t) = y(t - 1)$$

Applying the operator multiple times produces larger shifts:

$$L^2 y(t) = y(t - 2), \quad L^3 y(t) = y(t - 3), \text{ and so on.}$$

In general,

$$L^k y(t) = y(t - k)$$

Using this notation, a one-step difference can be written compactly as:

$$\Delta y(t) = y(t) - y(t - 1) = y(t) - Ly(t) = (1 - L)y(t)$$

while a second difference becomes:

$$\Delta^2 y(t) = (1 - L)[(1 - L)y(t)] = (1 - L)^2 y(t) = y(t) - 2y(t - 1) + y(t - 2)$$

More generally, applying differencing d times gives:

$$\Delta^d y(t) = (1 - L)^d y(t)$$

This operation approximates the **d-th derivative** of the series and removes **polynomial trends up to degree d**. The hyperparameter d is called the **order of integration**. This idea lies at the core of the **Autoregressive Integrated Moving Average (ARIMA)** model. Before modeling temporal dependencies using an ARMA model, the series is first "flattened" through differencing so that its statistical properties remain stable over time:

$$z(t) = \Delta^d y(t) = (1 - L)^d y(t)$$

An ARMA model is then applied to the differenced series:

$$\hat{z}(t) = \sum_{i=1}^p \alpha_i z(t - i) + \sum_{i=1}^q \beta_i \epsilon(t - i)$$

with prediction error

$$\epsilon(t) = z(t) - \hat{z}(t)$$

In summary, ARIMA is characterized by three hyperparameters: p and q , as in ARMA, and the additional parameter d , which controls the degree of differencing applied to the series before fitting the ARMA model.

1.2.3 Seasonal ARIMA (SARIMA)

A natural final extension of ARIMA is the **Seasonal ARIMA (SARIMA)** model, which explicitly accounts for **seasonal patterns**. The key idea is that many time series evolve on **two time scales simultaneously**:

- a **short-term scale**, capturing local dynamics from one time step to the next;
- a **seasonal scale**, capturing patterns that repeat every s time steps.

While ARIMA models only the short-term structure, SARIMA augments it by introducing additional autoregressive and moving-average components that operate at the **seasonal lag**. To do

this, the model introduces three **seasonal hyperparameters** P, D, and Q, together with the **seasonal period** s. After applying both regular and seasonal differencing, the transformed series is defined as:

$$z(t) = \Delta^d \Delta_s^D y(t) = (1 - L)^d (1 - L_s)^D y(t)$$

An ARMA model is then applied to this stationary series:

$$\hat{z}(t) = \underbrace{\sum_{i=1}^p \alpha_i z(t-i)}_{\text{non-seasonal AR}} + \underbrace{\sum_{i=1}^P A_i z(t-i \cdot s)}_{\text{seasonal AR}} + \underbrace{\sum_{i=1}^q \beta_i \epsilon(t-i)}_{\text{non-seasonal MA}} + \underbrace{\sum_{i=1}^Q B_i \epsilon(t-i \cdot s)}_{\text{seasonal MA}}$$

with prediction error:

$$\epsilon(t) = z(t) - \hat{z}(t)$$

In this formulation, the non-seasonal terms model short-term dependencies, while the seasonal terms capture repeating patterns occurring every s time steps. Together, these components allow SARIMA to represent both **local dynamics** and **long-range seasonal structure** within a unified framework.

The **statsmodels** library is a Python package for classical statistical modeling and time series analysis. It provides implementations of ARMA, ARIMA, and SARIMA models, along with a wide range of statistical diagnostics and hypothesis tests. Unlike many machine learning libraries, statsmodels emphasizes **explicit model assumptions**, **parameter estimation**, and **interpretability**. We can exploit this library to fit a SARIMA model to the rail time series and generate forecasts that explicitly account for its seasonal behavior:

```
from statsmodels.tsa.arima.model import ARIMA

# Select the rail time series for a specific period
origin, today = "2019-01-01", "2019-05-31"
rail_series = df.loc[origin:today]["rail"].asfreq("D")

# Define the ARIMA model parameters
order = (1, 0, 0) # p=1, d=0, q=0
seasonal_order = (0, 1, 1, 7) # P=0, D=1, Q=1, and s=7

# Fit the SARIMA model to the rail time series
model = ARIMA(rail_series, order=order, seasonal_order=seasonal_order)
model = model.fit()
```

We can use the model to generate a forecast for the next day

```
# Use the model to generate a forecast for the next day
y_pred = model.forecast()
print("SARIMA prediction: ", y_pred.values[0])

# Compare with the actual value
print("Actual value: ", df["rail"].loc["2019-06-01"])

# Compare with the naive prediction
print("Naive prediction: ", df["rail"].loc["2019-05-25"])
```

```
SARIMA prediction: 427758.6263023059
Actual value: 379044
Naive prediction: 426932
```

The forecast for this particular day is quite poor, similar or even worse than the naive baseline. However, judging a model based on a single prediction can be misleading. To properly assess its performance, we need to generate forecasts over a longer horizon and evaluate them using an aggregate metric, such as the Mean Absolute Error computed over the entire period:

```
# Earliest date used for training
origin = "2019-01-01"

# Start and end of the evaluation period
start_date = "2019-03-01"
end_date = "2019-05-31"

# Create a daily date range for evaluation
time_period = pd.date_range(start_date, end_date)

# Extract the rail time series and enforce daily frequency
rail_series = df.loc[origin:end_date]["rail"].asfreq("D")

# List to store one-step-ahead predictions
y_preds = []

# Rolling (walk-forward) forecasting
for today in time_period.shift(-1):

    # Train the model using all data available up to "today"
    train_series = rail_series.loc[origin:today]
```

```

# Fit the SARIMA model
model = ARIMA(
    train_series,
    order=order,
    seasonal_order=seasonal_order
)
results = model.fit()

# Forecast the next day (one-step ahead)
y_pred = results.forecast(steps=1).iloc[0]

# Store the prediction
y_preds.append(y_pred)

# Convert predictions to a time-indexed Series
y_preds_sarima = pd.Series(y_preds, index=time_period)

# Compute SARIMA MAE
mae_sarima = (y_preds_sarima - rail_series.loc[time_period]).abs().mean()

# Naive forecast
y_preds_naive = rail_series.shift(7).loc[time_period]

# Compute Naive MAE
mae_naive = (y_preds_naive - rail_series.loc[time_period]).abs().mean()

print("SARIMA MAE:", mae_sarima)
print("Naive MAE:", mae_naive)

```

SARIMA MAE: 32040.720095473138

Naive MAE: 42143.27173913043

Although the SARIMA model is not perfect, it consistently outperforms the naive baseline by a wide margin on average. To better understand this improvement, we can visualize the SARIMA forecasts alongside the actual observed values:

```

fig, ax = plt.subplots(figsize=(8, 3))

ax.plot(rail_series.loc[time_period].index,
        rail_series.loc[time_period].values,
        marker=".", label="True",

```

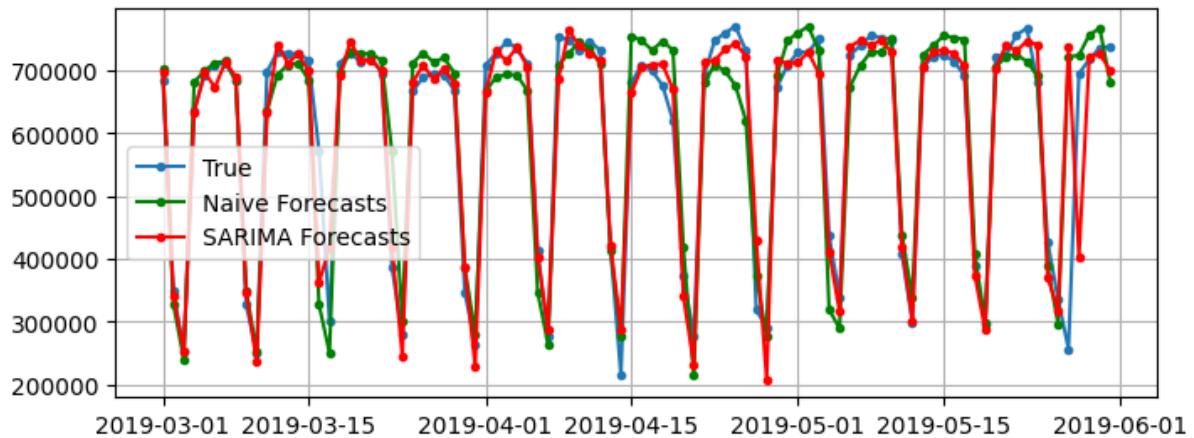
```

)
ax.plot(y_preds_naive.index,
         y_preds_naive.values,
         color="green", marker=".",
         label="Naive Forecasts"
)

ax.plot(y_preds_sarima.index,
         y_preds_sarima.values,
         color="r", marker=".",
         label="SARIMA Forecasts"
)

ax.grid(True)
ax.legend()
plt.show()

```



The plot shows that both models capture the strong weekly seasonality, but the SARIMA forecasts track the true series more closely than the naive baseline, especially around weekday peaks and weekend drops, resulting in a consistently lower average error.

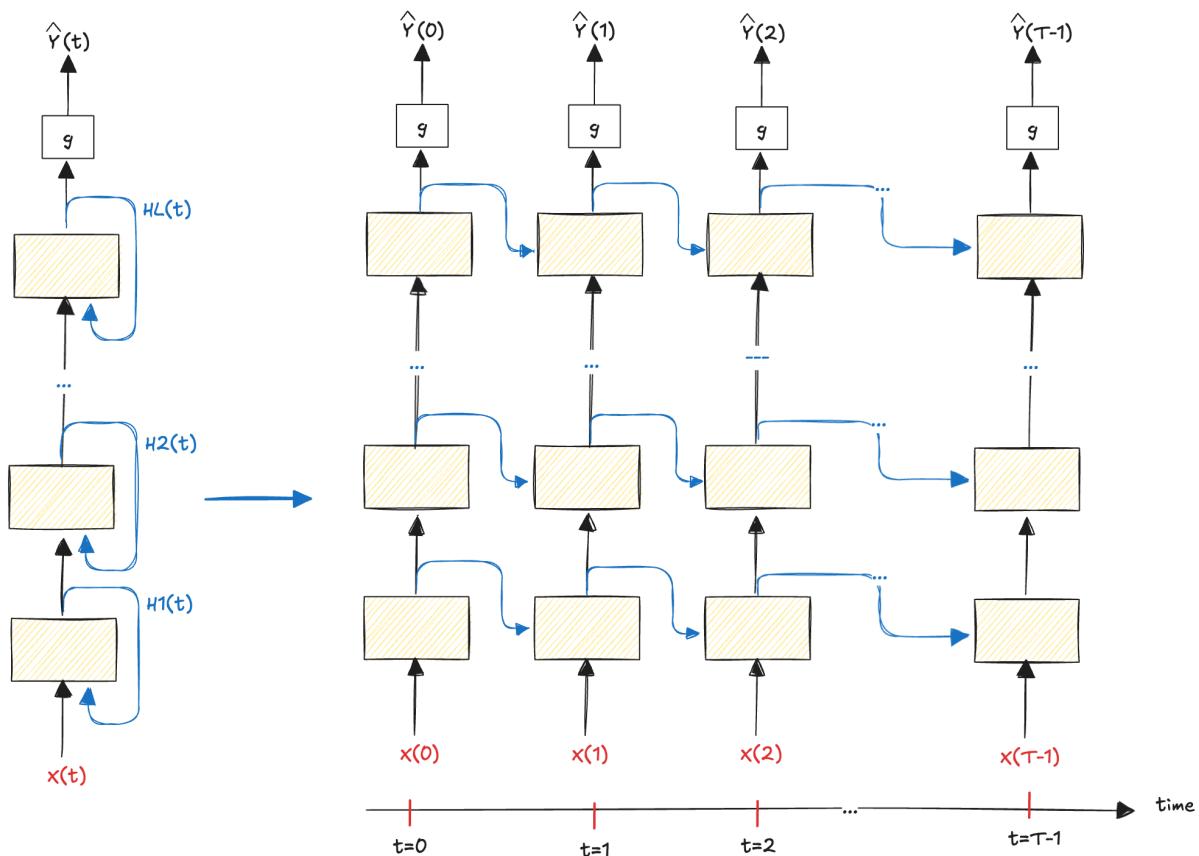
1.3 Recurrent Neural Network (RNN)

In a feedforward neural network, information flows strictly in one direction, from the input layer to the output layer. **Predictions are made independently**, without any memory of past inputs. This makes feedforward networks poorly suited for time series, where the current value often depends on what happened before. A **Recurrent Neural Network (RNN)** addresses this limitation by introducing **feedback connections**: the network feeds part of its output back into itself.

This creates an internal **state** that acts as a memory, allowing the model to carry information forward in time.

1.3.1 Recurrent Neuron

We begin with the simplest possible recurrent neural network: a single recurrent neuron. Unlike a standard feedforward neuron, which processes each input independently, a recurrent neuron maintains an **internal state that evolves over time**. This internal state allows the neuron to **retain information about past inputs** and to use it when processing new data:



At each time step t , the neuron receives two sources of information. The first is the current input vector:

$$x(t) \in \mathbb{R}^d$$

The second is the hidden state:

$$h(t-1) \in \mathbb{R}$$

which is carried over from the previous time step and acts as a **learned summary of the past**

inputs. These two signals are combined to produce a new hidden state $h(t)$, which is then fed back into the neuron and becomes part of the input at the next time step.

Because the hidden state is scalar in this setting, the neuron has two distinct sets of parameters: an input weight vector:

$$w_x \in \mathbb{R}^d$$

which processes the current input, and a recurrent weight

$$w_h \in \mathbb{R}$$

which controls how strongly past information influences the present. A bias term

$$b \in \mathbb{R}$$

is also included. The state update equation is therefore:

$$h(t) = \phi(w_x^\top x(t) + w_h h(t-1) + b)$$

where $\phi(\cdot)$ is a nonlinear activation function, such as ReLU. To fully define the recurrence, an initial hidden state must be specified, and is typically set to zero or treated as a learnable parameter.

The output of the neuron at time t is obtained by applying an output function $g()$ to the hidden state:

$$\hat{y}(t) = g(h(t))$$

This simple recurrent neuron already captures the key idea behind recurrent neural networks: **information flows not only forward from input to output, but also through time, allowing the model to build and update an internal memory as new data arrive.**

For example, in the context of daily ridership data, the hidden state can be interpreted as a compact summary of past travel patterns (weekly seasonality) which is then combined with the current day's observation to produce a prediction for the next day's ridership. We can implement this idea in Python:

```
import numpy as np

# Phi activation function (ReLU)
def phi(z):
    return np.maximum(0.0, z)

# g output function (identity)
def g(h):
    return h
```

```
def rnn_cell(x_t, h_prev, w_x, w_h, b):

    # Compute the linear combination of inputs and previous hidden state
    z = w_x @ x_t + w_h * h_prev + b

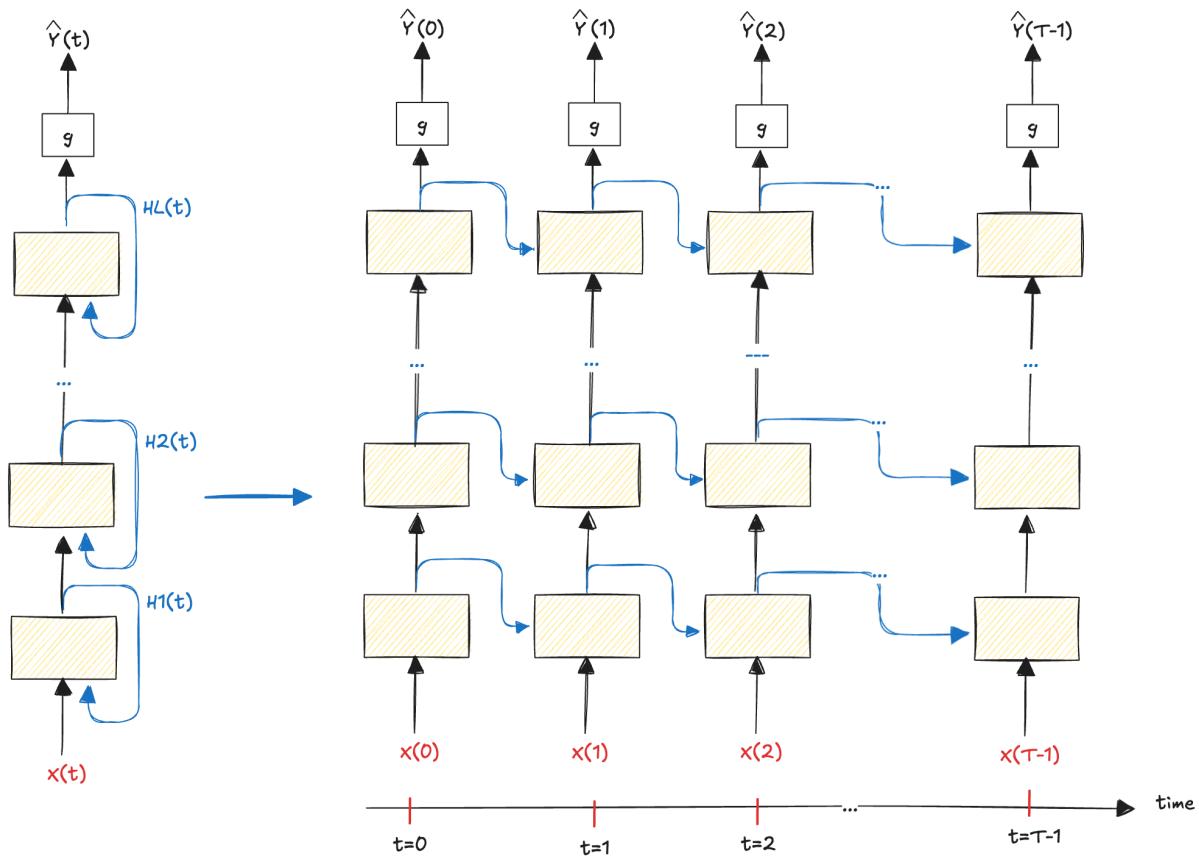
    # Apply activation function
    h_t = phi(z)

    # Compute output
    y_t = g(h_t)

    # Return the new hidden state and output
    return h_t, y_t
```

The recurrent neuron maintains an **internal state** that acts as a **compact memory** of the input sequence, enabling the network to retain, update, and exploit **information about the past**. As a result, observations are not processed in isolation: the model explicitly captures temporal dependencies, which makes it particularly well suited for sequential and time-series data. In this sense, the hidden state functions as a learned representation of past information that is continuously updated over time.

To better understand how this mechanism works, we can **conceptually unroll the recurrent connection along the time axis**. Consider an input sequence of length T , indexed by time steps $t=0,1,\dots,T-1$. At each time step, the recurrent neuron receives both the current input $x(t)$ and its own hidden state $h(t-1)$ from the previous time step. Since no past state exists at the initial time step $t=0$, the hidden state is typically initialized to zero.



In this **unfolded view**, the recurrent neuron appears as a **small network replicated across time**, with the same parameters shared at every step. This representation, known as **unrolling** the network through time, makes explicit how information flows forward from one time step to the next and how the hidden state evolves over the entire sequence of length T . We can unroll the previous Python implementation as follows:

```
def rnn_cell_unroll(x, w_x, w_h, b):

    # Get the number of time steps
    T = len(x)

    # Initialize the hidden state for the first time step
    h = 0.0

    # Lists to store hidden states and outputs
    H = []
    Y_hat = []

    # Unroll the RNN over time steps
```

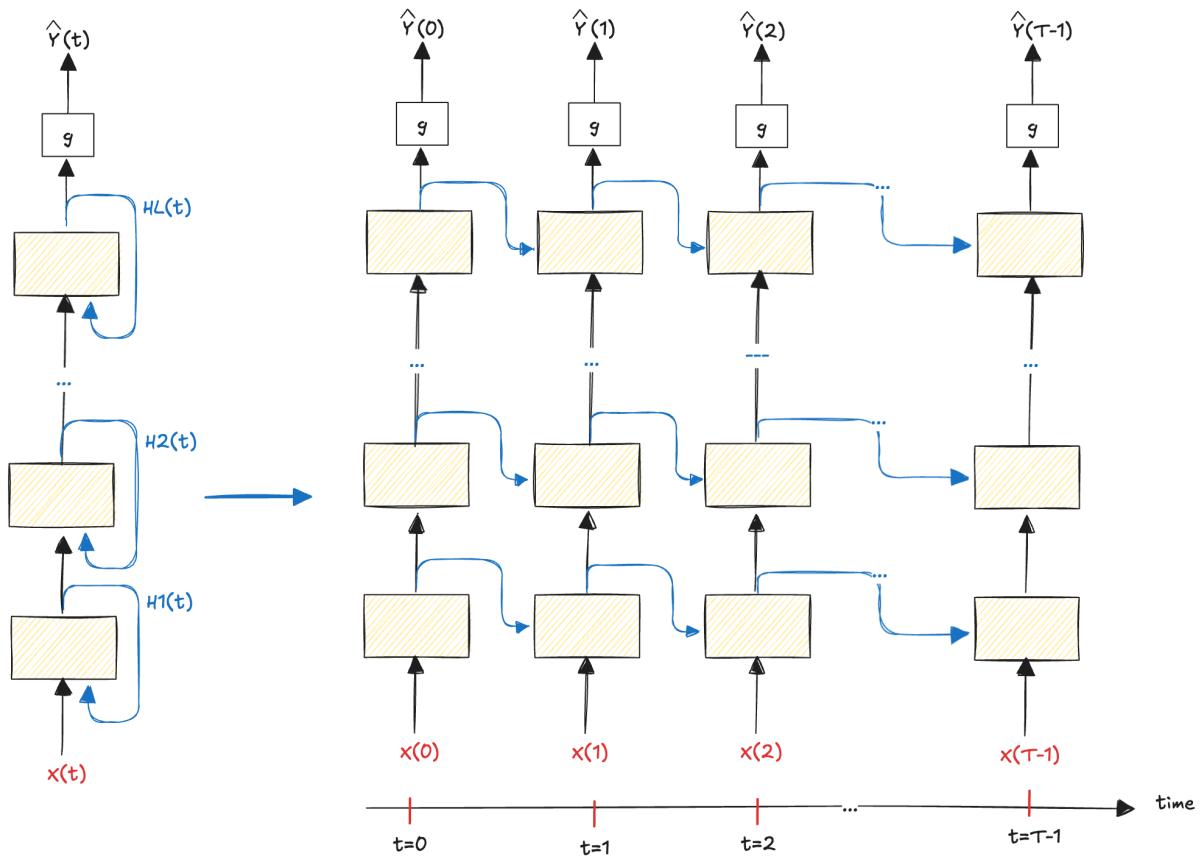
```
for t in range(T):
    # Compute the new hidden state and output
    h, y = rnn_cell(X[t], h, w_x, w_h, b)

    # Append to the lists
    H.append(h)
    Y_hat.append(y)

# return the arrays of hidden states and outputs
return np.array(H), np.array(Y_hat)
```

1.3.2 Recurrent Layer

A single recurrent neuron has very limited representational power. For this reason, practical recurrent neural networks use not just one recurrent neuron, but a **layer of recurrent neurons operating in parallel**. The layer consists of m recurrent neurons, all processing the **same input sequence**:



At each time step t , the input vector $x(t)$ is provided simultaneously to all neurons in the layer. Each neuron maintains its own hidden activation, and together these activations form the hidden state vector:

$$h(t) = [h_0(t), h_1(t), \dots, h_{m-1}(t)]^\top \in \mathbb{R}^m$$

Each neuron combines the shared input with the **entire hidden state vector from the previous time step**, not just with its own past activation. As a result, every neuron has its own input weight vector and its own recurrent weight vector:

$$h_i(t) = \phi(w_{x,i}^\top x(t) + w_{h,i}^\top h(t-1) + b_i)$$

where:

$$w_{x,i} \in \mathbb{R}^d$$

$$w_{h,i} \in \mathbb{R}^m$$

$$b_i \in \mathbb{R}$$

This formulation shows that the hidden state represents a **distributed memory**: information about the past is not stored in a single unit, but is spread across multiple neurons, each capturing different aspects of the history.

ing different aspects of the temporal structure (such as trends, periodic patterns, or short-term fluctuations).

The output at time t is obtained by applying a mapping function g to the entire hidden state vector:

$$\hat{y}(t) = g(h(t))$$

rather than to individual hidden units.

```
import numpy as np

def init_rnn_layer(input_dim, hidden_dim, output_dim):

    params = {
        # Input weights:
        "W_x": np.random.normal(0.0, 0.1, size=(input_dim, hidden_dim)),

        # Recurrent weights:
        "W_h": np.random.normal(0.0, 0.1, size=(hidden_dim, hidden_dim)),

        # Bias for each recurrent neuron
        "b": np.zeros(hidden_dim),

        # Output weights:
        "W_y": np.random.normal(0.0, 0.1, size=(hidden_dim, output_dim)),

        # Output bias
        "b_y": np.zeros(output_dim)
    }

    return params
```

In practice, as for feedforward networks, recurrent neural networks are almost always trained using **mini-batches**, rather than processing a single sequence at a time. Thus, instead of working with a single input sequence:

$$x(0), x(1), \dots, x(T-1)$$

we work with a mini-batch of B sequences, each of length T :

$$x^{(1)}(0), \dots, x^{(1)}(T-1)$$

$$x^{(2)}(0), \dots, x^{(2)}(T-1)$$

...

$$x^{(B)}(0), \dots, x^{(B)}(T-1)$$

At a fixed time step t , the input to the recurrent layer can be represented as a matrix

$$X(t) = \begin{bmatrix} x^{(1)}(t) \\ x^{(2)}(t) \\ \vdots \\ x^{(B)}(t) \end{bmatrix} \in \mathbb{R}^{B \times d}$$

where each row corresponds to the input vector of one sequence at time t . Importantly, $X(t)$ **does not represent a sequence by itself; the temporal structure arises from the ordered collection of such matrices over time**. Similarly, the hidden state at time t is represented as

$$H(t) = \begin{bmatrix} h^{(1)}(t) \\ h^{(2)}(t) \\ \vdots \\ h^{(B)}(t) \end{bmatrix} \in \mathbb{R}^{B \times m}$$

where each row is the hidden state of one sequence, and m is the hidden dimension of the recurrent layer.

At each time step, each sequence in the mini-batch **evolves independently**, but all sequences **share the same parameters**. For the b -th sequence in the mini-batch, the hidden state update is given by:

$$h^{(b)}(t) = \phi(W_x x^{(b)}(t) + W_h h^{(b)}(t-1) + b), \quad b = 1, \dots, B$$

Explicitly, this corresponds to

$$h^{(1)}(t) = \phi(W_x x^{(1)}(t) + W_h h^{(1)}(t-1) + b),$$

$$h^{(2)}(t) = \phi(W_x x^{(2)}(t) + W_h h^{(2)}(t-1) + b),$$

\vdots

$$h^{(B)}(t) = \phi(W_x x^{(B)}(t) + W_h h^{(B)}(t-1) + b).$$

Stacking these equations for all sequences in the mini-batch leads to the vectorized formulation:

$$H(t) = \phi(X(t)W_x + H(t-1)W_h + b)$$

where

$$W_x \in \mathbb{R}^{d \times m}$$

$$W_h \in \mathbb{R}^{m \times m}$$

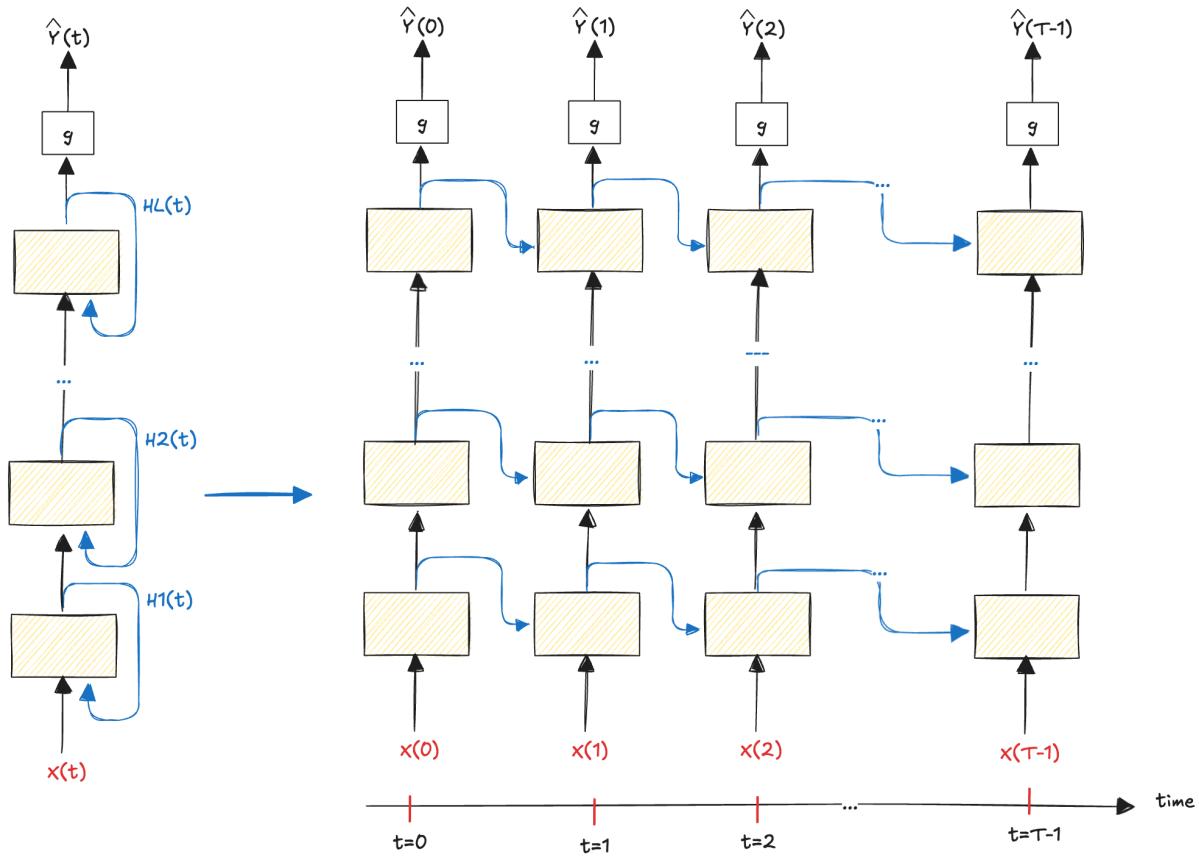
$$b \in \mathbb{R}^m$$

The output at time t is then obtained by applying an output mapping to the hidden states:

$$\hat{Y}(t) = g(H(t))$$

The same weight matrices are shared across all sequences in the mini-batch and across all time steps, ensuring that the model learns a single, consistent temporal transformation. Mini-batch processing significantly **improves computational efficiency** and **stabilizes training**, while fully preserving the temporal dynamics learned by the recurrent layer.

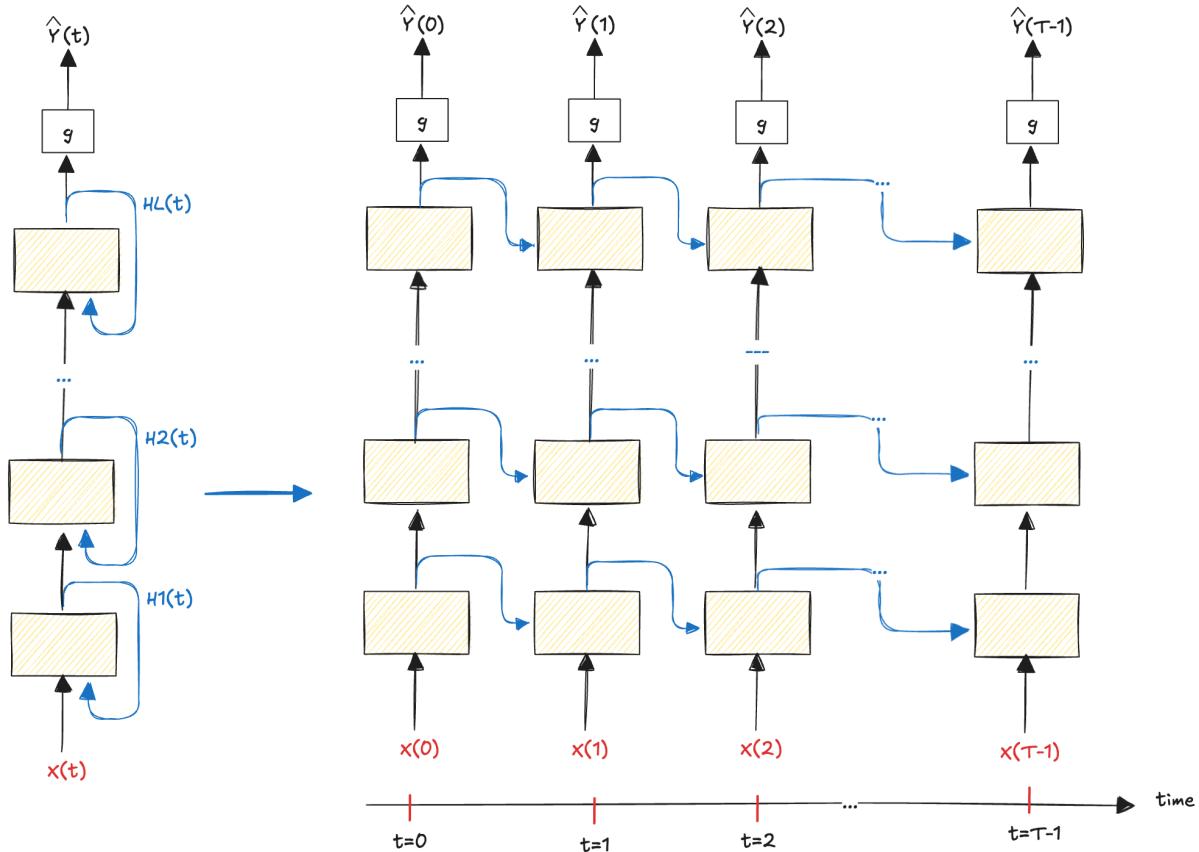
As with a single recurrent neuron, a recurrent layer can be **unrolled through time**. Each block represents the same recurrent layer evaluated at a different time step, with **shared parameters** across all copies. At time t , the layer receives the current input and the hidden state from the previous step, and produces a new hidden state and an output:



The same layer (with the same parameters) is replicated at each time step. This makes explicit how the hidden state evolves and how information flows forward in time, while emphasizing a crucial property of recurrent networks: **parameter sharing across time**. The weight matrices and the bias are reused at every time step, allowing the model to process sequences of arbitrary length with a fixed number of parameters.

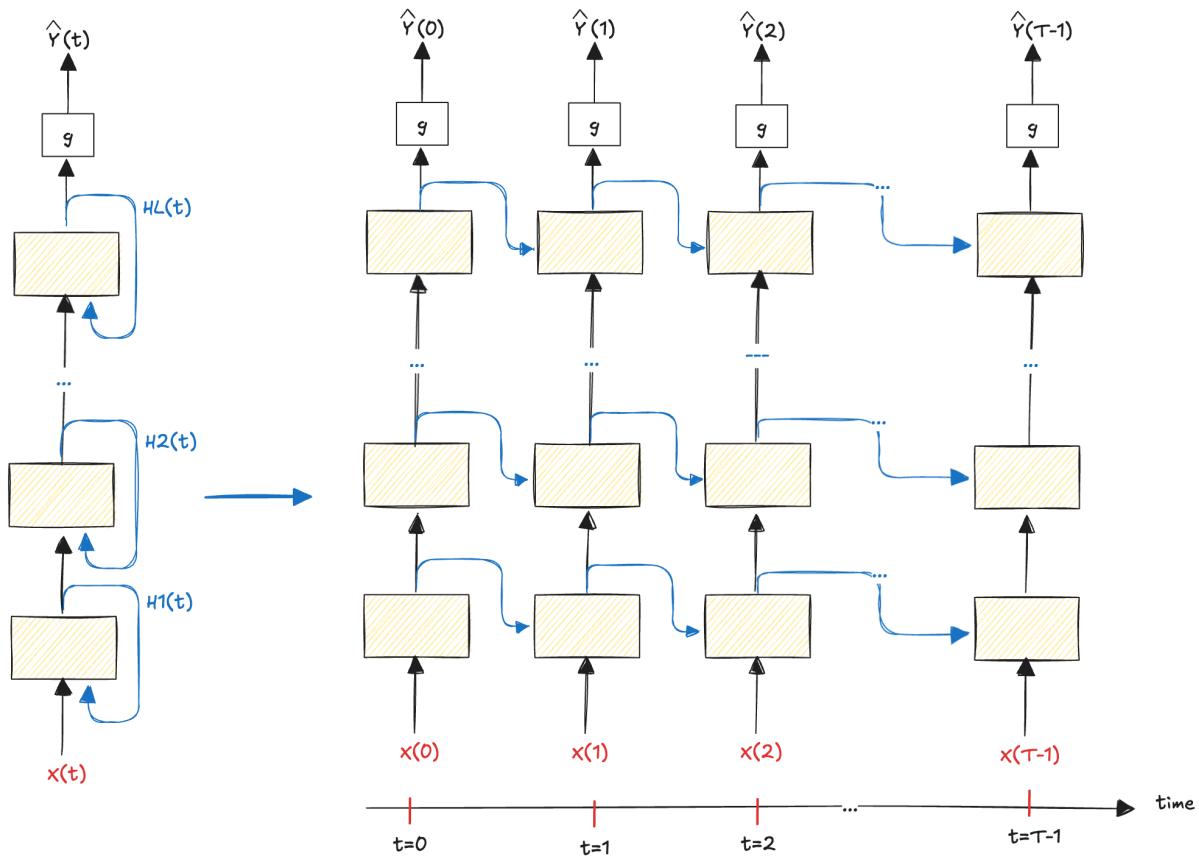
1.3.3 Recurrent Architectures

An RNN can process an entire sequence and produce a **corresponding sequence of outputs**. This architecture is known as a **sequence-to-sequence network**:



For example, the input consists of the last T observed values of household power consumption, and train it to predict the same sequence shifted **one day into the future**, effectively learning to forecast the next day's consumption at each time step.

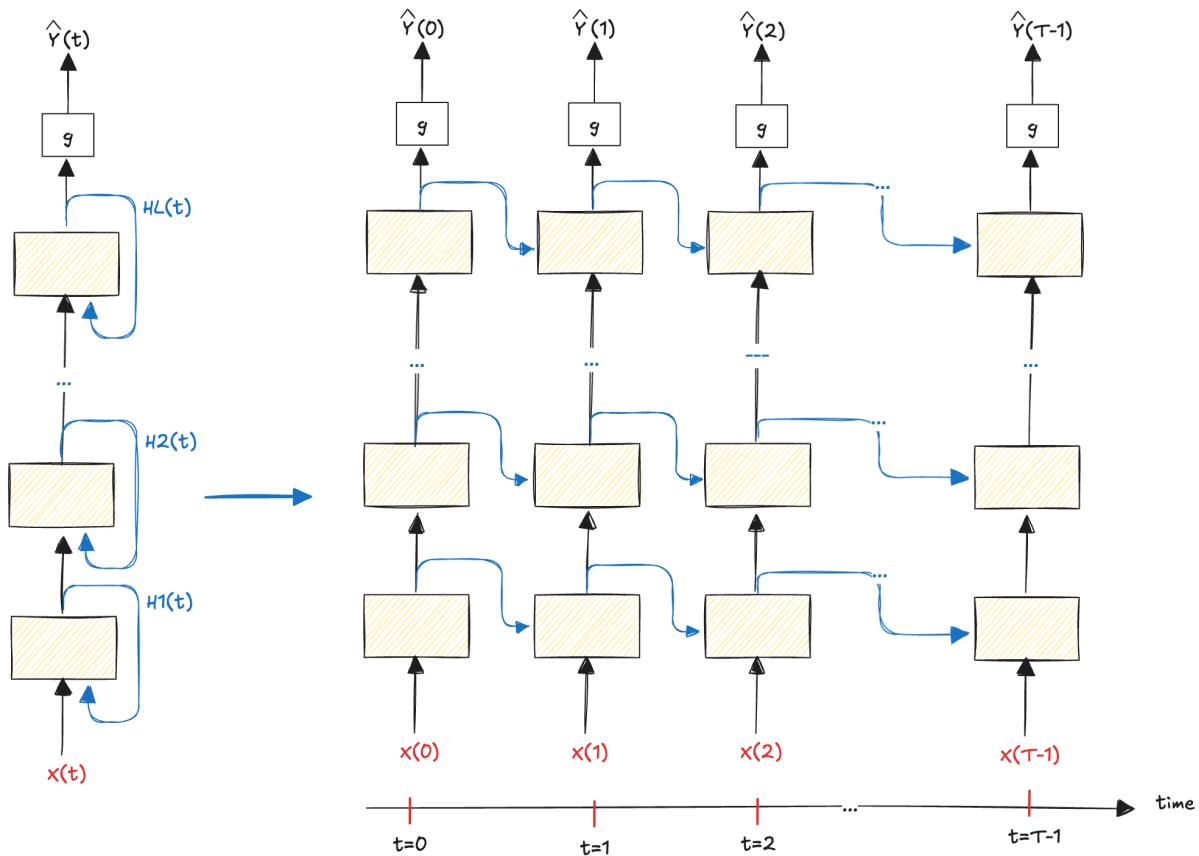
Alternatively, we can feed the network a **sequence of inputs** and use **only the final output**, obtaining a **sequence-to-vector** architecture:



In this case, the network compresses the entire input sequence into a single representation, which can then be used for tasks such classification. For example, we can input a sequence of words forming a movie review and have the network output a single sentiment score summarizing the overall opinion expressed in the text.

Although both architectures can be used to predict values one step into the future, **they differ in how supervision is applied**. In a sequence-to-sequence model, the network is trained to make accurate predictions at every time step, whereas in a sequence-to-vector model only the final prediction is supervised. As a result, sequence-to-sequence training provides a denser learning signal and encourages accurate intermediate predictions, while sequence-to-vector training forces the network to compress the entire sequence into a single representation optimized for the final output. The use case for sequence-to-sequence models is typically forecasting tasks, while sequence-to-vector models are often employed for classification or regression tasks based on sequential data.

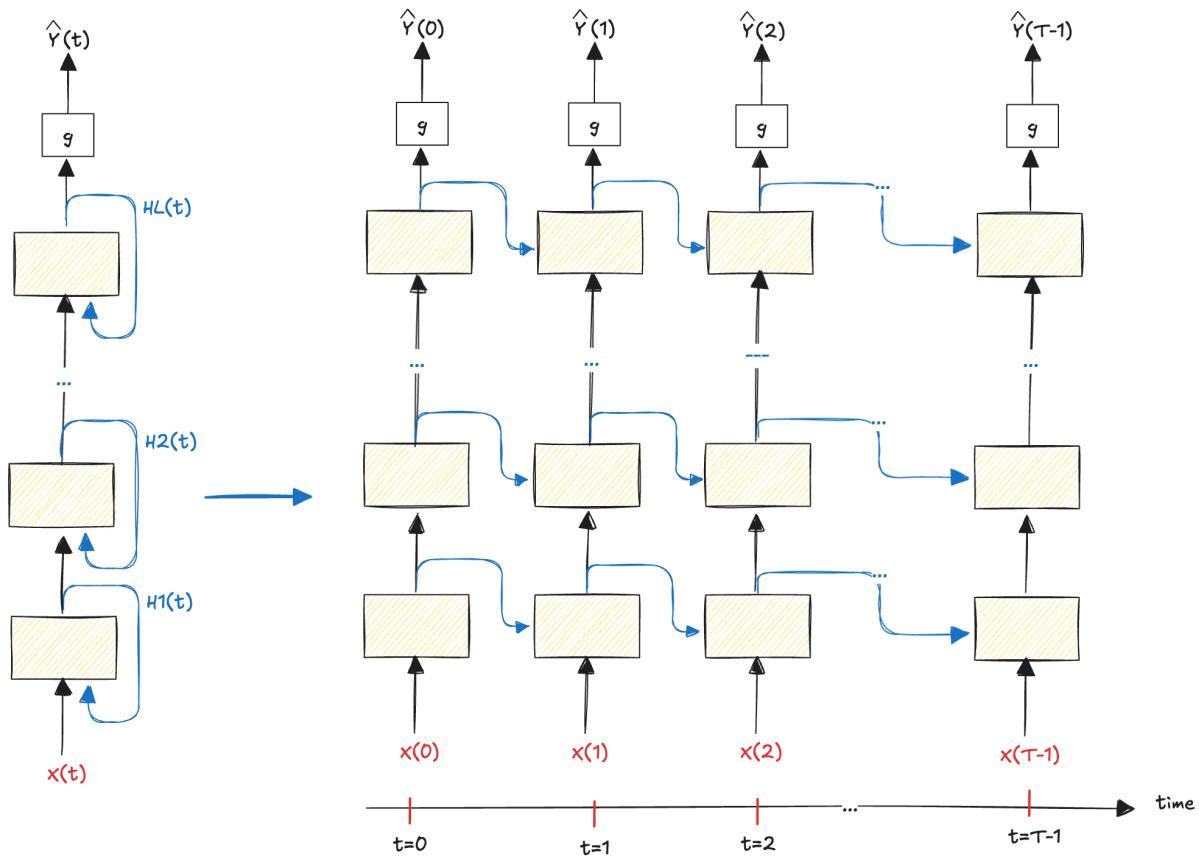
We can also provide the network with a **single input vector** and reuse it at every time step, allowing the network to generate a **sequence of outputs**. This architecture is known as a **vector-to-sequence** network:



In this setting, the same input vector is provided to the recurrent layer at each time step, while the hidden state evolves over time and produces a sequence of outputs. Alternatively, the input vector can be used to initialize the hidden state, with subsequent outputs generated purely through the recurrent dynamics; both formulations are commonly used in practice.

A common example is image captioning, where an image (or a feature vector extracted by a convolutional neural network) is given as input, and the recurrent network generates a sequence of words describing the image.

Lastly, we can consider an architecture composed of a sequence-to-vector network, called an **encoder**, followed by a vector-to-sequence network, called a **decoder**. This **encoder-decoder structure** allows the model to transform an input sequence into an output sequence of potentially different length.



The encoder processes the input sequence and compresses it into a single vector representation, given by its final hidden state. Intermediate encoder outputs are ignored. The final encoder hidden state is then used to initialize the hidden state of the decoder. The decoder generates an output sequence, one element at a time, using the encoded representation and its own recurrent dynamics. In practice, the decoder may also receive previous output tokens as inputs (**teacher forcing**). For example, this architecture can be used for translating a sentence from one language to another. We feed the network a sentence in one language, the encoder converts it into a single vector representation, and then the decoder decodes the vector into a sentence in another language.

1.3.4 Backpropagation Through Time

When working with sequential data, each prediction produced by a recurrent neural network depends not only on the current input, but also on the internal state built from all previous inputs. An intuitive way to think about this is to view the hidden state as a running summary of the past. If the network makes a poor prediction at time t , the problem may not lie only in the current input, but in how the hidden state was constructed at earlier time steps. Perhaps the

network failed to retain relevant information from the past, or assigned too much importance to irrelevant inputs. In this sense, an error observed at time t reveals that something went wrong earlier in the sequence, when the hidden state was being built.

During training, we therefore want the network to "learn from the future". When an error is observed at a later time step, we propagate that error backward through the sequence in order to adjust the parameters that were used at earlier steps. This allows the network to modify how it processes early inputs so that the hidden states they produce will be more useful for making accurate predictions later on.

Backpropagation Through Time (BPTT) formalizes this intuition. By unrolling the network through time, we obtain a feedforward computation network in which the same parameters appear at each time step. As in any feedforward network, we can then apply the chain rule to compute gradients. However, because the parameters are shared across time, the gradient with respect to each parameter is obtained by summing its contributions across all time steps in the unrolled network. In this way, the error at later outputs can be traced back through all intermediate hidden states, assigning responsibility to earlier computations. The network thus learns not only to predict well at each step, but to build internal representations that support accurate predictions in the future. Given an input sequence:

$$X(0), X(1), \dots, X(T-1)$$

the network performs a forward pass through the unrolled structure. At each time step, the hidden state is updated by combining the current input with the previous hidden state:

$$H(t) = \phi(X(t)W_x + H(t-1)W_h + b), \quad H(-1) = H_0$$

From the hidden state, the network produces an output. Here we explicitly include an linear implementation of the $g()$ function, defined by an output weight matrix and bias term:

$$\hat{Y}(t) = g(H(t)) = H(t)W_y + b_y$$

As the sequence is processed, this forward pass produces a sequence of hidden states:

$$H(0), H(1), \dots, H(T-1)$$

and a corresponding sequence of outputs

$$\hat{Y}(0), \hat{Y}(1), \dots, \hat{Y}(T-1)$$

Here the implementation of the forward pass through time:

```
# Phi activation function (ReLU)
def phi(z):
    return np.maximum(0.0, z)
```

```
# g output function (linear)
def g(H, W_y, b_y):
    return H @ W_y + b_y

def rnn_layer_forward(X, params):

    # Get dimensions (T: time steps, B: batch size)
    T, B, _ = X.shape

    # Get hidden dimension
    hidden_dim = params["b"].shape[0]

    # Get output dimension
    output_dim = params["b_y"].shape[0]

    # Initialize hidden states
    H = np.zeros((T, B, hidden_dim))

    # Initialize pre-activations
    Z = np.zeros_like(H)

    # Initialize outputs
    Y_hat = np.zeros((T, B, output_dim))

    # Initialize previous hidden state
    H_prev = np.zeros((B, hidden_dim))

    # Unroll through time
    for t in range(T):

        # Compute pre-activation
        Z[t] = (
            np.dot(X[t], params["W_x"]) +
            np.dot(H_prev, params["W_h"]) +
            params["b"]
        )
        # Apply activation function (ReLU)
        H[t] = phi(Z[t])

        # Compute output
        Y_hat[t] = g(H[t], params["W_y"], params["b_y"])
```

```

# Update previous hidden state
H_prev = H[t]

# Return hidden states, pre-activations, and outputs
return H, Z, Y_hat

```

The predictions produced by the recurrent network should be **evaluated using a loss function**.

Let

$$Y(0), Y(1), \dots, Y(T-1)$$

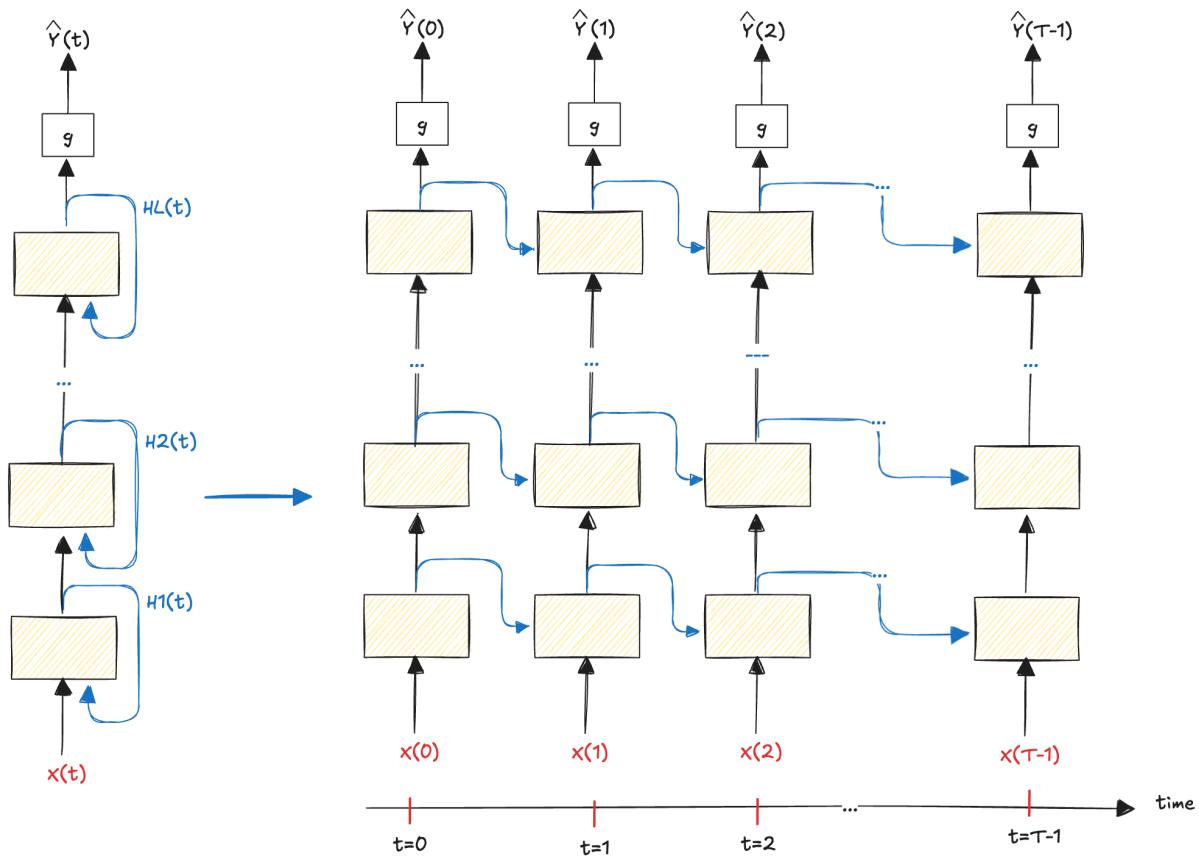
denote the target sequence associated with the input sequence. This sequence should represent **what we want the network to predict** at each time step. While the input sequence describes the information available to the model, the target sequence specifies the desired behavior of the model over time. In many problems, Y is naturally aligned with the input sequence. For example, in time-series forecasting, $Y(t)$ may represent the future value of the signal relative to the input at time t . In a one-step-ahead forecasting task, this corresponds to

$$Y(t) = X(t+1)$$

so that each prediction is trained to anticipate what will happen next. More generally, Y can be interpreted as a **supervisory signal** that tells the network what information extracted from the past should be useful. Depending on the task, the loss may involve all time steps (as in sequence-to-sequence problems) or only a subset of them (for example, only the final output in sequence-to-vector models). Let M denote the number of time steps over which supervision is applied. The training objective can then be written as

$$\mathcal{L} = \sum_{t=T-M}^{T-1} \ell(Y(t), \hat{Y}(t))$$

where $\ell()$ is the loss evaluated at each time step (for example, mean squared error for regression or cross-entropy for classification) and M is the number of supervised time steps. In the following figure, for example, the early part of the sequence is used to build up a meaningful hidden state, while only the later part of the sequence is used to evaluate prediction quality:



We can implement the loss function using the mean squared error between the predicted and target sequences, computed only over the last M supervised time steps, while earlier predictions are ignored:

```
def mse_loss(Y_hat, Y, M):
    # M = T → full sequence-to-sequence
    # M = 1 → sequence-to-vector
    # 1 < M < T → partial supervision

    T = Y_hat.shape[0]
    return 0.5 * np.sum((Y_hat[T-M:T] - Y[T-M:T]) ** 2)
```

To train the model, we now need analyze **how the loss accumulated over time depends on parameters that are reused at every time step**, requiring a form of backpropagation that operates across time. The gradient with respect to a parameter can be written as:

$$\frac{\partial \mathcal{L}}{\partial \omega} = \sum_{t=T-M}^{T-1} \frac{\partial \ell(Y(t), \hat{Y}(t))}{\partial \omega}$$

Applying the chain rule at a given time step t , we obtain:

$$\frac{\partial \ell(Y(t), \hat{Y}(t))}{\partial \omega} = \frac{\partial \ell(Y(t), \hat{Y}(t))}{\partial \hat{Y}(t)} \frac{\partial \hat{Y}(t)}{\partial H(t)} \frac{\partial H(t)}{\partial \omega}$$

The nontrivial term in this expression is the derivative of the hidden state with respect to the parameters:

$$\frac{\partial H(t)}{\partial \omega}$$

Unlike in feedforward networks, the hidden state at time t is not a simple function of the parameters. Because of the recurrence:

$$H(t) = \phi(X(t)W_x + H(t-1)W_h + b)$$

The hidden state depends on the parameters both **directly** at time t and **indirectly** through all previous hidden states, which themselves depend on the parameters. This temporal dependency must therefore be explicitly accounted for. Applying the chain rule yields the recursion:

$$\frac{\partial H(t)}{\partial \omega} = \frac{\partial \phi(\cdot)}{\partial \omega} + \frac{\partial \phi(\cdot)}{\partial H(t-1)} \frac{\partial H(t-1)}{\partial \omega}$$

where the first term captures the direct contribution of the parameters at the current time step, and the second term captures the indirect contribution propagated from the past through the recurrent connection. By recursively applying this relation backward in time, the derivative can be expressed as a sum of contributions from all earlier time steps. This recursive accumulation of gradients is precisely what is implemented by **Backpropagation Through Time (BPTT)** algorithm.

```
def rnn_layer_backward(X, Y, H, Z, Y_hat, params, M):

    # Get dimensions (T: time steps, B: batch size)
    T, B, D = X.shape

    # Get hidden and output dimension
    hidden_dim = params["W_h"].shape[0]
    output_dim = params["W_y"].shape[1]

    # Initialize gradients
    grads = {k: np.zeros_like(v) for k, v in params.items()}

    # Gradient flowing from the next time step
    dH_next = np.zeros((B, hidden_dim))

    # Backward pass through time
```

```

for t in reversed(range(T)):

    # dL/dY_hat_t for 0.5 * ||Y_hat - Y||^2 is just (Y_hat - Y)
    # Inject gradient from the loss only on the last M steps
    if t >= T - M:
        dY = Y_hat[t] - Y[t]
    else:
        dY = np.zeros((B, output_dim))

    # Output layer: Y_hat_t = H_t @ W_y + b_y
    grads["W_y"] += H[t].T @ dY
    grads["b_y"] += dY.sum(axis=0)

    # Backpropagate into hidden state
    dH = dY @ params["W_y"].T + dH_next

    # ReLU backward: H_t = ReLU(Z_t)
    dZ = dH * (Z[t] > 0)

    # Recurrent part: Z_t = X_t @ W_x + H_{t-1} @ W_h + b
    grads["b"] += dZ.sum(axis=0)
    grads["W_x"] += X[t].T @ dZ

    if t > 0:
        grads["W_h"] += H[t-1].T @ dZ

    # Gradient to propagate to previous time step: dL/dH_{t-1}
    dH_next = dZ @ params["W_h"].T

return grads

```

After computing the gradients, the model parameters are updated using Stochastic Gradient Descent, moving them in the direction of decreasing loss using the learning rate to control the step size:

$$w \leftarrow w - \eta \nabla_w L$$

```

def sgd_step(params, grads, lr=1e-3):
    for k in params:
        params[k] -= lr * grads[k]

```

Finally, we define a function that performs **a single training step**. Given an input sequence,

it first runs a forward pass through time to compute the hidden states and the corresponding predictions at each time step. These predictions are then compared with the target sequence to evaluate the loss over the supervised time window. Next, BPTT is used to compute the gradients of the loss with respect to all model parameters, accounting for the fact that the same parameters are reused across time. The parameters are then updated with a single stochastic gradient descent step. The function returns the loss value, which provides a convenient scalar metric for monitoring training progress.

```
def rnn_layer_train_step(X, Y, params, M, lr=1e-3):

    # Forward pass through time
    H, Z, Y_hat = rnn_layer_forward(X, params)

    # Compute loss over the last M time steps
    loss = mse_loss(Y_hat, Y, M)

    # Backward pass through time (BPTT)
    grads = rnn_layer_backward(X, Y, H, Z, Y_hat, params, M)

    # Update parameters with SGD
    sgd_step(params, grads, lr)

    # Return scalar loss
    return loss
```

In practice, additional complications arise because **input sequences can be very long**. This creates challenges both from a **computational** standpoint, due to the large amount of memory required to store intermediate activations, and from an **optimization** standpoint, due to numerical instabilities in the gradient computation. Information originating from early time steps must pass through many successive matrix multiplications before influencing later outputs, and an equally long chain of matrix products appears when gradients are propagated backward through time. As a result, gradients may either **vanish** or **explode**, making learning difficult.

To mitigate these issues, it is common practice to truncate the unrolling of the network during backpropagation. Rather than propagating gradients all the way back to the beginning of the sequence, **the backward pass is limited to a fixed number of time steps**. This hyperparameter, known as the **truncation length**, controls how far into the past error signals are allowed to propagate.

A larger **truncation length** enables the model to capture longer-range temporal dependencies,

but comes at the cost of increased memory usage and computational complexity. Conversely, a **smaller truncation length** reduces resource requirements and improves numerical stability, but may prevent the network from learning from distant past information. In practice, the truncation length is treated as a hyperparameter and chosen based on the characteristics of the data and the task, **balancing representational power against computational efficiency**.

1.3.5 Applying RNNs to Time Series Forecasting

We can use the RNN architecture to build a model for time-series forecasting, where the objective is to predict future values of a sequence based on its past observations. In this setting, we adopt a **sequence-to-sequence formulation**: the network receives a window of past ridership values as input and produces a sequence of **one-step-ahead predictions**, so that at each time step it learns to forecast the next value in the series.

To ensure a fair comparison with the previously discussed naive baseline and SARIMA model, we use the same evaluation period as a test set, while earlier data are split into training and validation periods. Before training the model, the ridership values are **normalized** by applying a simple rescaling. This brings the data into a numerically convenient range, which improves the stability of training, reduces the risk of excessively large gradients, and allows the optimization to focus on learning temporal patterns rather than compensating for large magnitudes. Importantly, this normalization does not change the temporal structure of the series and can be inverted after forecasting, so that predictions can be reported in the original units.

```
start_date = "2019-03-01"
end_date = "2019-05-31"

# Test data
rail_test = df["rail"][start_date:end_date] / 1e6

# Split other data into training and validation periods
rail_train = df["rail"]["2016-01-01":"2018-12-31"] / 1e6
rail_valid = df["rail"]["2019-01-01":"2019-02-28"] / 1e6
```

To apply the RNN model, the time series must first be transformed into **input-target sequences**. Each input sequence consists of the ridership values observed over the past T days, while the corresponding target sequence consists of the same window shifted one step into the future. In other words, for an input sequence:

$$x(t), x(t+1), \dots, x(t+T-1)$$

the target sequence is:

$x(t+1), x(t+2), \dots, x(t+T)$

In this formulation, the RNN processes a window of historical observations and learns to predict the next value at every time step, providing dense supervision across the entire sequence.

The choice of T determines how much historical information the model can exploit. It should be large enough to capture the dominant temporal patterns in the data, such as weekly seasonality, but not so large as to introduce unnecessary model complexity or unstable gradient propagation during training. In practice, a reasonable starting point for T is one or two full seasonal cycles (for example, $T = 7$ or $T = 14$ days for weekly patterns), and the optimal value can be selected based on validation performance.

The following function illustrates how a time series can be translated into a collection of sequence-to-sequence input-target pairs suitable for training an RNN:

```
import numpy as np

def make_sequences(series, T):

    # Ensure series is a numpy array
    series = np.asarray(series)

    # Sequence and target lists
    X, Y = [], []

    # Create input sequences and corresponding target value
    for i in range(len(series) - T):

        # Input: T past values
        X.append(series[i:i+T])

        # Target: next values
        Y.append(series[i+1:i+T+1])

    # Convert lists to numpy arrays
    X = np.array(X)[..., None]  # (num_sequences, N, 1)
    Y = np.array(Y)[..., None]  # (num_sequences, 1)

    # Return input sequences and targets
    return X, Y
```

We use this function to transform the time series into **training**, **validation**, and **test datasets**, each composed of **input sequences of past observations** and their corresponding **target values**

to be predicted:

```
# Define sequence length
T = 14

# Create input-target sequences for training, validation, and test sets
X_train, Y_train = make_sequences(rail_train, T)
X_valid, Y_valid = make_sequences(rail_valid, T)
X_test, Y_test = make_sequences(rail_test, T)
```

We use a recurrent neural network composed of **a single recurrent layer** followed by a linear output layer. Since the ridership series is one-dimensional, each time step provides a single scalar input and the input dimension is therefore **d=1**. At each time step, the recurrent layer updates a hidden state of dimension **m=20**, which acts as a distributed memory summarizing the relevant information extracted from past observations. The output dimension is **p=1**, because the model produces a scalar prediction at each time step.

In the sequence-to-sequence formulation adopted here, this forecast represents the next-day ridership value at each time step within the input window. All model parameters (input-to-hidden weights, recurrent weights, recurrent bias, output weights, and output bias) are initialized randomly and then learned during training.

```
# Define RNN model dimensions
input_dim = 1
hidden_dim = 20
output_dim = 1

# Initialize RNN layer parameters
params = init_rnn_layer(input_dim, hidden_dim, output_dim)
```

We implement a **mini-batch iterator** to train the RNN efficiently. Given a set of input sequences and their corresponding target sequences, the iterator first constructs an array of sample indices and **randomly shuffles** it in order to randomize the order of the sequences seen during training. Importantly, this shuffling is performed **across sequences**, not within sequences: the temporal ordering inside each individual sequence is preserved, since maintaining the correct time structure is essential for recurrent models.

After shuffling, the dataset is partitioned into consecutive mini-batches of a specified size. For each mini-batch, the iterator yields a pair consisting of a batch of input sequences and the corresponding batch of target sequences, which can then be processed in parallel during the forward and backward passes.

```
def iterate_minibatches(X, Y, batch_size):

    # Create an array of indices and shuffle them
    indices = np.arange(len(X))
    np.random.shuffle(indices)

    # Generate mini-batches
    for start in range(0, len(X), batch_size):
        batch_idx = indices[start:start+batch_size]
        yield X[batch_idx], Y[batch_idx]
```

Finally we can implement the **training loop**. The model is trained for a fixed number of epochs, where each epoch corresponds to one full pass over the training dataset. At each epoch, the training data are randomly partitioned into mini-batches of fixed size. For each mini-batch, the input and target tensors are transposed into the format expected by the RNN. A forward pass through the unrolled network is then performed to compute predictions at all time steps, the loss is evaluated over the supervised portion of the sequence (corresponding to the last M time steps), and the model parameters are updated by applying BPTT followed by a stochastic gradient descent step. In the sequence-to-sequence setting considered here, the network produces a prediction at every time step, and supervision can be applied either at all time steps or only over the final part of the sequence, depending on the choice of M.

```
# Number of training epochs
epochs = 500

# Minibatch size
batch_size = 32

# Learning rate
lr = 1e-3

# Sequence length and loss window
# (full sequence-to-sequence supervision)
T = X_train.shape[1]
M = T

# List to store losses for the epoch
losses = []

# Training loop
for epoch in range(epochs):
```

```

# Iterate over mini-batches
for Xb, Yb in iterate_minibatches(X_train, Y_train, batch_size):

    # Transpose to time-major format: (T, B, 1)
    Xb = Xb.transpose(1, 0, 2)
    Yb = Yb.transpose(1, 0, 2)

    # Perform a training step
    loss = rnn_layer_train_step(Xb, Yb, params, M, lr=lr)

    losses.append(loss)

# Print epoch loss
print(f"\rEpoch {epoch+1:02d}/{epochs} | Train loss: {np.mean(losses):.6f}",
      end="", flush=True)

```

Epoch 500/500 | Train loss: 2.228523

We evaluate the trained RNN on the held-out test period by generating one-step-ahead forecasts for every test time and comparing them with the corresponding ground truth. Forecasting is performed in a rolling fashion: for each test time t we build an input window containing the T observations immediately preceding t , feed that window through the trained network, and retain the model's last output $\hat{y}(t)$ as the one-step forecast. Because each forecast uses the preceding T days, the very first test forecast requires the T days immediately before the test start; forecasting therefore cannot start from the first test day without historical context. This procedure guarantees that every prediction is produced using only information that would have been available at decision time.

```

# Full rail series up to end_date for rolling predictions (as numpy array)
rail_full = (df["rail"][:end_date] / 1e6).values

# Predictions list
y_preds_rnn = []

for t in range(len(rail_test)):

    # Index in the full series corresponding to this test time
    idx = len(rail_full) - len(rail_test) + t

```

```

# Extract last T observations before time t
x_input = rail_full[idx - T:idx]

# Shape to (T, B=1, input_dim=1)
x_input = x_input.reshape(T, 1, 1)

# Forward pass
_, _, Y_hat = rnn_layer_forward(x_input, params)

# One-step-ahead prediction
y_pred = Y_hat[-1, 0, 0]
y_preds_rnn.append(y_pred)

```

We compute the MAE between the RNN forecasts and the true ridership values on the test set in order to assess the model's predictive accuracy and to compare its performance with the previously discussed naive baseline and SARIMA models. All forecasts are first mapped back to the original (**unnormalized**) scale before computing error statistics:

```

# Use the test index
test_index = df.loc[start_date:end_date].index

# Convert predictions to a Pandas Series and unnormalize
y_preds_rnn = pd.Series(y_preds_rnn, index=test_index)
y_preds_rnn_unscaled = y_preds_rnn * 1e6

# True values as a Pandas Series
y_true_unscaled = df.loc[start_date:end_date, "rail"]

# Compute MAE
mae_rnn = np.mean(np.abs(y_preds_rnn_unscaled - y_true_unscaled))

# Print MAE
print("Naive MAE:", mae_naive)
print("SARIMA MAE:", mae_sarima)
print("RNN MAE:", mae_rnn)

```

```

Naive MAE: 42143.27173913043
SARIMA MAE: 32040.720095473138
RNN MAE: 31045.809812056777

```

The rail ridership time series exhibits a very strong and highly regular weekly seasonality, com-

bined with smooth dynamics and limited nonlinear structure. In such settings, **simple models that explicitly encode this regularity are often extremely difficult to outperform**. The naive forecasting strategy, which directly copies the value from one week earlier, effectively hard-codes the correct inductive bias about the data-generating process. In contrast, the RNN must **learn** this behavior implicitly from data through gradient-based optimization. Even though the RNN is, in principle, a far more expressive model, this makes **the learning problem substantially harder**. Moreover, in this experiment, we focus on demonstrating the modeling and training procedure, and **we do not explicitly address potential overfitting**; in practice, techniques such as validation-based early stopping, regularization, or dropout would be required to control model complexity and improve generalization.

It is worth noting that the performance of the RNN **depends on the choice of the sequence length T and on the loss window M** used during training. Varying T changes how much historical context is provided to the model, while varying M determines how many of the model's predictions are directly supervised. In principle, increasing T allows the network to access longer-term information, whereas reducing M can focus learning on the most recent predictions. However, in practice, these choices often involve a **delicate trade-off between representational capacity, optimization stability, and effective gradient propagation**. Exploring different values of T and M can therefore provide useful insight into the model's behavior.

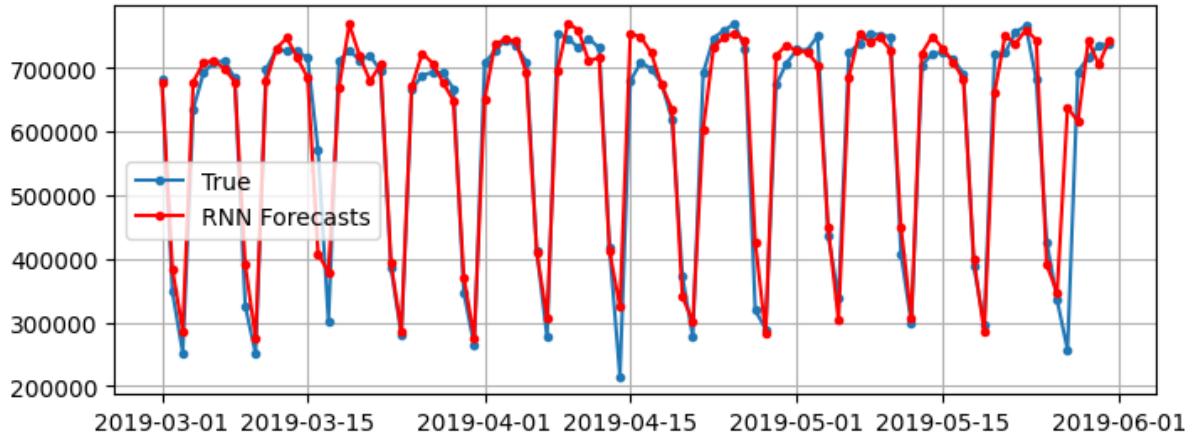
We can plot the forecasts produced by the RNN alongside the actual observed values to visually assess its performance:

```
fig, ax = plt.subplots(figsize=(8, 3))

ax.plot(rail_series.loc[time_period].index,
        rail_series.loc[time_period].values,
        marker=".", label="True",
        )

ax.plot(y_preds_rnn_unscaled.index,
        y_preds_rnn_unscaled.values,
        color="red", marker=".", label="RNN Forecasts"
        )

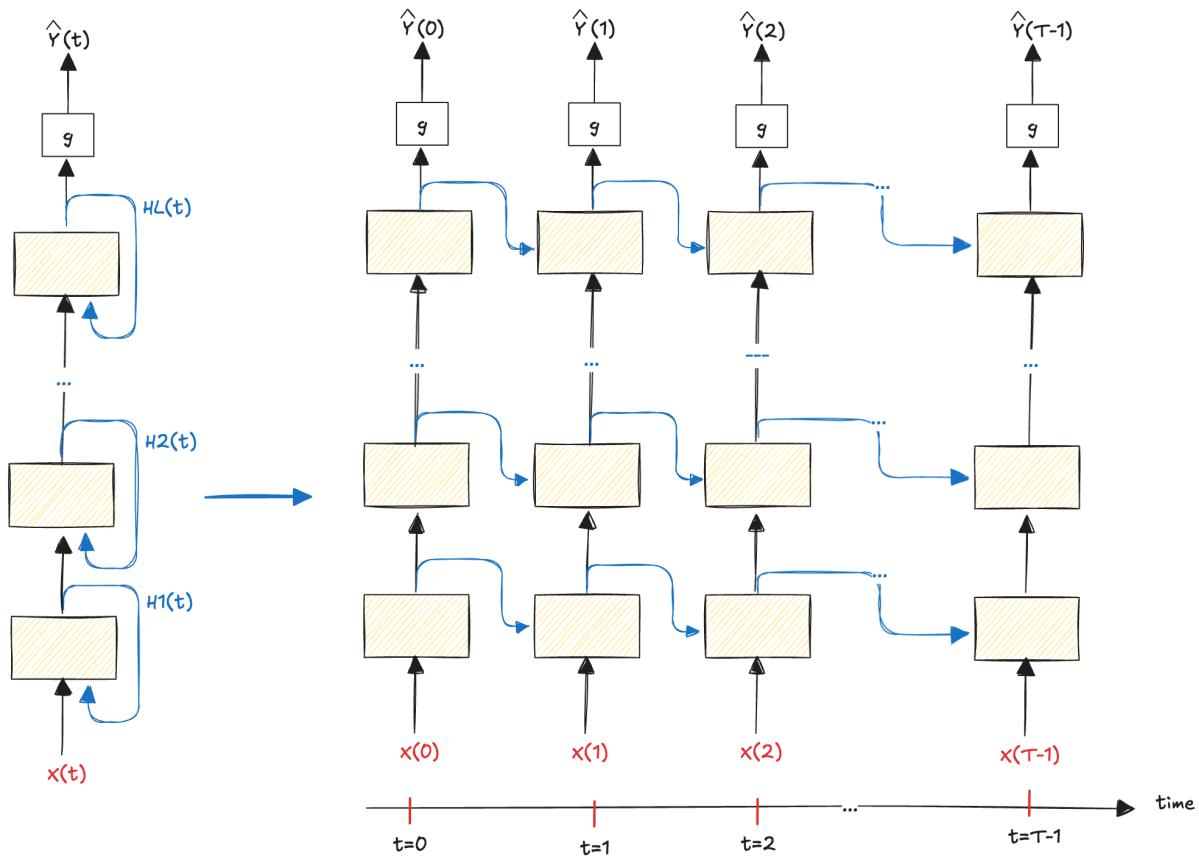
ax.grid(True)
ax.legend()
plt.show()
```



From a modeling perspective, the network used in this experiment is a **vanilla recurrent neural network**. Its practical ability to exploit long-term dependencies is limited by the well-known **vanishing gradient problem**. As gradients are propagated backward through time, they are repeatedly multiplied by the recurrent weight matrix and by the derivative of the activation function, which causes their magnitude to **shrink rapidly**. As a consequence, the effective temporal memory of the network is often much shorter than the nominal sequence length, making it **difficult for a simple RNN to reliably capture seasonal patterns** that extend over many time steps.

1.3.6 Deep networks

It is common to increase the representational capacity of recurrent networks by **stacking multiple recurrent layers**. In a stacked RNN, each layer contains its own hidden state and is unrolled through time in the same way as a single-layer model:



At a given time step, the first recurrent layer processes the input and produces a hidden representation. This representation is then fed as input to the second recurrent layer, producing a second hidden representation, and so on up to the top layer. In other words, information flows forward in time within each layer and upward across layers at each time step, while the parameters within each layer are shared across all time steps. Formally, the computation can be written as:

$$H^{(\ell)}(t) = \phi(H^{(\ell-1)}(t)W_x^{(\ell)} + H^{(\ell)}(t-1)W_h^{(\ell)} + b^{(\ell)})$$

with the output typically computed from the top-layer state:

$$\hat{Y}(t) = g(H^{(L)}(t))$$

This stacked structure allows higher layers to build **progressively more abstract temporal features**, often improving performance when the underlying dynamics are complex.

We now implement the stacked recurrent neural network using **PyTorch**, which provides built-in support for this architecture through the **nn.RNN module**, where the number of stacked layers is controlled by a single parameter:

```
import torch
import torch.nn as nn

# Input size, just the rail ridership
input_size = 1

# Number of hidden units in each RNN layer
hidden_size = 32

# Number of recurrent layers
num_layers = 3

# Output size. just the rail ridership prediction
output_size = 1

# Define the recurrent layers
deep_rnn = nn.RNN(input_size=input_size,
                   hidden_size=hidden_size,
                   num_layers=num_layers,
                   nonlinearity="tanh",
                   batch_first=True      # inputs: (B, T, D)
)

# Define the fully connected layer
fc = nn.Linear(hidden_size, output_size)

# Define the optimizer
optimizer = torch.optim.Adam(list(deep_rnn.parameters()) +
                             list(fc.parameters()),
                             lr=1e-3)

# Define the forward function
def forward(x):
    out, _ = deep_rnn(x)
    y_hat = fc(out)
    return y_hat
```

We need to write the loss we already defined using the PyTorch tensors instead of NumPy arrays:

```
def mse_loss_torch(y_hat, y, M):
    return 0.5 * ((y_hat[:, -M:, :] - y[:, -M:, :]) ** 2).sum()
```

We now define a function that performs a single training step for the deep network: given a mini-batch of input-target sequences, it executes a forward pass, computes the loss over the supervised time steps, propagates gradients backward through time, and updates the model parameters:

```
def train_step(Xb, Yb, deep, fc, M):

    # Set the model to training mode
    deep.train()
    fc.train()

    # Prepare mini-batch tensors
    xb = torch.as_tensor(Xb, dtype=torch.float32)
    yb = torch.as_tensor(Yb, dtype=torch.float32)

    # Clear previous gradients
    optimizer.zero_grad()

    # Forward pass
    y_hat = forward(xb)

    # Compute loss over the last M time steps
    loss = mse_loss_torch(y_hat, yb, M)

    # Backward pass
    loss.backward()

    # Update parameters
    optimizer.step()

    # Return scalar loss
    return float(loss.detach())
```

We can now train the stacked RNN by iterating over the training data for multiple epochs and applying the training step to each mini-batch, monitoring the average training loss at the end of each epoch:

```

# Number of training epochs
epochs = 80

# Minibatch size
batch_size = 32

# Loss window
M = T

# List to store losses
losses = []

# Training loop
for epoch in range(1, epochs + 1):

    for Xb, Yb in iterate_minibatches(X_train, Y_train, batch_size):
        loss = train_step(Xb, Yb, deep_rnn, fc, M)
        losses.append(loss)

    # Print epoch loss
    print(f"\rEpoch {epoch+1:02d}/{epochs} | Train loss: {np.mean(losses):.6f}",
          end="", flush=True)

```

Epoch 81/80 | Train loss: 2.8855038

We now evaluate the trained deep RNN on the test set by rolling the model forward to generate one-step-ahead forecasts at each test time step and computing the MAE on the original ridership scale, in direct comparison with the naive, SARIMA and simple RNN baselines:

```

# Predictions list
y_preds_deep_rnn = []

# Prepare the model for evaluation
deep_rnn.eval()
fc.eval()

with torch.no_grad():
    for t in range(len(rail_test)):

        # Index of current test time in rail_full
        idx = len(rail_full) - len(rail_test) + t

```

```

# # Extract last T observations before time t
x_input = rail_full[idx - T:idx]

# To tensor (B=1, T, 1)
x_tensor = torch.tensor(x_input.reshape(1, T, 1), dtype=torch.float32)

# Forward pass
Y_hat = forward(x_tensor)

# One-step-ahead prediction
y_pred = Y_hat[0, -1, 0].item()
y_preds_deep_rnn.append(y_pred)

# Convert to numpy and unnormalize to original units
y_preds_deep_rnn = np.array(y_preds_deep_rnn)
y_preds_deep_rnn_unscaled = y_preds_deep_rnn * 1e6

# Compute MAE
mae_deep_rnn = np.mean(np.abs(y_preds_deep_rnn_unscaled - y_true_unscaled))

# Print MAE
print("Naive MAE : ", mae_naive)
print("SARIMA MAE: ", mae_sarima)
print("RNN MAE   : ", mae_rnn)
print("Deep RNN MAE: ", mae_deep_rnn)

```

```

Naive MAE : 42143.27173913043
SARIMA MAE: 32040.720095473138
RNN MAE   : 31045.809812056777
Deep RNN MAE: 30394.18942018177

```

Stacking multiple recurrent layers does not outperform SARIMA or the simple RNN, highlighting that increased model complexity alone is not sufficient to surpass well-specified statistical models when the underlying data structure is simple and highly regular.

Overall, these results reinforce a key lesson in time-series forecasting: **model choice should be guided by the structure of the data rather than by model complexity alone**. While deep learning models offer great flexibility, classical approaches that explicitly encode known regularities can remain superior in settings dominated by strong seasonality and stable dynamics.

1.4 Extending to Multivariate and Multi-Step Forecasting

So far, we have focused on forecasting a univariate time series using one-step-ahead predictions. While this setting is useful for introducing recurrent neural networks and for understanding their strengths and limitations, many **real-world forecasting problems** are more complex. In practice, predictions often depend on **multiple correlated time series**, and the goal is frequently to **forecast several future time steps simultaneously** rather than a single value.

1.4.1 Multivariate Input

A **multivariate time series** consists of **multiple interrelated sequences** observed over time. At each time step, the data are represented by a **vector** of observations, where each component corresponds to a different variable or feature. These variables often capture **complementary aspects** of the underlying process and may provide additional predictive power when **modeled jointly**.

For example, in the context of the public transportation ridership example, we may observe separate time series for bus and rail usage. While each series reflects distinct travel behaviors, they are clearly related and may influence one another over time.

From a modeling perspective, multivariate time series can be handled with minimal changes to the recurrent network architecture. The main difference lies in the input representation: **instead of feeding a single scalar at each time step, the network receives a vector of features**.

To illustrate this setting, we consider forecasting the rail ridership time series while using both bus and rail ridership as inputs to the model. This allows the RNN to **exploit cross-dependencies** between variables when generating forecasts.

```
df_mulvar = df[["bus", "rail"]] / 1e6 # use both bus & rail series as input
```

```
df_mulvar.head()
```

	bus	rail
date		
2001-01-01	0.297192	0.126455
2001-01-02	0.780827	0.501952
2001-01-03	0.824923	0.536432
2001-01-04	0.870021	0.550011
2001-01-05	0.890426	0.557917

We split the multivariate time series into training, validation, and test sets using the same chronological boundaries adopted in the univariate experiments, ensuring a fair and consistent evaluation of the multivariate forecasting model:

```
mulvar_train = df_mulvar["2016-01":"2018-12-31"]
mulvar_valid = df_mulvar["2019-01-01":"2019-02-28"]
mulvar_test = df_mulvar["2019-03-01":"2019-05-31"]
```

We now define a **stacked recurrent neural network for multivariate forecasting**, where each time step is represented by a vector of inputs (bus and rail ridership), and the network produces a sequence of predictions for the target variable through a linear readout applied to the hidden states:

```
# Multivariate input: [bus, rail]
input_size = 2

# Number of hidden units in each RNN layer
hidden_size = 32

# Number of recurrent layers
num_layers = 3

# Output size: we predict rail only
output_size = 1

# Define the recurrent layers
deep_multi_rnn = nn.RNN(
    input_size=input_size,
    hidden_size=hidden_size,
    num_layers=num_layers,
    nonlinearity="tanh",
    batch_first=True      # (B, T, D)
)

# Define the fully connected layer
fc = nn.Linear(hidden_size, output_size)

# Define the optimizer
optimizer = torch.optim.Adam(list(deep_multi_rnn.parameters()) +
                             list(fc.parameters()),
                             lr=1e-3)
```

```
# Define the forward function
def forward_multi(x):
    out, _ = deep_multi_rnn(x)
    y_hat = fc(out)
    return y_hat
```

Notice that the only difference with respect to the univariate model is the **input dimensionality**. At each time step, the model now receives two input features instead of one: one for bus ridership and one for rail ridership. In that case we output only the rail ridership forecast, however the model could easily be extended to **produce forecasts for both series simultaneously** by adjusting the output layer accordingly. The train step is equivalent to the previous one, since the architecture is the same and

```
# Number of training epochs
epochs = 80

# Minibatch size
batch_size = 32

# Loss window
M = T

# List to store losses
losses = []

# Training loop
for epoch in range(1, epochs + 1):

    for Xb, Yb in iterate_minibatches(mulvar_train, Y_train, batch_size):
        loss = train_step(Xb, Yb, deep_multi_rnn, fc, M)
        losses.append(loss)

    # Print epoch loss (overwrite same line)
    print(
        f"\rEpoch {epoch:03d}/{epochs} | Train loss: {np.mean(losses):.6f}",
        end="",
        flush=True
    )
```

KeyError
Cell In[2188], line 16

Traceback (most recent call last)

```

13 # Training loop
14 for epoch in range(1, epochs + 1):
---> 16     for Xb, Yb in iterate_minibatches(mulvar_train, Y_train, batch_size):
17         loss = train_step(Xb, Yb, deep_multi_rnn, fc, M)
18         losses.append(loss)

```

```

Cell In[2177], line 10, in iterate_minibatches(X, Y, batch_size)
    8 for start in range(0, len(X), batch_size):
    9     batch_idx = indices[start:start+batch_size]
---> 10     yield X[batch_idx], Y[batch_idx]

```

```

File /opt/homebrew/Caskroom/miniconda/base/envs/machine-learning/lib/python3.13/site-
packages/pandas/core/frame.py:4113, in DataFrame.__getitem__(self, key)
4111     if is_iterator(key):
4112         key = list(key)
-> 4113     indexer = self.columns._get_indexer_strict(key, "columns")[1]
4115 # take() does not accept boolean indexers
4116 if getattr(indexer, "dtype", None) == bool:

```

```

File /opt/homebrew/Caskroom/miniconda/base/envs/machine-learning/lib/python3.13/site-
packages/pandas/core/indexes/base.py:6212, in Index._get_indexer_strict(self, key, axis_name)
6209 else:
6210     keyarr, indexer, new_indexer = self._reindex_non_unique(keyarr)
-> 6212 self._raise_if_missing(keyarr, indexer, axis_name)
6214 keyarr = self.take(indexer)
6215 if isinstance(key, Index):
6216     # GH 42790 - Preserve name from an Index

```

```

File /opt/homebrew/Caskroom/miniconda/base/envs/machine-learning/lib/python3.13/site-
packages/pandas/core/indexes/base.py:6261, in Index._raise_if_missing(self, key, indexer, ax
6259 if nmissing:
6260     if nmissing == len(indexer):
-> 6261         raise KeyError(f"None of [{key}] are in the [{axis_name}]")
6263     not_found = list(ensure_index(key)[missing_mask.nonzero()[0]].unique())
6264     raise KeyError(f"{not_found} not in index")

```

KeyError: "None of [Index([87, 192, 1058, 722, 890, 328, 664, 82, 812, 476, 925, 789, \

In general, using a single model for multiple related tasks often can results in better performance than using a separate model for each task, since features learned for one task may be useful for the other tasks, and also because having to perform well across multiple tasks prevents the model from overfitting (it's a form of regularization). Not in that case.

1.4.2 Forecasting Several Steps Ahead

So far we have only predicted the value at the next time step, and we could just as easily predict the value several steps ahead by changing the target appropriately (for example, to predict the ridership 2 weeks from now, we could just change the target to be the value 14 days ahead instead of 1 day ahead). But what if we want to predict the next N values?

The first option is called **recursive multi-step forecasting**: take the sequence-to-vector model, make it predict the next value, add that value to the inputs, and acting as if the predicted value had actually occurred, finally, use the model again to predict the following value, and so on:

```
import numpy as np

X = rail_valid.to_numpy()[:, np.newaxis, :, np.newaxis]
for step_ahead in range(14):
    y_pred_one = model.predict(X)
    X = np.concatenate([X, y_pred_one.reshape(1, 1, 1)], axis=1)
```

```
1/1 [=====] - 0s 300ms/step
1/1 [=====] - 0s 280ms/step
1/1 [=====] - 0s 24ms/step
1/1 [=====] - 0s 22ms/step
1/1 [=====] - 0s 24ms/step
1/1 [=====] - 0s 23ms/step
1/1 [=====] - 0s 26ms/step
1/1 [=====] - 0s 25ms/step
1/1 [=====] - 0s 29ms/step
1/1 [=====] - 0s 23ms/step
1/1 [=====] - 0s 25ms/step
```

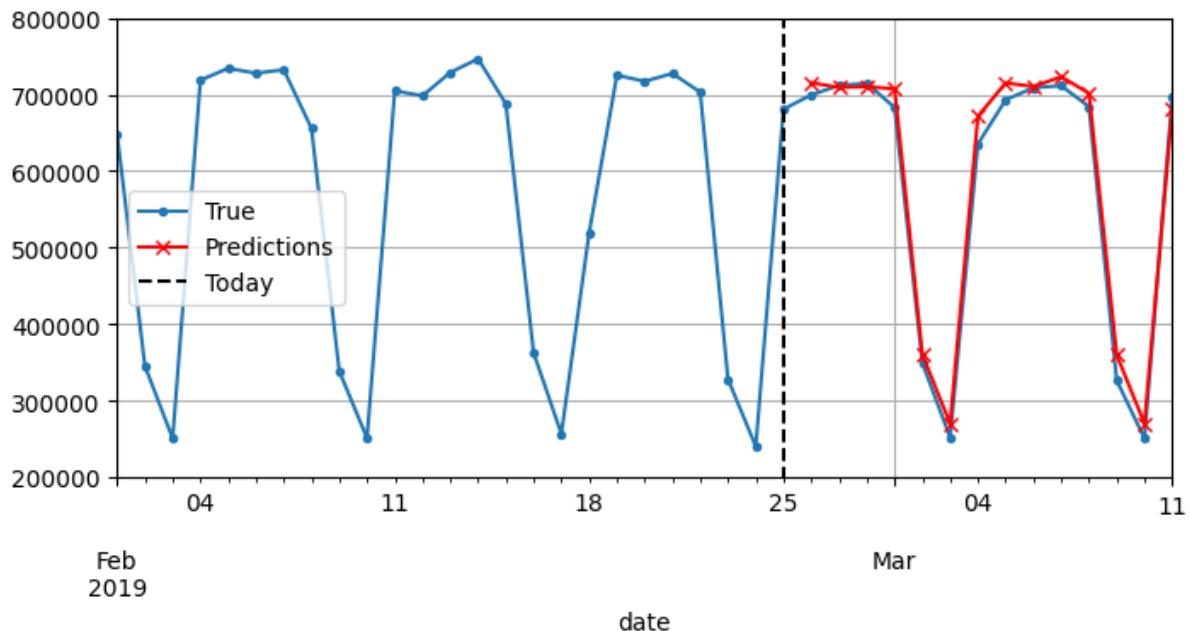
```
1/1 [=====] - 0s 24ms/step
1/1 [=====] - 0s 23ms/step
1/1 [=====] - 0s 24ms/step
```

We can plot the resulting forecasts, for example for the rail:

```
# Forecasts from 2019-02-26 (the 57th day of 2019) to 2019-03-11 (14 days in total)
Y_pred = pd.Series(X[0, -14:, 0], index=pd.date_range("2019-02-26", "2019-03-11"))

fig, ax = plt.subplots(figsize=(8, 3.5))
(rail_valid * 1e6)[["2019-02-01": "2019-03-11"].plot(label="True", marker=".", ax=ax)
(Y_pred * 1e6).plot(label="Predictions", grid=True, marker="x", color="r", ax=ax)
ax.vlines("2019-02-25", 0, 1e6, color="k", linestyle="--", label="Today")
ax.set_ylim([200_000, 800_000])
plt.legend(loc="center left")

plt.show()
```



Notice that if the model makes an error at one time step, then the forecasts for the following time steps are impacted as well, so the **errors tend to accumulate**. So, it's preferable to use this technique only for a small number of steps.

The second option is the **direct multi-step forecasting**: we can still use a sequence-to-vector model, but we train the RNN to predict the next N values in one shot:

However, we first **need to change the dataset in order to have targets as vectors** containing the next N values. To do this, we can use again `timeseries_dataset_from_array()`, but this time asking it to create datasets with longer sequences (`seq_length + N`) and without targets (`targets=None`). Then we can use `map()` method to apply a custom function to each batch of sequences, splitting them into inputs (`seq_length`) and targets (`N`).

```
N = 14

def split_inputs_and_targets(mulvar_series, ahead=14, target_col=1):
    return mulvar_series[:, :-ahead], mulvar_series[:, -ahead:, target_col]

ahead_train_ds = tf.keras.utils.timeseries_dataset_from_array(
    mulvar_train.to_numpy(),
    targets=None,
    sequence_length=seq_length + 14,
    batch_size=32,
    shuffle=True,
).map(split_inputs_and_targets)

ahead_valid_ds = tf.keras.utils.timeseries_dataset_from_array(
    mulvar_valid.to_numpy(),
    targets=None,
    sequence_length=seq_length + 14,
    batch_size=32
).map(split_inputs_and_targets)
```

Now we just need the output layer to have N units instead of 1:

```
ahead_model = tf.keras.Sequential([
    tf.keras.layers.SimpleRNN(32, input_shape=[None, 2]),
    tf.keras.layers.Dense(14)
])

opt = tf.keras.optimizers.SGD(learning_rate=0.02, momentum=0.9)
ahead_model.compile(loss=tf.keras.losses.Huber(), optimizer=opt, metrics=["mae"])

history = ahead_model.fit(ahead_train_ds, validation_data=ahead_valid_ds, epochs=500,
                           callbacks=[early_stopping_cb])
```

Epoch 1/500

```
33/33 [=====] - 1s 13ms/step - loss: 0.0628 -  
mae: 0.2620 - val_loss: 0.0236 - val_mae: 0.1694
```

```
2023-07-10 13:51:46.505127: I tensorflow/core/common_runtime/executor.cc:1197] [/device:CPU  
[{{node Placeholder/_0}}]]  
2023-07-10 13:51:46.505515: I tensorflow/core/common_runtime/executor.cc:1197] [/device:CPU  
[{{node Placeholder/_13}}]]
```

Epoch 2/500

```
33/33 [=====] - 0s 8ms/step - loss: 0.0193 -  
mae: 0.1656 - val_loss: 0.0177 - val_mae: 0.1474
```

Epoch 3/500

```
33/33 [=====] - 0s 8ms/step - loss: 0.0171 -  
mae: 0.1563 - val_loss: 0.0162 - val_mae: 0.1426
```

Epoch 4/500

```
33/33 [=====] - 0s 8ms/step - loss: 0.0159 -  
mae: 0.1504 - val_loss: 0.0152 - val_mae: 0.1335
```

Epoch 5/500

```
33/33 [=====] - 0s 8ms/step - loss: 0.0148 -  
mae: 0.1440 - val_loss: 0.0136 - val_mae: 0.1299
```

Epoch 6/500

```
33/33 [=====] - 0s 8ms/step - loss: 0.0137 -  
mae: 0.1375 - val_loss: 0.0124 - val_mae: 0.1252
```

Epoch 7/500

```
33/33 [=====] - 0s 8ms/step - loss: 0.0127 -  
mae: 0.1311 - val_loss: 0.0113 - val_mae: 0.1172
```

Epoch 8/500

```
33/33 [=====] - 0s 7ms/step - loss: 0.0117 -  
mae: 0.1248 - val_loss: 0.0103 - val_mae: 0.1114
```

Epoch 9/500

```
33/33 [=====] - 0s 8ms/step - loss: 0.0108 -  
mae: 0.1184 - val_loss: 0.0092 - val_mae: 0.1066
```

Epoch 10/500

```
33/33 [=====] - 0s 8ms/step - loss: 0.0099 -  
mae: 0.1123 - val_loss: 0.0083 - val_mae: 0.1008
```

Epoch 11/500

```
33/33 [=====] - 0s 7ms/step - loss: 0.0092 -  
mae: 0.1067 - val_loss: 0.0073 - val_mae: 0.0962
```

```
Epoch 12/500
33/33  [=====] - 0s 8ms/step - loss: 0.0085 -
mae: 0.1017 - val_loss: 0.0066 - val_mae: 0.0915
Epoch 13/500
33/33  [=====] - 0s 8ms/step - loss: 0.0079 -
mae: 0.0974 - val_loss: 0.0062 - val_mae: 0.0880
Epoch 14/500
33/33  [=====] - 0s 8ms/step - loss: 0.0075 -
mae: 0.0938 - val_loss: 0.0055 - val_mae: 0.0833
Epoch 15/500
33/33  [=====] - 0s 7ms/step - loss: 0.0071 -
mae: 0.0899 - val_loss: 0.0051 - val_mae: 0.0799
Epoch 16/500
33/33  [=====] - 0s 8ms/step - loss: 0.0067 -
mae: 0.0870 - val_loss: 0.0046 - val_mae: 0.0764
Epoch 17/500
33/33  [=====] - 0s 8ms/step - loss: 0.0065 -
mae: 0.0845 - val_loss: 0.0044 - val_mae: 0.0747
Epoch 18/500
33/33  [=====] - 0s 8ms/step - loss: 0.0062 -
mae: 0.0822 - val_loss: 0.0042 - val_mae: 0.0723
Epoch 19/500
33/33  [=====] - 0s 8ms/step - loss: 0.0060 -
mae: 0.0801 - val_loss: 0.0039 - val_mae: 0.0692
Epoch 20/500
33/33  [=====] - 0s 8ms/step - loss: 0.0058 -
mae: 0.0783 - val_loss: 0.0038 - val_mae: 0.0679
Epoch 21/500
33/33  [=====] - 0s 8ms/step - loss: 0.0057 -
mae: 0.0769 - val_loss: 0.0037 - val_mae: 0.0672
Epoch 22/500
33/33  [=====] - 0s 8ms/step - loss: 0.0055 -
mae: 0.0752 - val_loss: 0.0036 - val_mae: 0.0662
Epoch 23/500
33/33  [=====] - 0s 8ms/step - loss: 0.0054 -
mae: 0.0739 - val_loss: 0.0033 - val_mae: 0.0637
Epoch 24/500
33/33  [=====] - 0s 8ms/step - loss: 0.0053 -
```

```
mae: 0.0725 - val_loss: 0.0032 - val_mae: 0.0629
Epoch 25/500
33/33 [=====] - 0s 8ms/step - loss: 0.0052 -
mae: 0.0715 - val_loss: 0.0030 - val_mae: 0.0602
Epoch 26/500
33/33 [=====] - 0s 7ms/step - loss: 0.0051 -
mae: 0.0706 - val_loss: 0.0030 - val_mae: 0.0596
Epoch 27/500
33/33 [=====] - 0s 7ms/step - loss: 0.0050 -
mae: 0.0696 - val_loss: 0.0029 - val_mae: 0.0582
Epoch 28/500
33/33 [=====] - 0s 7ms/step - loss: 0.0049 -
mae: 0.0687 - val_loss: 0.0028 - val_mae: 0.0579
Epoch 29/500
33/33 [=====] - 0s 8ms/step - loss: 0.0049 -
mae: 0.0678 - val_loss: 0.0028 - val_mae: 0.0581
Epoch 30/500
33/33 [=====] - 0s 7ms/step - loss: 0.0048 -
mae: 0.0671 - val_loss: 0.0026 - val_mae: 0.0557
Epoch 31/500
33/33 [=====] - 0s 7ms/step - loss: 0.0047 -
mae: 0.0663 - val_loss: 0.0026 - val_mae: 0.0550
Epoch 32/500
33/33 [=====] - 0s 8ms/step - loss: 0.0047 -
mae: 0.0661 - val_loss: 0.0025 - val_mae: 0.0543
Epoch 33/500
33/33 [=====] - 0s 8ms/step - loss: 0.0046 -
mae: 0.0653 - val_loss: 0.0025 - val_mae: 0.0537
Epoch 34/500
33/33 [=====] - 0s 7ms/step - loss: 0.0046 -
mae: 0.0645 - val_loss: 0.0025 - val_mae: 0.0545
Epoch 35/500
33/33 [=====] - 0s 8ms/step - loss: 0.0045 -
mae: 0.0640 - val_loss: 0.0024 - val_mae: 0.0530
Epoch 36/500
33/33 [=====] - 0s 8ms/step - loss: 0.0045 -
mae: 0.0637 - val_loss: 0.0024 - val_mae: 0.0524
Epoch 37/500
```

```
33/33  [=====] - 0s 7ms/step - loss: 0.0044 -
mae: 0.0630 - val_loss: 0.0023 - val_mae: 0.0517
Epoch 38/500
33/33  [=====] - 0s 8ms/step - loss: 0.0044 -
mae: 0.0626 - val_loss: 0.0023 - val_mae: 0.0508
Epoch 39/500
33/33  [=====] - 0s 8ms/step - loss: 0.0044 -
mae: 0.0620 - val_loss: 0.0023 - val_mae: 0.0505
Epoch 40/500
33/33  [=====] - 0s 7ms/step - loss: 0.0043 -
mae: 0.0616 - val_loss: 0.0022 - val_mae: 0.0498
Epoch 41/500
33/33  [=====] - 0s 8ms/step - loss: 0.0043 -
mae: 0.0617 - val_loss: 0.0022 - val_mae: 0.0504
Epoch 42/500
33/33  [=====] - 0s 8ms/step - loss: 0.0043 -
mae: 0.0612 - val_loss: 0.0022 - val_mae: 0.0491
Epoch 43/500
33/33  [=====] - 0s 7ms/step - loss: 0.0042 -
mae: 0.0607 - val_loss: 0.0022 - val_mae: 0.0504
Epoch 44/500
33/33  [=====] - 0s 7ms/step - loss: 0.0042 -
mae: 0.0603 - val_loss: 0.0022 - val_mae: 0.0496
Epoch 45/500
33/33  [=====] - 0s 7ms/step - loss: 0.0042 -
mae: 0.0600 - val_loss: 0.0023 - val_mae: 0.0511
Epoch 46/500
33/33  [=====] - 0s 7ms/step - loss: 0.0042 -
mae: 0.0598 - val_loss: 0.0021 - val_mae: 0.0482
Epoch 47/500
33/33  [=====] - 0s 8ms/step - loss: 0.0042 -
mae: 0.0599 - val_loss: 0.0022 - val_mae: 0.0498
Epoch 48/500
33/33  [=====] - 0s 8ms/step - loss: 0.0041 -
mae: 0.0593 - val_loss: 0.0020 - val_mae: 0.0474
Epoch 49/500
33/33  [=====] - 0s 8ms/step - loss: 0.0041 -
mae: 0.0592 - val_loss: 0.0021 - val_mae: 0.0488
```

```
Epoch 50/500
33/33  [=====] - 0s 8ms/step - loss: 0.0041 -
mae: 0.0589 - val_loss: 0.0020 - val_mae: 0.0467
Epoch 51/500
33/33  [=====] - 0s 8ms/step - loss: 0.0041 -
mae: 0.0586 - val_loss: 0.0020 - val_mae: 0.0479
Epoch 52/500
33/33  [=====] - 0s 7ms/step - loss: 0.0041 -
mae: 0.0583 - val_loss: 0.0021 - val_mae: 0.0490
Epoch 53/500
33/33  [=====] - 0s 8ms/step - loss: 0.0040 -
mae: 0.0584 - val_loss: 0.0020 - val_mae: 0.0475
Epoch 54/500
33/33  [=====] - 0s 8ms/step - loss: 0.0040 -
mae: 0.0579 - val_loss: 0.0019 - val_mae: 0.0460
Epoch 55/500
33/33  [=====] - 0s 7ms/step - loss: 0.0040 -
mae: 0.0587 - val_loss: 0.0021 - val_mae: 0.0488
Epoch 56/500
33/33  [=====] - 0s 8ms/step - loss: 0.0040 -
mae: 0.0576 - val_loss: 0.0019 - val_mae: 0.0465
Epoch 57/500
33/33  [=====] - 0s 7ms/step - loss: 0.0040 -
mae: 0.0576 - val_loss: 0.0019 - val_mae: 0.0463
Epoch 58/500
33/33  [=====] - 0s 8ms/step - loss: 0.0040 -
mae: 0.0575 - val_loss: 0.0019 - val_mae: 0.0459
Epoch 59/500
33/33  [=====] - 0s 8ms/step - loss: 0.0040 -
mae: 0.0572 - val_loss: 0.0020 - val_mae: 0.0466
Epoch 60/500
33/33  [=====] - 0s 8ms/step - loss: 0.0040 -
mae: 0.0573 - val_loss: 0.0020 - val_mae: 0.0468
Epoch 61/500
33/33  [=====] - 0s 8ms/step - loss: 0.0040 -
mae: 0.0571 - val_loss: 0.0019 - val_mae: 0.0460
Epoch 62/500
33/33  [=====] - 0s 8ms/step - loss: 0.0039 -
```

```
mae: 0.0567 - val_loss: 0.0019 - val_mae: 0.0461
Epoch 63/500
33/33 [=====] - 0s 8ms/step - loss: 0.0039 -
mae: 0.0568 - val_loss: 0.0019 - val_mae: 0.0450
Epoch 64/500
33/33 [=====] - 0s 8ms/step - loss: 0.0039 -
mae: 0.0569 - val_loss: 0.0019 - val_mae: 0.0449
Epoch 65/500
33/33 [=====] - 0s 8ms/step - loss: 0.0039 -
mae: 0.0569 - val_loss: 0.0019 - val_mae: 0.0456
Epoch 66/500
33/33 [=====] - 0s 7ms/step - loss: 0.0039 -
mae: 0.0564 - val_loss: 0.0020 - val_mae: 0.0473
Epoch 67/500
33/33 [=====] - 0s 8ms/step - loss: 0.0039 -
mae: 0.0563 - val_loss: 0.0019 - val_mae: 0.0461
Epoch 68/500
33/33 [=====] - 0s 8ms/step - loss: 0.0039 -
mae: 0.0560 - val_loss: 0.0018 - val_mae: 0.0443
Epoch 69/500
33/33 [=====] - 0s 8ms/step - loss: 0.0039 -
mae: 0.0560 - val_loss: 0.0018 - val_mae: 0.0444
Epoch 70/500
33/33 [=====] - 0s 8ms/step - loss: 0.0039 -
mae: 0.0566 - val_loss: 0.0019 - val_mae: 0.0447
Epoch 71/500
33/33 [=====] - 0s 8ms/step - loss: 0.0039 -
mae: 0.0559 - val_loss: 0.0019 - val_mae: 0.0450
Epoch 72/500
33/33 [=====] - 0s 8ms/step - loss: 0.0039 -
mae: 0.0558 - val_loss: 0.0019 - val_mae: 0.0455
Epoch 73/500
33/33 [=====] - 0s 8ms/step - loss: 0.0038 -
mae: 0.0557 - val_loss: 0.0019 - val_mae: 0.0453
Epoch 74/500
33/33 [=====] - 0s 8ms/step - loss: 0.0038 -
mae: 0.0557 - val_loss: 0.0018 - val_mae: 0.0440
Epoch 75/500
```

```
33/33  [=====] - 0s 8ms/step - loss: 0.0038 -
mae: 0.0553 - val_loss: 0.0018 - val_mae: 0.0444
Epoch 76/500
33/33  [=====] - 0s 8ms/step - loss: 0.0038 -
mae: 0.0553 - val_loss: 0.0019 - val_mae: 0.0455
Epoch 77/500
33/33  [=====] - 0s 8ms/step - loss: 0.0038 -
mae: 0.0553 - val_loss: 0.0020 - val_mae: 0.0465
Epoch 78/500
33/33  [=====] - 0s 8ms/step - loss: 0.0038 -
mae: 0.0551 - val_loss: 0.0019 - val_mae: 0.0452
Epoch 79/500
33/33  [=====] - 0s 9ms/step - loss: 0.0038 -
mae: 0.0553 - val_loss: 0.0019 - val_mae: 0.0444
Epoch 80/500
33/33  [=====] - 0s 8ms/step - loss: 0.0038 -
mae: 0.0552 - val_loss: 0.0018 - val_mae: 0.0436
Epoch 81/500
33/33  [=====] - 0s 8ms/step - loss: 0.0038 -
mae: 0.0551 - val_loss: 0.0018 - val_mae: 0.0441
Epoch 82/500
33/33  [=====] - 0s 8ms/step - loss: 0.0038 -
mae: 0.0549 - val_loss: 0.0017 - val_mae: 0.0428
Epoch 83/500
33/33  [=====] - 0s 8ms/step - loss: 0.0038 -
mae: 0.0556 - val_loss: 0.0018 - val_mae: 0.0443
Epoch 84/500
33/33  [=====] - 0s 8ms/step - loss: 0.0038 -
mae: 0.0548 - val_loss: 0.0017 - val_mae: 0.0429
Epoch 85/500
33/33  [=====] - 0s 8ms/step - loss: 0.0038 -
mae: 0.0554 - val_loss: 0.0018 - val_mae: 0.0433
Epoch 86/500
33/33  [=====] - 0s 9ms/step - loss: 0.0038 -
mae: 0.0545 - val_loss: 0.0018 - val_mae: 0.0437
Epoch 87/500
33/33  [=====] - 0s 8ms/step - loss: 0.0037 -
mae: 0.0545 - val_loss: 0.0018 - val_mae: 0.0437
```

```
Epoch 88/500
33/33  [=====] - 0s 8ms/step - loss: 0.0038 -
mae: 0.0548 - val_loss: 0.0018 - val_mae: 0.0435
Epoch 89/500
33/33  [=====] - 0s 8ms/step - loss: 0.0037 -
mae: 0.0545 - val_loss: 0.0020 - val_mae: 0.0461
Epoch 90/500
33/33  [=====] - 0s 8ms/step - loss: 0.0037 -
mae: 0.0542 - val_loss: 0.0018 - val_mae: 0.0441
Epoch 91/500
33/33  [=====] - 0s 8ms/step - loss: 0.0037 -
mae: 0.0542 - val_loss: 0.0019 - val_mae: 0.0451
Epoch 92/500
33/33  [=====] - 0s 8ms/step - loss: 0.0037 -
mae: 0.0543 - val_loss: 0.0017 - val_mae: 0.0427
Epoch 93/500
33/33  [=====] - 0s 8ms/step - loss: 0.0037 -
mae: 0.0543 - val_loss: 0.0019 - val_mae: 0.0454
Epoch 94/500
33/33  [=====] - 0s 8ms/step - loss: 0.0037 -
mae: 0.0542 - val_loss: 0.0018 - val_mae: 0.0429
Epoch 95/500
33/33  [=====] - 0s 8ms/step - loss: 0.0037 -
mae: 0.0539 - val_loss: 0.0018 - val_mae: 0.0441
Epoch 96/500
33/33  [=====] - 0s 9ms/step - loss: 0.0037 -
mae: 0.0537 - val_loss: 0.0017 - val_mae: 0.0425
Epoch 97/500
33/33  [=====] - 0s 8ms/step - loss: 0.0037 -
mae: 0.0544 - val_loss: 0.0018 - val_mae: 0.0442
Epoch 98/500
33/33  [=====] - 0s 8ms/step - loss: 0.0037 -
mae: 0.0539 - val_loss: 0.0018 - val_mae: 0.0443
Epoch 99/500
33/33  [=====] - 0s 8ms/step - loss: 0.0037 -
mae: 0.0537 - val_loss: 0.0017 - val_mae: 0.0423
Epoch 100/500
33/33  [=====] - 0s 8ms/step - loss: 0.0037 -
```

```
mae: 0.0535 - val_loss: 0.0017 - val_mae: 0.0422
Epoch 101/500
33/33 [=====] - 0s 8ms/step - loss: 0.0037 -
mae: 0.0532 - val_loss: 0.0018 - val_mae: 0.0438
Epoch 102/500
33/33 [=====] - 0s 8ms/step - loss: 0.0037 -
mae: 0.0537 - val_loss: 0.0018 - val_mae: 0.0438
Epoch 103/500
33/33 [=====] - 0s 8ms/step - loss: 0.0037 -
mae: 0.0535 - val_loss: 0.0017 - val_mae: 0.0425
Epoch 104/500
33/33 [=====] - 0s 8ms/step - loss: 0.0037 -
mae: 0.0532 - val_loss: 0.0017 - val_mae: 0.0426
Epoch 105/500
33/33 [=====] - 0s 8ms/step - loss: 0.0037 -
mae: 0.0533 - val_loss: 0.0017 - val_mae: 0.0427
Epoch 106/500
33/33 [=====] - 0s 9ms/step - loss: 0.0037 -
mae: 0.0532 - val_loss: 0.0017 - val_mae: 0.0413
Epoch 107/500
33/33 [=====] - 0s 8ms/step - loss: 0.0037 -
mae: 0.0533 - val_loss: 0.0018 - val_mae: 0.0434
Epoch 108/500
33/33 [=====] - 0s 8ms/step - loss: 0.0037 -
mae: 0.0533 - val_loss: 0.0017 - val_mae: 0.0415
Epoch 109/500
33/33 [=====] - 0s 8ms/step - loss: 0.0036 -
mae: 0.0532 - val_loss: 0.0017 - val_mae: 0.0420
Epoch 110/500
33/33 [=====] - 0s 8ms/step - loss: 0.0036 -
mae: 0.0530 - val_loss: 0.0017 - val_mae: 0.0422
Epoch 111/500
33/33 [=====] - 0s 8ms/step - loss: 0.0036 -
mae: 0.0530 - val_loss: 0.0017 - val_mae: 0.0410
Epoch 112/500
33/33 [=====] - 0s 8ms/step - loss: 0.0036 -
mae: 0.0528 - val_loss: 0.0017 - val_mae: 0.0421
Epoch 113/500
```

```
33/33  [=====] - 0s 9ms/step - loss: 0.0036 -  
mae: 0.0526 - val_loss: 0.0017 - val_mae: 0.0414  
Epoch 114/500  
33/33  [=====] - 0s 8ms/step - loss: 0.0036 -  
mae: 0.0529 - val_loss: 0.0017 - val_mae: 0.0419  
Epoch 115/500  
33/33  [=====] - 0s 9ms/step - loss: 0.0036 -  
mae: 0.0530 - val_loss: 0.0018 - val_mae: 0.0436  
Epoch 116/500  
33/33  [=====] - 0s 8ms/step - loss: 0.0037 -  
mae: 0.0530 - val_loss: 0.0017 - val_mae: 0.0415  
Epoch 117/500  
33/33  [=====] - 0s 8ms/step - loss: 0.0036 -  
mae: 0.0525 - val_loss: 0.0016 - val_mae: 0.0401  
Epoch 118/500  
33/33  [=====] - 0s 14ms/step - loss: 0.0036 -  
mae: 0.0537 - val_loss: 0.0016 - val_mae: 0.0407  
Epoch 119/500  
33/33  [=====] - 0s 8ms/step - loss: 0.0036 -  
mae: 0.0529 - val_loss: 0.0018 - val_mae: 0.0428  
Epoch 120/500  
33/33  [=====] - 0s 8ms/step - loss: 0.0036 -  
mae: 0.0524 - val_loss: 0.0016 - val_mae: 0.0401  
Epoch 121/500  
33/33  [=====] - 0s 8ms/step - loss: 0.0036 -  
mae: 0.0526 - val_loss: 0.0017 - val_mae: 0.0413  
Epoch 122/500  
33/33  [=====] - 0s 8ms/step - loss: 0.0036 -  
mae: 0.0528 - val_loss: 0.0018 - val_mae: 0.0438  
Epoch 123/500  
33/33  [=====] - 0s 8ms/step - loss: 0.0036 -  
mae: 0.0525 - val_loss: 0.0017 - val_mae: 0.0416  
Epoch 124/500  
33/33  [=====] - 0s 8ms/step - loss: 0.0036 -  
mae: 0.0521 - val_loss: 0.0017 - val_mae: 0.0411  
Epoch 125/500  
33/33  [=====] - 0s 8ms/step - loss: 0.0036 -  
mae: 0.0518 - val_loss: 0.0016 - val_mae: 0.0407
```

```
Epoch 126/500
33/33  [=====] - 0s 8ms/step - loss: 0.0036 -
mae: 0.0520 - val_loss: 0.0016 - val_mae: 0.0399
Epoch 127/500
33/33  [=====] - 0s 8ms/step - loss: 0.0036 -
mae: 0.0521 - val_loss: 0.0016 - val_mae: 0.0396
Epoch 128/500
33/33  [=====] - 0s 9ms/step - loss: 0.0036 -
mae: 0.0518 - val_loss: 0.0017 - val_mae: 0.0411
Epoch 129/500
33/33  [=====] - 0s 10ms/step - loss: 0.0036 -
mae: 0.0519 - val_loss: 0.0016 - val_mae: 0.0398
Epoch 130/500
33/33  [=====] - 0s 8ms/step - loss: 0.0036 -
mae: 0.0523 - val_loss: 0.0017 - val_mae: 0.0410
Epoch 131/500
33/33  [=====] - 0s 8ms/step - loss: 0.0036 -
mae: 0.0517 - val_loss: 0.0018 - val_mae: 0.0422
Epoch 132/500
33/33  [=====] - 0s 8ms/step - loss: 0.0036 -
mae: 0.0519 - val_loss: 0.0016 - val_mae: 0.0403
Epoch 133/500
33/33  [=====] - 0s 8ms/step - loss: 0.0036 -
mae: 0.0516 - val_loss: 0.0016 - val_mae: 0.0396
Epoch 134/500
33/33  [=====] - 0s 9ms/step - loss: 0.0035 -
mae: 0.0513 - val_loss: 0.0016 - val_mae: 0.0398
Epoch 135/500
33/33  [=====] - 0s 9ms/step - loss: 0.0035 -
mae: 0.0513 - val_loss: 0.0016 - val_mae: 0.0393
Epoch 136/500
33/33  [=====] - 0s 8ms/step - loss: 0.0035 -
mae: 0.0513 - val_loss: 0.0016 - val_mae: 0.0397
Epoch 137/500
33/33  [=====] - 0s 8ms/step - loss: 0.0035 -
mae: 0.0513 - val_loss: 0.0016 - val_mae: 0.0397
Epoch 138/500
33/33  [=====] - 0s 8ms/step - loss: 0.0035 -
```

```
mae: 0.0513 - val_loss: 0.0017 - val_mae: 0.0412
Epoch 139/500
33/33 [=====] - 0s 8ms/step - loss: 0.0035 -
mae: 0.0511 - val_loss: 0.0018 - val_mae: 0.0432
Epoch 140/500
33/33 [=====] - 0s 9ms/step - loss: 0.0036 -
mae: 0.0519 - val_loss: 0.0016 - val_mae: 0.0401
Epoch 141/500
33/33 [=====] - 0s 8ms/step - loss: 0.0036 -
mae: 0.0516 - val_loss: 0.0016 - val_mae: 0.0395
Epoch 142/500
33/33 [=====] - 0s 8ms/step - loss: 0.0035 -
mae: 0.0513 - val_loss: 0.0017 - val_mae: 0.0415
Epoch 143/500
33/33 [=====] - 0s 8ms/step - loss: 0.0035 -
mae: 0.0513 - val_loss: 0.0016 - val_mae: 0.0395
Epoch 144/500
33/33 [=====] - 0s 8ms/step - loss: 0.0036 -
mae: 0.0521 - val_loss: 0.0016 - val_mae: 0.0397
Epoch 145/500
33/33 [=====] - 0s 8ms/step - loss: 0.0036 -
mae: 0.0519 - val_loss: 0.0016 - val_mae: 0.0394
Epoch 146/500
33/33 [=====] - 0s 8ms/step - loss: 0.0035 -
mae: 0.0509 - val_loss: 0.0015 - val_mae: 0.0382
Epoch 147/500
33/33 [=====] - 0s 7ms/step - loss: 0.0035 -
mae: 0.0507 - val_loss: 0.0016 - val_mae: 0.0388
Epoch 148/500
33/33 [=====] - 0s 8ms/step - loss: 0.0035 -
mae: 0.0509 - val_loss: 0.0016 - val_mae: 0.0392
Epoch 149/500
33/33 [=====] - 0s 8ms/step - loss: 0.0035 -
mae: 0.0506 - val_loss: 0.0016 - val_mae: 0.0401
Epoch 150/500
33/33 [=====] - 0s 8ms/step - loss: 0.0035 -
mae: 0.0507 - val_loss: 0.0015 - val_mae: 0.0375
Epoch 151/500
```

```
33/33  [=====] - 0s 9ms/step - loss: 0.0035 -  
mae: 0.0515 - val_loss: 0.0016 - val_mae: 0.0395  
Epoch 152/500  
33/33  [=====] - 0s 8ms/step - loss: 0.0035 -  
mae: 0.0506 - val_loss: 0.0015 - val_mae: 0.0385  
Epoch 153/500  
33/33  [=====] - 0s 8ms/step - loss: 0.0035 -  
mae: 0.0508 - val_loss: 0.0015 - val_mae: 0.0381  
Epoch 154/500  
33/33  [=====] - 0s 8ms/step - loss: 0.0035 -  
mae: 0.0506 - val_loss: 0.0015 - val_mae: 0.0371  
Epoch 155/500  
33/33  [=====] - 0s 8ms/step - loss: 0.0035 -  
mae: 0.0511 - val_loss: 0.0016 - val_mae: 0.0389  
Epoch 156/500  
33/33  [=====] - 0s 8ms/step - loss: 0.0035 -  
mae: 0.0506 - val_loss: 0.0016 - val_mae: 0.0390  
Epoch 157/500  
33/33  [=====] - 0s 8ms/step - loss: 0.0035 -  
mae: 0.0505 - val_loss: 0.0016 - val_mae: 0.0391  
Epoch 158/500  
33/33  [=====] - 0s 8ms/step - loss: 0.0035 -  
mae: 0.0504 - val_loss: 0.0016 - val_mae: 0.0397  
Epoch 159/500  
33/33  [=====] - 0s 8ms/step - loss: 0.0035 -  
mae: 0.0504 - val_loss: 0.0016 - val_mae: 0.0390  
Epoch 160/500  
33/33  [=====] - 0s 8ms/step - loss: 0.0035 -  
mae: 0.0501 - val_loss: 0.0016 - val_mae: 0.0394  
Epoch 161/500  
33/33  [=====] - 0s 8ms/step - loss: 0.0035 -  
mae: 0.0501 - val_loss: 0.0015 - val_mae: 0.0374  
Epoch 162/500  
33/33  [=====] - 0s 8ms/step - loss: 0.0035 -  
mae: 0.0505 - val_loss: 0.0015 - val_mae: 0.0383  
Epoch 163/500  
33/33  [=====] - 0s 8ms/step - loss: 0.0035 -  
mae: 0.0502 - val_loss: 0.0015 - val_mae: 0.0368
```

```
Epoch 164/500
33/33  [=====] - 0s 8ms/step - loss: 0.0035 -
mae: 0.0514 - val_loss: 0.0015 - val_mae: 0.0366
Epoch 165/500
33/33  [=====] - 0s 8ms/step - loss: 0.0035 -
mae: 0.0510 - val_loss: 0.0015 - val_mae: 0.0364
Epoch 166/500
33/33  [=====] - 0s 9ms/step - loss: 0.0036 -
mae: 0.0519 - val_loss: 0.0015 - val_mae: 0.0366
Epoch 167/500
33/33  [=====] - 0s 8ms/step - loss: 0.0035 -
mae: 0.0505 - val_loss: 0.0015 - val_mae: 0.0366
Epoch 168/500
33/33  [=====] - 0s 8ms/step - loss: 0.0034 -
mae: 0.0499 - val_loss: 0.0016 - val_mae: 0.0386
Epoch 169/500
33/33  [=====] - 0s 8ms/step - loss: 0.0035 -
mae: 0.0500 - val_loss: 0.0017 - val_mae: 0.0404
Epoch 170/500
33/33  [=====] - 0s 8ms/step - loss: 0.0035 -
mae: 0.0502 - val_loss: 0.0016 - val_mae: 0.0395
Epoch 171/500
33/33  [=====] - 0s 8ms/step - loss: 0.0035 -
mae: 0.0503 - val_loss: 0.0016 - val_mae: 0.0394
Epoch 172/500
33/33  [=====] - 0s 7ms/step - loss: 0.0035 -
mae: 0.0497 - val_loss: 0.0015 - val_mae: 0.0379
Epoch 173/500
33/33  [=====] - 0s 8ms/step - loss: 0.0035 -
mae: 0.0503 - val_loss: 0.0015 - val_mae: 0.0378
Epoch 174/500
33/33  [=====] - 0s 8ms/step - loss: 0.0034 -
mae: 0.0497 - val_loss: 0.0015 - val_mae: 0.0372
Epoch 175/500
33/33  [=====] - 0s 7ms/step - loss: 0.0034 -
mae: 0.0500 - val_loss: 0.0016 - val_mae: 0.0397
Epoch 176/500
33/33  [=====] - 0s 8ms/step - loss: 0.0035 -
```

```
mae: 0.0498 - val_loss: 0.0014 - val_mae: 0.0360
Epoch 177/500
33/33 [=====] - 0s 8ms/step - loss: 0.0034 -
mae: 0.0496 - val_loss: 0.0014 - val_mae: 0.0361
Epoch 178/500
33/33 [=====] - 0s 7ms/step - loss: 0.0034 -
mae: 0.0497 - val_loss: 0.0015 - val_mae: 0.0374
Epoch 179/500
33/33 [=====] - 0s 8ms/step - loss: 0.0034 -
mae: 0.0492 - val_loss: 0.0016 - val_mae: 0.0387
Epoch 180/500
33/33 [=====] - 0s 8ms/step - loss: 0.0034 -
mae: 0.0494 - val_loss: 0.0014 - val_mae: 0.0362
Epoch 181/500
33/33 [=====] - 0s 9ms/step - loss: 0.0034 -
mae: 0.0492 - val_loss: 0.0014 - val_mae: 0.0358
Epoch 182/500
33/33 [=====] - 0s 8ms/step - loss: 0.0034 -
mae: 0.0495 - val_loss: 0.0014 - val_mae: 0.0358
Epoch 183/500
33/33 [=====] - 0s 8ms/step - loss: 0.0034 -
mae: 0.0493 - val_loss: 0.0016 - val_mae: 0.0383
Epoch 184/500
33/33 [=====] - 0s 8ms/step - loss: 0.0034 -
mae: 0.0493 - val_loss: 0.0015 - val_mae: 0.0376
Epoch 185/500
33/33 [=====] - 0s 7ms/step - loss: 0.0034 -
mae: 0.0491 - val_loss: 0.0014 - val_mae: 0.0357
Epoch 186/500
33/33 [=====] - 0s 8ms/step - loss: 0.0034 -
mae: 0.0501 - val_loss: 0.0014 - val_mae: 0.0359
Epoch 187/500
33/33 [=====] - 0s 8ms/step - loss: 0.0034 -
mae: 0.0491 - val_loss: 0.0015 - val_mae: 0.0362
Epoch 188/500
33/33 [=====] - 0s 8ms/step - loss: 0.0034 -
mae: 0.0495 - val_loss: 0.0014 - val_mae: 0.0351
Epoch 189/500
```

```
33/33  [=====] - 0s 7ms/step - loss: 0.0034 -
mae: 0.0496 - val_loss: 0.0014 - val_mae: 0.0359
Epoch 190/500
33/33  [=====] - 0s 7ms/step - loss: 0.0034 -
mae: 0.0494 - val_loss: 0.0016 - val_mae: 0.0381
Epoch 191/500
33/33  [=====] - 0s 7ms/step - loss: 0.0034 -
mae: 0.0490 - val_loss: 0.0014 - val_mae: 0.0357
Epoch 192/500
33/33  [=====] - 0s 7ms/step - loss: 0.0034 -
mae: 0.0488 - val_loss: 0.0014 - val_mae: 0.0360
Epoch 193/500
33/33  [=====] - 0s 7ms/step - loss: 0.0034 -
mae: 0.0489 - val_loss: 0.0014 - val_mae: 0.0357
Epoch 194/500
33/33  [=====] - 0s 7ms/step - loss: 0.0034 -
mae: 0.0489 - val_loss: 0.0015 - val_mae: 0.0362
Epoch 195/500
33/33  [=====] - 0s 7ms/step - loss: 0.0034 -
mae: 0.0492 - val_loss: 0.0015 - val_mae: 0.0371
Epoch 196/500
33/33  [=====] - 0s 8ms/step - loss: 0.0034 -
mae: 0.0491 - val_loss: 0.0015 - val_mae: 0.0363
Epoch 197/500
33/33  [=====] - 0s 8ms/step - loss: 0.0034 -
mae: 0.0496 - val_loss: 0.0015 - val_mae: 0.0375
Epoch 198/500
33/33  [=====] - 0s 8ms/step - loss: 0.0034 -
mae: 0.0487 - val_loss: 0.0015 - val_mae: 0.0363
Epoch 199/500
33/33  [=====] - 0s 7ms/step - loss: 0.0034 -
mae: 0.0491 - val_loss: 0.0015 - val_mae: 0.0366
Epoch 200/500
33/33  [=====] - 0s 8ms/step - loss: 0.0034 -
mae: 0.0488 - val_loss: 0.0014 - val_mae: 0.0356
Epoch 201/500
33/33  [=====] - 0s 7ms/step - loss: 0.0034 -
mae: 0.0500 - val_loss: 0.0014 - val_mae: 0.0354
```

```
Epoch 202/500
33/33 [=====] - 0s 7ms/step - loss: 0.0034 -
mae: 0.0495 - val_loss: 0.0014 - val_mae: 0.0346
Epoch 203/500
33/33 [=====] - 0s 7ms/step - loss: 0.0034 -
mae: 0.0497 - val_loss: 0.0016 - val_mae: 0.0387
Epoch 204/500
33/33 [=====] - 0s 7ms/step - loss: 0.0034 -
mae: 0.0488 - val_loss: 0.0015 - val_mae: 0.0375
Epoch 205/500
33/33 [=====] - 0s 7ms/step - loss: 0.0034 -
mae: 0.0490 - val_loss: 0.0016 - val_mae: 0.0388
Epoch 206/500
33/33 [=====] - 0s 7ms/step - loss: 0.0034 -
mae: 0.0485 - val_loss: 0.0014 - val_mae: 0.0357
Epoch 207/500
33/33 [=====] - 0s 7ms/step - loss: 0.0034 -
mae: 0.0485 - val_loss: 0.0014 - val_mae: 0.0346
Epoch 208/500
33/33 [=====] - 0s 7ms/step - loss: 0.0034 -
mae: 0.0486 - val_loss: 0.0015 - val_mae: 0.0362
Epoch 209/500
33/33 [=====] - 0s 7ms/step - loss: 0.0034 -
mae: 0.0485 - val_loss: 0.0014 - val_mae: 0.0358
Epoch 210/500
33/33 [=====] - 0s 7ms/step - loss: 0.0034 -
mae: 0.0484 - val_loss: 0.0014 - val_mae: 0.0351
Epoch 211/500
33/33 [=====] - 0s 8ms/step - loss: 0.0034 -
mae: 0.0488 - val_loss: 0.0014 - val_mae: 0.0346
Epoch 212/500
33/33 [=====] - 0s 7ms/step - loss: 0.0034 -
mae: 0.0485 - val_loss: 0.0015 - val_mae: 0.0376
Epoch 213/500
33/33 [=====] - 0s 7ms/step - loss: 0.0034 -
mae: 0.0484 - val_loss: 0.0015 - val_mae: 0.0365
Epoch 214/500
33/33 [=====] - 0s 7ms/step - loss: 0.0034 -
```

```
mae: 0.0486 - val_loss: 0.0014 - val_mae: 0.0348
Epoch 215/500
33/33 [=====] - 0s 7ms/step - loss: 0.0034 -
mae: 0.0483 - val_loss: 0.0014 - val_mae: 0.0344
Epoch 216/500
33/33 [=====] - 0s 8ms/step - loss: 0.0034 -
mae: 0.0483 - val_loss: 0.0014 - val_mae: 0.0342
Epoch 217/500
33/33 [=====] - 0s 9ms/step - loss: 0.0033 -
mae: 0.0479 - val_loss: 0.0014 - val_mae: 0.0344
Epoch 218/500
33/33 [=====] - 0s 7ms/step - loss: 0.0033 -
mae: 0.0485 - val_loss: 0.0014 - val_mae: 0.0360
Epoch 219/500
33/33 [=====] - 0s 7ms/step - loss: 0.0033 -
mae: 0.0479 - val_loss: 0.0014 - val_mae: 0.0356
Epoch 220/500
33/33 [=====] - 0s 7ms/step - loss: 0.0033 -
mae: 0.0482 - val_loss: 0.0014 - val_mae: 0.0354
Epoch 221/500
33/33 [=====] - 0s 7ms/step - loss: 0.0034 -
mae: 0.0484 - val_loss: 0.0014 - val_mae: 0.0354
Epoch 222/500
33/33 [=====] - 0s 7ms/step - loss: 0.0034 -
mae: 0.0485 - val_loss: 0.0014 - val_mae: 0.0341
Epoch 223/500
33/33 [=====] - 0s 7ms/step - loss: 0.0033 -
mae: 0.0483 - val_loss: 0.0014 - val_mae: 0.0357
Epoch 224/500
33/33 [=====] - 0s 7ms/step - loss: 0.0033 -
mae: 0.0479 - val_loss: 0.0014 - val_mae: 0.0343
Epoch 225/500
33/33 [=====] - 0s 7ms/step - loss: 0.0033 -
mae: 0.0480 - val_loss: 0.0015 - val_mae: 0.0372
Epoch 226/500
33/33 [=====] - 0s 8ms/step - loss: 0.0034 -
mae: 0.0482 - val_loss: 0.0016 - val_mae: 0.0391
Epoch 227/500
```

```
33/33  [=====] - 0s 7ms/step - loss: 0.0034 -
mae: 0.0484 - val_loss: 0.0014 - val_mae: 0.0348
Epoch 228/500
33/33  [=====] - 0s 7ms/step - loss: 0.0034 -
mae: 0.0487 - val_loss: 0.0014 - val_mae: 0.0343
Epoch 229/500
33/33  [=====] - 0s 7ms/step - loss: 0.0033 -
mae: 0.0479 - val_loss: 0.0015 - val_mae: 0.0367
Epoch 230/500
33/33  [=====] - 0s 7ms/step - loss: 0.0033 -
mae: 0.0482 - val_loss: 0.0015 - val_mae: 0.0361
Epoch 231/500
33/33  [=====] - 0s 7ms/step - loss: 0.0033 -
mae: 0.0477 - val_loss: 0.0016 - val_mae: 0.0376
Epoch 232/500
33/33  [=====] - 0s 7ms/step - loss: 0.0033 -
mae: 0.0477 - val_loss: 0.0014 - val_mae: 0.0346
Epoch 233/500
33/33  [=====] - 0s 7ms/step - loss: 0.0033 -
mae: 0.0482 - val_loss: 0.0015 - val_mae: 0.0373
Epoch 234/500
33/33  [=====] - 0s 7ms/step - loss: 0.0034 -
mae: 0.0485 - val_loss: 0.0015 - val_mae: 0.0374
Epoch 235/500
33/33  [=====] - 0s 7ms/step - loss: 0.0034 -
mae: 0.0485 - val_loss: 0.0014 - val_mae: 0.0355
Epoch 236/500
33/33  [=====] - 0s 7ms/step - loss: 0.0033 -
mae: 0.0478 - val_loss: 0.0015 - val_mae: 0.0365
Epoch 237/500
33/33  [=====] - 0s 7ms/step - loss: 0.0033 -
mae: 0.0475 - val_loss: 0.0015 - val_mae: 0.0368
Epoch 238/500
33/33  [=====] - 0s 7ms/step - loss: 0.0034 -
mae: 0.0486 - val_loss: 0.0016 - val_mae: 0.0380
Epoch 239/500
33/33  [=====] - 0s 7ms/step - loss: 0.0033 -
mae: 0.0482 - val_loss: 0.0016 - val_mae: 0.0381
```

```
Epoch 240/500
33/33  [=====] - 0s 7ms/step - loss: 0.0033 -
mae: 0.0478 - val_loss: 0.0014 - val_mae: 0.0348
Epoch 241/500
33/33  [=====] - 0s 7ms/step - loss: 0.0033 -
mae: 0.0477 - val_loss: 0.0013 - val_mae: 0.0335
Epoch 242/500
33/33  [=====] - 0s 7ms/step - loss: 0.0033 -
mae: 0.0483 - val_loss: 0.0015 - val_mae: 0.0356
Epoch 243/500
33/33  [=====] - 0s 8ms/step - loss: 0.0033 -
mae: 0.0479 - val_loss: 0.0015 - val_mae: 0.0367
Epoch 244/500
33/33  [=====] - 0s 8ms/step - loss: 0.0033 -
mae: 0.0476 - val_loss: 0.0015 - val_mae: 0.0365
Epoch 245/500
33/33  [=====] - 0s 7ms/step - loss: 0.0034 -
mae: 0.0483 - val_loss: 0.0014 - val_mae: 0.0346
Epoch 246/500
33/33  [=====] - 0s 7ms/step - loss: 0.0033 -
mae: 0.0478 - val_loss: 0.0014 - val_mae: 0.0355
Epoch 247/500
33/33  [=====] - 0s 7ms/step - loss: 0.0033 -
mae: 0.0473 - val_loss: 0.0014 - val_mae: 0.0340
Epoch 248/500
33/33  [=====] - 0s 7ms/step - loss: 0.0033 -
mae: 0.0479 - val_loss: 0.0013 - val_mae: 0.0334
Epoch 249/500
33/33  [=====] - 0s 7ms/step - loss: 0.0033 -
mae: 0.0481 - val_loss: 0.0015 - val_mae: 0.0362
Epoch 250/500
33/33  [=====] - 0s 7ms/step - loss: 0.0033 -
mae: 0.0478 - val_loss: 0.0015 - val_mae: 0.0365
Epoch 251/500
33/33  [=====] - 0s 7ms/step - loss: 0.0033 -
mae: 0.0476 - val_loss: 0.0014 - val_mae: 0.0345
Epoch 252/500
33/33  [=====] - 0s 7ms/step - loss: 0.0033 -
```

```
mae: 0.0474 - val_loss: 0.0016 - val_mae: 0.0380
Epoch 253/500
33/33 [=====] - 0s 7ms/step - loss: 0.0034 -
mae: 0.0483 - val_loss: 0.0014 - val_mae: 0.0354
Epoch 254/500
33/33 [=====] - 0s 7ms/step - loss: 0.0033 -
mae: 0.0475 - val_loss: 0.0015 - val_mae: 0.0363
Epoch 255/500
33/33 [=====] - 0s 7ms/step - loss: 0.0033 -
mae: 0.0476 - val_loss: 0.0016 - val_mae: 0.0377
Epoch 256/500
33/33 [=====] - 0s 8ms/step - loss: 0.0033 -
mae: 0.0476 - val_loss: 0.0014 - val_mae: 0.0343
Epoch 257/500
33/33 [=====] - 0s 7ms/step - loss: 0.0033 -
mae: 0.0482 - val_loss: 0.0013 - val_mae: 0.0336
Epoch 258/500
33/33 [=====] - 0s 7ms/step - loss: 0.0033 -
mae: 0.0474 - val_loss: 0.0014 - val_mae: 0.0342
Epoch 259/500
33/33 [=====] - 0s 7ms/step - loss: 0.0033 -
mae: 0.0473 - val_loss: 0.0013 - val_mae: 0.0340
Epoch 260/500
33/33 [=====] - 0s 7ms/step - loss: 0.0033 -
mae: 0.0474 - val_loss: 0.0016 - val_mae: 0.0377
Epoch 261/500
33/33 [=====] - 0s 7ms/step - loss: 0.0033 -
mae: 0.0476 - val_loss: 0.0015 - val_mae: 0.0365
Epoch 262/500
33/33 [=====] - 0s 7ms/step - loss: 0.0033 -
mae: 0.0473 - val_loss: 0.0014 - val_mae: 0.0346
Epoch 263/500
33/33 [=====] - 0s 7ms/step - loss: 0.0033 -
mae: 0.0476 - val_loss: 0.0014 - val_mae: 0.0342
Epoch 264/500
33/33 [=====] - 0s 7ms/step - loss: 0.0033 -
mae: 0.0476 - val_loss: 0.0013 - val_mae: 0.0332
Epoch 265/500
```

```
33/33  [=====] - 0s 7ms/step - loss: 0.0033 -
mae: 0.0476 - val_loss: 0.0014 - val_mae: 0.0353
Epoch 266/500
33/33  [=====] - 0s 7ms/step - loss: 0.0033 -
mae: 0.0476 - val_loss: 0.0015 - val_mae: 0.0367
Epoch 267/500
33/33  [=====] - 0s 7ms/step - loss: 0.0033 -
mae: 0.0477 - val_loss: 0.0014 - val_mae: 0.0352
Epoch 268/500
33/33  [=====] - 0s 8ms/step - loss: 0.0033 -
mae: 0.0475 - val_loss: 0.0014 - val_mae: 0.0348
Epoch 269/500
33/33  [=====] - 0s 7ms/step - loss: 0.0033 -
mae: 0.0470 - val_loss: 0.0015 - val_mae: 0.0355
Epoch 270/500
33/33  [=====] - 0s 7ms/step - loss: 0.0033 -
mae: 0.0475 - val_loss: 0.0013 - val_mae: 0.0336
Epoch 271/500
33/33  [=====] - 0s 8ms/step - loss: 0.0033 -
mae: 0.0472 - val_loss: 0.0013 - val_mae: 0.0333
Epoch 272/500
33/33  [=====] - 0s 7ms/step - loss: 0.0033 -
mae: 0.0471 - val_loss: 0.0014 - val_mae: 0.0350
Epoch 273/500
33/33  [=====] - 0s 7ms/step - loss: 0.0033 -
mae: 0.0475 - val_loss: 0.0016 - val_mae: 0.0381
Epoch 274/500
33/33  [=====] - 0s 7ms/step - loss: 0.0033 -
mae: 0.0472 - val_loss: 0.0013 - val_mae: 0.0334
Epoch 275/500
33/33  [=====] - 0s 7ms/step - loss: 0.0033 -
mae: 0.0471 - val_loss: 0.0013 - val_mae: 0.0335
Epoch 276/500
33/33  [=====] - 0s 8ms/step - loss: 0.0033 -
mae: 0.0472 - val_loss: 0.0013 - val_mae: 0.0337
Epoch 277/500
33/33  [=====] - 0s 7ms/step - loss: 0.0033 -
mae: 0.0471 - val_loss: 0.0015 - val_mae: 0.0356
```

```
Epoch 278/500
33/33  [=====] - 0s 7ms/step - loss: 0.0033 -
mae: 0.0476 - val_loss: 0.0014 - val_mae: 0.0342
Epoch 279/500
33/33  [=====] - 0s 7ms/step - loss: 0.0033 -
mae: 0.0474 - val_loss: 0.0014 - val_mae: 0.0355
Epoch 280/500
33/33  [=====] - 0s 8ms/step - loss: 0.0033 -
mae: 0.0474 - val_loss: 0.0014 - val_mae: 0.0343
Epoch 281/500
33/33  [=====] - 0s 7ms/step - loss: 0.0033 -
mae: 0.0474 - val_loss: 0.0014 - val_mae: 0.0353
Epoch 282/500
33/33  [=====] - 0s 7ms/step - loss: 0.0033 -
mae: 0.0471 - val_loss: 0.0015 - val_mae: 0.0364
Epoch 283/500
33/33  [=====] - 0s 7ms/step - loss: 0.0033 -
mae: 0.0473 - val_loss: 0.0014 - val_mae: 0.0347
Epoch 284/500
33/33  [=====] - 0s 8ms/step - loss: 0.0033 -
mae: 0.0468 - val_loss: 0.0014 - val_mae: 0.0339
Epoch 285/500
33/33  [=====] - 0s 7ms/step - loss: 0.0033 -
mae: 0.0470 - val_loss: 0.0013 - val_mae: 0.0332
Epoch 286/500
33/33  [=====] - 0s 7ms/step - loss: 0.0033 -
mae: 0.0468 - val_loss: 0.0014 - val_mae: 0.0338
Epoch 287/500
33/33  [=====] - 0s 7ms/step - loss: 0.0033 -
mae: 0.0472 - val_loss: 0.0013 - val_mae: 0.0339
Epoch 288/500
33/33  [=====] - 0s 7ms/step - loss: 0.0033 -
mae: 0.0471 - val_loss: 0.0014 - val_mae: 0.0340
Epoch 289/500
33/33  [=====] - 0s 7ms/step - loss: 0.0033 -
mae: 0.0469 - val_loss: 0.0014 - val_mae: 0.0343
Epoch 290/500
33/33  [=====] - 0s 7ms/step - loss: 0.0033 -
```

```
mae: 0.0470 - val_loss: 0.0014 - val_mae: 0.0350
Epoch 291/500
33/33 [=====] - 0s 8ms/step - loss: 0.0033 -
mae: 0.0469 - val_loss: 0.0013 - val_mae: 0.0323
Epoch 292/500
33/33 [=====] - 0s 8ms/step - loss: 0.0033 -
mae: 0.0478 - val_loss: 0.0013 - val_mae: 0.0325
Epoch 293/500
33/33 [=====] - 0s 8ms/step - loss: 0.0033 -
mae: 0.0469 - val_loss: 0.0014 - val_mae: 0.0340
Epoch 294/500
33/33 [=====] - 0s 8ms/step - loss: 0.0033 -
mae: 0.0472 - val_loss: 0.0014 - val_mae: 0.0348
Epoch 295/500
33/33 [=====] - 0s 8ms/step - loss: 0.0033 -
mae: 0.0473 - val_loss: 0.0013 - val_mae: 0.0323
Epoch 296/500
33/33 [=====] - 0s 8ms/step - loss: 0.0033 -
mae: 0.0473 - val_loss: 0.0013 - val_mae: 0.0336
Epoch 297/500
33/33 [=====] - 0s 8ms/step - loss: 0.0033 -
mae: 0.0469 - val_loss: 0.0013 - val_mae: 0.0327
Epoch 298/500
33/33 [=====] - 0s 8ms/step - loss: 0.0032 -
mae: 0.0463 - val_loss: 0.0013 - val_mae: 0.0334
Epoch 299/500
33/33 [=====] - 0s 7ms/step - loss: 0.0033 -
mae: 0.0468 - val_loss: 0.0013 - val_mae: 0.0330
Epoch 300/500
33/33 [=====] - 0s 8ms/step - loss: 0.0033 -
mae: 0.0476 - val_loss: 0.0013 - val_mae: 0.0336
Epoch 301/500
33/33 [=====] - 0s 9ms/step - loss: 0.0033 -
mae: 0.0472 - val_loss: 0.0017 - val_mae: 0.0406
Epoch 302/500
33/33 [=====] - 0s 8ms/step - loss: 0.0033 -
mae: 0.0472 - val_loss: 0.0013 - val_mae: 0.0321
Epoch 303/500
```

```
33/33  [=====] - 0s 8ms/step - loss: 0.0033 -
mae: 0.0470 - val_loss: 0.0014 - val_mae: 0.0339
Epoch 304/500
33/33  [=====] - 0s 8ms/step - loss: 0.0032 -
mae: 0.0467 - val_loss: 0.0013 - val_mae: 0.0338
Epoch 305/500
33/33  [=====] - 0s 8ms/step - loss: 0.0033 -
mae: 0.0470 - val_loss: 0.0014 - val_mae: 0.0345
Epoch 306/500
33/33  [=====] - 0s 8ms/step - loss: 0.0033 -
mae: 0.0473 - val_loss: 0.0015 - val_mae: 0.0370
Epoch 307/500
33/33  [=====] - 0s 8ms/step - loss: 0.0033 -
mae: 0.0468 - val_loss: 0.0013 - val_mae: 0.0337
Epoch 308/500
33/33  [=====] - 0s 8ms/step - loss: 0.0033 -
mae: 0.0468 - val_loss: 0.0013 - val_mae: 0.0332
Epoch 309/500
33/33  [=====] - 0s 8ms/step - loss: 0.0032 -
mae: 0.0464 - val_loss: 0.0013 - val_mae: 0.0334
Epoch 310/500
33/33  [=====] - 0s 9ms/step - loss: 0.0032 -
mae: 0.0468 - val_loss: 0.0016 - val_mae: 0.0376
Epoch 311/500
33/33  [=====] - 0s 8ms/step - loss: 0.0033 -
mae: 0.0469 - val_loss: 0.0013 - val_mae: 0.0329
Epoch 312/500
33/33  [=====] - 0s 8ms/step - loss: 0.0032 -
mae: 0.0467 - val_loss: 0.0013 - val_mae: 0.0337
Epoch 313/500
33/33  [=====] - 0s 8ms/step - loss: 0.0033 -
mae: 0.0468 - val_loss: 0.0013 - val_mae: 0.0326
Epoch 314/500
33/33  [=====] - 0s 8ms/step - loss: 0.0033 -
mae: 0.0472 - val_loss: 0.0013 - val_mae: 0.0333
Epoch 315/500
33/33  [=====] - 0s 10ms/step - loss: 0.0032 -
mae: 0.0467 - val_loss: 0.0013 - val_mae: 0.0323
```

```
Epoch 316/500
33/33  [=====] - 0s 9ms/step - loss: 0.0032 -
mae: 0.0466 - val_loss: 0.0013 - val_mae: 0.0328
Epoch 317/500
33/33  [=====] - 0s 9ms/step - loss: 0.0032 -
mae: 0.0467 - val_loss: 0.0013 - val_mae: 0.0337
Epoch 318/500
33/33  [=====] - 0s 9ms/step - loss: 0.0033 -
mae: 0.0480 - val_loss: 0.0014 - val_mae: 0.0347
Epoch 319/500
33/33  [=====] - 0s 9ms/step - loss: 0.0033 -
mae: 0.0477 - val_loss: 0.0013 - val_mae: 0.0327
Epoch 320/500
33/33  [=====] - 0s 9ms/step - loss: 0.0032 -
mae: 0.0471 - val_loss: 0.0014 - val_mae: 0.0345
Epoch 321/500
33/33  [=====] - 0s 8ms/step - loss: 0.0033 -
mae: 0.0467 - val_loss: 0.0013 - val_mae: 0.0334
Epoch 322/500
33/33  [=====] - 0s 9ms/step - loss: 0.0032 -
mae: 0.0466 - val_loss: 0.0013 - val_mae: 0.0337
Epoch 323/500
33/33  [=====] - 0s 9ms/step - loss: 0.0033 -
mae: 0.0470 - val_loss: 0.0013 - val_mae: 0.0323
Epoch 324/500
33/33  [=====] - 0s 9ms/step - loss: 0.0033 -
mae: 0.0469 - val_loss: 0.0014 - val_mae: 0.0339
Epoch 325/500
33/33  [=====] - 0s 9ms/step - loss: 0.0032 -
mae: 0.0467 - val_loss: 0.0013 - val_mae: 0.0334
Epoch 326/500
33/33  [=====] - 0s 9ms/step - loss: 0.0032 -
mae: 0.0466 - val_loss: 0.0015 - val_mae: 0.0361
Epoch 327/500
33/33  [=====] - 0s 9ms/step - loss: 0.0032 -
mae: 0.0468 - val_loss: 0.0013 - val_mae: 0.0327
Epoch 328/500
33/33  [=====] - 0s 10ms/step - loss: 0.0032 -
```

```
mae: 0.0467 - val_loss: 0.0013 - val_mae: 0.0323
Epoch 329/500
33/33 [=====] - 0s 10ms/step - loss: 0.0032 -
mae: 0.0469 - val_loss: 0.0014 - val_mae: 0.0353
Epoch 330/500
33/33 [=====] - 0s 10ms/step - loss: 0.0033 -
mae: 0.0468 - val_loss: 0.0013 - val_mae: 0.0329
Epoch 331/500
33/33 [=====] - 0s 10ms/step - loss: 0.0033 -
mae: 0.0469 - val_loss: 0.0013 - val_mae: 0.0320
Epoch 332/500
33/33 [=====] - 0s 10ms/step - loss: 0.0033 -
mae: 0.0475 - val_loss: 0.0014 - val_mae: 0.0340
Epoch 333/500
33/33 [=====] - 0s 11ms/step - loss: 0.0033 -
mae: 0.0465 - val_loss: 0.0013 - val_mae: 0.0322
Epoch 334/500
33/33 [=====] - 0s 11ms/step - loss: 0.0032 -
mae: 0.0470 - val_loss: 0.0014 - val_mae: 0.0347
Epoch 335/500
33/33 [=====] - 0s 13ms/step - loss: 0.0032 -
mae: 0.0466 - val_loss: 0.0014 - val_mae: 0.0354
Epoch 336/500
33/33 [=====] - 0s 14ms/step - loss: 0.0032 -
mae: 0.0462 - val_loss: 0.0013 - val_mae: 0.0329
Epoch 337/500
33/33 [=====] - 0s 11ms/step - loss: 0.0033 -
mae: 0.0474 - val_loss: 0.0013 - val_mae: 0.0335
Epoch 338/500
33/33 [=====] - 0s 10ms/step - loss: 0.0032 -
mae: 0.0463 - val_loss: 0.0014 - val_mae: 0.0337
Epoch 339/500
33/33 [=====] - 0s 10ms/step - loss: 0.0032 -
mae: 0.0464 - val_loss: 0.0013 - val_mae: 0.0322
Epoch 340/500
33/33 [=====] - 0s 10ms/step - loss: 0.0032 -
mae: 0.0465 - val_loss: 0.0014 - val_mae: 0.0340
Epoch 341/500
```

```
33/33  [=====] - 0s 10ms/step - loss: 0.0033 -  
mae: 0.0472 - val_loss: 0.0013 - val_mae: 0.0323  
Epoch 342/500  
33/33  [=====] - 0s 10ms/step - loss: 0.0032 -  
mae: 0.0465 - val_loss: 0.0013 - val_mae: 0.0338  
Epoch 343/500  
33/33  [=====] - 0s 11ms/step - loss: 0.0033 -  
mae: 0.0472 - val_loss: 0.0013 - val_mae: 0.0326  
Epoch 344/500  
33/33  [=====] - 0s 12ms/step - loss: 0.0032 -  
mae: 0.0467 - val_loss: 0.0013 - val_mae: 0.0327  
Epoch 345/500  
33/33  [=====] - 0s 12ms/step - loss: 0.0032 -  
mae: 0.0465 - val_loss: 0.0013 - val_mae: 0.0328  
Epoch 346/500  
33/33  [=====] - 0s 12ms/step - loss: 0.0032 -  
mae: 0.0463 - val_loss: 0.0013 - val_mae: 0.0323  
Epoch 347/500  
33/33  [=====] - 0s 11ms/step - loss: 0.0033 -  
mae: 0.0474 - val_loss: 0.0013 - val_mae: 0.0325  
Epoch 348/500  
33/33  [=====] - 0s 14ms/step - loss: 0.0033 -  
mae: 0.0476 - val_loss: 0.0013 - val_mae: 0.0338  
Epoch 349/500  
33/33  [=====] - 0s 12ms/step - loss: 0.0032 -  
mae: 0.0462 - val_loss: 0.0013 - val_mae: 0.0330  
Epoch 350/500  
33/33  [=====] - 0s 13ms/step - loss: 0.0032 -  
mae: 0.0468 - val_loss: 0.0015 - val_mae: 0.0363  
Epoch 351/500  
33/33  [=====] - 0s 14ms/step - loss: 0.0033 -  
mae: 0.0472 - val_loss: 0.0013 - val_mae: 0.0332  
Epoch 352/500  
33/33  [=====] - 0s 12ms/step - loss: 0.0032 -  
mae: 0.0467 - val_loss: 0.0014 - val_mae: 0.0341  
Epoch 353/500  
33/33  [=====] - 0s 12ms/step - loss: 0.0032 -  
mae: 0.0463 - val_loss: 0.0014 - val_mae: 0.0348
```

```
Epoch 354/500
33/33 [=====] - 0s 12ms/step - loss: 0.0032 -
mae: 0.0462 - val_loss: 0.0013 - val_mae: 0.0320
Epoch 355/500
33/33 [=====] - 0s 12ms/step - loss: 0.0032 -
mae: 0.0466 - val_loss: 0.0013 - val_mae: 0.0327
Epoch 356/500
33/33 [=====] - 0s 13ms/step - loss: 0.0032 -
mae: 0.0463 - val_loss: 0.0013 - val_mae: 0.0325
Epoch 357/500
33/33 [=====] - 0s 12ms/step - loss: 0.0033 -
mae: 0.0476 - val_loss: 0.0014 - val_mae: 0.0338
Epoch 358/500
33/33 [=====] - 0s 12ms/step - loss: 0.0032 -
mae: 0.0463 - val_loss: 0.0014 - val_mae: 0.0339
Epoch 359/500
33/33 [=====] - 0s 12ms/step - loss: 0.0032 -
mae: 0.0466 - val_loss: 0.0013 - val_mae: 0.0324
Epoch 360/500
33/33 [=====] - 0s 12ms/step - loss: 0.0032 -
mae: 0.0464 - val_loss: 0.0014 - val_mae: 0.0345
Epoch 361/500
33/33 [=====] - 0s 12ms/step - loss: 0.0032 -
mae: 0.0464 - val_loss: 0.0013 - val_mae: 0.0337
Epoch 362/500
33/33 [=====] - 0s 12ms/step - loss: 0.0032 -
mae: 0.0472 - val_loss: 0.0015 - val_mae: 0.0357
Epoch 363/500
33/33 [=====] - 0s 13ms/step - loss: 0.0032 -
mae: 0.0463 - val_loss: 0.0013 - val_mae: 0.0325
Epoch 364/500
33/33 [=====] - 0s 12ms/step - loss: 0.0032 -
mae: 0.0475 - val_loss: 0.0013 - val_mae: 0.0325
Epoch 365/500
33/33 [=====] - 0s 12ms/step - loss: 0.0032 -
mae: 0.0466 - val_loss: 0.0013 - val_mae: 0.0331
Epoch 366/500
33/33 [=====] - 0s 12ms/step - loss: 0.0032 -
```

```
mae: 0.0462 - val_loss: 0.0013 - val_mae: 0.0324
Epoch 367/500
33/33 [=====] - 0s 12ms/step - loss: 0.0032 -
mae: 0.0463 - val_loss: 0.0013 - val_mae: 0.0326
Epoch 368/500
33/33 [=====] - 0s 12ms/step - loss: 0.0032 -
mae: 0.0464 - val_loss: 0.0014 - val_mae: 0.0346
Epoch 369/500
33/33 [=====] - 0s 12ms/step - loss: 0.0033 -
mae: 0.0469 - val_loss: 0.0014 - val_mae: 0.0348
Epoch 370/500
33/33 [=====] - 0s 13ms/step - loss: 0.0032 -
mae: 0.0464 - val_loss: 0.0013 - val_mae: 0.0319
Epoch 371/500
33/33 [=====] - 0s 12ms/step - loss: 0.0032 -
mae: 0.0463 - val_loss: 0.0014 - val_mae: 0.0347
Epoch 372/500
33/33 [=====] - 0s 12ms/step - loss: 0.0033 -
mae: 0.0475 - val_loss: 0.0016 - val_mae: 0.0376
Epoch 373/500
33/33 [=====] - 0s 12ms/step - loss: 0.0032 -
mae: 0.0467 - val_loss: 0.0013 - val_mae: 0.0325
Epoch 374/500
33/33 [=====] - 0s 12ms/step - loss: 0.0032 -
mae: 0.0459 - val_loss: 0.0013 - val_mae: 0.0324
Epoch 375/500
33/33 [=====] - 0s 12ms/step - loss: 0.0032 -
mae: 0.0463 - val_loss: 0.0013 - val_mae: 0.0326
Epoch 376/500
33/33 [=====] - 0s 12ms/step - loss: 0.0032 -
mae: 0.0462 - val_loss: 0.0013 - val_mae: 0.0325
Epoch 377/500
33/33 [=====] - 0s 13ms/step - loss: 0.0032 -
mae: 0.0461 - val_loss: 0.0013 - val_mae: 0.0333
Epoch 378/500
33/33 [=====] - 0s 12ms/step - loss: 0.0032 -
mae: 0.0463 - val_loss: 0.0014 - val_mae: 0.0343
Epoch 379/500
```

```
33/33  [=====] - 0s 12ms/step - loss: 0.0032 -  
mae: 0.0465 - val_loss: 0.0013 - val_mae: 0.0332  
Epoch 380/500  
33/33  [=====] - 0s 12ms/step - loss: 0.0032 -  
mae: 0.0462 - val_loss: 0.0013 - val_mae: 0.0326  
Epoch 381/500  
33/33  [=====] - 0s 12ms/step - loss: 0.0032 -  
mae: 0.0463 - val_loss: 0.0015 - val_mae: 0.0362  
Epoch 382/500  
33/33  [=====] - 0s 12ms/step - loss: 0.0032 -  
mae: 0.0461 - val_loss: 0.0013 - val_mae: 0.0321  
Epoch 383/500  
33/33  [=====] - 0s 12ms/step - loss: 0.0032 -  
mae: 0.0464 - val_loss: 0.0013 - val_mae: 0.0322  
Epoch 384/500  
33/33  [=====] - 0s 14ms/step - loss: 0.0032 -  
mae: 0.0463 - val_loss: 0.0012 - val_mae: 0.0317  
Epoch 385/500  
33/33  [=====] - 0s 12ms/step - loss: 0.0032 -  
mae: 0.0466 - val_loss: 0.0014 - val_mae: 0.0345  
Epoch 386/500  
33/33  [=====] - 0s 12ms/step - loss: 0.0032 -  
mae: 0.0466 - val_loss: 0.0013 - val_mae: 0.0323  
Epoch 387/500  
33/33  [=====] - 1s 14ms/step - loss: 0.0033 -  
mae: 0.0485 - val_loss: 0.0013 - val_mae: 0.0323  
Epoch 388/500  
33/33  [=====] - 0s 14ms/step - loss: 0.0033 -  
mae: 0.0471 - val_loss: 0.0012 - val_mae: 0.0320  
Epoch 389/500  
33/33  [=====] - 0s 12ms/step - loss: 0.0032 -  
mae: 0.0467 - val_loss: 0.0015 - val_mae: 0.0355  
Epoch 390/500  
33/33  [=====] - 1s 16ms/step - loss: 0.0032 -  
mae: 0.0469 - val_loss: 0.0016 - val_mae: 0.0371  
Epoch 391/500  
33/33  [=====] - 0s 12ms/step - loss: 0.0033 -  
mae: 0.0472 - val_loss: 0.0013 - val_mae: 0.0327
```

```
Epoch 392/500
33/33 [=====] - 0s 12ms/step - loss: 0.0032 -
mae: 0.0461 - val_loss: 0.0013 - val_mae: 0.0334
Epoch 393/500
33/33 [=====] - 0s 12ms/step - loss: 0.0032 -
mae: 0.0463 - val_loss: 0.0014 - val_mae: 0.0345
Epoch 394/500
33/33 [=====] - 0s 12ms/step - loss: 0.0032 -
mae: 0.0462 - val_loss: 0.0013 - val_mae: 0.0324
Epoch 395/500
33/33 [=====] - 0s 12ms/step - loss: 0.0032 -
mae: 0.0463 - val_loss: 0.0014 - val_mae: 0.0338
Epoch 396/500
33/33 [=====] - 0s 12ms/step - loss: 0.0032 -
mae: 0.0462 - val_loss: 0.0013 - val_mae: 0.0327
Epoch 397/500
33/33 [=====] - 0s 13ms/step - loss: 0.0032 -
mae: 0.0469 - val_loss: 0.0016 - val_mae: 0.0377
Epoch 398/500
33/33 [=====] - 0s 12ms/step - loss: 0.0032 -
mae: 0.0460 - val_loss: 0.0013 - val_mae: 0.0330
Epoch 399/500
33/33 [=====] - 0s 12ms/step - loss: 0.0032 -
mae: 0.0469 - val_loss: 0.0013 - val_mae: 0.0327
Epoch 400/500
33/33 [=====] - 0s 12ms/step - loss: 0.0032 -
mae: 0.0463 - val_loss: 0.0013 - val_mae: 0.0328
Epoch 401/500
33/33 [=====] - 0s 14ms/step - loss: 0.0032 -
mae: 0.0462 - val_loss: 0.0013 - val_mae: 0.0323
Epoch 402/500
33/33 [=====] - 0s 12ms/step - loss: 0.0032 -
mae: 0.0461 - val_loss: 0.0012 - val_mae: 0.0320
Epoch 403/500
33/33 [=====] - 0s 13ms/step - loss: 0.0032 -
mae: 0.0468 - val_loss: 0.0015 - val_mae: 0.0368
Epoch 404/500
33/33 [=====] - 0s 12ms/step - loss: 0.0033 -
```

```
mae: 0.0469 - val_loss: 0.0014 - val_mae: 0.0339
Epoch 405/500
33/33 [=====] - 0s 12ms/step - loss: 0.0032 -
mae: 0.0461 - val_loss: 0.0013 - val_mae: 0.0322
Epoch 406/500
33/33 [=====] - 0s 12ms/step - loss: 0.0032 -
mae: 0.0461 - val_loss: 0.0013 - val_mae: 0.0320
Epoch 407/500
33/33 [=====] - 0s 12ms/step - loss: 0.0032 -
mae: 0.0463 - val_loss: 0.0013 - val_mae: 0.0324
Epoch 408/500
33/33 [=====] - 0s 12ms/step - loss: 0.0032 -
mae: 0.0460 - val_loss: 0.0013 - val_mae: 0.0320
Epoch 409/500
33/33 [=====] - 0s 12ms/step - loss: 0.0032 -
mae: 0.0462 - val_loss: 0.0013 - val_mae: 0.0330
Epoch 410/500
33/33 [=====] - 0s 13ms/step - loss: 0.0032 -
mae: 0.0467 - val_loss: 0.0012 - val_mae: 0.0320
Epoch 411/500
33/33 [=====] - 0s 12ms/step - loss: 0.0032 -
mae: 0.0461 - val_loss: 0.0013 - val_mae: 0.0329
Epoch 412/500
33/33 [=====] - 0s 12ms/step - loss: 0.0032 -
mae: 0.0465 - val_loss: 0.0013 - val_mae: 0.0327
Epoch 413/500
33/33 [=====] - 0s 12ms/step - loss: 0.0032 -
mae: 0.0460 - val_loss: 0.0013 - val_mae: 0.0330
Epoch 414/500
33/33 [=====] - 0s 12ms/step - loss: 0.0032 -
mae: 0.0468 - val_loss: 0.0012 - val_mae: 0.0318
Epoch 415/500
33/33 [=====] - 0s 12ms/step - loss: 0.0032 -
mae: 0.0468 - val_loss: 0.0013 - val_mae: 0.0324
Epoch 416/500
33/33 [=====] - 0s 12ms/step - loss: 0.0032 -
mae: 0.0462 - val_loss: 0.0013 - val_mae: 0.0320
Epoch 417/500
```

```
33/33  [=====] - 0s 12ms/step - loss: 0.0032 -  
mae: 0.0465 - val_loss: 0.0014 - val_mae: 0.0342  
Epoch 418/500  
33/33  [=====] - 0s 12ms/step - loss: 0.0032 -  
mae: 0.0463 - val_loss: 0.0012 - val_mae: 0.0321  
Epoch 419/500  
33/33  [=====] - 0s 12ms/step - loss: 0.0032 -  
mae: 0.0462 - val_loss: 0.0013 - val_mae: 0.0323  
Epoch 420/500  
33/33  [=====] - 0s 12ms/step - loss: 0.0032 -  
mae: 0.0462 - val_loss: 0.0014 - val_mae: 0.0337  
Epoch 421/500  
33/33  [=====] - 0s 12ms/step - loss: 0.0032 -  
mae: 0.0460 - val_loss: 0.0013 - val_mae: 0.0322  
Epoch 422/500  
33/33  [=====] - 0s 12ms/step - loss: 0.0032 -  
mae: 0.0465 - val_loss: 0.0014 - val_mae: 0.0347  
Epoch 423/500  
33/33  [=====] - 0s 13ms/step - loss: 0.0032 -  
mae: 0.0464 - val_loss: 0.0013 - val_mae: 0.0323  
Epoch 424/500  
33/33  [=====] - 0s 12ms/step - loss: 0.0032 -  
mae: 0.0462 - val_loss: 0.0013 - val_mae: 0.0321  
Epoch 425/500  
33/33  [=====] - 0s 12ms/step - loss: 0.0032 -  
mae: 0.0458 - val_loss: 0.0012 - val_mae: 0.0322  
Epoch 426/500  
33/33  [=====] - 0s 12ms/step - loss: 0.0032 -  
mae: 0.0465 - val_loss: 0.0013 - val_mae: 0.0334  
Epoch 427/500  
33/33  [=====] - 0s 12ms/step - loss: 0.0032 -  
mae: 0.0473 - val_loss: 0.0013 - val_mae: 0.0321  
Epoch 428/500  
33/33  [=====] - 0s 12ms/step - loss: 0.0032 -  
mae: 0.0458 - val_loss: 0.0013 - val_mae: 0.0327  
Epoch 429/500  
33/33  [=====] - 0s 13ms/step - loss: 0.0032 -  
mae: 0.0461 - val_loss: 0.0012 - val_mae: 0.0321
```

```

Epoch 430/500
33/33 [=====] - 0s 12ms/step - loss: 0.0032 - 
mae: 0.0462 - val_loss: 0.0013 - val_mae: 0.0320
Epoch 431/500
33/33 [=====] - 0s 12ms/step - loss: 0.0032 - 
mae: 0.0459 - val_loss: 0.0013 - val_mae: 0.0326
Epoch 432/500
33/33 [=====] - 0s 12ms/step - loss: 0.0032 - 
mae: 0.0465 - val_loss: 0.0013 - val_mae: 0.0333
Epoch 433/500
33/33 [=====] - 0s 12ms/step - loss: 0.0032 - 
mae: 0.0460 - val_loss: 0.0013 - val_mae: 0.0323
Epoch 434/500
33/33 [=====] - 0s 12ms/step - loss: 0.0032 - 
mae: 0.0463 - val_loss: 0.0014 - val_mae: 0.0335

valid_loss, valid_mae = ahead_model.evaluate(ahead_valid_ds)
print("Validation MAE: ", valid_mae * 1e6)

```

```

3/3 [=====] - 0s 3ms/step - loss: 0.0012 - mae: 0.0317
Validation MAE: 31744.718551635742

```

After training this model, we can predict the next N values at once:

```

X = mulvar_valid.to_numpy()[np.newaxis, :seq_length] # shape [1, 56, 5]
Y_pred = ahead_model.predict(X) # shape [1, 14]

print(Y_pred)

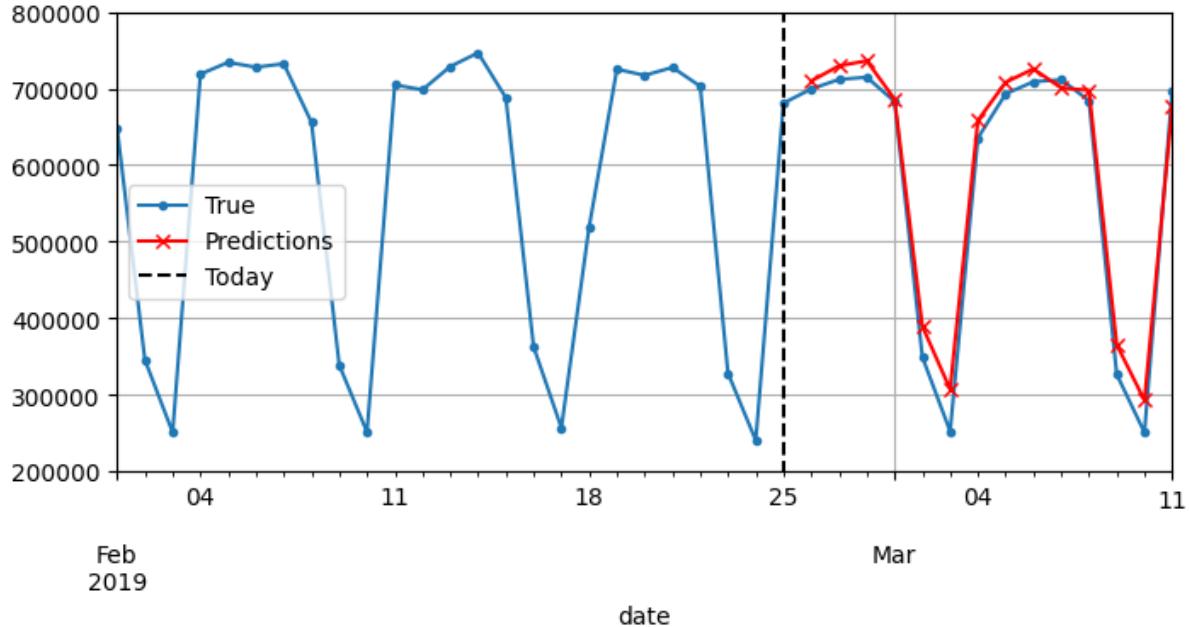
1/1 [=====] - 0s 19ms/step
[[0.70980793 0.7296526 0.7361691 0.68608063 0.389592 0.3071083
 0.6583105 0.7080724 0.72561973 0.70070535 0.697616 0.3654149
 0.29267445 0.67664796]]

Y_pred = pd.Series(Y_pred[0], index=pd.date_range("2019-02-26", "2019-03-11"))
fig, ax = plt.subplots(figsize=(8, 3.5))
(mulvar_valid.rail * 1e6)[["2019-02-01":"2019-03-11"].plot(label="True", marker=".",
    ↳ ax=ax)

```

```
(Y_pred * 1e6).plot(label="Predictions", grid=True, marker="x", color="r", ax=ax)
ax.vlines("2019-02-25", 0, 1e6, color="k", linestyle="--", label="Today")
ax.set_ylim([200_000, 800_000])
plt.legend(loc="center left")

plt.show()
```



The forecasts of this approach for the next day are obviously better than its forecasts for N days into the future, but it doesn't accumulate errors like the previous one.

The third option is **using a sequence-to-sequence model**: instead of training the model to forecast the next N values only at the very last time step, we can train it to forecast the next N values at each and every time step.

The advantage of this technique is that the loss will contain a term for the output of the RNN at each and every time step, not just for the output at the last time step. This means there will be many more error gradients flowing through the model, and they won't have to flow through time as much since they will come from the output of each time step, not just the last one. This will both stabilize and speed up training.

To be clear, at time step 0 the model will output a vector containing the forecasts for time steps 1 to N, then at time step 1 the model will forecast time steps 2 to N+1, and so on. In other words, the targets are sequences of consecutive windows, shifted by one time step at each time step. The target is not a vector anymore, but a sequence of the same length as the inputs, containing

a N-dimensional vector at each step.

Preparing the datasets is not trivial, since each instance has a window as input and a sequence of windows as output. A way to get the result is to use the `window()` method that returns a dataset of window datasets:

```
data = tf.data.Dataset.range(7)

dummy_dataset = data.window(3, shift=1)

for window in dummy_dataset:
    for element in window:
        print(f"{element}", end=" ")
    print()
```

```
0 1 2
1 2 3
2 3 4
3 4 5
4 5 6
5 6
6
```

```
2023-07-10 14:03:50.360099: W tensorflow/core/framework/dataset.cc:807] Input of Window will
2023-07-10 14:03:50.360322: I tensorflow/core/common_runtime/executor.cc:1197] [/device:CPU
[[{{node Placeholder/_0}}]]
2023-07-10 14:03:50.363667: I tensorflow/core/common_runtime/executor.cc:1197] [/device:CPU
[[{{node Placeholder/_0}}]]
2023-07-10 14:03:50.366503: I tensorflow/core/common_runtime/executor.cc:1197] [/device:CPU
[[{{node Placeholder/_0}}]]
2023-07-10 14:03:50.369415: I tensorflow/core/common_runtime/executor.cc:1197] [/device:CPU
[[{{node Placeholder/_0}}]]
2023-07-10 14:03:50.372394: I tensorflow/core/common_runtime/executor.cc:1197] [/device:CPU
[[{{node Placeholder/_0}}]]
2023-07-10 14:03:50.375200: I tensorflow/core/common_runtime/executor.cc:1197] [/device:CPU
[[{{node Placeholder/_0}}]]
2023-07-10 14:03:50.378053: I tensorflow/core/common_runtime/executor.cc:1197] [/device:CPU
[[{{node Placeholder/_0}}]]
```

In the example, the dataset contains six windows, each shifted by one step compared to the previous one, and the last three windows are smaller because they've reached the end of the series. In general, we can get rid of these smaller windows by passing `drop_remainder=True` to the method.

However, `window()` returns a nested dataset (a list of lists), useful when we want to transform each window by calling its methods, but we cannot use it directly for training: our model expects tensors as not datasets. Therefore, we can use `flat_map()` method, which converts a nested dataset into a flat one:

```
dummy_dataset = data.window(3, shift=1, drop_remainder=True)
dummy_dataset = dummy_dataset.flat_map(lambda window: window.batch(3))
list(dummy_dataset)
```

```
[<tf.Tensor: shape=(3,), dtype=int64, numpy=array([0, 1, 2])>,
 <tf.Tensor: shape=(3,), dtype=int64, numpy=array([1, 2, 3])>,
 <tf.Tensor: shape=(3,), dtype=int64, numpy=array([2, 3, 4])>,
 <tf.Tensor: shape=(3,), dtype=int64, numpy=array([3, 4, 5])>,
 <tf.Tensor: shape=(3,), dtype=int64, numpy=array([4, 5, 6])>]
```

We now have a dataset containing consecutive windows represented as tensors. We can repeat the same procedure to get windows of consecutive windows:

```
dummy_dataset = dummy_dataset.window(4, shift=1, drop_remainder=True)
dummy_dataset = dummy_dataset.flat_map(lambda window: window.batch(4))
list(dummy_dataset)
```

```
[<tf.Tensor: shape=(4, 3), dtype=int64, numpy=
 array([[0, 1, 2],
        [1, 2, 3],
        [2, 3, 4],
        [3, 4, 5]])>,
 <tf.Tensor: shape=(4, 3), dtype=int64, numpy=
 array([[1, 2, 3],
        [2, 3, 4],
        [3, 4, 5],
        [4, 5, 6]])>]
```

Finally, we can use the `map()` method to split these windows of windows into inputs and targets:

```
dummy_dataset = dummy_dataset.map(lambda s: (s[:, 0], s[:, 1:]))
list(dummy_dataset)

[(<tf.Tensor: shape=(4,), dtype=int64, numpy=array([0, 1, 2, 3])>,
 <tf.Tensor: shape=(4, 2), dtype=int64, numpy=
 array([[1, 2],
        [2, 3],
        [3, 4],
        [4, 5]]>),
 (<tf.Tensor: shape=(4,), dtype=int64, numpy=array([1, 2, 3, 4])>,
 <tf.Tensor: shape=(4, 2), dtype=int64, numpy=
 array([[2, 3],
        [3, 4],
        [4, 5],
        [5, 6]])>)]
```

Now the dataset contains sequences of length 4 as inputs, and the targets are sequences containing the next two steps, for each time step. For example, the first input sequence is $[0, 1, 2, 3]$ has a corresponding target as $[[1, 2], [2, 3], [3, 4], [4, 5]]$, which are the next two values for each time step.

Now we can create an utility function to prepare the datasets for sequence-to-sequence models:

```
def to_seq2seq_dataset(series, seq_length=56, ahead=14, target_col=1, batch_size=32):
    ds = tf.data.Dataset.from_tensor_slices(series)
    ds = ds.window(ahead + 1, shift=1, drop_remainder=True)
    ds = ds.flat_map(lambda window: window.batch(ahead + 1))
    ds = ds.window(seq_length, shift=1, drop_remainder=True)
    ds = ds.flat_map(lambda window: window.batch(seq_length))
    ds = ds.map(lambda s: (s[:, 0], s[:, 1:, 1]))
    return ds.batch(batch_size)
```

Now we can create the datasets for our running example:

```
seq2seq_train = to_seq2seq_dataset(mulvar_train)
seq2seq_valid = to_seq2seq_dataset(mulvar_valid)
```

And lastly, we can build the sequence-to-sequence model

```
seq2seq_model = tf.keras.Sequential([
    tf.keras.layers.SimpleRNN(32, return_sequences=True, input_shape=[None, 2]),
    tf.keras.layers.Dense(14)
])
```

It is almost identical to previous model: the only difference is that we set `return_sequences=True` in the RNN layer. This way, it will output a sequence of vectors, instead of outputting a single vector at the last time step:

```
opt = tf.keras.optimizers.SGD(learning_rate=0.02, momentum=0.9)
seq2seq_model.compile(loss=tf.keras.losses.Huber(), optimizer=opt, metrics=["mae"])

history = seq2seq_model.fit(seq2seq_train, validation_data=seq2seq_valid, epochs=500,
                           callbacks=[early_stopping_cb])
```

```
Epoch 1/1000
1/Unknown - 0s 17ms/step - loss: 0.0021 - mae: 0.0440

33/33  [=====] - 0s 13ms/step - loss: 0.0037 -
mae: 0.0526 - val_loss: 0.0039 - val_mae: 0.0485
Epoch 2/1000
33/33  [=====] - 0s 9ms/step - loss: 0.0037 -
mae: 0.0526 - val_loss: 0.0039 - val_mae: 0.0485
Epoch 3/1000
33/33  [=====] - 0s 9ms/step - loss: 0.0037 -
mae: 0.0526 - val_loss: 0.0039 - val_mae: 0.0485
Epoch 4/1000
33/33  [=====] - 0s 9ms/step - loss: 0.0037 -
mae: 0.0526 - val_loss: 0.0039 - val_mae: 0.0485
Epoch 5/1000
33/33  [=====] - 0s 9ms/step - loss: 0.0037 -
mae: 0.0526 - val_loss: 0.0039 - val_mae: 0.0485
```

```
Epoch 6/1000
33/33  [=====] - 0s 9ms/step - loss: 0.0037 -
mae: 0.0526 - val_loss: 0.0039 - val_mae: 0.0485
Epoch 7/1000
33/33  [=====] - 0s 9ms/step - loss: 0.0037 -
mae: 0.0526 - val_loss: 0.0039 - val_mae: 0.0485
Epoch 8/1000
33/33  [=====] - 0s 9ms/step - loss: 0.0037 -
mae: 0.0526 - val_loss: 0.0039 - val_mae: 0.0485
Epoch 9/1000
33/33  [=====] - 0s 9ms/step - loss: 0.0037 -
mae: 0.0526 - val_loss: 0.0039 - val_mae: 0.0485
Epoch 10/1000
33/33  [=====] - 0s 9ms/step - loss: 0.0037 -
mae: 0.0526 - val_loss: 0.0039 - val_mae: 0.0485
Epoch 11/1000
33/33  [=====] - 0s 9ms/step - loss: 0.0037 -
mae: 0.0526 - val_loss: 0.0039 - val_mae: 0.0485
Epoch 12/1000
33/33  [=====] - 0s 9ms/step - loss: 0.0037 -
mae: 0.0526 - val_loss: 0.0039 - val_mae: 0.0485
Epoch 13/1000
33/33  [=====] - 0s 9ms/step - loss: 0.0037 -
mae: 0.0526 - val_loss: 0.0039 - val_mae: 0.0485
Epoch 14/1000
33/33  [=====] - 0s 9ms/step - loss: 0.0037 -
mae: 0.0526 - val_loss: 0.0039 - val_mae: 0.0485
Epoch 15/1000
33/33  [=====] - 0s 9ms/step - loss: 0.0037 -
mae: 0.0525 - val_loss: 0.0039 - val_mae: 0.0485
Epoch 16/1000
33/33  [=====] - 0s 9ms/step - loss: 0.0037 -
mae: 0.0525 - val_loss: 0.0039 - val_mae: 0.0485
Epoch 17/1000
33/33  [=====] - 0s 9ms/step - loss: 0.0037 -
mae: 0.0525 - val_loss: 0.0039 - val_mae: 0.0485
Epoch 18/1000
33/33  [=====] - 0s 10ms/step - loss: 0.0037 -
```

```
mae: 0.0525 - val_loss: 0.0039 - val_mae: 0.0485
Epoch 19/1000
33/33 [=====] - 0s 9ms/step - loss: 0.0037 -
mae: 0.0525 - val_loss: 0.0039 - val_mae: 0.0485
Epoch 20/1000
33/33 [=====] - 0s 10ms/step - loss: 0.0037 -
mae: 0.0525 - val_loss: 0.0039 - val_mae: 0.0485
Epoch 21/1000
33/33 [=====] - 0s 10ms/step - loss: 0.0037 -
mae: 0.0525 - val_loss: 0.0039 - val_mae: 0.0485
Epoch 22/1000
33/33 [=====] - 0s 10ms/step - loss: 0.0037 -
mae: 0.0525 - val_loss: 0.0039 - val_mae: 0.0485
Epoch 23/1000
33/33 [=====] - 0s 9ms/step - loss: 0.0037 -
mae: 0.0525 - val_loss: 0.0039 - val_mae: 0.0485
Epoch 24/1000
33/33 [=====] - 0s 9ms/step - loss: 0.0037 -
mae: 0.0525 - val_loss: 0.0039 - val_mae: 0.0485
Epoch 25/1000
33/33 [=====] - 0s 9ms/step - loss: 0.0037 -
mae: 0.0525 - val_loss: 0.0039 - val_mae: 0.0485
Epoch 26/1000
33/33 [=====] - 0s 9ms/step - loss: 0.0037 -
mae: 0.0525 - val_loss: 0.0039 - val_mae: 0.0485
Epoch 27/1000
33/33 [=====] - 0s 9ms/step - loss: 0.0037 -
mae: 0.0525 - val_loss: 0.0039 - val_mae: 0.0485
Epoch 28/1000
33/33 [=====] - 0s 9ms/step - loss: 0.0037 -
mae: 0.0525 - val_loss: 0.0039 - val_mae: 0.0485
Epoch 29/1000
33/33 [=====] - 0s 9ms/step - loss: 0.0037 -
mae: 0.0525 - val_loss: 0.0039 - val_mae: 0.0485
Epoch 30/1000
33/33 [=====] - 0s 9ms/step - loss: 0.0037 -
mae: 0.0525 - val_loss: 0.0039 - val_mae: 0.0485
Epoch 31/1000
```

```
33/33  [=====] - 0s 9ms/step - loss: 0.0037 -  
mae: 0.0525 - val_loss: 0.0039 - val_mae: 0.0485  
Epoch 32/1000  
33/33  [=====] - 0s 9ms/step - loss: 0.0037 -  
mae: 0.0525 - val_loss: 0.0039 - val_mae: 0.0485  
Epoch 33/1000  
33/33  [=====] - 0s 9ms/step - loss: 0.0037 -  
mae: 0.0525 - val_loss: 0.0039 - val_mae: 0.0485  
Epoch 34/1000  
33/33  [=====] - 0s 9ms/step - loss: 0.0037 -  
mae: 0.0525 - val_loss: 0.0039 - val_mae: 0.0485  
Epoch 35/1000  
33/33  [=====] - 0s 9ms/step - loss: 0.0037 -  
mae: 0.0525 - val_loss: 0.0039 - val_mae: 0.0485  
Epoch 36/1000  
33/33  [=====] - 0s 10ms/step - loss: 0.0037 -  
mae: 0.0525 - val_loss: 0.0039 - val_mae: 0.0485  
Epoch 37/1000  
33/33  [=====] - 0s 9ms/step - loss: 0.0037 -  
mae: 0.0525 - val_loss: 0.0039 - val_mae: 0.0485  
Epoch 38/1000  
33/33  [=====] - 0s 10ms/step - loss: 0.0037 -  
mae: 0.0525 - val_loss: 0.0039 - val_mae: 0.0485  
Epoch 39/1000  
33/33  [=====] - 0s 11ms/step - loss: 0.0037 -  
mae: 0.0525 - val_loss: 0.0039 - val_mae: 0.0485  
Epoch 40/1000  
33/33  [=====] - 0s 9ms/step - loss: 0.0037 -  
mae: 0.0525 - val_loss: 0.0039 - val_mae: 0.0485  
Epoch 41/1000  
33/33  [=====] - 0s 9ms/step - loss: 0.0037 -  
mae: 0.0525 - val_loss: 0.0039 - val_mae: 0.0485  
Epoch 42/1000  
33/33  [=====] - 0s 9ms/step - loss: 0.0037 -  
mae: 0.0525 - val_loss: 0.0039 - val_mae: 0.0485  
Epoch 43/1000  
33/33  [=====] - 0s 9ms/step - loss: 0.0037 -  
mae: 0.0525 - val_loss: 0.0039 - val_mae: 0.0485
```

```
Epoch 44/1000
33/33  [=====] - 0s 9ms/step - loss: 0.0037 -
mae: 0.0525 - val_loss: 0.0039 - val_mae: 0.0485
Epoch 45/1000
33/33  [=====] - 0s 9ms/step - loss: 0.0037 -
mae: 0.0525 - val_loss: 0.0039 - val_mae: 0.0485
Epoch 46/1000
33/33  [=====] - 0s 9ms/step - loss: 0.0037 -
mae: 0.0525 - val_loss: 0.0039 - val_mae: 0.0485
Epoch 47/1000
33/33  [=====] - 0s 9ms/step - loss: 0.0037 -
mae: 0.0525 - val_loss: 0.0039 - val_mae: 0.0485
Epoch 48/1000
33/33  [=====] - 0s 9ms/step - loss: 0.0037 -
mae: 0.0525 - val_loss: 0.0039 - val_mae: 0.0485
Epoch 49/1000
33/33  [=====] - 0s 9ms/step - loss: 0.0037 -
mae: 0.0525 - val_loss: 0.0039 - val_mae: 0.0485
Epoch 50/1000
33/33  [=====] - 0s 9ms/step - loss: 0.0037 -
mae: 0.0525 - val_loss: 0.0039 - val_mae: 0.0485
Epoch 51/1000
33/33  [=====] - 0s 9ms/step - loss: 0.0037 -
mae: 0.0525 - val_loss: 0.0039 - val_mae: 0.0485
Epoch 52/1000
33/33  [=====] - 0s 9ms/step - loss: 0.0037 -
mae: 0.0525 - val_loss: 0.0039 - val_mae: 0.0485
Epoch 53/1000
33/33  [=====] - 0s 9ms/step - loss: 0.0037 -
mae: 0.0525 - val_loss: 0.0039 - val_mae: 0.0485
Epoch 54/1000
33/33  [=====] - 0s 9ms/step - loss: 0.0037 -
mae: 0.0524 - val_loss: 0.0039 - val_mae: 0.0484
Epoch 55/1000
33/33  [=====] - 0s 9ms/step - loss: 0.0037 -
mae: 0.0524 - val_loss: 0.0039 - val_mae: 0.0484
Epoch 56/1000
33/33  [=====] - 0s 10ms/step - loss: 0.0037 -
```

```
mae: 0.0524 - val_loss: 0.0039 - val_mae: 0.0484
Epoch 57/1000
33/33 [=====] - 0s 9ms/step - loss: 0.0037 -
mae: 0.0524 - val_loss: 0.0039 - val_mae: 0.0484
Epoch 58/1000
33/33 [=====] - 0s 9ms/step - loss: 0.0037 -
mae: 0.0524 - val_loss: 0.0039 - val_mae: 0.0484
Epoch 59/1000
33/33 [=====] - 0s 9ms/step - loss: 0.0037 -
mae: 0.0524 - val_loss: 0.0039 - val_mae: 0.0484
Epoch 60/1000
33/33 [=====] - 0s 9ms/step - loss: 0.0037 -
mae: 0.0524 - val_loss: 0.0039 - val_mae: 0.0484
Epoch 61/1000
33/33 [=====] - 0s 9ms/step - loss: 0.0037 -
mae: 0.0524 - val_loss: 0.0039 - val_mae: 0.0484
Epoch 62/1000
33/33 [=====] - 0s 9ms/step - loss: 0.0037 -
mae: 0.0524 - val_loss: 0.0039 - val_mae: 0.0484
Epoch 63/1000
33/33 [=====] - 0s 9ms/step - loss: 0.0037 -
mae: 0.0524 - val_loss: 0.0039 - val_mae: 0.0484
Epoch 64/1000
33/33 [=====] - 0s 9ms/step - loss: 0.0037 -
mae: 0.0524 - val_loss: 0.0039 - val_mae: 0.0484
Epoch 65/1000
33/33 [=====] - 0s 10ms/step - loss: 0.0037 -
mae: 0.0524 - val_loss: 0.0039 - val_mae: 0.0484
Epoch 66/1000
33/33 [=====] - 0s 10ms/step - loss: 0.0037 -
mae: 0.0524 - val_loss: 0.0039 - val_mae: 0.0484
Epoch 67/1000
33/33 [=====] - 0s 11ms/step - loss: 0.0037 -
mae: 0.0524 - val_loss: 0.0039 - val_mae: 0.0484
Epoch 68/1000
33/33 [=====] - 0s 9ms/step - loss: 0.0037 -
mae: 0.0524 - val_loss: 0.0039 - val_mae: 0.0484
Epoch 69/1000
```

```
33/33  [=====] - 0s 9ms/step - loss: 0.0037 -
mae: 0.0524 - val_loss: 0.0039 - val_mae: 0.0484
Epoch 70/1000
33/33  [=====] - 0s 9ms/step - loss: 0.0037 -
mae: 0.0524 - val_loss: 0.0039 - val_mae: 0.0484
Epoch 71/1000
33/33  [=====] - 0s 9ms/step - loss: 0.0037 -
mae: 0.0524 - val_loss: 0.0039 - val_mae: 0.0484
Epoch 72/1000
33/33  [=====] - 0s 10ms/step - loss: 0.0037 -
mae: 0.0524 - val_loss: 0.0039 - val_mae: 0.0484
Epoch 73/1000
33/33  [=====] - 0s 9ms/step - loss: 0.0037 -
mae: 0.0524 - val_loss: 0.0039 - val_mae: 0.0484
Epoch 74/1000
33/33  [=====] - 0s 9ms/step - loss: 0.0037 -
mae: 0.0524 - val_loss: 0.0039 - val_mae: 0.0484
Epoch 75/1000
33/33  [=====] - 0s 9ms/step - loss: 0.0037 -
mae: 0.0524 - val_loss: 0.0039 - val_mae: 0.0484
Epoch 76/1000
33/33  [=====] - 0s 9ms/step - loss: 0.0037 -
mae: 0.0524 - val_loss: 0.0039 - val_mae: 0.0484
Epoch 77/1000
33/33  [=====] - 0s 9ms/step - loss: 0.0037 -
mae: 0.0524 - val_loss: 0.0039 - val_mae: 0.0484
Epoch 78/1000
33/33  [=====] - 0s 9ms/step - loss: 0.0037 -
mae: 0.0524 - val_loss: 0.0039 - val_mae: 0.0484
Epoch 79/1000
33/33  [=====] - 0s 12ms/step - loss: 0.0037 -
mae: 0.0524 - val_loss: 0.0039 - val_mae: 0.0484
Epoch 80/1000
33/33  [=====] - 0s 9ms/step - loss: 0.0037 -
mae: 0.0524 - val_loss: 0.0039 - val_mae: 0.0484
Epoch 81/1000
33/33  [=====] - 0s 9ms/step - loss: 0.0037 -
mae: 0.0524 - val_loss: 0.0039 - val_mae: 0.0484
```

```
Epoch 82/1000
33/33  [=====] - 0s 9ms/step - loss: 0.0037 -
mae: 0.0524 - val_loss: 0.0039 - val_mae: 0.0484
Epoch 83/1000
33/33  [=====] - 0s 9ms/step - loss: 0.0037 -
mae: 0.0524 - val_loss: 0.0039 - val_mae: 0.0484
Epoch 84/1000
33/33  [=====] - 0s 9ms/step - loss: 0.0037 -
mae: 0.0524 - val_loss: 0.0039 - val_mae: 0.0484
Epoch 85/1000
33/33  [=====] - 0s 9ms/step - loss: 0.0037 -
mae: 0.0524 - val_loss: 0.0039 - val_mae: 0.0484
Epoch 86/1000
33/33  [=====] - 0s 9ms/step - loss: 0.0037 -
mae: 0.0524 - val_loss: 0.0039 - val_mae: 0.0484
Epoch 87/1000
33/33  [=====] - 0s 9ms/step - loss: 0.0037 -
mae: 0.0524 - val_loss: 0.0039 - val_mae: 0.0484
Epoch 88/1000
33/33  [=====] - 0s 9ms/step - loss: 0.0037 -
mae: 0.0524 - val_loss: 0.0039 - val_mae: 0.0484
Epoch 89/1000
33/33  [=====] - 0s 9ms/step - loss: 0.0037 -
mae: 0.0524 - val_loss: 0.0039 - val_mae: 0.0484
Epoch 90/1000
33/33  [=====] - 0s 9ms/step - loss: 0.0037 -
mae: 0.0524 - val_loss: 0.0039 - val_mae: 0.0484
Epoch 91/1000
33/33  [=====] - 0s 9ms/step - loss: 0.0037 -
mae: 0.0524 - val_loss: 0.0039 - val_mae: 0.0484
Epoch 92/1000
33/33  [=====] - 0s 9ms/step - loss: 0.0037 -
mae: 0.0524 - val_loss: 0.0039 - val_mae: 0.0484
Epoch 93/1000
33/33  [=====] - 0s 9ms/step - loss: 0.0037 -
mae: 0.0524 - val_loss: 0.0039 - val_mae: 0.0484
Epoch 94/1000
33/33  [=====] - 0s 9ms/step - loss: 0.0037 -
```

```
mae: 0.0524 - val_loss: 0.0039 - val_mae: 0.0484
Epoch 95/1000
33/33 [=====] - 0s 10ms/step - loss: 0.0037 -
mae: 0.0523 - val_loss: 0.0039 - val_mae: 0.0484
Epoch 96/1000
33/33 [=====] - 0s 9ms/step - loss: 0.0037 -
mae: 0.0523 - val_loss: 0.0039 - val_mae: 0.0484
Epoch 97/1000
33/33 [=====] - 0s 9ms/step - loss: 0.0037 -
mae: 0.0523 - val_loss: 0.0039 - val_mae: 0.0484
Epoch 98/1000
33/33 [=====] - 0s 9ms/step - loss: 0.0037 -
mae: 0.0523 - val_loss: 0.0039 - val_mae: 0.0484
Epoch 99/1000
33/33 [=====] - 0s 9ms/step - loss: 0.0037 -
mae: 0.0523 - val_loss: 0.0039 - val_mae: 0.0484
Epoch 100/1000
33/33 [=====] - 0s 9ms/step - loss: 0.0037 -
mae: 0.0523 - val_loss: 0.0039 - val_mae: 0.0484
Epoch 101/1000
33/33 [=====] - 0s 9ms/step - loss: 0.0037 -
mae: 0.0523 - val_loss: 0.0039 - val_mae: 0.0484
Epoch 102/1000
33/33 [=====] - 0s 9ms/step - loss: 0.0037 -
mae: 0.0523 - val_loss: 0.0039 - val_mae: 0.0484
Epoch 103/1000
33/33 [=====] - 0s 9ms/step - loss: 0.0037 -
mae: 0.0523 - val_loss: 0.0039 - val_mae: 0.0484
Epoch 104/1000
33/33 [=====] - 0s 9ms/step - loss: 0.0037 -
mae: 0.0523 - val_loss: 0.0039 - val_mae: 0.0484
Epoch 105/1000
33/33 [=====] - 0s 9ms/step - loss: 0.0037 -
mae: 0.0523 - val_loss: 0.0039 - val_mae: 0.0484
Epoch 106/1000
33/33 [=====] - 0s 9ms/step - loss: 0.0037 -
mae: 0.0523 - val_loss: 0.0039 - val_mae: 0.0484
Epoch 107/1000
```

```
33/33  [=====] - 0s 9ms/step - loss: 0.0037 -  
mae: 0.0523 - val_loss: 0.0039 - val_mae: 0.0484  
Epoch 108/1000  
33/33  [=====] - 0s 9ms/step - loss: 0.0037 -  
mae: 0.0523 - val_loss: 0.0039 - val_mae: 0.0484  
Epoch 109/1000  
33/33  [=====] - 0s 9ms/step - loss: 0.0037 -  
mae: 0.0523 - val_loss: 0.0039 - val_mae: 0.0484  
Epoch 110/1000  
33/33  [=====] - 0s 9ms/step - loss: 0.0037 -  
mae: 0.0523 - val_loss: 0.0039 - val_mae: 0.0484  
Epoch 111/1000  
33/33  [=====] - 0s 9ms/step - loss: 0.0037 -  
mae: 0.0523 - val_loss: 0.0039 - val_mae: 0.0484  
Epoch 112/1000  
33/33  [=====] - 0s 10ms/step - loss: 0.0037 -  
mae: 0.0523 - val_loss: 0.0039 - val_mae: 0.0484  
Epoch 113/1000  
33/33  [=====] - 0s 9ms/step - loss: 0.0037 -  
mae: 0.0523 - val_loss: 0.0039 - val_mae: 0.0484  
Epoch 114/1000  
33/33  [=====] - 0s 9ms/step - loss: 0.0037 -  
mae: 0.0523 - val_loss: 0.0039 - val_mae: 0.0484  
Epoch 115/1000  
33/33  [=====] - 0s 9ms/step - loss: 0.0037 -  
mae: 0.0523 - val_loss: 0.0039 - val_mae: 0.0484  
Epoch 116/1000  
33/33  [=====] - 0s 9ms/step - loss: 0.0037 -  
mae: 0.0523 - val_loss: 0.0039 - val_mae: 0.0484  
Epoch 117/1000  
33/33  [=====] - 0s 9ms/step - loss: 0.0037 -  
mae: 0.0523 - val_loss: 0.0039 - val_mae: 0.0484  
Epoch 118/1000  
33/33  [=====] - 0s 9ms/step - loss: 0.0037 -  
mae: 0.0523 - val_loss: 0.0039 - val_mae: 0.0484  
Epoch 119/1000  
33/33  [=====] - 0s 9ms/step - loss: 0.0037 -  
mae: 0.0523 - val_loss: 0.0039 - val_mae: 0.0484
```

```
Epoch 120/1000
33/33  [=====] - 0s 9ms/step - loss: 0.0037 -
mae: 0.0523 - val_loss: 0.0039 - val_mae: 0.0484
Epoch 121/1000
33/33  [=====] - 0s 9ms/step - loss: 0.0037 -
mae: 0.0523 - val_loss: 0.0039 - val_mae: 0.0484
Epoch 122/1000
33/33  [=====] - 0s 9ms/step - loss: 0.0037 -
mae: 0.0523 - val_loss: 0.0039 - val_mae: 0.0484
Epoch 123/1000
33/33  [=====] - 0s 9ms/step - loss: 0.0037 -
mae: 0.0523 - val_loss: 0.0039 - val_mae: 0.0484
Epoch 124/1000
33/33  [=====] - 0s 9ms/step - loss: 0.0037 -
mae: 0.0523 - val_loss: 0.0039 - val_mae: 0.0484
Epoch 125/1000
33/33  [=====] - 0s 9ms/step - loss: 0.0037 -
mae: 0.0523 - val_loss: 0.0039 - val_mae: 0.0484
Epoch 126/1000
33/33  [=====] - 0s 9ms/step - loss: 0.0037 -
mae: 0.0523 - val_loss: 0.0039 - val_mae: 0.0484
Epoch 127/1000
33/33  [=====] - 0s 9ms/step - loss: 0.0037 -
mae: 0.0523 - val_loss: 0.0039 - val_mae: 0.0484
Epoch 128/1000
33/33  [=====] - 0s 10ms/step - loss: 0.0037 -
mae: 0.0523 - val_loss: 0.0039 - val_mae: 0.0484
Epoch 129/1000
33/33  [=====] - 0s 9ms/step - loss: 0.0037 -
mae: 0.0523 - val_loss: 0.0039 - val_mae: 0.0484
Epoch 130/1000
33/33  [=====] - 0s 9ms/step - loss: 0.0037 -
mae: 0.0523 - val_loss: 0.0039 - val_mae: 0.0484
Epoch 131/1000
33/33  [=====] - 0s 9ms/step - loss: 0.0037 -
mae: 0.0523 - val_loss: 0.0039 - val_mae: 0.0484
Epoch 132/1000
33/33  [=====] - 0s 9ms/step - loss: 0.0037 -
```

```
mae: 0.0523 - val_loss: 0.0039 - val_mae: 0.0484
Epoch 133/1000
33/33 [=====] - 0s 9ms/step - loss: 0.0037 -
mae: 0.0523 - val_loss: 0.0039 - val_mae: 0.0484
Epoch 134/1000
33/33 [=====] - 0s 9ms/step - loss: 0.0037 -
mae: 0.0523 - val_loss: 0.0039 - val_mae: 0.0484
Epoch 135/1000
33/33 [=====] - 0s 9ms/step - loss: 0.0037 -
mae: 0.0523 - val_loss: 0.0039 - val_mae: 0.0484
Epoch 136/1000
33/33 [=====] - 0s 9ms/step - loss: 0.0037 -
mae: 0.0523 - val_loss: 0.0039 - val_mae: 0.0484
Epoch 137/1000
33/33 [=====] - 0s 9ms/step - loss: 0.0037 -
mae: 0.0523 - val_loss: 0.0039 - val_mae: 0.0484
Epoch 138/1000
33/33 [=====] - 0s 9ms/step - loss: 0.0037 -
mae: 0.0522 - val_loss: 0.0039 - val_mae: 0.0484
Epoch 139/1000
33/33 [=====] - 0s 9ms/step - loss: 0.0037 -
mae: 0.0522 - val_loss: 0.0039 - val_mae: 0.0484
Epoch 140/1000
33/33 [=====] - 0s 9ms/step - loss: 0.0037 -
mae: 0.0522 - val_loss: 0.0039 - val_mae: 0.0484
Epoch 141/1000
33/33 [=====] - 0s 9ms/step - loss: 0.0037 -
mae: 0.0522 - val_loss: 0.0039 - val_mae: 0.0484
Epoch 142/1000
33/33 [=====] - 0s 9ms/step - loss: 0.0037 -
mae: 0.0522 - val_loss: 0.0039 - val_mae: 0.0484
Epoch 143/1000
33/33 [=====] - 0s 11ms/step - loss: 0.0037 -
mae: 0.0522 - val_loss: 0.0039 - val_mae: 0.0484
Epoch 144/1000
33/33 [=====] - 0s 9ms/step - loss: 0.0037 -
mae: 0.0522 - val_loss: 0.0039 - val_mae: 0.0484
Epoch 145/1000
```

```
33/33  [=====] - 0s 9ms/step - loss: 0.0037 -
mae: 0.0522 - val_loss: 0.0039 - val_mae: 0.0484
Epoch 146/1000
33/33  [=====] - 0s 9ms/step - loss: 0.0037 -
mae: 0.0522 - val_loss: 0.0039 - val_mae: 0.0484
Epoch 147/1000
33/33  [=====] - 0s 9ms/step - loss: 0.0037 -
mae: 0.0522 - val_loss: 0.0039 - val_mae: 0.0484
Epoch 148/1000
33/33  [=====] - 0s 9ms/step - loss: 0.0037 -
mae: 0.0522 - val_loss: 0.0039 - val_mae: 0.0484
Epoch 149/1000
33/33  [=====] - 0s 9ms/step - loss: 0.0037 -
mae: 0.0522 - val_loss: 0.0039 - val_mae: 0.0484
Epoch 150/1000
33/33  [=====] - 0s 9ms/step - loss: 0.0037 -
mae: 0.0522 - val_loss: 0.0039 - val_mae: 0.0484
Epoch 151/1000
33/33  [=====] - 0s 9ms/step - loss: 0.0037 -
mae: 0.0522 - val_loss: 0.0039 - val_mae: 0.0484
Epoch 152/1000
33/33  [=====] - 0s 9ms/step - loss: 0.0037 -
mae: 0.0522 - val_loss: 0.0039 - val_mae: 0.0484
Epoch 153/1000
33/33  [=====] - 0s 9ms/step - loss: 0.0037 -
mae: 0.0522 - val_loss: 0.0039 - val_mae: 0.0484
Epoch 154/1000
33/33  [=====] - 0s 9ms/step - loss: 0.0037 -
mae: 0.0522 - val_loss: 0.0039 - val_mae: 0.0484
Epoch 155/1000
33/33  [=====] - 0s 9ms/step - loss: 0.0037 -
mae: 0.0522 - val_loss: 0.0039 - val_mae: 0.0484
Epoch 156/1000
33/33  [=====] - 0s 9ms/step - loss: 0.0037 -
mae: 0.0522 - val_loss: 0.0039 - val_mae: 0.0484
Epoch 157/1000
33/33  [=====] - 0s 11ms/step - loss: 0.0037 -
mae: 0.0522 - val_loss: 0.0039 - val_mae: 0.0484
```

```
Epoch 158/1000
33/33  [=====] - 0s 9ms/step - loss: 0.0037 -
mae: 0.0522 - val_loss: 0.0039 - val_mae: 0.0484
Epoch 159/1000
33/33  [=====] - 0s 9ms/step - loss: 0.0037 -
mae: 0.0522 - val_loss: 0.0039 - val_mae: 0.0484
Epoch 160/1000
33/33  [=====] - 0s 9ms/step - loss: 0.0037 -
mae: 0.0522 - val_loss: 0.0039 - val_mae: 0.0484
Epoch 161/1000
33/33  [=====] - 0s 9ms/step - loss: 0.0037 -
mae: 0.0522 - val_loss: 0.0039 - val_mae: 0.0484
Epoch 162/1000
33/33  [=====] - 0s 9ms/step - loss: 0.0037 -
mae: 0.0522 - val_loss: 0.0039 - val_mae: 0.0484
Epoch 163/1000
33/33  [=====] - 0s 9ms/step - loss: 0.0037 -
mae: 0.0522 - val_loss: 0.0039 - val_mae: 0.0484
Epoch 164/1000
33/33  [=====] - 0s 9ms/step - loss: 0.0037 -
mae: 0.0522 - val_loss: 0.0039 - val_mae: 0.0484
Epoch 165/1000
33/33  [=====] - 0s 9ms/step - loss: 0.0037 -
mae: 0.0522 - val_loss: 0.0039 - val_mae: 0.0484
Epoch 166/1000
33/33  [=====] - 0s 9ms/step - loss: 0.0037 -
mae: 0.0522 - val_loss: 0.0039 - val_mae: 0.0484
Epoch 167/1000
33/33  [=====] - 0s 9ms/step - loss: 0.0037 -
mae: 0.0522 - val_loss: 0.0039 - val_mae: 0.0484
Epoch 168/1000
33/33  [=====] - 0s 9ms/step - loss: 0.0037 -
mae: 0.0522 - val_loss: 0.0039 - val_mae: 0.0484
Epoch 169/1000
33/33  [=====] - 0s 9ms/step - loss: 0.0037 -
mae: 0.0522 - val_loss: 0.0039 - val_mae: 0.0484
Epoch 170/1000
33/33  [=====] - 0s 9ms/step - loss: 0.0037 -
```

```
mae: 0.0522 - val_loss: 0.0039 - val_mae: 0.0484
Epoch 171/1000
33/33 [=====] - 0s 11ms/step - loss: 0.0037 -
mae: 0.0522 - val_loss: 0.0039 - val_mae: 0.0484
Epoch 172/1000
33/33 [=====] - 0s 10ms/step - loss: 0.0037 -
mae: 0.0522 - val_loss: 0.0039 - val_mae: 0.0484
Epoch 173/1000
33/33 [=====] - 0s 9ms/step - loss: 0.0037 -
mae: 0.0522 - val_loss: 0.0039 - val_mae: 0.0484
Epoch 174/1000
33/33 [=====] - 0s 9ms/step - loss: 0.0037 -
mae: 0.0522 - val_loss: 0.0039 - val_mae: 0.0484
Epoch 175/1000
33/33 [=====] - 0s 9ms/step - loss: 0.0037 -
mae: 0.0522 - val_loss: 0.0039 - val_mae: 0.0484
Epoch 176/1000
33/33 [=====] - 0s 10ms/step - loss: 0.0037 -
mae: 0.0522 - val_loss: 0.0039 - val_mae: 0.0484
Epoch 177/1000
33/33 [=====] - 0s 9ms/step - loss: 0.0037 -
mae: 0.0522 - val_loss: 0.0039 - val_mae: 0.0484
Epoch 178/1000
33/33 [=====] - 0s 9ms/step - loss: 0.0037 -
mae: 0.0522 - val_loss: 0.0039 - val_mae: 0.0484
Epoch 179/1000
33/33 [=====] - 0s 9ms/step - loss: 0.0037 -
mae: 0.0522 - val_loss: 0.0039 - val_mae: 0.0484
Epoch 180/1000
33/33 [=====] - 0s 10ms/step - loss: 0.0037 -
mae: 0.0522 - val_loss: 0.0039 - val_mae: 0.0484
Epoch 181/1000
33/33 [=====] - 0s 10ms/step - loss: 0.0037 -
mae: 0.0522 - val_loss: 0.0039 - val_mae: 0.0484
Epoch 182/1000
33/33 [=====] - 0s 11ms/step - loss: 0.0037 -
mae: 0.0521 - val_loss: 0.0039 - val_mae: 0.0484
Epoch 183/1000
```

```
33/33 [=====] - 0s 10ms/step - loss: 0.0037 -
mae: 0.0521 - val_loss: 0.0039 - val_mae: 0.0484
Epoch 184/1000
33/33 [=====] - 0s 11ms/step - loss: 0.0037 -
mae: 0.0521 - val_loss: 0.0039 - val_mae: 0.0484
Epoch 185/1000
33/33 [=====] - 0s 9ms/step - loss: 0.0037 -
mae: 0.0521 - val_loss: 0.0039 - val_mae: 0.0484
Epoch 186/1000
33/33 [=====] - 0s 10ms/step - loss: 0.0037 -
mae: 0.0521 - val_loss: 0.0039 - val_mae: 0.0484
Epoch 187/1000
33/33 [=====] - 0s 10ms/step - loss: 0.0037 -
mae: 0.0521 - val_loss: 0.0039 - val_mae: 0.0484
Epoch 188/1000
33/33 [=====] - 0s 10ms/step - loss: 0.0037 -
mae: 0.0521 - val_loss: 0.0039 - val_mae: 0.0484
Epoch 189/1000
33/33 [=====] - 0s 10ms/step - loss: 0.0037 -
mae: 0.0521 - val_loss: 0.0039 - val_mae: 0.0484
Epoch 190/1000
33/33 [=====] - 0s 10ms/step - loss: 0.0037 -
mae: 0.0521 - val_loss: 0.0039 - val_mae: 0.0484
Epoch 191/1000
33/33 [=====] - 0s 11ms/step - loss: 0.0037 -
mae: 0.0521 - val_loss: 0.0039 - val_mae: 0.0484
Epoch 192/1000
33/33 [=====] - 0s 10ms/step - loss: 0.0037 -
mae: 0.0521 - val_loss: 0.0039 - val_mae: 0.0484
Epoch 193/1000
33/33 [=====] - 0s 11ms/step - loss: 0.0037 -
mae: 0.0521 - val_loss: 0.0039 - val_mae: 0.0484
Epoch 194/1000
33/33 [=====] - 0s 11ms/step - loss: 0.0037 -
mae: 0.0521 - val_loss: 0.0039 - val_mae: 0.0484
Epoch 195/1000
33/33 [=====] - 0s 12ms/step - loss: 0.0037 -
mae: 0.0521 - val_loss: 0.0039 - val_mae: 0.0484
```

```
Epoch 196/1000
33/33 [=====] - 0s 11ms/step - loss: 0.0037 -
mae: 0.0521 - val_loss: 0.0039 - val_mae: 0.0484
Epoch 197/1000
33/33 [=====] - 0s 11ms/step - loss: 0.0037 -
mae: 0.0521 - val_loss: 0.0039 - val_mae: 0.0484
Epoch 198/1000
33/33 [=====] - 0s 10ms/step - loss: 0.0037 -
mae: 0.0521 - val_loss: 0.0039 - val_mae: 0.0484
Epoch 199/1000
33/33 [=====] - 0s 10ms/step - loss: 0.0037 -
mae: 0.0521 - val_loss: 0.0039 - val_mae: 0.0483
Epoch 200/1000
33/33 [=====] - 0s 10ms/step - loss: 0.0037 -
mae: 0.0521 - val_loss: 0.0039 - val_mae: 0.0483
Epoch 201/1000
33/33 [=====] - 0s 9ms/step - loss: 0.0037 -
mae: 0.0521 - val_loss: 0.0039 - val_mae: 0.0483
Epoch 202/1000
33/33 [=====] - 0s 10ms/step - loss: 0.0037 -
mae: 0.0521 - val_loss: 0.0039 - val_mae: 0.0483
Epoch 203/1000
33/33 [=====] - 0s 11ms/step - loss: 0.0037 -
mae: 0.0521 - val_loss: 0.0039 - val_mae: 0.0483
Epoch 204/1000
33/33 [=====] - 0s 10ms/step - loss: 0.0037 -
mae: 0.0521 - val_loss: 0.0039 - val_mae: 0.0483
Epoch 205/1000
33/33 [=====] - 0s 10ms/step - loss: 0.0037 -
mae: 0.0521 - val_loss: 0.0039 - val_mae: 0.0483
Epoch 206/1000
33/33 [=====] - 0s 12ms/step - loss: 0.0037 -
mae: 0.0521 - val_loss: 0.0039 - val_mae: 0.0483
Epoch 207/1000
33/33 [=====] - 0s 10ms/step - loss: 0.0037 -
mae: 0.0521 - val_loss: 0.0039 - val_mae: 0.0483
Epoch 208/1000
33/33 [=====] - 0s 10ms/step - loss: 0.0037 -
```

```
mae: 0.0521 - val_loss: 0.0039 - val_mae: 0.0483
Epoch 209/1000
33/33 [=====] - 0s 10ms/step - loss: 0.0037 -
mae: 0.0521 - val_loss: 0.0039 - val_mae: 0.0483
Epoch 210/1000
33/33 [=====] - 0s 10ms/step - loss: 0.0037 -
mae: 0.0521 - val_loss: 0.0039 - val_mae: 0.0483
Epoch 211/1000
33/33 [=====] - 0s 10ms/step - loss: 0.0037 -
mae: 0.0521 - val_loss: 0.0039 - val_mae: 0.0483
Epoch 212/1000
33/33 [=====] - 0s 10ms/step - loss: 0.0037 -
mae: 0.0521 - val_loss: 0.0039 - val_mae: 0.0483
Epoch 213/1000
33/33 [=====] - 0s 12ms/step - loss: 0.0037 -
mae: 0.0521 - val_loss: 0.0039 - val_mae: 0.0483
Epoch 214/1000
33/33 [=====] - 0s 10ms/step - loss: 0.0037 -
mae: 0.0521 - val_loss: 0.0039 - val_mae: 0.0483
Epoch 215/1000
33/33 [=====] - 0s 13ms/step - loss: 0.0037 -
mae: 0.0521 - val_loss: 0.0039 - val_mae: 0.0483
Epoch 216/1000
33/33 [=====] - 0s 12ms/step - loss: 0.0037 -
mae: 0.0521 - val_loss: 0.0039 - val_mae: 0.0483
Epoch 217/1000
33/33 [=====] - 0s 10ms/step - loss: 0.0037 -
mae: 0.0521 - val_loss: 0.0039 - val_mae: 0.0483
Epoch 218/1000
33/33 [=====] - 0s 10ms/step - loss: 0.0037 -
mae: 0.0521 - val_loss: 0.0039 - val_mae: 0.0483
Epoch 219/1000
33/33 [=====] - 0s 10ms/step - loss: 0.0037 -
mae: 0.0521 - val_loss: 0.0039 - val_mae: 0.0483
Epoch 220/1000
33/33 [=====] - 0s 11ms/step - loss: 0.0037 -
mae: 0.0521 - val_loss: 0.0039 - val_mae: 0.0483
Epoch 221/1000
```

```
33/33  [=====] - 0s 10ms/step - loss: 0.0037 -  
mae: 0.0521 - val_loss: 0.0039 - val_mae: 0.0483  
Epoch 222/1000  
33/33  [=====] - 0s 12ms/step - loss: 0.0037 -  
mae: 0.0521 - val_loss: 0.0039 - val_mae: 0.0483  
Epoch 223/1000  
33/33  [=====] - 0s 11ms/step - loss: 0.0037 -  
mae: 0.0521 - val_loss: 0.0039 - val_mae: 0.0483  
Epoch 224/1000  
33/33  [=====] - 0s 10ms/step - loss: 0.0037 -  
mae: 0.0521 - val_loss: 0.0039 - val_mae: 0.0483  
Epoch 225/1000  
33/33  [=====] - 0s 10ms/step - loss: 0.0037 -  
mae: 0.0521 - val_loss: 0.0039 - val_mae: 0.0483  
Epoch 226/1000  
33/33  [=====] - 0s 11ms/step - loss: 0.0037 -  
mae: 0.0521 - val_loss: 0.0039 - val_mae: 0.0483  
Epoch 227/1000  
33/33  [=====] - 0s 10ms/step - loss: 0.0037 -  
mae: 0.0521 - val_loss: 0.0039 - val_mae: 0.0483  
Epoch 228/1000  
33/33  [=====] - 0s 12ms/step - loss: 0.0037 -  
mae: 0.0520 - val_loss: 0.0039 - val_mae: 0.0483  
Epoch 229/1000  
33/33  [=====] - 0s 10ms/step - loss: 0.0037 -  
mae: 0.0520 - val_loss: 0.0039 - val_mae: 0.0483  
Epoch 230/1000  
33/33  [=====] - 0s 10ms/step - loss: 0.0037 -  
mae: 0.0520 - val_loss: 0.0039 - val_mae: 0.0483  
Epoch 231/1000  
33/33  [=====] - 0s 10ms/step - loss: 0.0037 -  
mae: 0.0520 - val_loss: 0.0039 - val_mae: 0.0483  
Epoch 232/1000  
33/33  [=====] - 0s 10ms/step - loss: 0.0037 -  
mae: 0.0520 - val_loss: 0.0039 - val_mae: 0.0483  
Epoch 233/1000  
33/33  [=====] - 0s 11ms/step - loss: 0.0037 -  
mae: 0.0520 - val_loss: 0.0039 - val_mae: 0.0483
```

```
Epoch 234/1000
33/33 [=====] - 0s 10ms/step - loss: 0.0037 -
mae: 0.0520 - val_loss: 0.0039 - val_mae: 0.0483
Epoch 235/1000
33/33 [=====] - 0s 10ms/step - loss: 0.0037 -
mae: 0.0520 - val_loss: 0.0039 - val_mae: 0.0483
Epoch 236/1000
33/33 [=====] - 0s 10ms/step - loss: 0.0037 -
mae: 0.0520 - val_loss: 0.0039 - val_mae: 0.0483
Epoch 237/1000
33/33 [=====] - 0s 12ms/step - loss: 0.0037 -
mae: 0.0520 - val_loss: 0.0039 - val_mae: 0.0483
Epoch 238/1000
33/33 [=====] - 0s 10ms/step - loss: 0.0037 -
mae: 0.0520 - val_loss: 0.0039 - val_mae: 0.0483
Epoch 239/1000
33/33 [=====] - 0s 10ms/step - loss: 0.0037 -
mae: 0.0520 - val_loss: 0.0039 - val_mae: 0.0483
Epoch 240/1000
33/33 [=====] - 0s 10ms/step - loss: 0.0037 -
mae: 0.0520 - val_loss: 0.0039 - val_mae: 0.0483
Epoch 241/1000
33/33 [=====] - 0s 10ms/step - loss: 0.0037 -
mae: 0.0520 - val_loss: 0.0039 - val_mae: 0.0483
Epoch 242/1000
33/33 [=====] - 0s 10ms/step - loss: 0.0037 -
mae: 0.0520 - val_loss: 0.0039 - val_mae: 0.0483
Epoch 243/1000
33/33 [=====] - 0s 10ms/step - loss: 0.0037 -
mae: 0.0520 - val_loss: 0.0039 - val_mae: 0.0483
Epoch 244/1000
33/33 [=====] - 0s 10ms/step - loss: 0.0037 -
mae: 0.0520 - val_loss: 0.0039 - val_mae: 0.0483
Epoch 245/1000
33/33 [=====] - 0s 12ms/step - loss: 0.0037 -
mae: 0.0520 - val_loss: 0.0039 - val_mae: 0.0483
Epoch 246/1000
33/33 [=====] - 0s 10ms/step - loss: 0.0037 -
```

```
mae: 0.0520 - val_loss: 0.0039 - val_mae: 0.0483
Epoch 247/1000
33/33 [=====] - 0s 11ms/step - loss: 0.0037 -
mae: 0.0520 - val_loss: 0.0039 - val_mae: 0.0483
Epoch 248/1000
33/33 [=====] - 0s 10ms/step - loss: 0.0037 -
mae: 0.0520 - val_loss: 0.0039 - val_mae: 0.0483
Epoch 249/1000
33/33 [=====] - 0s 10ms/step - loss: 0.0037 -
mae: 0.0520 - val_loss: 0.0039 - val_mae: 0.0483
Epoch 250/1000
33/33 [=====] - 0s 10ms/step - loss: 0.0037 -
mae: 0.0520 - val_loss: 0.0039 - val_mae: 0.0483
Epoch 251/1000
33/33 [=====] - 0s 10ms/step - loss: 0.0037 -
mae: 0.0520 - val_loss: 0.0039 - val_mae: 0.0483
Epoch 252/1000
33/33 [=====] - 0s 12ms/step - loss: 0.0037 -
mae: 0.0520 - val_loss: 0.0039 - val_mae: 0.0483
Epoch 253/1000
33/33 [=====] - 0s 9ms/step - loss: 0.0037 -
mae: 0.0520 - val_loss: 0.0039 - val_mae: 0.0483
Epoch 254/1000
33/33 [=====] - 0s 11ms/step - loss: 0.0037 -
mae: 0.0520 - val_loss: 0.0039 - val_mae: 0.0483
Epoch 255/1000
33/33 [=====] - 0s 10ms/step - loss: 0.0037 -
mae: 0.0520 - val_loss: 0.0039 - val_mae: 0.0483
Epoch 256/1000
33/33 [=====] - 0s 10ms/step - loss: 0.0037 -
mae: 0.0520 - val_loss: 0.0039 - val_mae: 0.0483
Epoch 257/1000
33/33 [=====] - 0s 10ms/step - loss: 0.0037 -
mae: 0.0520 - val_loss: 0.0039 - val_mae: 0.0483
Epoch 258/1000
33/33 [=====] - 0s 10ms/step - loss: 0.0037 -
mae: 0.0520 - val_loss: 0.0039 - val_mae: 0.0483
Epoch 259/1000
```

```
33/33  [=====] - 0s 10ms/step - loss: 0.0037 -  
mae: 0.0520 - val_loss: 0.0039 - val_mae: 0.0483  
Epoch 260/1000  
33/33  [=====] - 0s 10ms/step - loss: 0.0037 -  
mae: 0.0520 - val_loss: 0.0039 - val_mae: 0.0483  
Epoch 261/1000  
33/33  [=====] - 0s 11ms/step - loss: 0.0037 -  
mae: 0.0520 - val_loss: 0.0039 - val_mae: 0.0483  
Epoch 262/1000  
33/33  [=====] - 0s 10ms/step - loss: 0.0037 -  
mae: 0.0520 - val_loss: 0.0039 - val_mae: 0.0483  
Epoch 263/1000  
33/33  [=====] - 0s 11ms/step - loss: 0.0037 -  
mae: 0.0520 - val_loss: 0.0039 - val_mae: 0.0483  
Epoch 264/1000  
33/33  [=====] - 0s 10ms/step - loss: 0.0037 -  
mae: 0.0520 - val_loss: 0.0039 - val_mae: 0.0483  
Epoch 265/1000  
33/33  [=====] - 0s 10ms/step - loss: 0.0037 -  
mae: 0.0520 - val_loss: 0.0039 - val_mae: 0.0483  
Epoch 266/1000  
33/33  [=====] - 0s 10ms/step - loss: 0.0037 -  
mae: 0.0520 - val_loss: 0.0039 - val_mae: 0.0483  
Epoch 267/1000  
33/33  [=====] - 0s 10ms/step - loss: 0.0037 -  
mae: 0.0520 - val_loss: 0.0039 - val_mae: 0.0483  
Epoch 268/1000  
33/33  [=====] - 0s 9ms/step - loss: 0.0037 -  
mae: 0.0520 - val_loss: 0.0039 - val_mae: 0.0483  
Epoch 269/1000  
33/33  [=====] - 0s 10ms/step - loss: 0.0037 -  
mae: 0.0520 - val_loss: 0.0039 - val_mae: 0.0483  
Epoch 270/1000  
33/33  [=====] - 0s 11ms/step - loss: 0.0037 -  
mae: 0.0520 - val_loss: 0.0039 - val_mae: 0.0483  
Epoch 271/1000  
33/33  [=====] - 0s 10ms/step - loss: 0.0037 -  
mae: 0.0520 - val_loss: 0.0039 - val_mae: 0.0483
```

```
Epoch 272/1000
33/33 [=====] - 0s 10ms/step - loss: 0.0037 -
mae: 0.0520 - val_loss: 0.0039 - val_mae: 0.0483
Epoch 273/1000
33/33 [=====] - 0s 10ms/step - loss: 0.0037 -
mae: 0.0520 - val_loss: 0.0039 - val_mae: 0.0483
Epoch 274/1000
33/33 [=====] - 0s 10ms/step - loss: 0.0037 -
mae: 0.0520 - val_loss: 0.0039 - val_mae: 0.0483
Epoch 275/1000
33/33 [=====] - 0s 11ms/step - loss: 0.0037 -
mae: 0.0520 - val_loss: 0.0039 - val_mae: 0.0483
Epoch 276/1000
33/33 [=====] - 0s 14ms/step - loss: 0.0037 -
mae: 0.0520 - val_loss: 0.0039 - val_mae: 0.0483
Epoch 277/1000
33/33 [=====] - 0s 11ms/step - loss: 0.0037 -
mae: 0.0520 - val_loss: 0.0039 - val_mae: 0.0483
Epoch 278/1000
33/33 [=====] - 0s 11ms/step - loss: 0.0037 -
mae: 0.0519 - val_loss: 0.0039 - val_mae: 0.0483
Epoch 279/1000
33/33 [=====] - 0s 11ms/step - loss: 0.0037 -
mae: 0.0519 - val_loss: 0.0039 - val_mae: 0.0483
Epoch 280/1000
33/33 [=====] - 0s 11ms/step - loss: 0.0037 -
mae: 0.0519 - val_loss: 0.0039 - val_mae: 0.0483
Epoch 281/1000
33/33 [=====] - 0s 11ms/step - loss: 0.0037 -
mae: 0.0519 - val_loss: 0.0039 - val_mae: 0.0483
Epoch 282/1000
33/33 [=====] - 0s 11ms/step - loss: 0.0037 -
mae: 0.0519 - val_loss: 0.0039 - val_mae: 0.0483
Epoch 283/1000
33/33 [=====] - 0s 10ms/step - loss: 0.0037 -
mae: 0.0519 - val_loss: 0.0039 - val_mae: 0.0483
Epoch 284/1000
33/33 [=====] - 0s 10ms/step - loss: 0.0037 -
```

```
mae: 0.0519 - val_loss: 0.0039 - val_mae: 0.0483
Epoch 285/1000
33/33 [=====] - 0s 10ms/step - loss: 0.0037 -
mae: 0.0519 - val_loss: 0.0039 - val_mae: 0.0483
Epoch 286/1000
33/33 [=====] - 0s 10ms/step - loss: 0.0037 -
mae: 0.0519 - val_loss: 0.0039 - val_mae: 0.0483
Epoch 287/1000
33/33 [=====] - 0s 9ms/step - loss: 0.0037 -
mae: 0.0519 - val_loss: 0.0039 - val_mae: 0.0483
Epoch 288/1000
33/33 [=====] - 0s 10ms/step - loss: 0.0037 -
mae: 0.0519 - val_loss: 0.0039 - val_mae: 0.0483
Epoch 289/1000
33/33 [=====] - 0s 11ms/step - loss: 0.0037 -
mae: 0.0519 - val_loss: 0.0039 - val_mae: 0.0483
Epoch 290/1000
33/33 [=====] - 0s 10ms/step - loss: 0.0037 -
mae: 0.0519 - val_loss: 0.0039 - val_mae: 0.0483
Epoch 291/1000
33/33 [=====] - 0s 10ms/step - loss: 0.0037 -
mae: 0.0519 - val_loss: 0.0039 - val_mae: 0.0483
Epoch 292/1000
33/33 [=====] - 0s 10ms/step - loss: 0.0037 -
mae: 0.0519 - val_loss: 0.0039 - val_mae: 0.0483
Epoch 293/1000
33/33 [=====] - 0s 10ms/step - loss: 0.0037 -
mae: 0.0519 - val_loss: 0.0039 - val_mae: 0.0483
Epoch 294/1000
33/33 [=====] - 0s 10ms/step - loss: 0.0037 -
mae: 0.0519 - val_loss: 0.0039 - val_mae: 0.0483
Epoch 295/1000
33/33 [=====] - 0s 10ms/step - loss: 0.0037 -
mae: 0.0519 - val_loss: 0.0039 - val_mae: 0.0483
Epoch 296/1000
33/33 [=====] - 0s 11ms/step - loss: 0.0037 -
mae: 0.0519 - val_loss: 0.0039 - val_mae: 0.0483
Epoch 297/1000
```

```
33/33  [=====] - 0s 10ms/step - loss: 0.0037 -  
mae: 0.0519 - val_loss: 0.0039 - val_mae: 0.0483  
Epoch 298/1000  
33/33  [=====] - 0s 10ms/step - loss: 0.0037 -  
mae: 0.0519 - val_loss: 0.0039 - val_mae: 0.0483  
Epoch 299/1000  
33/33  [=====] - 0s 11ms/step - loss: 0.0037 -  
mae: 0.0519 - val_loss: 0.0039 - val_mae: 0.0483  
Epoch 300/1000  
33/33  [=====] - 0s 10ms/step - loss: 0.0037 -  
mae: 0.0519 - val_loss: 0.0039 - val_mae: 0.0483  
Epoch 301/1000  
33/33  [=====] - 0s 10ms/step - loss: 0.0037 -  
mae: 0.0519 - val_loss: 0.0039 - val_mae: 0.0483  
Epoch 302/1000  
33/33  [=====] - 0s 10ms/step - loss: 0.0037 -  
mae: 0.0519 - val_loss: 0.0039 - val_mae: 0.0483  
Epoch 303/1000  
33/33  [=====] - 0s 10ms/step - loss: 0.0037 -  
mae: 0.0519 - val_loss: 0.0039 - val_mae: 0.0483  
Epoch 304/1000  
33/33  [=====] - 0s 10ms/step - loss: 0.0037 -  
mae: 0.0519 - val_loss: 0.0039 - val_mae: 0.0483  
Epoch 305/1000  
33/33  [=====] - 0s 10ms/step - loss: 0.0037 -  
mae: 0.0519 - val_loss: 0.0039 - val_mae: 0.0483  
Epoch 306/1000  
33/33  [=====] - 0s 10ms/step - loss: 0.0037 -  
mae: 0.0519 - val_loss: 0.0039 - val_mae: 0.0483  
Epoch 307/1000  
33/33  [=====] - 0s 11ms/step - loss: 0.0037 -  
mae: 0.0519 - val_loss: 0.0039 - val_mae: 0.0483  
Epoch 308/1000  
33/33  [=====] - 0s 12ms/step - loss: 0.0037 -  
mae: 0.0519 - val_loss: 0.0039 - val_mae: 0.0483  
Epoch 309/1000  
33/33  [=====] - 0s 11ms/step - loss: 0.0037 -  
mae: 0.0519 - val_loss: 0.0039 - val_mae: 0.0483
```

```
Epoch 310/1000
33/33 [=====] - 0s 11ms/step - loss: 0.0037 -
mae: 0.0519 - val_loss: 0.0039 - val_mae: 0.0483
Epoch 311/1000
33/33 [=====] - 0s 11ms/step - loss: 0.0037 -
mae: 0.0519 - val_loss: 0.0039 - val_mae: 0.0483
Epoch 312/1000
33/33 [=====] - 0s 11ms/step - loss: 0.0037 -
mae: 0.0519 - val_loss: 0.0039 - val_mae: 0.0483
Epoch 313/1000
33/33 [=====] - 0s 11ms/step - loss: 0.0037 -
mae: 0.0519 - val_loss: 0.0039 - val_mae: 0.0483
Epoch 314/1000
33/33 [=====] - 0s 11ms/step - loss: 0.0037 -
mae: 0.0519 - val_loss: 0.0039 - val_mae: 0.0483
Epoch 315/1000
33/33 [=====] - 0s 11ms/step - loss: 0.0037 -
mae: 0.0519 - val_loss: 0.0039 - val_mae: 0.0483
Epoch 316/1000
33/33 [=====] - 0s 11ms/step - loss: 0.0037 -
mae: 0.0519 - val_loss: 0.0039 - val_mae: 0.0483
Epoch 317/1000
33/33 [=====] - 0s 13ms/step - loss: 0.0037 -
mae: 0.0519 - val_loss: 0.0039 - val_mae: 0.0483
Epoch 318/1000
33/33 [=====] - 0s 10ms/step - loss: 0.0037 -
mae: 0.0519 - val_loss: 0.0039 - val_mae: 0.0483
Epoch 319/1000
33/33 [=====] - 0s 10ms/step - loss: 0.0037 -
mae: 0.0519 - val_loss: 0.0039 - val_mae: 0.0483
Epoch 320/1000
33/33 [=====] - 0s 10ms/step - loss: 0.0037 -
mae: 0.0519 - val_loss: 0.0039 - val_mae: 0.0483
Epoch 321/1000
33/33 [=====] - 0s 10ms/step - loss: 0.0037 -
mae: 0.0519 - val_loss: 0.0039 - val_mae: 0.0483
Epoch 322/1000
33/33 [=====] - 0s 10ms/step - loss: 0.0037 -
```

```
mae: 0.0519 - val_loss: 0.0039 - val_mae: 0.0483
Epoch 323/1000
33/33 [=====] - 0s 10ms/step - loss: 0.0037 -
mae: 0.0519 - val_loss: 0.0039 - val_mae: 0.0483
Epoch 324/1000
33/33 [=====] - 0s 10ms/step - loss: 0.0037 -
mae: 0.0519 - val_loss: 0.0039 - val_mae: 0.0483
Epoch 325/1000
33/33 [=====] - 0s 10ms/step - loss: 0.0037 -
mae: 0.0519 - val_loss: 0.0039 - val_mae: 0.0483
Epoch 326/1000
33/33 [=====] - 0s 12ms/step - loss: 0.0037 -
mae: 0.0519 - val_loss: 0.0039 - val_mae: 0.0483
Epoch 327/1000
33/33 [=====] - 0s 10ms/step - loss: 0.0037 -
mae: 0.0519 - val_loss: 0.0039 - val_mae: 0.0483
Epoch 328/1000
33/33 [=====] - 0s 10ms/step - loss: 0.0037 -
mae: 0.0519 - val_loss: 0.0039 - val_mae: 0.0483
Epoch 329/1000
33/33 [=====] - 0s 10ms/step - loss: 0.0037 -
mae: 0.0518 - val_loss: 0.0039 - val_mae: 0.0483
Epoch 330/1000
33/33 [=====] - 0s 10ms/step - loss: 0.0037 -
mae: 0.0518 - val_loss: 0.0039 - val_mae: 0.0483
Epoch 331/1000
33/33 [=====] - 0s 10ms/step - loss: 0.0037 -
mae: 0.0518 - val_loss: 0.0039 - val_mae: 0.0483
Epoch 332/1000
33/33 [=====] - 0s 10ms/step - loss: 0.0037 -
mae: 0.0518 - val_loss: 0.0039 - val_mae: 0.0483
Epoch 333/1000
33/33 [=====] - 0s 10ms/step - loss: 0.0037 -
mae: 0.0518 - val_loss: 0.0039 - val_mae: 0.0483
Epoch 334/1000
33/33 [=====] - 0s 10ms/step - loss: 0.0037 -
mae: 0.0518 - val_loss: 0.0039 - val_mae: 0.0483
Epoch 335/1000
```

```
33/33  [=====] - 0s 12ms/step - loss: 0.0037 -  
mae: 0.0518 - val_loss: 0.0039 - val_mae: 0.0483  
Epoch 336/1000  
33/33  [=====] - 0s 10ms/step - loss: 0.0037 -  
mae: 0.0518 - val_loss: 0.0039 - val_mae: 0.0483  
Epoch 337/1000  
33/33  [=====] - 0s 10ms/step - loss: 0.0037 -  
mae: 0.0518 - val_loss: 0.0039 - val_mae: 0.0483  
Epoch 338/1000  
33/33  [=====] - 0s 10ms/step - loss: 0.0037 -  
mae: 0.0518 - val_loss: 0.0039 - val_mae: 0.0483  
Epoch 339/1000  
33/33  [=====] - 0s 10ms/step - loss: 0.0037 -  
mae: 0.0518 - val_loss: 0.0039 - val_mae: 0.0483  
Epoch 340/1000  
33/33  [=====] - 0s 10ms/step - loss: 0.0037 -  
mae: 0.0518 - val_loss: 0.0039 - val_mae: 0.0483  
Epoch 341/1000  
33/33  [=====] - 0s 10ms/step - loss: 0.0037 -  
mae: 0.0518 - val_loss: 0.0039 - val_mae: 0.0483  
Epoch 342/1000  
33/33  [=====] - 0s 10ms/step - loss: 0.0037 -  
mae: 0.0518 - val_loss: 0.0039 - val_mae: 0.0483  
Epoch 343/1000  
33/33  [=====] - 0s 10ms/step - loss: 0.0037 -  
mae: 0.0518 - val_loss: 0.0039 - val_mae: 0.0483  
Epoch 344/1000  
33/33  [=====] - 0s 12ms/step - loss: 0.0037 -  
mae: 0.0518 - val_loss: 0.0039 - val_mae: 0.0483  
Epoch 345/1000  
33/33  [=====] - 0s 10ms/step - loss: 0.0037 -  
mae: 0.0518 - val_loss: 0.0039 - val_mae: 0.0483  
Epoch 346/1000  
33/33  [=====] - 0s 10ms/step - loss: 0.0037 -  
mae: 0.0518 - val_loss: 0.0039 - val_mae: 0.0483  
Epoch 347/1000  
33/33  [=====] - 0s 10ms/step - loss: 0.0037 -  
mae: 0.0518 - val_loss: 0.0039 - val_mae: 0.0483
```

```
Epoch 348/1000
33/33 [=====] - 0s 10ms/step - loss: 0.0037 -
mae: 0.0518 - val_loss: 0.0039 - val_mae: 0.0483
Epoch 349/1000
33/33 [=====] - 0s 10ms/step - loss: 0.0037 -
mae: 0.0518 - val_loss: 0.0039 - val_mae: 0.0483
Epoch 350/1000
33/33 [=====] - 0s 10ms/step - loss: 0.0037 -
mae: 0.0518 - val_loss: 0.0039 - val_mae: 0.0483
Epoch 351/1000
33/33 [=====] - 0s 10ms/step - loss: 0.0037 -
mae: 0.0518 - val_loss: 0.0039 - val_mae: 0.0483
Epoch 352/1000
33/33 [=====] - 0s 10ms/step - loss: 0.0037 -
mae: 0.0518 - val_loss: 0.0039 - val_mae: 0.0483
Epoch 353/1000
33/33 [=====] - 0s 12ms/step - loss: 0.0037 -
mae: 0.0518 - val_loss: 0.0039 - val_mae: 0.0483
Epoch 354/1000
33/33 [=====] - 0s 10ms/step - loss: 0.0037 -
mae: 0.0518 - val_loss: 0.0039 - val_mae: 0.0483
Epoch 355/1000
33/33 [=====] - 0s 10ms/step - loss: 0.0037 -
mae: 0.0518 - val_loss: 0.0039 - val_mae: 0.0483
Epoch 356/1000
33/33 [=====] - 0s 10ms/step - loss: 0.0037 -
mae: 0.0518 - val_loss: 0.0039 - val_mae: 0.0483
Epoch 357/1000
33/33 [=====] - 0s 10ms/step - loss: 0.0037 -
mae: 0.0518 - val_loss: 0.0039 - val_mae: 0.0483
Epoch 358/1000
33/33 [=====] - 0s 11ms/step - loss: 0.0037 -
mae: 0.0518 - val_loss: 0.0039 - val_mae: 0.0483
Epoch 359/1000
33/33 [=====] - 0s 10ms/step - loss: 0.0037 -
mae: 0.0518 - val_loss: 0.0039 - val_mae: 0.0483
Epoch 360/1000
33/33 [=====] - 0s 10ms/step - loss: 0.0037 -
```

```
mae: 0.0518 - val_loss: 0.0039 - val_mae: 0.0483
Epoch 361/1000
33/33 [=====] - 0s 12ms/step - loss: 0.0037 -
mae: 0.0518 - val_loss: 0.0039 - val_mae: 0.0483
Epoch 362/1000
33/33 [=====] - 0s 10ms/step - loss: 0.0037 -
mae: 0.0518 - val_loss: 0.0039 - val_mae: 0.0483
Epoch 363/1000
33/33 [=====] - 0s 10ms/step - loss: 0.0037 -
mae: 0.0518 - val_loss: 0.0039 - val_mae: 0.0483
Epoch 364/1000
33/33 [=====] - 0s 10ms/step - loss: 0.0037 -
mae: 0.0518 - val_loss: 0.0039 - val_mae: 0.0483
Epoch 365/1000
33/33 [=====] - 0s 10ms/step - loss: 0.0037 -
mae: 0.0518 - val_loss: 0.0039 - val_mae: 0.0483
Epoch 366/1000
33/33 [=====] - 0s 10ms/step - loss: 0.0037 -
mae: 0.0518 - val_loss: 0.0039 - val_mae: 0.0483
Epoch 367/1000
33/33 [=====] - 0s 10ms/step - loss: 0.0037 -
mae: 0.0518 - val_loss: 0.0039 - val_mae: 0.0483
Epoch 368/1000
33/33 [=====] - 0s 9ms/step - loss: 0.0037 -
mae: 0.0518 - val_loss: 0.0039 - val_mae: 0.0483
Epoch 369/1000
33/33 [=====] - 0s 10ms/step - loss: 0.0037 -
mae: 0.0518 - val_loss: 0.0039 - val_mae: 0.0483
Epoch 370/1000
33/33 [=====] - 0s 10ms/step - loss: 0.0037 -
mae: 0.0518 - val_loss: 0.0039 - val_mae: 0.0483
Epoch 371/1000
33/33 [=====] - 0s 9ms/step - loss: 0.0037 -
mae: 0.0518 - val_loss: 0.0039 - val_mae: 0.0483
Epoch 372/1000
33/33 [=====] - 0s 10ms/step - loss: 0.0037 -
mae: 0.0518 - val_loss: 0.0039 - val_mae: 0.0483
Epoch 373/1000
```

```
33/33  [=====] - 0s 10ms/step - loss: 0.0037 -  
mae: 0.0518 - val_loss: 0.0039 - val_mae: 0.0483  
Epoch 374/1000  
33/33  [=====] - 0s 10ms/step - loss: 0.0037 -  
mae: 0.0518 - val_loss: 0.0039 - val_mae: 0.0483  
Epoch 375/1000  
33/33  [=====] - 0s 10ms/step - loss: 0.0037 -  
mae: 0.0518 - val_loss: 0.0039 - val_mae: 0.0483  
Epoch 376/1000  
33/33  [=====] - 0s 10ms/step - loss: 0.0037 -  
mae: 0.0518 - val_loss: 0.0039 - val_mae: 0.0483  
Epoch 377/1000  
33/33  [=====] - 0s 10ms/step - loss: 0.0037 -  
mae: 0.0518 - val_loss: 0.0039 - val_mae: 0.0483  
Epoch 378/1000  
33/33  [=====] - 0s 11ms/step - loss: 0.0037 -  
mae: 0.0518 - val_loss: 0.0039 - val_mae: 0.0483  
Epoch 379/1000  
33/33  [=====] - 0s 10ms/step - loss: 0.0037 -  
mae: 0.0518 - val_loss: 0.0039 - val_mae: 0.0483  
Epoch 380/1000  
33/33  [=====] - 0s 10ms/step - loss: 0.0037 -  
mae: 0.0518 - val_loss: 0.0039 - val_mae: 0.0483  
Epoch 381/1000  
33/33  [=====] - 0s 10ms/step - loss: 0.0037 -  
mae: 0.0518 - val_loss: 0.0039 - val_mae: 0.0483  
Epoch 382/1000  
33/33  [=====] - 0s 10ms/step - loss: 0.0037 -  
mae: 0.0518 - val_loss: 0.0039 - val_mae: 0.0483  
Epoch 383/1000  
33/33  [=====] - 0s 10ms/step - loss: 0.0037 -  
mae: 0.0517 - val_loss: 0.0039 - val_mae: 0.0483  
Epoch 384/1000  
33/33  [=====] - 0s 10ms/step - loss: 0.0037 -  
mae: 0.0517 - val_loss: 0.0039 - val_mae: 0.0483  
Epoch 385/1000  
33/33  [=====] - 0s 10ms/step - loss: 0.0037 -  
mae: 0.0517 - val_loss: 0.0039 - val_mae: 0.0483
```

```
Epoch 386/1000
33/33 [=====] - 0s 11ms/step - loss: 0.0037 -
mae: 0.0517 - val_loss: 0.0039 - val_mae: 0.0483
Epoch 387/1000
33/33 [=====] - 0s 10ms/step - loss: 0.0037 -
mae: 0.0517 - val_loss: 0.0039 - val_mae: 0.0483
Epoch 388/1000
33/33 [=====] - 0s 10ms/step - loss: 0.0037 -
mae: 0.0517 - val_loss: 0.0039 - val_mae: 0.0483
Epoch 389/1000
33/33 [=====] - 0s 10ms/step - loss: 0.0037 -
mae: 0.0517 - val_loss: 0.0039 - val_mae: 0.0483
Epoch 390/1000
33/33 [=====] - 0s 13ms/step - loss: 0.0037 -
mae: 0.0517 - val_loss: 0.0039 - val_mae: 0.0483
Epoch 391/1000
33/33 [=====] - 0s 10ms/step - loss: 0.0037 -
mae: 0.0517 - val_loss: 0.0039 - val_mae: 0.0483
Epoch 392/1000
33/33 [=====] - 0s 10ms/step - loss: 0.0037 -
mae: 0.0517 - val_loss: 0.0039 - val_mae: 0.0483
Epoch 393/1000
33/33 [=====] - 0s 10ms/step - loss: 0.0037 -
mae: 0.0517 - val_loss: 0.0039 - val_mae: 0.0483
Epoch 394/1000
33/33 [=====] - 0s 12ms/step - loss: 0.0037 -
mae: 0.0517 - val_loss: 0.0039 - val_mae: 0.0483
Epoch 395/1000
33/33 [=====] - 0s 10ms/step - loss: 0.0037 -
mae: 0.0517 - val_loss: 0.0039 - val_mae: 0.0483
Epoch 396/1000
33/33 [=====] - 0s 10ms/step - loss: 0.0037 -
mae: 0.0517 - val_loss: 0.0039 - val_mae: 0.0483
Epoch 397/1000
33/33 [=====] - 0s 10ms/step - loss: 0.0037 -
mae: 0.0517 - val_loss: 0.0039 - val_mae: 0.0483
Epoch 398/1000
33/33 [=====] - 0s 10ms/step - loss: 0.0037 -
```

```
mae: 0.0517 - val_loss: 0.0039 - val_mae: 0.0483
Epoch 399/1000
33/33 [=====] - 0s 10ms/step - loss: 0.0037 -
mae: 0.0517 - val_loss: 0.0039 - val_mae: 0.0483
Epoch 400/1000
33/33 [=====] - 0s 10ms/step - loss: 0.0037 -
mae: 0.0517 - val_loss: 0.0039 - val_mae: 0.0483
Epoch 401/1000
33/33 [=====] - 0s 10ms/step - loss: 0.0037 -
mae: 0.0517 - val_loss: 0.0039 - val_mae: 0.0483
Epoch 402/1000
33/33 [=====] - 0s 12ms/step - loss: 0.0037 -
mae: 0.0517 - val_loss: 0.0039 - val_mae: 0.0483
Epoch 403/1000
33/33 [=====] - 0s 10ms/step - loss: 0.0037 -
mae: 0.0517 - val_loss: 0.0039 - val_mae: 0.0483
Epoch 404/1000
33/33 [=====] - 0s 10ms/step - loss: 0.0037 -
mae: 0.0517 - val_loss: 0.0039 - val_mae: 0.0483
Epoch 405/1000
33/33 [=====] - 0s 10ms/step - loss: 0.0037 -
mae: 0.0517 - val_loss: 0.0039 - val_mae: 0.0483
Epoch 406/1000
33/33 [=====] - 0s 9ms/step - loss: 0.0037 -
mae: 0.0517 - val_loss: 0.0039 - val_mae: 0.0483
Epoch 407/1000
33/33 [=====] - 0s 10ms/step - loss: 0.0037 -
mae: 0.0517 - val_loss: 0.0039 - val_mae: 0.0483
Epoch 408/1000
33/33 [=====] - 0s 9ms/step - loss: 0.0037 -
mae: 0.0517 - val_loss: 0.0039 - val_mae: 0.0483
Epoch 409/1000
33/33 [=====] - 0s 9ms/step - loss: 0.0037 -
mae: 0.0517 - val_loss: 0.0039 - val_mae: 0.0483
Epoch 410/1000
33/33 [=====] - 0s 11ms/step - loss: 0.0037 -
mae: 0.0517 - val_loss: 0.0039 - val_mae: 0.0483
Epoch 411/1000
```

```
33/33  [=====] - 0s 9ms/step - loss: 0.0037 -  
mae: 0.0517 - val_loss: 0.0039 - val_mae: 0.0483  
Epoch 412/1000  
33/33  [=====] - 0s 10ms/step - loss: 0.0037 -  
mae: 0.0517 - val_loss: 0.0039 - val_mae: 0.0483  
Epoch 413/1000  
33/33  [=====] - 0s 9ms/step - loss: 0.0037 -  
mae: 0.0517 - val_loss: 0.0039 - val_mae: 0.0483  
Epoch 414/1000  
33/33  [=====] - 0s 9ms/step - loss: 0.0037 -  
mae: 0.0517 - val_loss: 0.0039 - val_mae: 0.0483  
Epoch 415/1000  
33/33  [=====] - 0s 9ms/step - loss: 0.0037 -  
mae: 0.0517 - val_loss: 0.0039 - val_mae: 0.0483  
Epoch 416/1000  
33/33  [=====] - 0s 10ms/step - loss: 0.0037 -  
mae: 0.0517 - val_loss: 0.0039 - val_mae: 0.0483  
Epoch 417/1000  
33/33  [=====] - 0s 9ms/step - loss: 0.0037 -  
mae: 0.0517 - val_loss: 0.0039 - val_mae: 0.0483  
Epoch 418/1000  
33/33  [=====] - 0s 10ms/step - loss: 0.0037 -  
mae: 0.0517 - val_loss: 0.0039 - val_mae: 0.0483  
Epoch 419/1000  
33/33  [=====] - 0s 10ms/step - loss: 0.0037 -  
mae: 0.0517 - val_loss: 0.0039 - val_mae: 0.0483  
Epoch 420/1000  
33/33  [=====] - 0s 10ms/step - loss: 0.0037 -  
mae: 0.0517 - val_loss: 0.0039 - val_mae: 0.0483  
Epoch 421/1000  
33/33  [=====] - 0s 10ms/step - loss: 0.0037 -  
mae: 0.0517 - val_loss: 0.0039 - val_mae: 0.0483  
Epoch 422/1000  
33/33  [=====] - 0s 10ms/step - loss: 0.0037 -  
mae: 0.0517 - val_loss: 0.0039 - val_mae: 0.0483  
Epoch 423/1000  
33/33  [=====] - 0s 10ms/step - loss: 0.0037 -  
mae: 0.0517 - val_loss: 0.0039 - val_mae: 0.0482
```

```
Epoch 424/1000
33/33 [=====] - 0s 10ms/step - loss: 0.0037 -
mae: 0.0517 - val_loss: 0.0039 - val_mae: 0.0482
Epoch 425/1000
33/33 [=====] - 0s 12ms/step - loss: 0.0037 -
mae: 0.0517 - val_loss: 0.0039 - val_mae: 0.0482
Epoch 426/1000
33/33 [=====] - 0s 10ms/step - loss: 0.0037 -
mae: 0.0517 - val_loss: 0.0039 - val_mae: 0.0482
Epoch 427/1000
33/33 [=====] - 0s 10ms/step - loss: 0.0037 -
mae: 0.0517 - val_loss: 0.0039 - val_mae: 0.0482
Epoch 428/1000
33/33 [=====] - 0s 10ms/step - loss: 0.0037 -
mae: 0.0517 - val_loss: 0.0039 - val_mae: 0.0482
Epoch 429/1000
33/33 [=====] - 0s 10ms/step - loss: 0.0037 -
mae: 0.0517 - val_loss: 0.0039 - val_mae: 0.0482
Epoch 430/1000
33/33 [=====] - 0s 10ms/step - loss: 0.0036 -
mae: 0.0517 - val_loss: 0.0039 - val_mae: 0.0482
Epoch 431/1000
33/33 [=====] - 0s 10ms/step - loss: 0.0036 -
mae: 0.0517 - val_loss: 0.0039 - val_mae: 0.0482
Epoch 432/1000
33/33 [=====] - 0s 10ms/step - loss: 0.0036 -
mae: 0.0517 - val_loss: 0.0039 - val_mae: 0.0482
Epoch 433/1000
33/33 [=====] - 0s 11ms/step - loss: 0.0036 -
mae: 0.0517 - val_loss: 0.0039 - val_mae: 0.0482
Epoch 434/1000
33/33 [=====] - 0s 10ms/step - loss: 0.0036 -
mae: 0.0517 - val_loss: 0.0039 - val_mae: 0.0482
Epoch 435/1000
33/33 [=====] - 0s 9ms/step - loss: 0.0036 -
mae: 0.0517 - val_loss: 0.0039 - val_mae: 0.0482
Epoch 436/1000
33/33 [=====] - 0s 9ms/step - loss: 0.0036 -
```

```
mae: 0.0517 - val_loss: 0.0039 - val_mae: 0.0482
Epoch 437/1000
33/33 [=====] - 0s 10ms/step - loss: 0.0036 -
mae: 0.0517 - val_loss: 0.0039 - val_mae: 0.0482
Epoch 438/1000
33/33 [=====] - 0s 9ms/step - loss: 0.0036 -
mae: 0.0517 - val_loss: 0.0039 - val_mae: 0.0482
Epoch 439/1000
33/33 [=====] - 0s 9ms/step - loss: 0.0036 -
mae: 0.0517 - val_loss: 0.0039 - val_mae: 0.0482
Epoch 440/1000
33/33 [=====] - 0s 10ms/step - loss: 0.0036 -
mae: 0.0516 - val_loss: 0.0039 - val_mae: 0.0482
Epoch 441/1000
33/33 [=====] - 0s 10ms/step - loss: 0.0036 -
mae: 0.0516 - val_loss: 0.0039 - val_mae: 0.0482
Epoch 442/1000
33/33 [=====] - 0s 9ms/step - loss: 0.0036 -
mae: 0.0516 - val_loss: 0.0039 - val_mae: 0.0482
Epoch 443/1000
33/33 [=====] - 0s 9ms/step - loss: 0.0036 -
mae: 0.0516 - val_loss: 0.0039 - val_mae: 0.0482
Epoch 444/1000
33/33 [=====] - 0s 10ms/step - loss: 0.0036 -
mae: 0.0516 - val_loss: 0.0039 - val_mae: 0.0482
Epoch 445/1000
33/33 [=====] - 0s 10ms/step - loss: 0.0036 -
mae: 0.0516 - val_loss: 0.0039 - val_mae: 0.0482
Epoch 446/1000
33/33 [=====] - 0s 10ms/step - loss: 0.0036 -
mae: 0.0516 - val_loss: 0.0039 - val_mae: 0.0482
Epoch 447/1000
33/33 [=====] - 0s 10ms/step - loss: 0.0036 -
mae: 0.0516 - val_loss: 0.0039 - val_mae: 0.0482
Epoch 448/1000
33/33 [=====] - 0s 11ms/step - loss: 0.0036 -
mae: 0.0516 - val_loss: 0.0039 - val_mae: 0.0482
Epoch 449/1000
```

```
33/33  [=====] - 0s 10ms/step - loss: 0.0036 -  
mae: 0.0516 - val_loss: 0.0039 - val_mae: 0.0482  
Epoch 450/1000  
33/33  [=====] - 0s 10ms/step - loss: 0.0036 -  
mae: 0.0516 - val_loss: 0.0039 - val_mae: 0.0482  
Epoch 451/1000  
33/33  [=====] - 0s 10ms/step - loss: 0.0036 -  
mae: 0.0516 - val_loss: 0.0039 - val_mae: 0.0482  
Epoch 452/1000  
33/33  [=====] - 0s 10ms/step - loss: 0.0036 -  
mae: 0.0516 - val_loss: 0.0039 - val_mae: 0.0482  
Epoch 453/1000  
33/33  [=====] - 0s 10ms/step - loss: 0.0036 -  
mae: 0.0516 - val_loss: 0.0039 - val_mae: 0.0482  
Epoch 454/1000  
33/33  [=====] - 0s 10ms/step - loss: 0.0036 -  
mae: 0.0516 - val_loss: 0.0039 - val_mae: 0.0482  
Epoch 455/1000  
33/33  [=====] - 0s 11ms/step - loss: 0.0036 -  
mae: 0.0516 - val_loss: 0.0039 - val_mae: 0.0482  
Epoch 456/1000  
33/33  [=====] - 0s 10ms/step - loss: 0.0036 -  
mae: 0.0516 - val_loss: 0.0039 - val_mae: 0.0482  
Epoch 457/1000  
33/33  [=====] - 0s 10ms/step - loss: 0.0036 -  
mae: 0.0516 - val_loss: 0.0039 - val_mae: 0.0482  
Epoch 458/1000  
33/33  [=====] - 0s 10ms/step - loss: 0.0036 -  
mae: 0.0516 - val_loss: 0.0039 - val_mae: 0.0482  
Epoch 459/1000  
33/33  [=====] - 0s 9ms/step - loss: 0.0036 -  
mae: 0.0516 - val_loss: 0.0039 - val_mae: 0.0482  
Epoch 460/1000  
33/33  [=====] - 0s 10ms/step - loss: 0.0036 -  
mae: 0.0516 - val_loss: 0.0039 - val_mae: 0.0482  
Epoch 461/1000  
33/33  [=====] - 0s 10ms/step - loss: 0.0036 -  
mae: 0.0516 - val_loss: 0.0039 - val_mae: 0.0482
```

```
Epoch 462/1000
33/33  [=====] - 0s 9ms/step - loss: 0.0036 -
mae: 0.0516 - val_loss: 0.0039 - val_mae: 0.0482
Epoch 463/1000
33/33  [=====] - 0s 10ms/step - loss: 0.0036 -
mae: 0.0516 - val_loss: 0.0039 - val_mae: 0.0482
Epoch 464/1000
33/33  [=====] - 0s 9ms/step - loss: 0.0036 -
mae: 0.0516 - val_loss: 0.0039 - val_mae: 0.0482
Epoch 465/1000
33/33  [=====] - 0s 9ms/step - loss: 0.0036 -
mae: 0.0516 - val_loss: 0.0039 - val_mae: 0.0482
Epoch 466/1000
33/33  [=====] - 0s 9ms/step - loss: 0.0036 -
mae: 0.0516 - val_loss: 0.0039 - val_mae: 0.0482
Epoch 467/1000
33/33  [=====] - 0s 9ms/step - loss: 0.0036 -
mae: 0.0516 - val_loss: 0.0039 - val_mae: 0.0482
Epoch 468/1000
33/33  [=====] - 0s 10ms/step - loss: 0.0036 -
mae: 0.0516 - val_loss: 0.0039 - val_mae: 0.0482
Epoch 469/1000
33/33  [=====] - 0s 9ms/step - loss: 0.0036 -
mae: 0.0516 - val_loss: 0.0039 - val_mae: 0.0482
Epoch 470/1000
33/33  [=====] - 0s 12ms/step - loss: 0.0036 -
mae: 0.0516 - val_loss: 0.0039 - val_mae: 0.0482
Epoch 471/1000
33/33  [=====] - 0s 10ms/step - loss: 0.0036 -
mae: 0.0516 - val_loss: 0.0039 - val_mae: 0.0482
Epoch 472/1000
33/33  [=====] - 0s 10ms/step - loss: 0.0036 -
mae: 0.0516 - val_loss: 0.0039 - val_mae: 0.0482
Epoch 473/1000
33/33  [=====] - 0s 10ms/step - loss: 0.0036 -
mae: 0.0516 - val_loss: 0.0039 - val_mae: 0.0482
Epoch 474/1000
33/33  [=====] - 0s 10ms/step - loss: 0.0036 -
```

```
mae: 0.0516 - val_loss: 0.0039 - val_mae: 0.0482
Epoch 475/1000
33/33 [=====] - 0s 10ms/step - loss: 0.0036 -
mae: 0.0516 - val_loss: 0.0039 - val_mae: 0.0482
Epoch 476/1000
33/33 [=====] - 0s 10ms/step - loss: 0.0036 -
mae: 0.0516 - val_loss: 0.0039 - val_mae: 0.0482
Epoch 477/1000
33/33 [=====] - 0s 10ms/step - loss: 0.0036 -
mae: 0.0516 - val_loss: 0.0039 - val_mae: 0.0482
Epoch 478/1000
33/33 [=====] - 0s 10ms/step - loss: 0.0036 -
mae: 0.0516 - val_loss: 0.0039 - val_mae: 0.0482
Epoch 479/1000
33/33 [=====] - 0s 10ms/step - loss: 0.0036 -
mae: 0.0516 - val_loss: 0.0039 - val_mae: 0.0482
Epoch 480/1000
33/33 [=====] - 0s 10ms/step - loss: 0.0036 -
mae: 0.0516 - val_loss: 0.0039 - val_mae: 0.0482
Epoch 481/1000
33/33 [=====] - 0s 10ms/step - loss: 0.0036 -
mae: 0.0516 - val_loss: 0.0039 - val_mae: 0.0482
Epoch 482/1000
33/33 [=====] - 0s 10ms/step - loss: 0.0036 -
mae: 0.0516 - val_loss: 0.0039 - val_mae: 0.0482
Epoch 483/1000
33/33 [=====] - 0s 10ms/step - loss: 0.0036 -
mae: 0.0516 - val_loss: 0.0039 - val_mae: 0.0482
Epoch 484/1000
33/33 [=====] - 0s 12ms/step - loss: 0.0036 -
mae: 0.0516 - val_loss: 0.0039 - val_mae: 0.0482
Epoch 485/1000
33/33 [=====] - 0s 10ms/step - loss: 0.0036 -
mae: 0.0516 - val_loss: 0.0039 - val_mae: 0.0482
Epoch 486/1000
33/33 [=====] - 0s 10ms/step - loss: 0.0036 -
mae: 0.0516 - val_loss: 0.0039 - val_mae: 0.0482
Epoch 487/1000
```

```
33/33  [=====] - 0s 10ms/step - loss: 0.0036 -  
mae: 0.0516 - val_loss: 0.0039 - val_mae: 0.0482  
Epoch 488/1000  
33/33  [=====] - 0s 10ms/step - loss: 0.0036 -  
mae: 0.0516 - val_loss: 0.0039 - val_mae: 0.0482  
Epoch 489/1000  
33/33  [=====] - 0s 10ms/step - loss: 0.0036 -  
mae: 0.0516 - val_loss: 0.0039 - val_mae: 0.0482  
Epoch 490/1000  
33/33  [=====] - 0s 9ms/step - loss: 0.0036 -  
mae: 0.0516 - val_loss: 0.0039 - val_mae: 0.0482  
Epoch 491/1000  
33/33  [=====] - 0s 10ms/step - loss: 0.0036 -  
mae: 0.0516 - val_loss: 0.0039 - val_mae: 0.0482  
Epoch 492/1000  
33/33  [=====] - 0s 9ms/step - loss: 0.0036 -  
mae: 0.0516 - val_loss: 0.0039 - val_mae: 0.0482  
Epoch 493/1000  
33/33  [=====] - 0s 9ms/step - loss: 0.0036 -  
mae: 0.0516 - val_loss: 0.0039 - val_mae: 0.0482  
Epoch 494/1000  
33/33  [=====] - 0s 10ms/step - loss: 0.0036 -  
mae: 0.0516 - val_loss: 0.0039 - val_mae: 0.0482  
Epoch 495/1000  
33/33  [=====] - 0s 9ms/step - loss: 0.0036 -  
mae: 0.0516 - val_loss: 0.0039 - val_mae: 0.0482  
Epoch 496/1000  
33/33  [=====] - 0s 9ms/step - loss: 0.0036 -  
mae: 0.0516 - val_loss: 0.0039 - val_mae: 0.0482  
Epoch 497/1000  
33/33  [=====] - 0s 10ms/step - loss: 0.0036 -  
mae: 0.0516 - val_loss: 0.0039 - val_mae: 0.0482  
Epoch 498/1000  
33/33  [=====] - 0s 10ms/step - loss: 0.0036 -  
mae: 0.0516 - val_loss: 0.0039 - val_mae: 0.0482  
Epoch 499/1000  
33/33  [=====] - 0s 10ms/step - loss: 0.0036 -  
mae: 0.0516 - val_loss: 0.0039 - val_mae: 0.0482
```

```
Epoch 500/1000
33/33 [=====] - 0s 10ms/step - loss: 0.0036 -
mae: 0.0516 - val_loss: 0.0039 - val_mae: 0.0482
Epoch 501/1000
33/33 [=====] - 0s 10ms/step - loss: 0.0036 -
mae: 0.0515 - val_loss: 0.0039 - val_mae: 0.0482
Epoch 502/1000
33/33 [=====] - 0s 10ms/step - loss: 0.0036 -
mae: 0.0515 - val_loss: 0.0039 - val_mae: 0.0482
Epoch 503/1000
33/33 [=====] - 0s 10ms/step - loss: 0.0036 -
mae: 0.0515 - val_loss: 0.0039 - val_mae: 0.0482
Epoch 504/1000
33/33 [=====] - 0s 12ms/step - loss: 0.0036 -
mae: 0.0515 - val_loss: 0.0039 - val_mae: 0.0482
Epoch 505/1000
33/33 [=====] - 0s 10ms/step - loss: 0.0036 -
mae: 0.0515 - val_loss: 0.0039 - val_mae: 0.0482
Epoch 506/1000
33/33 [=====] - 0s 10ms/step - loss: 0.0036 -
mae: 0.0515 - val_loss: 0.0039 - val_mae: 0.0482
Epoch 507/1000
33/33 [=====] - 0s 10ms/step - loss: 0.0036 -
mae: 0.0515 - val_loss: 0.0039 - val_mae: 0.0482
Epoch 508/1000
33/33 [=====] - 0s 10ms/step - loss: 0.0036 -
mae: 0.0515 - val_loss: 0.0039 - val_mae: 0.0482
Epoch 509/1000
33/33 [=====] - 0s 10ms/step - loss: 0.0036 -
mae: 0.0515 - val_loss: 0.0039 - val_mae: 0.0482
Epoch 510/1000
33/33 [=====] - 0s 10ms/step - loss: 0.0036 -
mae: 0.0515 - val_loss: 0.0039 - val_mae: 0.0482
Epoch 511/1000
33/33 [=====] - 0s 12ms/step - loss: 0.0036 -
mae: 0.0515 - val_loss: 0.0039 - val_mae: 0.0482
Epoch 512/1000
33/33 [=====] - 0s 10ms/step - loss: 0.0036 -
```

```
mae: 0.0515 - val_loss: 0.0039 - val_mae: 0.0482
Epoch 513/1000
33/33 [=====] - 0s 10ms/step - loss: 0.0036 -
mae: 0.0515 - val_loss: 0.0039 - val_mae: 0.0482
Epoch 514/1000
33/33 [=====] - 0s 10ms/step - loss: 0.0036 -
mae: 0.0515 - val_loss: 0.0039 - val_mae: 0.0482
Epoch 515/1000
33/33 [=====] - 0s 10ms/step - loss: 0.0036 -
mae: 0.0515 - val_loss: 0.0039 - val_mae: 0.0482
Epoch 516/1000
33/33 [=====] - 0s 10ms/step - loss: 0.0036 -
mae: 0.0515 - val_loss: 0.0039 - val_mae: 0.0482
Epoch 517/1000
33/33 [=====] - 0s 11ms/step - loss: 0.0036 -
mae: 0.0515 - val_loss: 0.0039 - val_mae: 0.0482
Epoch 518/1000
33/33 [=====] - 0s 10ms/step - loss: 0.0036 -
mae: 0.0515 - val_loss: 0.0039 - val_mae: 0.0482
Epoch 519/1000
33/33 [=====] - 0s 10ms/step - loss: 0.0036 -
mae: 0.0515 - val_loss: 0.0039 - val_mae: 0.0482
Epoch 520/1000
33/33 [=====] - 0s 10ms/step - loss: 0.0036 -
mae: 0.0515 - val_loss: 0.0039 - val_mae: 0.0482
Epoch 521/1000
33/33 [=====] - 0s 11ms/step - loss: 0.0036 -
mae: 0.0515 - val_loss: 0.0039 - val_mae: 0.0482
Epoch 522/1000
33/33 [=====] - 0s 11ms/step - loss: 0.0036 -
mae: 0.0515 - val_loss: 0.0039 - val_mae: 0.0482
Epoch 523/1000
33/33 [=====] - 0s 12ms/step - loss: 0.0036 -
mae: 0.0515 - val_loss: 0.0039 - val_mae: 0.0482
Epoch 524/1000
33/33 [=====] - 0s 11ms/step - loss: 0.0036 -
mae: 0.0515 - val_loss: 0.0039 - val_mae: 0.0482
Epoch 525/1000
```

```
33/33  [=====] - 0s 11ms/step - loss: 0.0036 -  
mae: 0.0515 - val_loss: 0.0039 - val_mae: 0.0482  
Epoch 526/1000  
33/33  [=====] - 0s 11ms/step - loss: 0.0036 -  
mae: 0.0515 - val_loss: 0.0039 - val_mae: 0.0482  
Epoch 527/1000  
33/33  [=====] - 0s 11ms/step - loss: 0.0036 -  
mae: 0.0515 - val_loss: 0.0039 - val_mae: 0.0482  
Epoch 528/1000  
33/33  [=====] - 0s 11ms/step - loss: 0.0036 -  
mae: 0.0515 - val_loss: 0.0039 - val_mae: 0.0482  
Epoch 529/1000  
33/33  [=====] - 0s 13ms/step - loss: 0.0036 -  
mae: 0.0515 - val_loss: 0.0039 - val_mae: 0.0482  
Epoch 530/1000  
33/33  [=====] - 0s 11ms/step - loss: 0.0036 -  
mae: 0.0515 - val_loss: 0.0039 - val_mae: 0.0482  
Epoch 531/1000  
33/33  [=====] - 0s 11ms/step - loss: 0.0036 -  
mae: 0.0515 - val_loss: 0.0039 - val_mae: 0.0482  
Epoch 532/1000  
33/33  [=====] - 0s 10ms/step - loss: 0.0036 -  
mae: 0.0515 - val_loss: 0.0039 - val_mae: 0.0482  
Epoch 533/1000  
33/33  [=====] - 0s 11ms/step - loss: 0.0036 -  
mae: 0.0515 - val_loss: 0.0039 - val_mae: 0.0482  
Epoch 534/1000  
33/33  [=====] - 0s 11ms/step - loss: 0.0036 -  
mae: 0.0515 - val_loss: 0.0039 - val_mae: 0.0482  
Epoch 535/1000  
33/33  [=====] - 0s 12ms/step - loss: 0.0036 -  
mae: 0.0515 - val_loss: 0.0039 - val_mae: 0.0482  
Epoch 536/1000  
33/33  [=====] - 0s 11ms/step - loss: 0.0036 -  
mae: 0.0515 - val_loss: 0.0039 - val_mae: 0.0482  
Epoch 537/1000  
33/33  [=====] - 0s 11ms/step - loss: 0.0036 -  
mae: 0.0515 - val_loss: 0.0039 - val_mae: 0.0482
```

```
Epoch 538/1000
33/33  [=====] - 0s 11ms/step - loss: 0.0036 -
mae: 0.0515 - val_loss: 0.0039 - val_mae: 0.0482
Epoch 539/1000
33/33  [=====] - 0s 11ms/step - loss: 0.0036 -
mae: 0.0515 - val_loss: 0.0039 - val_mae: 0.0482
Epoch 540/1000
33/33  [=====] - 0s 12ms/step - loss: 0.0036 -
mae: 0.0515 - val_loss: 0.0039 - val_mae: 0.0482
Epoch 541/1000
33/33  [=====] - 0s 11ms/step - loss: 0.0036 -
mae: 0.0515 - val_loss: 0.0039 - val_mae: 0.0482
Epoch 542/1000
33/33  [=====] - 0s 11ms/step - loss: 0.0036 -
mae: 0.0515 - val_loss: 0.0039 - val_mae: 0.0482
Epoch 543/1000
33/33  [=====] - 0s 11ms/step - loss: 0.0036 -
mae: 0.0515 - val_loss: 0.0039 - val_mae: 0.0482
Epoch 544/1000
33/33  [=====] - 0s 10ms/step - loss: 0.0036 -
mae: 0.0515 - val_loss: 0.0039 - val_mae: 0.0482
Epoch 545/1000
33/33  [=====] - 0s 10ms/step - loss: 0.0036 -
mae: 0.0515 - val_loss: 0.0039 - val_mae: 0.0482
Epoch 546/1000
33/33  [=====] - 0s 12ms/step - loss: 0.0036 -
mae: 0.0515 - val_loss: 0.0039 - val_mae: 0.0482
Epoch 547/1000
33/33  [=====] - 0s 10ms/step - loss: 0.0036 -
mae: 0.0515 - val_loss: 0.0039 - val_mae: 0.0482
Epoch 548/1000
33/33  [=====] - 0s 10ms/step - loss: 0.0036 -
mae: 0.0515 - val_loss: 0.0039 - val_mae: 0.0482
Epoch 549/1000
33/33  [=====] - 0s 10ms/step - loss: 0.0036 -
mae: 0.0515 - val_loss: 0.0039 - val_mae: 0.0482
Epoch 550/1000
33/33  [=====] - 0s 10ms/step - loss: 0.0036 -
```

```
mae: 0.0515 - val_loss: 0.0039 - val_mae: 0.0482
Epoch 551/1000
33/33 [=====] - 0s 10ms/step - loss: 0.0036 -
mae: 0.0515 - val_loss: 0.0039 - val_mae: 0.0482
Epoch 552/1000
33/33 [=====] - 0s 9ms/step - loss: 0.0036 -
mae: 0.0515 - val_loss: 0.0039 - val_mae: 0.0482
Epoch 553/1000
33/33 [=====] - 0s 12ms/step - loss: 0.0036 -
mae: 0.0515 - val_loss: 0.0039 - val_mae: 0.0482
Epoch 554/1000
33/33 [=====] - 0s 9ms/step - loss: 0.0036 -
mae: 0.0515 - val_loss: 0.0039 - val_mae: 0.0482
Epoch 555/1000
33/33 [=====] - 0s 10ms/step - loss: 0.0036 -
mae: 0.0515 - val_loss: 0.0039 - val_mae: 0.0482
Epoch 556/1000
33/33 [=====] - 0s 10ms/step - loss: 0.0036 -
mae: 0.0515 - val_loss: 0.0039 - val_mae: 0.0482
Epoch 557/1000
33/33 [=====] - 0s 9ms/step - loss: 0.0036 -
mae: 0.0515 - val_loss: 0.0039 - val_mae: 0.0482
Epoch 558/1000
33/33 [=====] - 0s 10ms/step - loss: 0.0036 -
mae: 0.0515 - val_loss: 0.0039 - val_mae: 0.0482
Epoch 559/1000
33/33 [=====] - 0s 12ms/step - loss: 0.0036 -
mae: 0.0515 - val_loss: 0.0039 - val_mae: 0.0482
Epoch 560/1000
33/33 [=====] - 0s 9ms/step - loss: 0.0036 -
mae: 0.0515 - val_loss: 0.0039 - val_mae: 0.0482
Epoch 561/1000
33/33 [=====] - 0s 10ms/step - loss: 0.0036 -
mae: 0.0515 - val_loss: 0.0039 - val_mae: 0.0482
Epoch 562/1000
33/33 [=====] - 0s 10ms/step - loss: 0.0036 -
mae: 0.0515 - val_loss: 0.0039 - val_mae: 0.0482
Epoch 563/1000
```

```
33/33  [=====] - 0s 9ms/step - loss: 0.0036 -  
mae: 0.0515 - val_loss: 0.0039 - val_mae: 0.0482  
Epoch 564/1000  
33/33  [=====] - 0s 10ms/step - loss: 0.0036 -  
mae: 0.0515 - val_loss: 0.0039 - val_mae: 0.0482  
Epoch 565/1000  
33/33  [=====] - 0s 12ms/step - loss: 0.0036 -  
mae: 0.0514 - val_loss: 0.0039 - val_mae: 0.0482  
Epoch 566/1000  
33/33  [=====] - 0s 10ms/step - loss: 0.0036 -  
mae: 0.0514 - val_loss: 0.0039 - val_mae: 0.0482  
Epoch 567/1000  
33/33  [=====] - 0s 10ms/step - loss: 0.0036 -  
mae: 0.0514 - val_loss: 0.0039 - val_mae: 0.0482  
Epoch 568/1000  
33/33  [=====] - 0s 10ms/step - loss: 0.0036 -  
mae: 0.0514 - val_loss: 0.0039 - val_mae: 0.0482  
Epoch 569/1000  
33/33  [=====] - 0s 10ms/step - loss: 0.0036 -  
mae: 0.0514 - val_loss: 0.0039 - val_mae: 0.0482  
Epoch 570/1000  
33/33  [=====] - 0s 10ms/step - loss: 0.0036 -  
mae: 0.0514 - val_loss: 0.0039 - val_mae: 0.0482  
Epoch 571/1000  
33/33  [=====] - 0s 12ms/step - loss: 0.0036 -  
mae: 0.0514 - val_loss: 0.0039 - val_mae: 0.0482  
Epoch 572/1000  
33/33  [=====] - 0s 10ms/step - loss: 0.0036 -  
mae: 0.0514 - val_loss: 0.0039 - val_mae: 0.0482  
Epoch 573/1000  
33/33  [=====] - 0s 10ms/step - loss: 0.0036 -  
mae: 0.0514 - val_loss: 0.0039 - val_mae: 0.0482  
Epoch 574/1000  
33/33  [=====] - 0s 10ms/step - loss: 0.0036 -  
mae: 0.0514 - val_loss: 0.0039 - val_mae: 0.0482  
Epoch 575/1000  
33/33  [=====] - 0s 10ms/step - loss: 0.0036 -  
mae: 0.0514 - val_loss: 0.0039 - val_mae: 0.0482
```

```
Epoch 576/1000
33/33  [=====] - 0s 10ms/step - loss: 0.0036 -
mae: 0.0514 - val_loss: 0.0039 - val_mae: 0.0482
Epoch 577/1000
33/33  [=====] - 0s 11ms/step - loss: 0.0036 -
mae: 0.0514 - val_loss: 0.0039 - val_mae: 0.0482
Epoch 578/1000
33/33  [=====] - 0s 11ms/step - loss: 0.0036 -
mae: 0.0514 - val_loss: 0.0039 - val_mae: 0.0482
Epoch 579/1000
33/33  [=====] - 0s 10ms/step - loss: 0.0036 -
mae: 0.0514 - val_loss: 0.0039 - val_mae: 0.0482
Epoch 580/1000
33/33  [=====] - 0s 10ms/step - loss: 0.0036 -
mae: 0.0514 - val_loss: 0.0039 - val_mae: 0.0482
Epoch 581/1000
33/33  [=====] - 0s 13ms/step - loss: 0.0036 -
mae: 0.0514 - val_loss: 0.0039 - val_mae: 0.0482
Epoch 582/1000
33/33  [=====] - 0s 10ms/step - loss: 0.0036 -
mae: 0.0514 - val_loss: 0.0039 - val_mae: 0.0482
Epoch 583/1000
33/33  [=====] - 0s 11ms/step - loss: 0.0036 -
mae: 0.0514 - val_loss: 0.0039 - val_mae: 0.0482
Epoch 584/1000
33/33  [=====] - 0s 10ms/step - loss: 0.0036 -
mae: 0.0514 - val_loss: 0.0039 - val_mae: 0.0482
Epoch 585/1000
33/33  [=====] - 0s 10ms/step - loss: 0.0036 -
mae: 0.0514 - val_loss: 0.0039 - val_mae: 0.0482
Epoch 586/1000
33/33  [=====] - 0s 10ms/step - loss: 0.0036 -
mae: 0.0514 - val_loss: 0.0039 - val_mae: 0.0482
Epoch 587/1000
33/33  [=====] - 0s 12ms/step - loss: 0.0036 -
mae: 0.0514 - val_loss: 0.0039 - val_mae: 0.0482
Epoch 588/1000
33/33  [=====] - 0s 10ms/step - loss: 0.0036 -
```

```
mae: 0.0514 - val_loss: 0.0039 - val_mae: 0.0482
Epoch 589/1000
33/33 [=====] - 0s 11ms/step - loss: 0.0036 -
mae: 0.0514 - val_loss: 0.0039 - val_mae: 0.0482
Epoch 590/1000
33/33 [=====] - 0s 10ms/step - loss: 0.0036 -
mae: 0.0514 - val_loss: 0.0039 - val_mae: 0.0482
Epoch 591/1000
33/33 [=====] - 0s 10ms/step - loss: 0.0036 -
mae: 0.0514 - val_loss: 0.0039 - val_mae: 0.0482
Epoch 592/1000
33/33 [=====] - 0s 11ms/step - loss: 0.0036 -
mae: 0.0514 - val_loss: 0.0039 - val_mae: 0.0482
Epoch 593/1000
33/33 [=====] - 0s 10ms/step - loss: 0.0036 -
mae: 0.0514 - val_loss: 0.0039 - val_mae: 0.0482
Epoch 594/1000
33/33 [=====] - 0s 10ms/step - loss: 0.0036 -
mae: 0.0514 - val_loss: 0.0039 - val_mae: 0.0482
Epoch 595/1000
33/33 [=====] - 0s 10ms/step - loss: 0.0036 -
mae: 0.0514 - val_loss: 0.0039 - val_mae: 0.0482
Epoch 596/1000
33/33 [=====] - 0s 10ms/step - loss: 0.0036 -
mae: 0.0514 - val_loss: 0.0039 - val_mae: 0.0482
Epoch 597/1000
33/33 [=====] - 0s 10ms/step - loss: 0.0036 -
mae: 0.0514 - val_loss: 0.0039 - val_mae: 0.0482
Epoch 598/1000
33/33 [=====] - 0s 11ms/step - loss: 0.0036 -
mae: 0.0514 - val_loss: 0.0039 - val_mae: 0.0482
Epoch 599/1000
33/33 [=====] - 0s 10ms/step - loss: 0.0036 -
mae: 0.0514 - val_loss: 0.0039 - val_mae: 0.0482
Epoch 600/1000
33/33 [=====] - 0s 10ms/step - loss: 0.0036 -
mae: 0.0514 - val_loss: 0.0039 - val_mae: 0.0482
Epoch 601/1000
```

```
33/33  [=====] - 0s 10ms/step - loss: 0.0036 -  
mae: 0.0514 - val_loss: 0.0039 - val_mae: 0.0482  
Epoch 602/1000  
33/33  [=====] - 0s 9ms/step - loss: 0.0036 -  
mae: 0.0514 - val_loss: 0.0039 - val_mae: 0.0482  
Epoch 603/1000  
33/33  [=====] - 0s 10ms/step - loss: 0.0036 -  
mae: 0.0514 - val_loss: 0.0039 - val_mae: 0.0482  
Epoch 604/1000  
33/33  [=====] - 0s 12ms/step - loss: 0.0036 -  
mae: 0.0514 - val_loss: 0.0039 - val_mae: 0.0482  
Epoch 605/1000  
33/33  [=====] - 0s 10ms/step - loss: 0.0036 -  
mae: 0.0514 - val_loss: 0.0039 - val_mae: 0.0482  
Epoch 606/1000  
33/33  [=====] - 0s 10ms/step - loss: 0.0036 -  
mae: 0.0514 - val_loss: 0.0039 - val_mae: 0.0482  
Epoch 607/1000  
33/33  [=====] - 0s 10ms/step - loss: 0.0036 -  
mae: 0.0514 - val_loss: 0.0039 - val_mae: 0.0482  
Epoch 608/1000  
33/33  [=====] - 0s 10ms/step - loss: 0.0036 -  
mae: 0.0514 - val_loss: 0.0039 - val_mae: 0.0482  
Epoch 609/1000  
33/33  [=====] - 0s 10ms/step - loss: 0.0036 -  
mae: 0.0514 - val_loss: 0.0039 - val_mae: 0.0482  
Epoch 610/1000  
33/33  [=====] - 0s 12ms/step - loss: 0.0036 -  
mae: 0.0514 - val_loss: 0.0039 - val_mae: 0.0482  
Epoch 611/1000  
33/33  [=====] - 0s 10ms/step - loss: 0.0036 -  
mae: 0.0514 - val_loss: 0.0039 - val_mae: 0.0482  
Epoch 612/1000  
33/33  [=====] - 0s 10ms/step - loss: 0.0036 -  
mae: 0.0514 - val_loss: 0.0039 - val_mae: 0.0482  
Epoch 613/1000  
33/33  [=====] - 0s 10ms/step - loss: 0.0036 -  
mae: 0.0514 - val_loss: 0.0039 - val_mae: 0.0482
```

```
Epoch 614/1000
33/33  [=====] - 0s 10ms/step - loss: 0.0036 -
mae: 0.0514 - val_loss: 0.0039 - val_mae: 0.0482
Epoch 615/1000
33/33  [=====] - 0s 10ms/step - loss: 0.0036 -
mae: 0.0514 - val_loss: 0.0039 - val_mae: 0.0482
Epoch 616/1000
33/33  [=====] - 0s 12ms/step - loss: 0.0036 -
mae: 0.0514 - val_loss: 0.0039 - val_mae: 0.0482
Epoch 617/1000
33/33  [=====] - 0s 10ms/step - loss: 0.0036 -
mae: 0.0514 - val_loss: 0.0039 - val_mae: 0.0482
Epoch 618/1000
33/33  [=====] - 0s 10ms/step - loss: 0.0036 -
mae: 0.0514 - val_loss: 0.0039 - val_mae: 0.0482
Epoch 619/1000
33/33  [=====] - 0s 10ms/step - loss: 0.0036 -
mae: 0.0514 - val_loss: 0.0039 - val_mae: 0.0482
Epoch 620/1000
33/33  [=====] - 0s 10ms/step - loss: 0.0036 -
mae: 0.0514 - val_loss: 0.0039 - val_mae: 0.0482
Epoch 621/1000
33/33  [=====] - 0s 10ms/step - loss: 0.0036 -
mae: 0.0514 - val_loss: 0.0039 - val_mae: 0.0482
Epoch 622/1000
33/33  [=====] - 0s 10ms/step - loss: 0.0036 -
mae: 0.0514 - val_loss: 0.0039 - val_mae: 0.0482
Epoch 623/1000
33/33  [=====] - 0s 10ms/step - loss: 0.0036 -
mae: 0.0514 - val_loss: 0.0039 - val_mae: 0.0482
Epoch 624/1000
33/33  [=====] - 0s 10ms/step - loss: 0.0036 -
mae: 0.0514 - val_loss: 0.0039 - val_mae: 0.0482
Epoch 625/1000
33/33  [=====] - 0s 10ms/step - loss: 0.0036 -
mae: 0.0514 - val_loss: 0.0039 - val_mae: 0.0482
Epoch 626/1000
33/33  [=====] - 0s 10ms/step - loss: 0.0036 -
```

```
mae: 0.0514 - val_loss: 0.0039 - val_mae: 0.0482
Epoch 627/1000
33/33 [=====] - 0s 11ms/step - loss: 0.0036 -
mae: 0.0514 - val_loss: 0.0039 - val_mae: 0.0482
Epoch 628/1000
33/33 [=====] - 0s 10ms/step - loss: 0.0036 -
mae: 0.0514 - val_loss: 0.0039 - val_mae: 0.0482
Epoch 629/1000
33/33 [=====] - 0s 10ms/step - loss: 0.0036 -
mae: 0.0514 - val_loss: 0.0039 - val_mae: 0.0482
Epoch 630/1000
33/33 [=====] - 0s 10ms/step - loss: 0.0036 -
mae: 0.0514 - val_loss: 0.0039 - val_mae: 0.0482
Epoch 631/1000
33/33 [=====] - 0s 10ms/step - loss: 0.0036 -
mae: 0.0514 - val_loss: 0.0039 - val_mae: 0.0482
Epoch 632/1000
33/33 [=====] - 0s 10ms/step - loss: 0.0036 -
mae: 0.0514 - val_loss: 0.0039 - val_mae: 0.0482
Epoch 633/1000
33/33 [=====] - 0s 12ms/step - loss: 0.0036 -
mae: 0.0513 - val_loss: 0.0039 - val_mae: 0.0482
Epoch 634/1000
33/33 [=====] - 0s 10ms/step - loss: 0.0036 -
mae: 0.0513 - val_loss: 0.0039 - val_mae: 0.0482
Epoch 635/1000
33/33 [=====] - 0s 10ms/step - loss: 0.0036 -
mae: 0.0513 - val_loss: 0.0039 - val_mae: 0.0482
Epoch 636/1000
33/33 [=====] - 0s 10ms/step - loss: 0.0036 -
mae: 0.0513 - val_loss: 0.0039 - val_mae: 0.0482
Epoch 637/1000
33/33 [=====] - 0s 10ms/step - loss: 0.0036 -
mae: 0.0513 - val_loss: 0.0039 - val_mae: 0.0482
Epoch 638/1000
33/33 [=====] - 0s 10ms/step - loss: 0.0036 -
mae: 0.0513 - val_loss: 0.0039 - val_mae: 0.0482
Epoch 639/1000
```

```
33/33  [=====] - 0s 10ms/step - loss: 0.0036 -  
mae: 0.0513 - val_loss: 0.0039 - val_mae: 0.0482  
Epoch 640/1000  
33/33  [=====] - 0s 10ms/step - loss: 0.0036 -  
mae: 0.0513 - val_loss: 0.0039 - val_mae: 0.0482  
Epoch 641/1000  
33/33  [=====] - 0s 10ms/step - loss: 0.0036 -  
mae: 0.0513 - val_loss: 0.0039 - val_mae: 0.0482  
Epoch 642/1000  
33/33  [=====] - 0s 10ms/step - loss: 0.0036 -  
mae: 0.0513 - val_loss: 0.0039 - val_mae: 0.0482  
Epoch 643/1000  
33/33  [=====] - 0s 10ms/step - loss: 0.0036 -  
mae: 0.0513 - val_loss: 0.0039 - val_mae: 0.0482  
Epoch 644/1000  
33/33  [=====] - 0s 12ms/step - loss: 0.0036 -  
mae: 0.0513 - val_loss: 0.0039 - val_mae: 0.0482  
Epoch 645/1000  
33/33  [=====] - 0s 9ms/step - loss: 0.0036 -  
mae: 0.0513 - val_loss: 0.0039 - val_mae: 0.0482  
Epoch 646/1000  
33/33  [=====] - 0s 10ms/step - loss: 0.0036 -  
mae: 0.0513 - val_loss: 0.0039 - val_mae: 0.0482  
Epoch 647/1000  
33/33  [=====] - 0s 10ms/step - loss: 0.0036 -  
mae: 0.0513 - val_loss: 0.0039 - val_mae: 0.0482  
Epoch 648/1000  
33/33  [=====] - 0s 10ms/step - loss: 0.0036 -  
mae: 0.0513 - val_loss: 0.0039 - val_mae: 0.0482  
Epoch 649/1000  
33/33  [=====] - 0s 9ms/step - loss: 0.0036 -  
mae: 0.0513 - val_loss: 0.0039 - val_mae: 0.0482  
Epoch 650/1000  
33/33  [=====] - 0s 12ms/step - loss: 0.0036 -  
mae: 0.0513 - val_loss: 0.0039 - val_mae: 0.0482  
Epoch 651/1000  
33/33  [=====] - 0s 9ms/step - loss: 0.0036 -  
mae: 0.0513 - val_loss: 0.0039 - val_mae: 0.0482
```

```
Epoch 652/1000
33/33 [=====] - 0s 10ms/step - loss: 0.0036 -
mae: 0.0513 - val_loss: 0.0039 - val_mae: 0.0482
Epoch 653/1000
33/33 [=====] - 0s 10ms/step - loss: 0.0036 -
mae: 0.0513 - val_loss: 0.0039 - val_mae: 0.0482
Epoch 654/1000
33/33 [=====] - 0s 9ms/step - loss: 0.0036 -
mae: 0.0513 - val_loss: 0.0039 - val_mae: 0.0482
Epoch 655/1000
33/33 [=====] - 0s 11ms/step - loss: 0.0036 -
mae: 0.0513 - val_loss: 0.0039 - val_mae: 0.0482
Epoch 656/1000
33/33 [=====] - 0s 10ms/step - loss: 0.0036 -
mae: 0.0513 - val_loss: 0.0039 - val_mae: 0.0482
Epoch 657/1000
33/33 [=====] - 0s 10ms/step - loss: 0.0036 -
mae: 0.0513 - val_loss: 0.0039 - val_mae: 0.0482
Epoch 658/1000
33/33 [=====] - 0s 10ms/step - loss: 0.0036 -
mae: 0.0513 - val_loss: 0.0039 - val_mae: 0.0482
Epoch 659/1000
33/33 [=====] - 0s 10ms/step - loss: 0.0036 -
mae: 0.0513 - val_loss: 0.0039 - val_mae: 0.0482
Epoch 660/1000
33/33 [=====] - 0s 10ms/step - loss: 0.0036 -
mae: 0.0513 - val_loss: 0.0039 - val_mae: 0.0482
Epoch 661/1000
33/33 [=====] - 0s 12ms/step - loss: 0.0036 -
mae: 0.0513 - val_loss: 0.0039 - val_mae: 0.0482
Epoch 662/1000
33/33 [=====] - 0s 10ms/step - loss: 0.0036 -
mae: 0.0513 - val_loss: 0.0039 - val_mae: 0.0482
Epoch 663/1000
33/33 [=====] - 0s 10ms/step - loss: 0.0036 -
mae: 0.0513 - val_loss: 0.0039 - val_mae: 0.0482
Epoch 664/1000
33/33 [=====] - 1s 29ms/step - loss: 0.0036 -
```

```
mae: 0.0513 - val_loss: 0.0039 - val_mae: 0.0482
Epoch 665/1000
33/33 [=====] - 0s 12ms/step - loss: 0.0036 -
mae: 0.0513 - val_loss: 0.0039 - val_mae: 0.0482
Epoch 666/1000
33/33 [=====] - 0s 14ms/step - loss: 0.0036 -
mae: 0.0513 - val_loss: 0.0039 - val_mae: 0.0482
Epoch 667/1000
33/33 [=====] - 0s 11ms/step - loss: 0.0036 -
mae: 0.0513 - val_loss: 0.0039 - val_mae: 0.0482
Epoch 668/1000
33/33 [=====] - 0s 11ms/step - loss: 0.0036 -
mae: 0.0513 - val_loss: 0.0039 - val_mae: 0.0482
Epoch 669/1000
33/33 [=====] - 0s 11ms/step - loss: 0.0036 -
mae: 0.0513 - val_loss: 0.0039 - val_mae: 0.0482
Epoch 670/1000
33/33 [=====] - 0s 10ms/step - loss: 0.0036 -
mae: 0.0513 - val_loss: 0.0039 - val_mae: 0.0482
Epoch 671/1000
33/33 [=====] - 0s 10ms/step - loss: 0.0036 -
mae: 0.0513 - val_loss: 0.0039 - val_mae: 0.0482
Epoch 672/1000
33/33 [=====] - 0s 10ms/step - loss: 0.0036 -
mae: 0.0513 - val_loss: 0.0039 - val_mae: 0.0482
Epoch 673/1000
33/33 [=====] - 0s 9ms/step - loss: 0.0036 -
mae: 0.0513 - val_loss: 0.0039 - val_mae: 0.0482
Epoch 674/1000
33/33 [=====] - 0s 10ms/step - loss: 0.0036 -
mae: 0.0513 - val_loss: 0.0039 - val_mae: 0.0482
Epoch 675/1000
33/33 [=====] - 0s 9ms/step - loss: 0.0036 -
mae: 0.0513 - val_loss: 0.0039 - val_mae: 0.0482
Epoch 676/1000
33/33 [=====] - 0s 9ms/step - loss: 0.0036 -
mae: 0.0513 - val_loss: 0.0039 - val_mae: 0.0482
Epoch 677/1000
```

```
33/33  [=====] - 0s 9ms/step - loss: 0.0036 -  
mae: 0.0513 - val_loss: 0.0039 - val_mae: 0.0482  
Epoch 678/1000  
33/33  [=====] - 0s 9ms/step - loss: 0.0036 -  
mae: 0.0513 - val_loss: 0.0039 - val_mae: 0.0482  
Epoch 679/1000  
33/33  [=====] - 0s 11ms/step - loss: 0.0036 -  
mae: 0.0513 - val_loss: 0.0039 - val_mae: 0.0482  
Epoch 680/1000  
33/33  [=====] - 1s 15ms/step - loss: 0.0036 -  
mae: 0.0513 - val_loss: 0.0039 - val_mae: 0.0482  
Epoch 681/1000  
33/33  [=====] - 0s 14ms/step - loss: 0.0036 -  
mae: 0.0513 - val_loss: 0.0039 - val_mae: 0.0482  
Epoch 682/1000  
33/33  [=====] - 1s 16ms/step - loss: 0.0036 -  
mae: 0.0513 - val_loss: 0.0039 - val_mae: 0.0482  
Epoch 683/1000  
33/33  [=====] - 1s 26ms/step - loss: 0.0036 -  
mae: 0.0513 - val_loss: 0.0039 - val_mae: 0.0482  
Epoch 684/1000  
33/33  [=====] - 1s 16ms/step - loss: 0.0036 -  
mae: 0.0513 - val_loss: 0.0039 - val_mae: 0.0482  
Epoch 685/1000  
33/33  [=====] - 0s 10ms/step - loss: 0.0036 -  
mae: 0.0513 - val_loss: 0.0039 - val_mae: 0.0482  
Epoch 686/1000  
33/33  [=====] - 0s 11ms/step - loss: 0.0036 -  
mae: 0.0513 - val_loss: 0.0039 - val_mae: 0.0482  
Epoch 687/1000  
33/33  [=====] - 0s 11ms/step - loss: 0.0036 -  
mae: 0.0513 - val_loss: 0.0039 - val_mae: 0.0482  
Epoch 688/1000  
33/33  [=====] - 0s 13ms/step - loss: 0.0036 -  
mae: 0.0513 - val_loss: 0.0039 - val_mae: 0.0482  
Epoch 689/1000  
33/33  [=====] - 0s 11ms/step - loss: 0.0036 -  
mae: 0.0513 - val_loss: 0.0039 - val_mae: 0.0482
```

```
Epoch 690/1000
33/33 [=====] - 0s 10ms/step - loss: 0.0036 -
mae: 0.0513 - val_loss: 0.0039 - val_mae: 0.0482
Epoch 691/1000
33/33 [=====] - 0s 10ms/step - loss: 0.0036 -
mae: 0.0513 - val_loss: 0.0039 - val_mae: 0.0482
Epoch 692/1000
33/33 [=====] - 0s 10ms/step - loss: 0.0036 -
mae: 0.0513 - val_loss: 0.0039 - val_mae: 0.0482
Epoch 693/1000
33/33 [=====] - 0s 11ms/step - loss: 0.0036 -
mae: 0.0513 - val_loss: 0.0039 - val_mae: 0.0482
Epoch 694/1000
33/33 [=====] - 0s 9ms/step - loss: 0.0036 -
mae: 0.0513 - val_loss: 0.0039 - val_mae: 0.0482
Epoch 695/1000
33/33 [=====] - 0s 9ms/step - loss: 0.0036 -
mae: 0.0513 - val_loss: 0.0039 - val_mae: 0.0482
Epoch 696/1000
33/33 [=====] - 0s 10ms/step - loss: 0.0036 -
mae: 0.0513 - val_loss: 0.0039 - val_mae: 0.0482
Epoch 697/1000
33/33 [=====] - 0s 9ms/step - loss: 0.0036 -
mae: 0.0513 - val_loss: 0.0039 - val_mae: 0.0482
Epoch 698/1000
33/33 [=====] - 0s 9ms/step - loss: 0.0036 -
mae: 0.0513 - val_loss: 0.0039 - val_mae: 0.0482
Epoch 699/1000
33/33 [=====] - 0s 10ms/step - loss: 0.0036 -
mae: 0.0513 - val_loss: 0.0039 - val_mae: 0.0482
Epoch 700/1000
33/33 [=====] - 0s 9ms/step - loss: 0.0036 -
mae: 0.0513 - val_loss: 0.0039 - val_mae: 0.0482
Epoch 701/1000
33/33 [=====] - 0s 9ms/step - loss: 0.0036 -
mae: 0.0513 - val_loss: 0.0039 - val_mae: 0.0482
Epoch 702/1000
33/33 [=====] - 0s 9ms/step - loss: 0.0036 -
```

```
mae: 0.0513 - val_loss: 0.0039 - val_mae: 0.0482
Epoch 703/1000
33/33 [=====] - 0s 9ms/step - loss: 0.0036 -
mae: 0.0513 - val_loss: 0.0039 - val_mae: 0.0482
Epoch 704/1000
33/33 [=====] - 0s 9ms/step - loss: 0.0036 -
mae: 0.0513 - val_loss: 0.0039 - val_mae: 0.0482
Epoch 705/1000
33/33 [=====] - 0s 10ms/step - loss: 0.0036 -
mae: 0.0513 - val_loss: 0.0039 - val_mae: 0.0482
Epoch 706/1000
33/33 [=====] - 0s 9ms/step - loss: 0.0036 -
mae: 0.0512 - val_loss: 0.0039 - val_mae: 0.0482
Epoch 707/1000
33/33 [=====] - 0s 9ms/step - loss: 0.0036 -
mae: 0.0512 - val_loss: 0.0039 - val_mae: 0.0482
Epoch 708/1000
33/33 [=====] - 0s 9ms/step - loss: 0.0036 -
mae: 0.0512 - val_loss: 0.0039 - val_mae: 0.0482
Epoch 709/1000
33/33 [=====] - 0s 9ms/step - loss: 0.0036 -
mae: 0.0512 - val_loss: 0.0039 - val_mae: 0.0482
Epoch 710/1000
33/33 [=====] - 0s 10ms/step - loss: 0.0036 -
mae: 0.0512 - val_loss: 0.0039 - val_mae: 0.0482
Epoch 711/1000
33/33 [=====] - 0s 9ms/step - loss: 0.0036 -
mae: 0.0512 - val_loss: 0.0039 - val_mae: 0.0482
Epoch 712/1000
33/33 [=====] - 0s 9ms/step - loss: 0.0036 -
mae: 0.0512 - val_loss: 0.0039 - val_mae: 0.0482
Epoch 713/1000
33/33 [=====] - 0s 9ms/step - loss: 0.0036 -
mae: 0.0512 - val_loss: 0.0039 - val_mae: 0.0482
Epoch 714/1000
33/33 [=====] - 0s 9ms/step - loss: 0.0036 -
mae: 0.0512 - val_loss: 0.0039 - val_mae: 0.0482
Epoch 715/1000
```

```
33/33  [=====] - 0s 9ms/step - loss: 0.0036 -
mae: 0.0512 - val_loss: 0.0039 - val_mae: 0.0482
Epoch 716/1000
33/33  [=====] - 0s 9ms/step - loss: 0.0036 -
mae: 0.0512 - val_loss: 0.0039 - val_mae: 0.0482
Epoch 717/1000
33/33  [=====] - 0s 9ms/step - loss: 0.0036 -
mae: 0.0512 - val_loss: 0.0039 - val_mae: 0.0482
Epoch 718/1000
33/33  [=====] - 0s 9ms/step - loss: 0.0036 -
mae: 0.0512 - val_loss: 0.0039 - val_mae: 0.0482
Epoch 719/1000
33/33  [=====] - 0s 9ms/step - loss: 0.0036 -
mae: 0.0512 - val_loss: 0.0039 - val_mae: 0.0482
Epoch 720/1000
33/33  [=====] - 0s 9ms/step - loss: 0.0036 -
mae: 0.0512 - val_loss: 0.0039 - val_mae: 0.0482
Epoch 721/1000
33/33  [=====] - 0s 10ms/step - loss: 0.0036 -
mae: 0.0512 - val_loss: 0.0039 - val_mae: 0.0482
Epoch 722/1000
33/33  [=====] - 0s 9ms/step - loss: 0.0036 -
mae: 0.0512 - val_loss: 0.0039 - val_mae: 0.0482
Epoch 723/1000
33/33  [=====] - 0s 10ms/step - loss: 0.0036 -
mae: 0.0512 - val_loss: 0.0039 - val_mae: 0.0482
Epoch 724/1000
33/33  [=====] - 0s 9ms/step - loss: 0.0036 -
mae: 0.0512 - val_loss: 0.0039 - val_mae: 0.0482
Epoch 725/1000
33/33  [=====] - 0s 9ms/step - loss: 0.0036 -
mae: 0.0512 - val_loss: 0.0039 - val_mae: 0.0482
Epoch 726/1000
33/33  [=====] - 0s 10ms/step - loss: 0.0036 -
mae: 0.0512 - val_loss: 0.0039 - val_mae: 0.0482
Epoch 727/1000
33/33  [=====] - 0s 9ms/step - loss: 0.0036 -
mae: 0.0512 - val_loss: 0.0039 - val_mae: 0.0482
```

```
Epoch 728/1000
33/33  [=====] - 0s 9ms/step - loss: 0.0036 -
mae: 0.0512 - val_loss: 0.0039 - val_mae: 0.0482
Epoch 729/1000
33/33  [=====] - 0s 9ms/step - loss: 0.0036 -
mae: 0.0512 - val_loss: 0.0039 - val_mae: 0.0482
Epoch 730/1000
33/33  [=====] - 0s 9ms/step - loss: 0.0036 -
mae: 0.0512 - val_loss: 0.0039 - val_mae: 0.0482
Epoch 731/1000
33/33  [=====] - 0s 10ms/step - loss: 0.0036 -
mae: 0.0512 - val_loss: 0.0039 - val_mae: 0.0482
Epoch 732/1000
33/33  [=====] - 0s 10ms/step - loss: 0.0036 -
mae: 0.0512 - val_loss: 0.0039 - val_mae: 0.0482
Epoch 733/1000
33/33  [=====] - 28s 880ms/step - loss: 0.0036 -
mae: 0.0512 - val_loss: 0.0039 - val_mae: 0.0482
Epoch 734/1000
33/33  [=====] - 0s 12ms/step - loss: 0.0036 -
mae: 0.0512 - val_loss: 0.0039 - val_mae: 0.0482
Epoch 735/1000
33/33  [=====] - 0s 10ms/step - loss: 0.0036 -
mae: 0.0512 - val_loss: 0.0039 - val_mae: 0.0482
Epoch 736/1000
33/33  [=====] - 0s 8ms/step - loss: 0.0036 -
mae: 0.0512 - val_loss: 0.0039 - val_mae: 0.0482
Epoch 737/1000
33/33  [=====] - 0s 9ms/step - loss: 0.0036 -
mae: 0.0512 - val_loss: 0.0039 - val_mae: 0.0482
Epoch 738/1000
33/33  [=====] - 0s 11ms/step - loss: 0.0036 -
mae: 0.0512 - val_loss: 0.0039 - val_mae: 0.0482
Epoch 739/1000
33/33  [=====] - 0s 9ms/step - loss: 0.0036 -
mae: 0.0512 - val_loss: 0.0039 - val_mae: 0.0482
Epoch 740/1000
33/33  [=====] - 0s 9ms/step - loss: 0.0036 -
```

```
mae: 0.0512 - val_loss: 0.0039 - val_mae: 0.0482
Epoch 741/1000
33/33 [=====] - 0s 9ms/step - loss: 0.0036 -
mae: 0.0512 - val_loss: 0.0039 - val_mae: 0.0482
Epoch 742/1000
33/33 [=====] - 2s 47ms/step - loss: 0.0036 -
mae: 0.0512 - val_loss: 0.0039 - val_mae: 0.0482
Epoch 743/1000
33/33 [=====] - 2s 55ms/step - loss: 0.0036 -
mae: 0.0512 - val_loss: 0.0039 - val_mae: 0.0482
Epoch 744/1000
33/33 [=====] - 2s 57ms/step - loss: 0.0036 -
mae: 0.0512 - val_loss: 0.0039 - val_mae: 0.0482
Epoch 745/1000
33/33 [=====] - 2s 53ms/step - loss: 0.0036 -
mae: 0.0512 - val_loss: 0.0039 - val_mae: 0.0482
Epoch 746/1000
33/33 [=====] - 2s 60ms/step - loss: 0.0036 -
mae: 0.0512 - val_loss: 0.0039 - val_mae: 0.0482
Epoch 747/1000
33/33 [=====] - 465s 15s/step - loss: 0.0036 -
mae: 0.0512 - val_loss: 0.0039 - val_mae: 0.0482
Epoch 748/1000
33/33 [=====] - 1s 15ms/step - loss: 0.0036 -
mae: 0.0512 - val_loss: 0.0039 - val_mae: 0.0482
Epoch 749/1000
33/33 [=====] - 0s 10ms/step - loss: 0.0036 -
mae: 0.0512 - val_loss: 0.0039 - val_mae: 0.0482
Epoch 750/1000
33/33 [=====] - 0s 11ms/step - loss: 0.0036 -
mae: 0.0512 - val_loss: 0.0039 - val_mae: 0.0482
Epoch 751/1000
33/33 [=====] - 0s 13ms/step - loss: 0.0036 -
mae: 0.0512 - val_loss: 0.0039 - val_mae: 0.0482
Epoch 752/1000
33/33 [=====] - 1s 21ms/step - loss: 0.0036 -
mae: 0.0512 - val_loss: 0.0039 - val_mae: 0.0482
Epoch 753/1000
```

```
33/33  [=====] - 0s 14ms/step - loss: 0.0036 -  
mae: 0.0512 - val_loss: 0.0039 - val_mae: 0.0482  
Epoch 754/1000  
33/33  [=====] - 1s 17ms/step - loss: 0.0036 -  
mae: 0.0512 - val_loss: 0.0039 - val_mae: 0.0482  
Epoch 755/1000  
33/33  [=====] - 0s 12ms/step - loss: 0.0036 -  
mae: 0.0512 - val_loss: 0.0039 - val_mae: 0.0482  
Epoch 756/1000  
33/33  [=====] - 0s 10ms/step - loss: 0.0036 -  
mae: 0.0512 - val_loss: 0.0039 - val_mae: 0.0482  
Epoch 757/1000  
33/33  [=====] - 0s 10ms/step - loss: 0.0036 -  
mae: 0.0512 - val_loss: 0.0039 - val_mae: 0.0482  
Epoch 758/1000  
33/33  [=====] - 0s 11ms/step - loss: 0.0036 -  
mae: 0.0512 - val_loss: 0.0039 - val_mae: 0.0482  
Epoch 759/1000  
33/33  [=====] - 0s 12ms/step - loss: 0.0036 -  
mae: 0.0512 - val_loss: 0.0039 - val_mae: 0.0482  
Epoch 760/1000  
33/33  [=====] - 0s 10ms/step - loss: 0.0036 -  
mae: 0.0512 - val_loss: 0.0039 - val_mae: 0.0482  
Epoch 761/1000  
33/33  [=====] - 0s 9ms/step - loss: 0.0036 -  
mae: 0.0512 - val_loss: 0.0039 - val_mae: 0.0482  
Epoch 762/1000  
33/33  [=====] - 0s 9ms/step - loss: 0.0036 -  
mae: 0.0512 - val_loss: 0.0039 - val_mae: 0.0482  
Epoch 763/1000  
33/33  [=====] - 0s 12ms/step - loss: 0.0036 -  
mae: 0.0512 - val_loss: 0.0039 - val_mae: 0.0482  
Epoch 764/1000  
33/33  [=====] - 0s 11ms/step - loss: 0.0036 -  
mae: 0.0512 - val_loss: 0.0039 - val_mae: 0.0482  
Epoch 765/1000  
33/33  [=====] - 0s 11ms/step - loss: 0.0036 -  
mae: 0.0512 - val_loss: 0.0039 - val_mae: 0.0482
```

```
Epoch 766/1000
33/33 [=====] - 0s 11ms/step - loss: 0.0036 -
mae: 0.0512 - val_loss: 0.0039 - val_mae: 0.0482
Epoch 767/1000
33/33 [=====] - 1s 20ms/step - loss: 0.0036 -
mae: 0.0512 - val_loss: 0.0039 - val_mae: 0.0482
Epoch 768/1000
33/33 [=====] - 0s 10ms/step - loss: 0.0036 -
mae: 0.0512 - val_loss: 0.0039 - val_mae: 0.0482
Epoch 769/1000
33/33 [=====] - 1s 16ms/step - loss: 0.0036 -
mae: 0.0512 - val_loss: 0.0039 - val_mae: 0.0482
Epoch 770/1000
33/33 [=====] - 1s 17ms/step - loss: 0.0036 -
mae: 0.0512 - val_loss: 0.0039 - val_mae: 0.0482
Epoch 771/1000
33/33 [=====] - 2s 57ms/step - loss: 0.0036 -
mae: 0.0512 - val_loss: 0.0039 - val_mae: 0.0482
Epoch 772/1000
33/33 [=====] - 2s 53ms/step - loss: 0.0036 -
mae: 0.0512 - val_loss: 0.0039 - val_mae: 0.0482
Epoch 773/1000
33/33 [=====] - 2s 61ms/step - loss: 0.0036 -
mae: 0.0512 - val_loss: 0.0039 - val_mae: 0.0482
Epoch 774/1000
33/33 [=====] - 2s 58ms/step - loss: 0.0036 -
mae: 0.0512 - val_loss: 0.0039 - val_mae: 0.0482
Epoch 775/1000
33/33 [=====] - 2s 54ms/step - loss: 0.0036 -
mae: 0.0512 - val_loss: 0.0039 - val_mae: 0.0482
Epoch 776/1000
33/33 [=====] - 2s 59ms/step - loss: 0.0036 -
mae: 0.0512 - val_loss: 0.0039 - val_mae: 0.0482
Epoch 777/1000
33/33 [=====] - 2s 56ms/step - loss: 0.0036 -
mae: 0.0512 - val_loss: 0.0039 - val_mae: 0.0482
Epoch 778/1000
33/33 [=====] - 2s 53ms/step - loss: 0.0036 -
```

```
mae: 0.0512 - val_loss: 0.0039 - val_mae: 0.0482
Epoch 779/1000
33/33 [=====] - 2s 57ms/step - loss: 0.0036 -
mae: 0.0512 - val_loss: 0.0039 - val_mae: 0.0482
Epoch 780/1000
33/33 [=====] - 2s 53ms/step - loss: 0.0036 -
mae: 0.0512 - val_loss: 0.0039 - val_mae: 0.0482
Epoch 781/1000
33/33 [=====] - 2s 56ms/step - loss: 0.0036 -
mae: 0.0512 - val_loss: 0.0039 - val_mae: 0.0482
Epoch 782/1000
33/33 [=====] - 2s 53ms/step - loss: 0.0036 -
mae: 0.0512 - val_loss: 0.0039 - val_mae: 0.0482
Epoch 783/1000
33/33 [=====] - 2s 53ms/step - loss: 0.0036 -
mae: 0.0512 - val_loss: 0.0039 - val_mae: 0.0482
Epoch 784/1000
33/33 [=====] - 2s 57ms/step - loss: 0.0036 -
mae: 0.0511 - val_loss: 0.0039 - val_mae: 0.0482
Epoch 785/1000
33/33 [=====] - 2s 54ms/step - loss: 0.0036 -
mae: 0.0511 - val_loss: 0.0039 - val_mae: 0.0482
Epoch 786/1000
33/33 [=====] - 2s 54ms/step - loss: 0.0036 -
mae: 0.0511 - val_loss: 0.0039 - val_mae: 0.0482
Epoch 787/1000
33/33 [=====] - 2s 55ms/step - loss: 0.0036 -
mae: 0.0511 - val_loss: 0.0039 - val_mae: 0.0482
Epoch 788/1000
33/33 [=====] - 2s 53ms/step - loss: 0.0036 -
mae: 0.0511 - val_loss: 0.0039 - val_mae: 0.0482
Epoch 789/1000
33/33 [=====] - 2s 55ms/step - loss: 0.0036 -
mae: 0.0511 - val_loss: 0.0039 - val_mae: 0.0482
Epoch 790/1000
33/33 [=====] - 2s 53ms/step - loss: 0.0036 -
mae: 0.0511 - val_loss: 0.0039 - val_mae: 0.0482
Epoch 791/1000
```

```
33/33  [=====] - 2s 55ms/step - loss: 0.0036 -  
mae: 0.0511 - val_loss: 0.0039 - val_mae: 0.0482  
Epoch 792/1000  
33/33  [=====] - 2s 55ms/step - loss: 0.0036 -  
mae: 0.0511 - val_loss: 0.0039 - val_mae: 0.0482  
Epoch 793/1000  
33/33  [=====] - 2s 57ms/step - loss: 0.0036 -  
mae: 0.0511 - val_loss: 0.0039 - val_mae: 0.0482  
Epoch 794/1000  
33/33  [=====] - 2s 56ms/step - loss: 0.0036 -  
mae: 0.0511 - val_loss: 0.0039 - val_mae: 0.0482  
Epoch 795/1000  
33/33  [=====] - 2s 54ms/step - loss: 0.0036 -  
mae: 0.0511 - val_loss: 0.0039 - val_mae: 0.0482  
Epoch 796/1000  
33/33  [=====] - 2s 57ms/step - loss: 0.0036 -  
mae: 0.0511 - val_loss: 0.0039 - val_mae: 0.0482  
Epoch 797/1000  
33/33  [=====] - 2s 55ms/step - loss: 0.0036 -  
mae: 0.0511 - val_loss: 0.0039 - val_mae: 0.0482  
Epoch 798/1000  
33/33  [=====] - 1s 31ms/step - loss: 0.0036 -  
mae: 0.0511 - val_loss: 0.0039 - val_mae: 0.0482  
Epoch 799/1000  
33/33  [=====] - 0s 9ms/step - loss: 0.0036 -  
mae: 0.0511 - val_loss: 0.0039 - val_mae: 0.0482  
Epoch 800/1000  
33/33  [=====] - 0s 8ms/step - loss: 0.0036 -  
mae: 0.0511 - val_loss: 0.0039 - val_mae: 0.0482  
Epoch 801/1000  
33/33  [=====] - 0s 9ms/step - loss: 0.0036 -  
mae: 0.0511 - val_loss: 0.0039 - val_mae: 0.0482  
Epoch 802/1000  
33/33  [=====] - 0s 10ms/step - loss: 0.0036 -  
mae: 0.0511 - val_loss: 0.0039 - val_mae: 0.0482  
Epoch 803/1000  
33/33  [=====] - 0s 9ms/step - loss: 0.0036 -  
mae: 0.0511 - val_loss: 0.0039 - val_mae: 0.0482
```

```
Epoch 804/1000
33/33 [=====] - 0s 9ms/step - loss: 0.0036 -
mae: 0.0511 - val_loss: 0.0039 - val_mae: 0.0482
Epoch 805/1000
33/33 [=====] - 0s 9ms/step - loss: 0.0036 -
mae: 0.0511 - val_loss: 0.0039 - val_mae: 0.0482
Epoch 806/1000
33/33 [=====] - 0s 10ms/step - loss: 0.0036 -
mae: 0.0511 - val_loss: 0.0039 - val_mae: 0.0482
Epoch 807/1000
33/33 [=====] - 0s 10ms/step - loss: 0.0036 -
mae: 0.0511 - val_loss: 0.0039 - val_mae: 0.0482
Epoch 808/1000
33/33 [=====] - 0s 9ms/step - loss: 0.0036 -
mae: 0.0511 - val_loss: 0.0039 - val_mae: 0.0482
Epoch 809/1000
33/33 [=====] - 0s 9ms/step - loss: 0.0036 -
mae: 0.0511 - val_loss: 0.0039 - val_mae: 0.0482
Epoch 810/1000
33/33 [=====] - 0s 9ms/step - loss: 0.0036 -
mae: 0.0511 - val_loss: 0.0039 - val_mae: 0.0482
Epoch 811/1000
33/33 [=====] - 0s 9ms/step - loss: 0.0036 -
mae: 0.0511 - val_loss: 0.0039 - val_mae: 0.0482
Epoch 812/1000
33/33 [=====] - 0s 10ms/step - loss: 0.0036 -
mae: 0.0511 - val_loss: 0.0039 - val_mae: 0.0482
Epoch 813/1000
33/33 [=====] - 0s 10ms/step - loss: 0.0036 -
mae: 0.0511 - val_loss: 0.0039 - val_mae: 0.0482
Epoch 814/1000
33/33 [=====] - 0s 9ms/step - loss: 0.0036 -
mae: 0.0511 - val_loss: 0.0039 - val_mae: 0.0482
Epoch 815/1000
33/33 [=====] - 0s 10ms/step - loss: 0.0036 -
mae: 0.0511 - val_loss: 0.0039 - val_mae: 0.0482
Epoch 816/1000
33/33 [=====] - 0s 10ms/step - loss: 0.0036 -
```

```
mae: 0.0511 - val_loss: 0.0039 - val_mae: 0.0482
Epoch 817/1000
33/33 [=====] - 0s 9ms/step - loss: 0.0036 -
mae: 0.0511 - val_loss: 0.0039 - val_mae: 0.0482
Epoch 818/1000
33/33 [=====] - 0s 9ms/step - loss: 0.0036 -
mae: 0.0511 - val_loss: 0.0039 - val_mae: 0.0482
Epoch 819/1000
33/33 [=====] - 0s 9ms/step - loss: 0.0036 -
mae: 0.0511 - val_loss: 0.0039 - val_mae: 0.0482
Epoch 820/1000
33/33 [=====] - 0s 10ms/step - loss: 0.0036 -
mae: 0.0511 - val_loss: 0.0039 - val_mae: 0.0482
Epoch 821/1000
33/33 [=====] - 0s 10ms/step - loss: 0.0036 -
mae: 0.0511 - val_loss: 0.0039 - val_mae: 0.0482
Epoch 822/1000
33/33 [=====] - 0s 9ms/step - loss: 0.0036 -
mae: 0.0511 - val_loss: 0.0039 - val_mae: 0.0482
Epoch 823/1000
33/33 [=====] - 0s 9ms/step - loss: 0.0036 -
mae: 0.0511 - val_loss: 0.0039 - val_mae: 0.0482
Epoch 824/1000
33/33 [=====] - 0s 9ms/step - loss: 0.0036 -
mae: 0.0511 - val_loss: 0.0039 - val_mae: 0.0482
Epoch 825/1000
33/33 [=====] - 0s 9ms/step - loss: 0.0036 -
mae: 0.0511 - val_loss: 0.0039 - val_mae: 0.0482
Epoch 826/1000
33/33 [=====] - 0s 11ms/step - loss: 0.0036 -
mae: 0.0511 - val_loss: 0.0039 - val_mae: 0.0482
Epoch 827/1000
33/33 [=====] - 0s 9ms/step - loss: 0.0036 -
mae: 0.0511 - val_loss: 0.0039 - val_mae: 0.0482
Epoch 828/1000
33/33 [=====] - 0s 9ms/step - loss: 0.0036 -
mae: 0.0511 - val_loss: 0.0039 - val_mae: 0.0482
Epoch 829/1000
```

```
33/33  [=====] - 0s 9ms/step - loss: 0.0036 -
mae: 0.0511 - val_loss: 0.0039 - val_mae: 0.0482
Epoch 830/1000
33/33  [=====] - 0s 10ms/step - loss: 0.0036 -
mae: 0.0511 - val_loss: 0.0039 - val_mae: 0.0482
Epoch 831/1000
33/33  [=====] - 0s 9ms/step - loss: 0.0036 -
mae: 0.0511 - val_loss: 0.0039 - val_mae: 0.0482
Epoch 832/1000
33/33  [=====] - 0s 9ms/step - loss: 0.0036 -
mae: 0.0511 - val_loss: 0.0039 - val_mae: 0.0482
Epoch 833/1000
33/33  [=====] - 0s 9ms/step - loss: 0.0036 -
mae: 0.0511 - val_loss: 0.0039 - val_mae: 0.0482
Epoch 834/1000
33/33  [=====] - 0s 9ms/step - loss: 0.0036 -
mae: 0.0511 - val_loss: 0.0039 - val_mae: 0.0482
Epoch 835/1000
33/33  [=====] - 0s 10ms/step - loss: 0.0036 -
mae: 0.0511 - val_loss: 0.0039 - val_mae: 0.0482
Epoch 836/1000
33/33  [=====] - 0s 9ms/step - loss: 0.0036 -
mae: 0.0511 - val_loss: 0.0039 - val_mae: 0.0482
Epoch 837/1000
33/33  [=====] - 0s 9ms/step - loss: 0.0036 -
mae: 0.0511 - val_loss: 0.0039 - val_mae: 0.0482
Epoch 838/1000
33/33  [=====] - 0s 9ms/step - loss: 0.0036 -
mae: 0.0511 - val_loss: 0.0039 - val_mae: 0.0482
Epoch 839/1000
33/33  [=====] - 0s 9ms/step - loss: 0.0036 -
mae: 0.0511 - val_loss: 0.0039 - val_mae: 0.0482
Epoch 840/1000
33/33  [=====] - 0s 11ms/step - loss: 0.0036 -
mae: 0.0511 - val_loss: 0.0039 - val_mae: 0.0482
Epoch 841/1000
33/33  [=====] - 0s 9ms/step - loss: 0.0036 -
mae: 0.0511 - val_loss: 0.0039 - val_mae: 0.0482
```

```

Epoch 842/1000
33/33 [=====] - 0s 9ms/step - loss: 0.0036 -
mae: 0.0511 - val_loss: 0.0039 - val_mae: 0.0482
Epoch 843/1000
33/33 [=====] - 0s 9ms/step - loss: 0.0036 -
mae: 0.0511 - val_loss: 0.0039 - val_mae: 0.0482
Epoch 844/1000
33/33 [=====] - 0s 10ms/step - loss: 0.0036 -
mae: 0.0511 - val_loss: 0.0039 - val_mae: 0.0482
Epoch 845/1000
33/33 [=====] - 0s 9ms/step - loss: 0.0036 -
mae: 0.0511 - val_loss: 0.0039 - val_mae: 0.0482

```

We can evaluate the model forecasts MAE for each steps:

```

Y_pred_valid = seq2seq_model.predict(seq2seq_valid)
for ahead in range(14):
    preds = pd.Series(Y_pred_valid[:-1, -1, ahead], index=mulvar_valid.index[56 +
    ↵ ahead : -14 + ahead])
    mae = (preds - mulvar_valid["rail"]).abs().mean() * 1e6
    print(f"MAE for +{ahead + 1}: {mae:.0f}")

```

```

3/3 [=====] - 0s 19ms/step
MAE for +1: 25,418
MAE for +2: 28,053
MAE for +3: 29,681
MAE for +4: 33,068
MAE for +5: 33,684
MAE for +6: 33,115
MAE for +7: 33,976
MAE for +8: 31,357
MAE for +9: 31,070
MAE for +10: 31,363
MAE for +11: 37,903
MAE for +12: 35,931
MAE for +13: 37,171
MAE for +14: 32,831

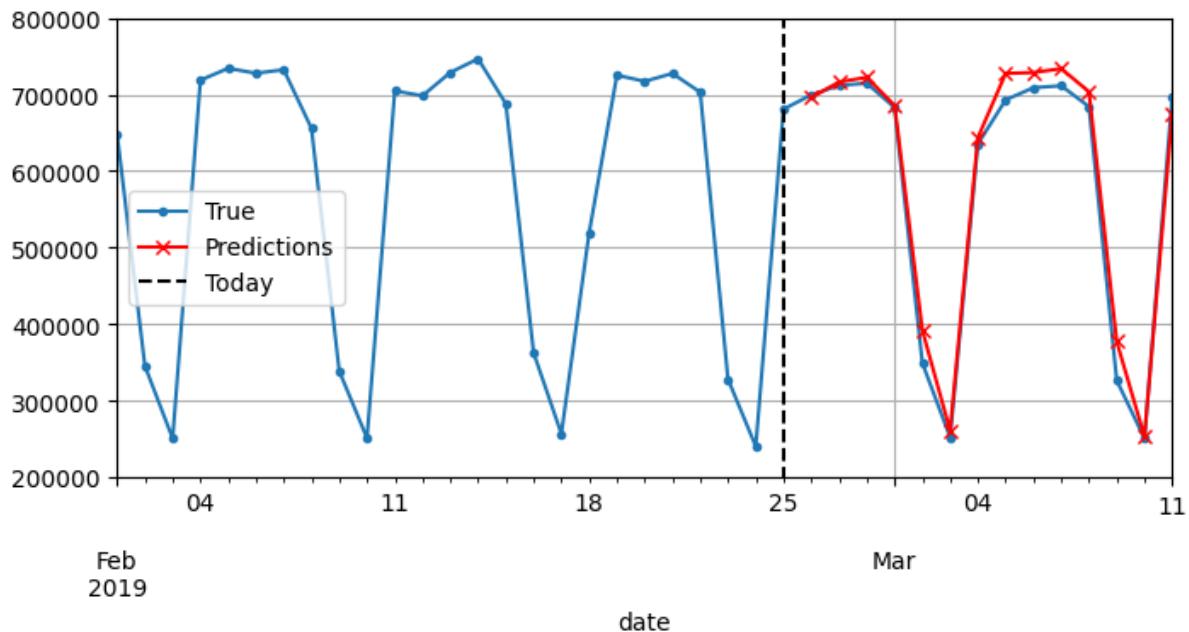
```

```
X = mulvar_valid.to_numpy()[np.newaxis, :seq_length] # shape [1, 56, 5]
Y_pred = seq2seq_model.predict(X)[0, -1] # shape [1, 14]
```

1/1 [=====] - 0s 19ms/step

```
Y_pred = pd.Series(Y_pred, index=pd.date_range("2019-02-26", "2019-03-11"))
fig, ax = plt.subplots(figsize=(8, 3.5))
(mulvar_valid.rail * 1e6)["2019-02-01":"2019-03-11"].plot(label="True", marker=".",
    ax=ax)
(Y_pred * 1e6).plot(label="Predictions", grid=True, marker="x", color="r", ax=ax)
ax.vlines("2019-02-25", 0, 1e6, color="k", linestyle="--", label="Today")
ax.set_ylim([200_000, 800_000])
plt.legend(loc="center left")

plt.show()
```



We can combine different approaches to forecasting multiple steps ahead. For example, we can train a model that forecasts N days ahead, then take its output and append it to the inputs, then run the model again to get forecasts for the following N days, and possibly repeat the process. However, simple RNNs can be quite good at forecasting short time series, but they do not perform as well on long time series.

1.5 Handling Long Sequences

xxxxxTo train an RNN on long sequences, we must run it over many time steps, making the unrolled RNN a very deep network. Just like any deep neural network it may suffer from the unstable gradients problem, discussed in Chapter 11: it may take forever to train, or training may be unstable. Moreover, when an RNN processes a long sequence, it will gradually forget the first inputs in the sequence. Let's look at both these problems, starting with the unstable gradients problem.xxx

1.5.1 Unstable Gradients Problem

On long sequences, we must run the network over many time steps, **making the unrolled RNN a very deep network**. Like any deep neural network, it may suffer from the unstable gradients problem: it may take forever to train, or training may be unstable. Many of the tricks used in deep nets can also be used for RNNs: good parameter initialization, faster optimizers, dropout, and so on.

However, **non-saturating activation functions (e.g., ReLU) may not help**. Suppose that gradient descent updates the weights in a way that increases the outputs slightly at the first time step. Because the same weights are used at every time step, the outputs at the second time step may also be slightly increased, and those at the third, and so on until the outputs explode, and a nonsaturating activation function does not prevent that. We can reduce this risk by using a smaller learning rate or we can use a saturating activation function (e.g. hyperbolic tangent). In the same way, gradients can explode. If we notice that training is unstable, we should monitor the size of the gradients and perhaps use gradient clipping.

Moreover, **batch normalization cannot be used as efficiently with RNNs**. We cannot use it between time steps, only between recurrent layers. To be more precise, it is technically possible to add a batch normalization layer to a memory cell, so that it will be applied at each time step. However, the same batch normalization layer will be used at each time step, with the same parameters, regardless of the actual scale and offset of the inputs and hidden state. In practice, this does not yield good results: as demonstrated in Batch Normalized Recurrent Neural Networks. It was slightly better than nothing when applied between recurrent layers, but not within recurrent layers.

Layer normalization often works better with RNNs: it is similar to batch normalization, but instead of normalizing across the batch dimension, it normalizes across the features dimension. One advantage is that it can compute the required statistics on the fly, at each time step, independently for each instance. This also means that it behaves the same way during training and

testing, and it does not need to use exponential moving averages to estimate the feature statistics across all instances in the training set. In an RNN, it is typically used right after the linear combination of the inputs and the hidden states.

We need to define a custom memory cell, which is just like a regular layer, except it will apply layer normalization at each time step, and its call() method takes two arguments: the inputs at the current time step and the hidden states from the previous time step.

```
class LNSimpleRNNCell(tf.keras.layers.Layer):
    def __init__(self, units, activation="tanh", **kwargs):
        super().__init__(**kwargs)
        self.state_size = units
        self.output_size = units

        # create a SimpleRNNCell with no activation function (because we
        # want to perform layer normalization after the linear operation but
        # before the activation function)
        self.simple_rnn_cell = tf.keras.layers.SimpleRNNCell(units, activation=None)
        self.layer_norm = tf.keras.layers.LayerNormalization()
        self.activation = tf.keras.activations.get(activation)

    def call(self, inputs, states):
        outputs, new_states = self.simple_rnn_cell(inputs, states)
        norm_outputs = self.activation(self.layer_norm(outputs))
        return norm_outputs, [norm_outputs]
```

Notice that the states argument is a list containing one or more tensors. In a simple RNN cell it contains a single tensor equal to the outputs of the previous time step, but other cells may have multiple state tensors. A cell must also have a state_size attribute and an output_size attribute. In a simple RNN, both are simply equal to the number of units.

```
custom_ln_model = tf.keras.Sequential([
    tf.keras.layers.RNN(LNSimpleRNNCell(32), return_sequences=True,
    ↵ input_shape=[None, 2]),
    tf.keras.layers.Dense(14)
])

opt = tf.keras.optimizers.SGD(learning_rate=0.02, momentum=0.9)
custom_ln_model.compile(loss=tf.keras.losses.Huber(), optimizer=opt, metrics=["mae"])
```

Just training for some epochs to show that it works:

```
history = custom_ln_model.fit(seq2seq_train, validation_data=seq2seq_valid,
→ epochs=500, callbacks=[early_stopping_cb])
```

Epoch 1/5

```
33/33 [=====] - 2s 25ms/step - loss: 0.1348 -  
mae: 0.3695 - val_loss: 0.0309 - val_mae: 0.1758
```

Epoch 2/5

```
33/33 [=====] - 1s 17ms/step - loss: 0.0214 -  
mae: 0.1760 - val_loss: 0.0219 - val_mae: 0.1611
```

Epoch 3/5

```
33/33 [=====] - 1s 17ms/step - loss: 0.0192 -  
mae: 0.1661 - val_loss: 0.0201 - val_mae: 0.1625
```

Epoch 4/5

```
33/33 [=====] - 1s 17ms/step - loss: 0.0183 -  
mae: 0.1644 - val_loss: 0.0194 - val_mae: 0.1598
```

Epoch 5/5

```
33/33 [=====] - 1s 16ms/step - loss: 0.0177 -  
mae: 0.1626 - val_loss: 0.0189 - val_mae: 0.1579
```

```
Y_pred_valid = custom_ln_model.predict(seq2seq_valid)
for ahead in range(14):
    preds = pd.Series(Y_pred_valid[::−1, −1, ahead], index=mulvar_valid.index[56 +
→ ahead : −14 + ahead])
    mae = (preds − mulvar_valid["rail"]).abs().mean() * 1e6
    print(f"MAE for +{ahead + 1}: {mae:.0f}")
```

3/3 [=====] - 0s 14ms/step

```
MAE for +1: 136,748
MAE for +2: 162,667
MAE for +3: 145,744
MAE for +4: 157,006
MAE for +5: 151,718
MAE for +6: 145,106
MAE for +7: 151,408
MAE for +8: 133,953
MAE for +9: 161,726
```

```
MAE for +10: 161,509  
MAE for +11: 154,425  
MAE for +12: 158,581  
MAE for +13: 166,979  
MAE for +14: 130,431
```

Similarly, we can create a custom cell to apply dropout between each time step. However, most recurrent layers in Keras have dropout hyperparameters. With these techniques, we can alleviate the unstable gradients problem and train an RNN much more efficiently.

1.6 Short-Term Memory Problem

One of the appeals of RNNs is the idea that they might be able to connect previous information to the present task, such as using previous video frames might inform the understanding of the present frame. Sometimes, we only need to look at recent information to perform the present task. For example, consider a language model trying to predict the next word based on the previous ones. If we are trying to predict the last word in "the clouds are in the" we don't need any further context, it's pretty obvious the next word is going to be "sky". In such cases, where the gap between the relevant information and the place that it's needed is small, RNNs can learn to use the past information. But there are also cases where we need **more context**. Consider trying to predict the last word in the text "I grew up in France... I speak fluent". Recent information suggests that the next word is probably the name of a language, but if we want to narrow down which language, we need the context of France, from further back. It's entirely possible for the gap between the relevant information and the point where it is needed to become very large. Unfortunately, as that gap grows, RNNs become unable to learn to connect the information.

Due to the transformations that the data goes through when traversing an RNN, some information is lost at each time step. After a while, the RNN's state contains virtually no trace of the first inputs. To tackle this problem, various types of cells with **long-term memory** have been introduced. They proven so successful that the basic cells are not used anymore.

The **long short-term memory (LSTM)** cell looks exactly like a regular cell from the input-output point of view (), however its state is split into two vectors: h_t (the short-term state) and c_t (the long-term state). The key idea is that **the network can learn what to store in the long-term state, what to throw away, and what to read from it**.

The long-term cell state c_t is like a conveyor belt: it runs straight down the entire chain, with only some minor linear interactions. It's very easy for information to just flow along it unchanged:

The LSTM cell has the ability to remove or add information to the long-term state, carefully regulated by structures called gates. They are composed of a sigmoid neural net layer and a pointwise multiplication operation. The sigmoid layer outputs numbers between zero and one, describing how much of each component should be let through. An LSTM cell has three of these gates, to protect and control the cell long-term state.

The first step is to decide what information we're going to throw away from the long-term cell state. This decision is made by a sigmoid layer called **the forget gate layer**:

$$f_t = \sigma(W_f \cdot [h_{t-1}, x_t] + b_f)$$

In the language model example, the cell long-term state might include the gender of the present subject, so that the correct pronouns can be used. When we see a new subject, we want to forget the gender of the old subject.

The next step is to decide what new information we're going to store in the cell long-term state:

$$i_t = \sigma(W_i \cdot [h_{t-1}, x_t] + b_i)$$

$$\hat{c}_t = \tanh(W_c \cdot [h_{t-1}, x_t] + b_c)$$

This has two parts. First, a sigmoid layer called **the input gate layer** decides which values we'll update. Next, a tanh layer creates a vector of new candidate values \hat{c}_t , that could be added to the state. In the next step, we'll combine these two to create an update to the state. In the example, we want to add the gender of the new subject to the cell state, to replace the old one we are forgetting.

It's now time to update the old long-term state C_{t-1} into the new long-term state C_t . The previous steps already decided what to do, we just need to actually do it:

$$c_t = f_t * C_{t-1} + i_t * \hat{c}_t$$

We multiply the old state by f_t (forgetting things we decided to forget earlier), then we add \hat{c}_t (the new candidate values), scaled by how much we decided to update each state value. In the example, this is where we actually drop the information about the old subject gender and add the new information.

Finally, we need to decide what we're going to output using the **output gate**. This output will be based on our long-term state, but will be a filtered version. First, we run a sigmoid layer which decides what parts of the long-term state we're going to output. Then, we put the long-term state through tanh (to push the values to be between -1 and 1) and multiply it by the output of the sigmoid gate, so that we only output the parts we decided to:

$$o_t = \sigma(W_o \cdot [h_{t-1}, x_t] + b_o)$$

$$h_t = o_t * \tanh(c_t)$$

In the example, since it just saw a subject, it might want to output information relevant to a verb, in case that's what is coming next. For example, it might output whether the subject is singular or plural, so that we know what form a verb should be conjugated into if that's what follows next. In short, an LSTM cell can learn to recognize an important input (that's the role of the input gate), store it in the long-term state, preserve it for as long as it is needed (that's the role of the forget gate), and extract it whenever it is needed.

In Keras, we can simply use the LSTM layer instead of the SimpleRNN layer:

```
lstm_model = tf.keras.models.Sequential([
    tf.keras.layers.LSTM(32, return_sequences=True, input_shape=[None, 2]),
    tf.keras.layers.Dense(14)
])
```

```
2023-07-11 13:02:30.223789: I tensorflow/core/common_runtime/executor.cc:1197] [/device:CPU
  [[{{node gradients/split_2_grad(concat/split_2/split_dim)}}]]
2023-07-11 13:02:30.225114: I tensorflow/core/common_runtime/executor.cc:1197] [/device:CPU
  [[{{node gradients/split_grad(concat/split/split_dim)}}]]]
2023-07-11 13:02:30.226109: I tensorflow/core/common_runtime/executor.cc:1197] [/device:CPU
  [[{{node gradients/split_1_grad(concat/split_1/split_dim)}}]]]

opt = tf.keras.optimizers.SGD(learning_rate=0.02, momentum=0.9)
lstm_model.compile(loss=tf.keras.losses.Huber(), optimizer=opt, metrics=["mae"])
```

Just training for 5 epochs to show that it works:

```
history = lstm_model.fit(seq2seq_train, validation_data=seq2seq_valid, epochs=5,
                         callbacks=[early_stopping_cb])
```

Epoch 1/5

```
33/33  [=====] - 1s 15ms/step - loss: 0.0076 -
mae: 0.0957 - val_loss: 0.0091 - val_mae: 0.1037
Epoch 2/5
33/33  [=====] - 0s 14ms/step - loss: 0.0076 -
mae: 0.0956 - val_loss: 0.0091 - val_mae: 0.1037
Epoch 3/5
```

```
33/33 [=====] - 0s 14ms/step - loss: 0.0076 -
mae: 0.0956 - val_loss: 0.0091 - val_mae: 0.1037
Epoch 4/5
33/33 [=====] - 0s 14ms/step - loss: 0.0076 -
mae: 0.0956 - val_loss: 0.0091 - val_mae: 0.1036
Epoch 5/5
33/33 [=====] - 0s 14ms/step - loss: 0.0076 -
mae: 0.0955 - val_loss: 0.0090 - val_mae: 0.1036
```

There are several variants of the LSTM cell. One particularly popular variant is the **Gated Recurrent Unit (GRU) cell**. It is a simplified version of the LSTM cell, and it seems to perform just as well:

$$\begin{aligned} z_t &= \sigma(W_z \cdot [h_{t-1}, x_t]) \\ r_t &= \sigma(W_r \cdot [h_{t-1}, x_t]) \\ \hat{h}_t &= \tanh(W \cdot [r_t * h_{t-1}, x_t]) \\ h_t &= (1 - z_t) * h_{t-1} + z_t * \hat{h}_t \end{aligned}$$

Both state vectors are merged into a single vector h_t . A single gate controller z controls both the forget gate and the input gate. If the gate controller outputs a 1, the forget gate is open (1) and the input gate is closed (0). If it outputs a 0, the opposite happens. In other words, whenever a memory must be stored, the location where it will be stored is erased first. There is no output gate: the full state vector is output at every time step. However, there is a new gate controller r that controls which part of the previous state will be shown to the main layer.

Keras provides a `tf.keras.layers.GRU` layer: using it is just a matter of replacing `SimpleRNN` or `LSTM` with `GRU`. It also provides a `tf.keras.layers.GRUCell`, in case you want to create a custom cell based on a GRU cell.

```
gru_model = tf.keras.Sequential([
    tf.keras.layers.GRU(32, return_sequences=True, input_shape=[None, 2]),
    tf.keras.layers.Dense(14)
])
```

```
2023-07-11 13:38:03.533281: I tensorflow/core/common_runtime/executor.cc:1197] [/device:CPU
[[{{node gradients/split_2_grad(concat/split_2/split_dim)}}]]
2023-07-11 13:38:03.534440: I tensorflow/core/common_runtime/executor.cc:1197] [/device:CPU
[[{{node gradients/split_grad(concat/split/split_dim)}}]]
2023-07-11 13:38:03.535775: I tensorflow/core/common_runtime/executor.cc:1197] [/device:CPU
```

```
[[{{node gradients/split_1_grad(concat/split_1/split_dim)}}]]
```

```
opt = tf.keras.optimizers.SGD(learning_rate=0.02, momentum=0.9)
gru_model.compile(loss=tf.keras.losses.Huber(), optimizer=opt, metrics=["mae"])
```

Just training for 5 epochs to show that it works:

```
history = gru_model.fit(seq2seq_train, validation_data=seq2seq_valid, epochs=5,
    callbacks=[early_stopping_cb])
```

Epoch 1/5

```
2023-07-11 13:38:06.653890: I tensorflow/core/common_runtime/executor.cc:1197] [/device:CPU
[[{{node gradients/split_2_grad(concat/split_2/split_dim)}}]]]
2023-07-11 13:38:06.655169: I tensorflow/core/common_runtime/executor.cc:1197] [/device:CPU
[[{{node gradients/split_grad(concat/split/split_dim)}}]]]
2023-07-11 13:38:06.656521: I tensorflow/core/common_runtime/executor.cc:1197] [/device:CPU
[[{{node gradients/split_1_grad(concat/split_1/split_dim)}}]]]
2023-07-11 13:38:07.227025: I tensorflow/core/common_runtime/executor.cc:1197] [/device:CPU
[[{{node gradients/split_2_grad(concat/split_2/split_dim)}}]]]
2023-07-11 13:38:07.228330: I tensorflow/core/common_runtime/executor.cc:1197] [/device:CPU
[[{{node gradients/split_grad(concat/split/split_dim)}}]]]
2023-07-11 13:38:07.229759: I tensorflow/core/common_runtime/executor.cc:1197] [/device:CPU
[[{{node gradients/split_1_grad(concat/split_1/split_dim)}}]]]
```

33/Unknown - 2s 14ms/step - loss: 0.1065 - mae: 0.3859

```
2023-07-11 13:38:08.810716: I tensorflow/core/common_runtime/executor.cc:1197] [/device:CPU
[[{{node gradients/split_2_grad(concat/split_2/split_dim)}}]]]
2023-07-11 13:38:08.812765: I tensorflow/core/common_runtime/executor.cc:1197] [/device:CPU
[[{{node gradients/split_grad(concat/split/split_dim)}}]]]
2023-07-11 13:38:08.813972: I tensorflow/core/common_runtime/executor.cc:1197] [/device:CPU
[[{{node gradients/split_1_grad(concat/split_1/split_dim)}}]]]
```

```
33/33 [=====] - 3s 27ms/step - loss: 0.1065 -
mae: 0.3859 - val_loss: 0.0226 - val_mae: 0.1683
Epoch 2/5
```

```

33/33  [=====] - 1s 16ms/step - loss: 0.0187 -
mae: 0.1533 - val_loss: 0.0191 - val_mae: 0.1520
Epoch 3/5
33/33  [=====] - 1s 15ms/step - loss: 0.0172 -
mae: 0.1594 - val_loss: 0.0186 - val_mae: 0.1534
Epoch 4/5
33/33  [=====] - 1s 15ms/step - loss: 0.0171 -
mae: 0.1585 - val_loss: 0.0185 - val_mae: 0.1525
Epoch 5/5
33/33  [=====] - 1s 15ms/step - loss: 0.0170 -
mae: 0.1580 - val_loss: 0.0184 - val_mae: 0.1520

```

These are only two of the most notable LSTM variants.

LSTM and GRU cells still have a **fairly limited short-term memory**, and they have a hard time learning long-term patterns in sequences of 100 time steps or more (such as audio samples, long time series, or long sentences). One way to solve this is to shorten the input sequences, for example, using **1D convolutional layers**.

Similarly two 2D convolutional, a 1D convolutional layer slides several kernels across a sequence, producing a 1D feature map per kernel. Each kernel will learn to detect a single very short sequential pattern (no longer than the kernel size).

This means that we can build a neural network composed of a mix of recurrent layers and 1D convolutional layers (or even 1D pooling layers). By shortening the sequences the convolutional layer may help the LSTM/GRU layers detect longer patterns.

For example, the following model is the same as earlier, except it starts with a 1D convolutional layer that **downsamples the input sequence** by a factor of 2 (using a stride of 2). The kernel size is larger than the stride, so all inputs will be used to compute the layer's output, and therefore the model can learn to preserve the useful information, dropping only the unimportant details.

```

conv_rnn_model = tf.keras.Sequential([
    tf.keras.layers.Conv1D(filters=32, kernel_size=4, strides=2, activation="relu",
    ↴ input_shape=[None, 2]),
    tf.keras.layers.GRU(32, return_sequences=True),
    tf.keras.layers.Dense(14)
])

```

```

2023-07-11 13:49:54.521924: I tensorflow/core/common_runtime/executor.cc:1197] [/device:CPU
[[{{node gradients/split_2_grad(concat/split_2/split_dim)}}]]

```

```
2023-07-11 13:49:54.523413: I tensorflow/core/common_runtime/executor.cc:1197] [/device:CPU
[[{{node gradients/split_grad(concat/split/split_dim)}}]]
2023-07-11 13:49:54.524512: I tensorflow/core/common_runtime/executor.cc:1197] [/device:CPU
[[{{node gradients/split_1_grad(concat/split_1/split_dim)}}]]
```

Notice that we must also crop off the first three time steps in the targets: indeed, the kernel's size is 4, so the first output of the convolutional layer will be based on the input time steps 0 to 3, and the first forecasts will be for time steps 4 to 17 (instead of time steps 1 to 14). Moreover, we must downsample the targets by a factor of 2, because of the stride:

```
longer_train = to_seq2seq_dataset(mulvar_train, seq_length=112)
longer_valid = to_seq2seq_dataset(mulvar_valid, seq_length=112)

downsampled_train = longer_train.map(lambda X, Y: (X, Y[:, 3::2]))
downsampled_valid = longer_valid.map(lambda X, Y: (X, Y[:, 3::2]))
```

Just training for 5 epochs to show that it works:

```
opt = tf.keras.optimizers.SGD(learning_rate=0.02, momentum=0.9)
conv_rnn_model.compile(loss=tf.keras.losses.Huber(), optimizer=opt, metrics=["mae"])

history = conv_rnn_model.fit(downsampled_train, validation_data=downsampled_valid,
                           epochs=5, callbacks=[early_stopping_cb])
```

Epoch 1/5

```
2023-07-11 13:50:16.688627: I tensorflow/core/common_runtime/executor.cc:1197] [/device:CPU
[[{{node Placeholder/_0}}]]
2023-07-11 13:50:16.688967: I tensorflow/core/common_runtime/executor.cc:1197] [/device:CPU
[[{{node Placeholder/_0}}]]
2023-07-11 13:50:16.935736: I tensorflow/core/common_runtime/executor.cc:1197] [/device:CPU
[[{{node gradients/split_2_grad(concat/split_2/split_dim)}}]]
2023-07-11 13:50:16.937288: I tensorflow/core/common_runtime/executor.cc:1197] [/device:CPU
[[{{node gradients/split_grad(concat/split/split_dim)}}]]
2023-07-11 13:50:16.939063: I tensorflow/core/common_runtime/executor.cc:1197] [/device:CPU
[[{{node gradients/split_1_grad(concat/split_1/split_dim)}}]]
2023-07-11 13:50:17.571753: I tensorflow/core/common_runtime/executor.cc:1197] [/device:CPU
[[{{node gradients/split_2_grad(concat/split_2/split_dim)}}]]
```

```

2023-07-11 13:50:17.573390: I tensorflow/core/common_runtime/executor.cc:1197] [/device:CPU
  [[{{node gradients/split_grad(concat/split/split_dim)}}]]
2023-07-11 13:50:17.574525: I tensorflow/core/common_runtime/executor.cc:1197] [/device:CPU
  [[{{node gradients/split_1_grad(concat/split_1/split_dim)}}]]]

29/Unknown - 2s 15ms/step - loss: 0.1116 - mae: 0.4002

2023-07-11 13:50:19.066413: I tensorflow/core/common_runtime/executor.cc:1197] [/device:CPU
  [[{{node Placeholder/_0}}]]]
2023-07-11 13:50:19.066746: I tensorflow/core/common_runtime/executor.cc:1197] [/device:CPU
  [[{{node Placeholder/_0}}]]]
2023-07-11 13:50:19.293512: I tensorflow/core/common_runtime/executor.cc:1197] [/device:CPU
  [[{{node gradients/split_2_grad(concat/split_2/split_dim)}}]]]
2023-07-11 13:50:19.295140: I tensorflow/core/common_runtime/executor.cc:1197] [/device:CPU
  [[{{node gradients/split_grad(concat/split/split_dim)}}]]]
2023-07-11 13:50:19.296585: I tensorflow/core/common_runtime/executor.cc:1197] [/device:CPU
  [[{{node gradients/split_1_grad(concat/split_1/split_dim)}}]]]

31/31 [=====] - 3s 31ms/step - loss: 0.1076 -
mae: 0.3904 - val_loss: 0.0214 - val_mae: 0.1614
Epoch 2/5
31/31 [=====] - 1s 17ms/step - loss: 0.0184 -
mae: 0.1495 - val_loss: 0.0181 - val_mae: 0.1595
Epoch 3/5
31/31 [=====] - 1s 17ms/step - loss: 0.0166 -
mae: 0.1582 - val_loss: 0.0181 - val_mae: 0.1553
Epoch 4/5
31/31 [=====] - 1s 17ms/step - loss: 0.0164 -
mae: 0.1554 - val_loss: 0.0179 - val_mae: 0.1551
Epoch 5/5
31/31 [=====] - 1s 19ms/step - loss: 0.0163 -
mae: 0.1551 - val_loss: 0.0178 - val_mae: 0.1546

```

If we train and evaluate this model, we will find that it outperforms the previous model (by a small margin). In fact, it is actually possible to use only 1D convolutional layers and drop the recurrent layers entirely! This is the **WaveNet architecture**. It stacked 1D convolutional layers, doubling the dilation rate (how spread apart each neuron's inputs are) at every layer: the first convolutional layer gets a glimpse of just two time steps at a time, while the next one sees four

time steps (its receptive field is four time steps long), the next one sees eight time steps, and so on.

This way, the lower layers learn short-term patterns, while the higher layers learn long-term patterns. Thanks to the doubling dilation rate, the network can process extremely large sequences very efficiently. The architecture also left-pads the input sequences with a number of zeros equal to the dilation rate before every layer, to preserve the same sequence length throughout the network.

Here is how to implement a simplified WaveNet to tackle the same sequences as earlier:

```
wavenet_model = tf.keras.Sequential()

# the model starts with an explicit input layer
wavenet_model.add(tf.keras.layers.InputLayer(input_shape=[None, 2]))

# then it continues with 1D convolutional layers using
# growing dilation rates: 1, 2, 4, 8, two times
# Thanks to the causal padding, every convolutional layer outputs a sequence of the
# same length as its input sequence, so the targets we use during training can
# be the full sequences: no need to crop them or downsample them.

for rate in (1, 2, 4, 8) * 2:
    wavenet_model.add(tf.keras.layers.Conv1D(filters=32, kernel_size=2,
    ↳ padding="causal",
                           activation="relu", dilation_rate=rate))

# finally, we add the output layer: a convolutional layer with 14 filters of size 1
# and
# without any activation function.
wavenet_model.add(tf.keras.layers.Conv1D(filters=14, kernel_size=1))
```

Just training for 5 epochs to show that it works:

```
opt = tf.keras.optimizers.SGD(learning_rate=0.02, momentum=0.9)
wavenet_model.compile(loss=tf.keras.losses.Huber(), optimizer=opt, metrics=["mae"])

history = wavenet_model.fit(longer_train, validation_data=longer_valid, epochs=5,
    ↳ callbacks=[early_stopping_cb])
```

Epoch 1/5

```
2023-07-11 14:08:41.717659: I tensorflow/core/common_runtime/executor.cc:1197] [/device:CPU
  [[{{node Placeholder/_0}}]]
2023-07-11 14:08:41.717984: I tensorflow/core/common_runtime/executor.cc:1197] [/device:CPU
  [[{{node Placeholder/_0}}]]]

31/Unknown - 2s 15ms/step - loss: 0.1528 - mae: 0.5097

2023-07-11 14:08:43.627238: I tensorflow/core/common_runtime/executor.cc:1197] [/device:CPU
  [[{{node Placeholder/_0}}]]
2023-07-11 14:08:43.627628: I tensorflow/core/common_runtime/executor.cc:1197] [/device:CPU
  [[{{node Placeholder/_0}}]]]

31/31  [=====] - 2s 31ms/step - loss: 0.1528 -
mae: 0.5097 - val_loss: 0.0578 - val_mae: 0.2978
Epoch 2/5
31/31  [=====] - 1s 18ms/step - loss: 0.0391 -
mae: 0.2358 - val_loss: 0.0213 - val_mae: 0.1669
Epoch 3/5
31/31  [=====] - 1s 18ms/step - loss: 0.0176 -
mae: 0.1623 - val_loss: 0.0186 - val_mae: 0.1524
Epoch 4/5
31/31  [=====] - 1s 18ms/step - loss: 0.0161 -
mae: 0.1574 - val_loss: 0.0186 - val_mae: 0.1508
Epoch 5/5
31/31  [=====] - 1s 22ms/step - loss: 0.0160 -
mae: 0.1564 - val_loss: 0.0186 - val_mae: 0.1505
```

The models discussed offer similar performance for the simple ridership forecasting task, but they may vary significantly depending on the task and the amount of available data.

1.7 Exercise

1 - Download the Bach chorales dataset and unzip it. It is composed of 382 chorales composed by Johann Sebastian Bach. Each chorale is 100 to 640 time steps long, and each time step contains 4 integers, where each integer corresponds to a note's index on a piano (except for the value 0, which means that no note is played).

```
url =
    "https://www.dropbox.com/scl/fi/naevqylv41x21zdiol1rq/jsb_chorales.zip?rlkey=zhsumvz5mfmezu1vs5"
    # dl=1 is important
u = urllib.request.urlopen(url)
data = u.read()
u.close()
with open("./data/jsb_chorales.zip", "wb") as f :
    f.write(data)

import zipfile

with zipfile.ZipFile("./data/jsb_chorales.zip", "r") as zip_ref:
    zip_ref.extractall("./data")

from pathlib import Path
jsb_chorales_dir = Path("data/jsb_chorales")

train_files = sorted(jsb_chorales_dir.glob("train/chorale_*.csv"))
valid_files = sorted(jsb_chorales_dir.glob("valid/chorale_*.csv"))
test_files = sorted(jsb_chorales_dir.glob("test/chorale_*.csv"))

def load_chorales(filepaths):
    return [pd.read_csv(filepath).values.tolist() for filepath in filepaths]

train_chorales = load_chorales(train_files)
valid_chorales = load_chorales(valid_files)
test_chorales = load_chorales(test_files)

train_chorales[0]

[[74, 70, 65, 58],
 [74, 70, 65, 58],
 [74, 70, 65, 58],
 [74, 70, 65, 58],
 [75, 70, 58, 55],
 [75, 70, 58, 55],
 [75, 70, 60, 55],
 [75, 70, 60, 55],
```

```
[77, 69, 62, 50],  
[77, 69, 62, 50],  
[77, 69, 62, 50],  
[77, 69, 62, 50],  
[77, 70, 62, 55],  
[77, 70, 62, 55],  
[77, 69, 62, 55],  
[77, 69, 62, 55],  
[75, 67, 63, 48],  
[75, 67, 63, 48],  
[75, 69, 63, 48],  
[75, 69, 63, 48],  
[74, 70, 65, 46],  
[74, 70, 65, 46],  
[74, 70, 65, 46],  
[72, 69, 65, 53],  
[72, 69, 65, 53],  
[72, 69, 65, 53],  
[72, 69, 65, 53],  
[72, 69, 65, 53],  
[72, 69, 65, 53],  
[72, 69, 65, 53],  
[74, 70, 65, 46],  
[74, 70, 65, 46],  
[74, 70, 65, 46],  
[74, 70, 65, 46],  
[75, 69, 63, 48],  
[75, 69, 63, 48],  
[75, 67, 63, 48],  
[75, 67, 63, 48],  
[77, 65, 62, 50],  
[77, 65, 62, 50],  
[77, 65, 60, 50],  
[77, 65, 60, 50],  
[74, 67, 58, 55],  
[74, 67, 58, 55],
```

```
[74, 67, 58, 53],  
[74, 67, 58, 53],  
[72, 67, 58, 51],  
[72, 67, 58, 51],  
[72, 67, 58, 51],  
[72, 67, 58, 51],  
[72, 65, 57, 53],  
[72, 65, 57, 53],  
[72, 65, 57, 53],  
[72, 65, 57, 53],  
[70, 65, 62, 46],  
[70, 65, 62, 46],  
[70, 65, 62, 46],  
[70, 65, 62, 46],  
[70, 65, 62, 46],  
[70, 65, 62, 46],  
[70, 65, 62, 46],  
[70, 65, 62, 46],  
[70, 65, 62, 46],  
[70, 69, 65, 53],  
[72, 69, 65, 53],  
[72, 69, 65, 53],  
[72, 69, 65, 53],  
[74, 71, 53, 50],  
[74, 71, 53, 50],  
[74, 71, 53, 50],  
[74, 71, 53, 50],  
[75, 72, 55, 48],  
[75, 72, 55, 48],  
[75, 72, 55, 50],  
[75, 72, 55, 50],  
[75, 67, 60, 51],  
[75, 67, 60, 51],  
[75, 67, 60, 53],  
[75, 67, 60, 53],  
[74, 67, 60, 55],  
[74, 67, 60, 55],  
[74, 67, 57, 55],  
[74, 67, 57, 55],
```

```
[74, 65, 59, 43],  
[74, 65, 59, 43],  
[72, 63, 59, 43],  
[72, 63, 59, 43],  
[72, 63, 55, 48],  
[72, 63, 55, 48],  
[72, 63, 55, 48],  
[72, 63, 55, 48],  
[72, 63, 55, 48],  
[72, 63, 55, 48],  
[72, 63, 55, 48],  
[72, 63, 55, 48],  
[72, 63, 55, 48],  
[72, 63, 55, 48],  
[72, 63, 55, 48],  
[72, 63, 55, 48],  
[72, 63, 55, 48],  
[72, 63, 55, 48],  
[72, 63, 55, 48],  
[72, 63, 55, 48],  
[75, 67, 60, 60],  
[75, 67, 60, 60],  
[75, 67, 60, 60],  
[75, 67, 60, 60],  
[77, 70, 62, 58],  
[77, 70, 62, 58],  
[77, 70, 62, 56],  
[77, 70, 62, 56],  
[79, 70, 62, 55],  
[79, 70, 62, 55],  
[79, 70, 62, 53],  
[79, 70, 62, 53],  
[79, 70, 63, 51],  
[79, 70, 63, 51],  
[79, 70, 63, 51],  
[79, 70, 63, 51],  
[77, 70, 63, 58],  
[77, 70, 63, 58],  
[77, 70, 60, 58],  
[77, 70, 60, 58],  
[77, 70, 62, 46],  
[77, 70, 62, 46],  
[77, 68, 62, 46],  
[75, 68, 62, 46],  
[75, 67, 58, 51],  
[75, 67, 58, 51],
```



```
[74, 65, 58, 58],  
[74, 65, 58, 58],  
[74, 65, 58, 58],  
[74, 65, 58, 58],  
[75, 67, 58, 57],  
[75, 67, 58, 57],  
[75, 67, 58, 55],  
[75, 67, 58, 55],  
[77, 65, 60, 57],  
[77, 65, 60, 57],  
[77, 65, 60, 53],  
[77, 65, 60, 53],  
[74, 65, 58, 58],  
[74, 65, 58, 58],  
[74, 65, 58, 58],  
[72, 67, 58, 51],  
[72, 67, 58, 51],  
[72, 67, 58, 51],  
[72, 67, 58, 51],  
[72, 65, 57, 53],  
[72, 65, 57, 53],  
[72, 65, 57, 53],  
[72, 65, 57, 53],  
[70, 65, 62, 46],  
[70, 65, 62, 46],  
[70, 65, 62, 46],  
[70, 65, 62, 46],  
[70, 65, 62, 46],  
[70, 65, 62, 46],  
[70, 65, 62, 46],  
[70, 65, 62, 46]]
```

```
notes = set()  
for chorales in (train_chorales, valid_chorales, test_chorales):  
    for chorale in chorales:  
        for chord in chorale:  
            notes |= set(chord)
```

```

n_notes = len(notes)
min_note = min(notes - {0})
max_note = max(notes)

assert min_note == 36
assert max_note == 81

```

Notice that values range from 36 (C1 = C on octave 1) to 81 (A5 = A on octave 5), plus 0 for silence. Let's write some functions to listen to chorales:

```

def notes_to_frequencies(notes):
    # Frequency doubles when you go up one octave; there are 12 semi-tones
    # per octave; Note A on octave 4 is 440 Hz, and it is note number 69.
    return 2 ** ((np.array(notes) - 69) / 12) * 440

```

```

def frequencies_to_samples(frequencies, tempo, sample_rate):
    note_duration = 60 / tempo # the tempo is measured in beats per minutes
    # To reduce click sound at every beat, we round the frequencies to try to
    # get the samples close to zero at the end of each note.
    frequencies = (note_duration * frequencies).round() / note_duration
    n_samples = int(note_duration * sample_rate)
    time = np.linspace(0, note_duration, n_samples)
    sine_waves = np.sin(2 * np.pi * frequencies.reshape(-1, 1) * time)
    # Removing all notes with frequencies ≤ 9 Hz (includes note 0 = silence)
    sine_waves *= (frequencies > 9.).reshape(-1, 1)
    return sine_waves.reshape(-1)

```

```

def chords_to_samples(chords, tempo, sample_rate):
    freqs = notes_to_frequencies(chords)
    freqs = np.r_[freqs, freqs[-1:]] # make last note a bit longer
    merged = np.mean([frequencies_to_samples(melody, tempo, sample_rate)
                     for melody in freqs.T], axis=0)
    n_fade_out_samples = sample_rate * 60 // tempo # fade out last note
    fade_out = np.linspace(1., 0., n_fade_out_samples)**2
    merged[-n_fade_out_samples:] *= fade_out
    return merged

```

```
from IPython.display import Audio

def play_chords(chords, tempo=160, amplitude=0.1, sample_rate=44100, filepath=None):
    samples = amplitude * chords_to_samples(chords, tempo, sample_rate)
    if filepath:
        from scipy.io import wavfile
        samples = (2**15 * samples).astype(np.int16)
        wavfile.write(filepath, sample_rate, samples)
        return display(Audio(filepath))
    else:
        return display(Audio(samples, rate=sample_rate))
```

Now let's listen to a few chorales:

```
for index in range(3):
    play_chords(train_chorales[index])
```

<IPython.lib.display.Audio object>

<IPython.lib.display.Audio object>

<IPython.lib.display.Audio object>

2 - Train a model (recurrent, convolutional, or both) that can predict the next time step (four notes), given a sequence of time steps from a chorale.

We need to train a model that can predict the next chord given all the previous chords. If we try to predict the next chord in one shot (all 4 notes at once), we get notes that don't go very well together (believe me, I tried). It's much better (and simpler) to predict one note at a time. So we will need to preprocess every chorale, turning each chord into an arpeggio (i.e., a sequence of notes rather than notes played simultaneously). So each chorale will be a long sequence of notes (rather than chords), and we can just train a model that can predict the next note given all the previous notes. We will also shift the values so that they range from 0 to 46, where 0 represents silence, and values 1 to 46 represent notes 36 (C1) to 81 (A5). And we will train the model on windows of 128 notes (i.e., 32 chords).

In this simple case, the dataset fits in memory, so we could preprocess the chorales in RAM using any Python code we like, but we will demonstrate here how to do all the preprocessing using tf.data, in order to be able to manage also dataset that doesn't fit in memory.

```

def create_target(batch):
    X = batch[:, :-1]
    Y = batch[:, 1:] # predict next note in each arpeggio, at each step
    return X, Y

def preprocess(window):
    window = tf.where(window == 0, window, window - min_note + 1) # shift values
    return tf.reshape(window, [-1]) # convert to arpeggio

def batch_dataset(chorales, batch_size=32, shuffle_buffer_size=None,
                  window_size=32, window_shift=16, cache=True):
    def batch_window(window):
        return window.batch(window_size + 1)

    def to_windows(chorale):
        dataset = tf.data.Dataset.from_tensor_slices(chorale)
        dataset = dataset.window(window_size + 1, window_shift, drop_remainder=True)
        return dataset.flat_map(batch_window)

    chorales = tf.ragged.constant(chorales, ragged_rank=1)
    dataset = tf.data.Dataset.from_tensor_slices(chorales)
    dataset = dataset.flat_map(to_windows).map(preprocess)
    if cache:
        dataset = dataset.cache()
    if shuffle_buffer_size:
        dataset = dataset.shuffle(shuffle_buffer_size)
    dataset = dataset.batch(batch_size)
    dataset = dataset.map(create_target)
    return dataset.prefetch(1)

```

Now let's create the training set, the validation set and the test set:

```

train_set = batch_dataset(train_chorales, shuffle_buffer_size=1000)
valid_set = batch_dataset(valid_chorales)
test_set = batch_dataset(test_chorales)

```

We will use a sequence-to-sequence approach, where we feed a window to the neural net, and it tries to predict that same window shifted one time step into the future. We could feed the note values directly to the model, as floats, but this would probably not give good results. The relationships between notes are not that simple: for example, if you replace a C3 with a C4, the

melody will still sound fine, even though these notes are 12 semi-tones apart (i.e., one octave). Conversely, if you replace a C3 with a C#3, it's very likely that the chord will sound horrible, despite these notes being just next to each other. So we will use an Embedding layer to convert each note to a small vector representation (we will see embeddings in another notebook). We will use 5-dimensional embeddings, so the output of this first layer will have a shape of [batch_size, window_size, 5]. We will then feed this data to a small WaveNet-like neural network, composed of a stack of 4 Conv1D layers with doubling dilation rates. We will intersperse these layers with BatchNormalization layers for faster better convergence. Then one LSTM layer to try to capture long-term patterns. And finally a Dense layer to produce the final note probabilities. It will predict one probability for each chorale in the batch, for each time step, and for each possible note (including silence). So the output shape will be [batch_size, window_size, 47].

```
n_embedding_dims = 5

model = tf.keras.Sequential([
    tf.keras.layers.Embedding(input_dim=n_notes, output_dim=n_embedding_dims,
                              input_shape=[None]),
    tf.keras.layers.Conv1D(32, kernel_size=2, padding="causal", activation="relu"),
    tf.keras.layers.BatchNormalization(),
    tf.keras.layers.Conv1D(48, kernel_size=2, padding="causal", activation="relu",
    ↵ dilation_rate=2),
    tf.keras.layers.BatchNormalization(),
    tf.keras.layers.Conv1D(64, kernel_size=2, padding="causal", activation="relu",
    ↵ dilation_rate=4),
    tf.keras.layers.BatchNormalization(),
    tf.keras.layers.Conv1D(96, kernel_size=2, padding="causal", activation="relu",
    ↵ dilation_rate=8),
    tf.keras.layers.BatchNormalization(),
    tf.keras.layers.LSTM(256, return_sequences=True),
    tf.keras.layers.Dense(n_notes, activation="softmax")
])

model.summary()
```

Model: "sequential"

Layer (type)	Output Shape	Param #
embedding (Embedding)	(None, None, 5)	235

conv1d (Conv1D)	(None, None, 32)	352
batch_normalization (BatchN ormalization)	(None, None, 32)	128
conv1d_1 (Conv1D)	(None, None, 48)	3120
batch_normalization_1 (BatchNormalization)	(None, None, 48)	192
conv1d_2 (Conv1D)	(None, None, 64)	6208
batch_normalization_2 (BatchNormalization)	(None, None, 64)	256
conv1d_3 (Conv1D)	(None, None, 96)	12384
batch_normalization_3 (BatchNormalization)	(None, None, 96)	384
lstm (LSTM)	(None, None, 256)	361472
dense (Dense)	(None, None, 47)	12079

Total params: 396,810

Trainable params: 396,330

Non-trainable params: 480

```
2023-07-14 13:00:24.617989: I tensorflow/core/common_runtime/executor.cc:1197] [/device:CPU
[[{{node gradients/split_2_grad(concat/split_2/split_dim)}}]]
2023-07-14 13:00:24.619568: I tensorflow/core/common_runtime/executor.cc:1197] [/device:CPU
[[{{node gradients/split_grad(concat/split/split_dim)}}]]
2023-07-14 13:00:24.620629: I tensorflow/core/common_runtime/executor.cc:1197] [/device:CPU
[[{{node gradients/split_1_grad(concat/split_1/split_dim)}}]]
```

Now we're ready to compile and train the model!

```
optimizer = tf.keras.optimizers.Nadam(learning_rate=1e-3)
model.compile(loss="sparse_categorical_crossentropy", optimizer=optimizer,
              metrics=["accuracy"])

model.fit(train_set, epochs=20, validation_data=valid_set)
```

Epoch 1/20

```
2023-07-14 13:00:31.782884: I tensorflow/core/common_runtime/executor.cc:1197] [/device:CPU
[[{{node Placeholder/_0}}]]
2023-07-14 13:00:31.783675: I tensorflow/core/common_runtime/executor.cc:1197] [/device:CPU
[[{{node Placeholder/_0}}]]
2023-07-14 13:00:32.282459: I tensorflow/core/common_runtime/executor.cc:1197] [/device:CPU
[[{{node gradients/split_2_grad(concat/split_2/split_dim)}}]]
2023-07-14 13:00:32.283612: I tensorflow/core/common_runtime/executor.cc:1197] [/device:CPU
[[{{node gradients/split_grad(concat/split/split_dim)}}]]
2023-07-14 13:00:32.284843: I tensorflow/core/common_runtime/executor.cc:1197] [/device:CPU
[[{{node gradients/split_1_grad(concat/split_1/split_dim)}}]]
2023-07-14 13:00:34.320783: I tensorflow/core/common_runtime/executor.cc:1197] [/device:CPU
[[{{node gradients/split_2_grad(concat/split_2/split_dim)}}]]
2023-07-14 13:00:34.324159: I tensorflow/core/common_runtime/executor.cc:1197] [/device:CPU
[[{{node gradients/split_grad(concat/split/split_dim)}}]]
2023-07-14 13:00:34.328429: I tensorflow/core/common_runtime/executor.cc:1197] [/device:CPU
[[{{node gradients/split_1_grad(concat/split_1/split_dim)}}]]
```

98/Unknown - 51s 460ms/step - loss: 1.8791 - accuracy: 0.5281

```
2023-07-14 13:01:22.455294: I tensorflow/core/common_runtime/executor.cc:1197] [/device:CPU
[[{{node Placeholder/_0}}]]
2023-07-14 13:01:22.837103: I tensorflow/core/common_runtime/executor.cc:1197] [/device:CPU
[[{{node gradients/split_2_grad(concat/split_2/split_dim)}}]]
2023-07-14 13:01:22.838266: I tensorflow/core/common_runtime/executor.cc:1197] [/device:CPU
[[{{node gradients/split_grad(concat/split/split_dim)}}]]
2023-07-14 13:01:22.839505: I tensorflow/core/common_runtime/executor.cc:1197] [/device:CPU
[[{{node gradients/split_1_grad(concat/split_1/split_dim)}}]]
```

98/98 [=====] - 61s 572ms/step - loss: 1.8791 -
accuracy: 0.5281 - val_loss: 3.5540 - val_accuracy: 0.0809

```
Epoch 2/20
98/98 [=====] - 80s 821ms/step - loss: 0.8585 -
accuracy: 0.7732 - val_loss: 3.7054 - val_accuracy: 0.0826
Epoch 3/20
98/98 [=====] - 62s 631ms/step - loss: 0.7213 -
accuracy: 0.7984 - val_loss: 3.5234 - val_accuracy: 0.1417
Epoch 4/20
98/98 [=====] - 173s 2s/step - loss: 0.6518 -
accuracy: 0.8131 - val_loss: 3.0466 - val_accuracy: 0.2125
Epoch 5/20
98/98 [=====] - 67s 688ms/step - loss: 0.6012 -
accuracy: 0.8238 - val_loss: 2.0218 - val_accuracy: 0.4137
Epoch 6/20
98/98 [=====] - 2073s 21s/step - loss: 0.5624 -
accuracy: 0.8325 - val_loss: 0.9131 - val_accuracy: 0.7403
Epoch 7/20
98/98 [=====] - 62s 634ms/step - loss: 0.5283 -
accuracy: 0.8410 - val_loss: 0.6834 - val_accuracy: 0.8016
Epoch 8/20
98/98 [=====] - 52s 534ms/step - loss: 0.4972 -
accuracy: 0.8490 - val_loss: 0.6265 - val_accuracy: 0.8183
Epoch 9/20
98/98 [=====] - 54s 550ms/step - loss: 0.4690 -
accuracy: 0.8563 - val_loss: 0.5987 - val_accuracy: 0.8247
Epoch 10/20
98/98 [=====] - 54s 556ms/step - loss: 0.4436 -
accuracy: 0.8636 - val_loss: 0.6081 - val_accuracy: 0.8222
Epoch 11/20
98/98 [=====] - 56s 568ms/step - loss: 0.4198 -
accuracy: 0.8702 - val_loss: 0.6039 - val_accuracy: 0.8244
Epoch 12/20
98/98 [=====] - 45s 463ms/step - loss: 0.3968 -
accuracy: 0.8775 - val_loss: 0.6061 - val_accuracy: 0.8223
Epoch 13/20
98/98 [=====] - 48s 489ms/step - loss: 0.3789 -
accuracy: 0.8823 - val_loss: 0.6176 - val_accuracy: 0.8212
Epoch 14/20
98/98 [=====] - 55s 558ms/step - loss: 0.3615 -
```

```
accuracy: 0.8877 - val_loss: 0.6122 - val_accuracy: 0.8237
Epoch 15/20
98/98 [=====] - 56s 566ms/step - loss: 0.3410 -
accuracy: 0.8934 - val_loss: 0.6133 - val_accuracy: 0.8230
Epoch 16/20
98/98 [=====] - 48s 488ms/step - loss: 0.3228 -
accuracy: 0.8995 - val_loss: 0.6269 - val_accuracy: 0.8217
Epoch 17/20
98/98 [=====] - 49s 505ms/step - loss: 0.3055 -
accuracy: 0.9047 - val_loss: 0.6305 - val_accuracy: 0.8189
Epoch 18/20
98/98 [=====] - 49s 503ms/step - loss: 0.2902 -
accuracy: 0.9096 - val_loss: 0.6416 - val_accuracy: 0.8214
Epoch 19/20
98/98 [=====] - 53s 544ms/step - loss: 0.2806 -
accuracy: 0.9124 - val_loss: 0.6492 - val_accuracy: 0.8167
Epoch 20/20
98/98 [=====] - 44s 448ms/step - loss: 0.2634 -
accuracy: 0.9184 - val_loss: 0.6530 - val_accuracy: 0.8193

<keras.callbacks.History at 0x14e089e50>
```

We have not done much hyperparameter search, so feel free to iterate on this model and try to optimize it. For example, we could try removing the LSTM layer and replacing it with Conv1D layers. You could also play with the number of layers, the learning rate, the optimizer, and so on. Once we're satisfied with the performance of the model on the validation set, we can save it and evaluate it one last time on the test set:

```
model.save("./data/bach_model", save_format="tf")
```

```
2023-07-14 13:57:25.172447: I tensorflow/core/common_runtime/executor.cc:1197] [/device:CPU
[[{{node gradients/split_2_grad(concat/split_2/split_dim)}}]]
2023-07-14 13:57:25.174101: I tensorflow/core/common_runtime/executor.cc:1197] [/device:CPU
[[{{node gradients/split_grad(concat/split/split_dim)}}]]
2023-07-14 13:57:25.175874: I tensorflow/core/common_runtime/executor.cc:1197] [/device:CPU
[[{{node gradients/split_1_grad(concat/split_1/split_dim)}}]]
2023-07-14 13:57:25.573083: I tensorflow/core/common_runtime/executor.cc:1197] [/device:CPU
[[{{node gradients/split_2_grad(concat/split_2/split_dim)}}]]
```

```
2023-07-14 13:57:25.574165: I tensorflow/core/common_runtime/executor.cc:1197] [/device:CPU
  [[{{node gradients/split_grad(concat/split/split_dim)}}]]
2023-07-14 13:57:25.575347: I tensorflow/core/common_runtime/executor.cc:1197] [/device:CPU
  [[{{node gradients/split_1_grad(concat/split_1/split_dim)}}]]
2023-07-14 13:57:25.782031: I tensorflow/core/common_runtime/executor.cc:1197] [/device:CPU
  [[{{node gradients/split_2_grad(concat/split_2/split_dim)}}]]
2023-07-14 13:57:25.783715: I tensorflow/core/common_runtime/executor.cc:1197] [/device:CPU
  [[{{node gradients/split_grad(concat/split/split_dim)}}]]
2023-07-14 13:57:25.784995: I tensorflow/core/common_runtime/executor.cc:1197] [/device:CPU
  [[{{node gradients/split_1_grad(concat/split_1/split_dim)}}]]
2023-07-14 13:57:25.813884: I tensorflow/core/common_runtime/executor.cc:1197] [/device:CPU
  [[{{node inputs}}]]
2023-07-14 13:57:25.823012: I tensorflow/core/common_runtime/executor.cc:1197] [/device:CPU
  [[{{node inputs}}]]
2023-07-14 13:57:26.169080: I tensorflow/core/common_runtime/executor.cc:1197] [/device:CPU
  [[{{node gradients/split_2_grad(concat/split_2/split_dim)}}]]
2023-07-14 13:57:26.170428: I tensorflow/core/common_runtime/executor.cc:1197] [/device:CPU
  [[{{node gradients/split_grad(concat/split/split_dim)}}]]
2023-07-14 13:57:26.171729: I tensorflow/core/common_runtime/executor.cc:1197] [/device:CPU
  [[{{node gradients/split_1_grad(concat/split_1/split_dim)}}]]
2023-07-14 13:57:26.413228: I tensorflow/core/common_runtime/executor.cc:1197] [/device:CPU
  [[{{node gradients/split_2_grad(concat/split_2/split_dim)}}]]
2023-07-14 13:57:26.414815: I tensorflow/core/common_runtime/executor.cc:1197] [/device:CPU
  [[{{node gradients/split_grad(concat/split/split_dim)}}]]
2023-07-14 13:57:26.416025: I tensorflow/core/common_runtime/executor.cc:1197] [/device:CPU
  [[{{node gradients/split_1_grad(concat/split_1/split_dim)}}]]
2023-07-14 13:57:26.607805: I tensorflow/core/common_runtime/executor.cc:1197] [/device:CPU
  [[{{node inputs}}]]
2023-07-14 13:57:26.614898: I tensorflow/core/common_runtime/executor.cc:1197] [/device:CPU
  [[{{node inputs}}]]
2023-07-14 13:57:26.622972: I tensorflow/core/common_runtime/executor.cc:1197] [/device:CPU
  [[{{node inputs}}]]
2023-07-14 13:57:26.632505: I tensorflow/core/common_runtime/executor.cc:1197] [/device:CPU
  [[{{node inputs}}]]
2023-07-14 13:57:26.641252: I tensorflow/core/common_runtime/executor.cc:1197] [/device:CPU
  [[{{node inputs}}]]
2023-07-14 13:57:26.646633: I tensorflow/core/common_runtime/executor.cc:1197] [/device:CPU
  [[{{node inputs}}]]
```

```
2023-07-14 13:57:27.146544: I tensorflow/core/common_runtime/executor.cc:1197] [/device:CPU
  [[{{node gradients/split_2_grad(concat/split_2/split_dim)}}]]
2023-07-14 13:57:27.147910: I tensorflow/core/common_runtime/executor.cc:1197] [/device:CPU
  [[{{node gradients/split_grad(concat/split/split_dim)}}]]
2023-07-14 13:57:27.149916: I tensorflow/core/common_runtime/executor.cc:1197] [/device:CPU
  [[{{node gradients/split_1_grad(concat/split_1/split_dim)}}]]
2023-07-14 13:57:27.516335: I tensorflow/core/common_runtime/executor.cc:1197] [/device:CPU
  [[{{node gradients/split_2_grad(concat/split_2/split_dim)}}]]
2023-07-14 13:57:27.517640: I tensorflow/core/common_runtime/executor.cc:1197] [/device:CPU
  [[{{node gradients/split_grad(concat/split/split_dim)}}]]
2023-07-14 13:57:27.519216: I tensorflow/core/common_runtime/executor.cc:1197] [/device:CPU
  [[{{node gradients/split_1_grad(concat/split_1/split_dim)}}]]
2023-07-14 13:57:27.650727: I tensorflow/core/common_runtime/executor.cc:1197] [/device:CPU
  [[{{node inputs}}]]
2023-07-14 13:57:27.724306: I tensorflow/core/common_runtime/executor.cc:1197] [/device:CPU
  [[{{node inputs}}]]
2023-07-14 13:57:27.806227: I tensorflow/core/common_runtime/executor.cc:1197] [/device:CPU
  [[{{node inputs}}]]
2023-07-14 13:57:27.881234: I tensorflow/core/common_runtime/executor.cc:1197] [/device:CPU
  [[{{node inputs_0}}]]
2023-07-14 13:57:27.888752: I tensorflow/core/common_runtime/executor.cc:1197] [/device:CPU
  [[{{node inputs_0}}]]
2023-07-14 13:57:27.895562: I tensorflow/core/common_runtime/executor.cc:1197] [/device:CPU
  [[{{node inputs}}]]
2023-07-14 13:57:27.903031: I tensorflow/core/common_runtime/executor.cc:1197] [/device:CPU
  [[{{node inputs}}]]
2023-07-14 13:57:28.062253: I tensorflow/core/common_runtime/executor.cc:1197] [/device:CPU
  [[{{node gradients/split_2_grad(concat/split_2/split_dim)}}]]
2023-07-14 13:57:28.063216: I tensorflow/core/common_runtime/executor.cc:1197] [/device:CPU
  [[{{node gradients/split_grad(concat/split/split_dim)}}]]
2023-07-14 13:57:28.064736: I tensorflow/core/common_runtime/executor.cc:1197] [/device:CPU
  [[{{node gradients/split_1_grad(concat/split_1/split_dim)}}]]
2023-07-14 13:57:28.547159: I tensorflow/core/common_runtime/executor.cc:1197] [/device:CPU
  [[{{node gradients/split_2_grad(concat/split_2/split_dim)}}]]
2023-07-14 13:57:28.548343: I tensorflow/core/common_runtime/executor.cc:1197] [/device:CPU
  [[{{node gradients/split_grad(concat/split/split_dim)}}]]
2023-07-14 13:57:28.549688: I tensorflow/core/common_runtime/executor.cc:1197] [/device:CPU
  [[{{node gradients/split_1_grad(concat/split_1/split_dim)}}]]
```

```
2023-07-14 13:57:28.719933: I tensorflow/core/common_runtime/executor.cc:1197] [/device:CPU
  [[{{node gradients/split_2_grad(concat/split_2/split_dim)}}]]
2023-07-14 13:57:28.721093: I tensorflow/core/common_runtime/executor.cc:1197] [/device:CPU
  [[{{node gradients/split_grad(concat/split/split_dim)}}]]
2023-07-14 13:57:28.721947: I tensorflow/core/common_runtime/executor.cc:1197] [/device:CPU
  [[{{node gradients/split_1_grad(concat/split_1/split_dim)}}]]
2023-07-14 13:57:28.900179: I tensorflow/core/common_runtime/executor.cc:1197] [/device:CPU
  [[{{node gradients/split_2_grad(concat/split_2/split_dim)}}]]
2023-07-14 13:57:28.901647: I tensorflow/core/common_runtime/executor.cc:1197] [/device:CPU
  [[{{node gradients/split_grad(concat/split/split_dim)}}]]
2023-07-14 13:57:28.902886: I tensorflow/core/common_runtime/executor.cc:1197] [/device:CPU
  [[{{node gradients/split_1_grad(concat/split_1/split_dim)}}]]
2023-07-14 13:57:28.937200: I tensorflow/core/common_runtime/executor.cc:1197] [/device:CPU
  [[{{node inputs}}]]]
WARNING:absl:Found untraced functions such as _jit_compiled_convolution_op, _jit_compiled_co

INFO:tensorflow:Assets written to: ./data/bach_model/assets

INFO:tensorflow:Assets written to: ./data/bach_model/assets

model.evaluate(test_set)

34/34  [=====] - 6s 174ms/step - loss: 0.6610 -
accuracy: 0.8178

[0.6609758734703064, 0.8177668452262878]
```

3 - Use the model to generate Bach-like music, one note at a time: you can do this by giving the model the start of a chorale and asking it to predict the next time step, then appending these time steps to the input sequence and asking the model for the next note, and so on.

There's no real need for a test set in this exercise, since we will perform the final evaluation by just listening to the music produced by the model. So if we want, you can add the test set to the train set, and train the model again, hopefully getting a slightly better model.

Now let's write a function that will generate a new chorale. We will give it a few seed chords, it will convert them to arpeggios (the format expected by the model), and use the model to predict the next note, then the next, and so on. In the end, it will group the notes 4 by 4 to create chords again, and return the resulting chorale.

```
def generate_chorale(model, seed_chords, length):
    arpeggio = preprocess(tf.constant(seed_chords, dtype=tf.int64))
    arpeggio = tf.reshape(arpeggio, [1, -1])
    for chord in range(length):
        for note in range(4):
            next_note = model.predict(arpeggio, verbose=0).argmax(axis=-1)[:1, -1:]
            arpeggio = tf.concat([arpeggio, next_note], axis=1)
    arpeggio = tf.where(arpeggio == 0, arpeggio, arpeggio + min_note - 1)
    return tf.reshape(arpeggio, shape=[-1, 4])
```

To test this function, we need some seed chords. Let's use the first 8 chords of one of the test chorales:

```
seed_chords = test_chorales[2][:8]
play_chords(seed_chords, amplitude=0.2)
```

<IPython.lib.display.Audio object>

Now we are ready to generate our first chorale. Let's ask the function to generate 56 more chords, for a total of 64 chords, i.e., 16 bars (assuming 4 chords per bar, i.e., a 4/4 signature):

```
new_chorale = generate_chorale(model, seed_chords, 56)
play_chords(new_chorale)
```

```
2023-07-14 13:58:00.800734: I tensorflow/core/common_runtime/executor.cc:1197] [/device:CPU
[[{{node gradients/split_2_grad(concat/split_2/split_dim)}}]]
2023-07-14 13:58:00.801876: I tensorflow/core/common_runtime/executor.cc:1197] [/device:CPU
[[{{node gradients/split_grad(concat/split/split_dim)}}]]
2023-07-14 13:58:00.803183: I tensorflow/core/common_runtime/executor.cc:1197] [/device:CPU
[[{{node gradients/split_1_grad(concat/split_1/split_dim)}}]]
2023-07-14 13:58:01.376297: I tensorflow/core/common_runtime/executor.cc:1197] [/device:CPU
[[{{node gradients/split_2_grad(concat/split_2/split_dim)}}]]
2023-07-14 13:58:01.378218: I tensorflow/core/common_runtime/executor.cc:1197] [/device:CPU
[[{{node gradients/split_grad(concat/split/split_dim)}}]]
2023-07-14 13:58:01.379676: I tensorflow/core/common_runtime/executor.cc:1197] [/device:CPU
[[{{node gradients/split_1_grad(concat/split_1/split_dim)}}]]
```

<IPython.lib.display.Audio object>

This approach has one major flaw: it is often too conservative. The model will not take any risk, it will always choose the note with the highest score, and since repeating the previous note generally sounds good enough, it's the least risky option, so the algorithm will tend to make notes last longer and longer. Pretty boring. Plus, if you run the model multiple times, it will always generate the same melody. So, instead of always picking the note with the highest score, we will pick the next note randomly, according to the predicted probabilities. For example, if the model predicts a C3 with 75% probability, and a G3 with a 25% probability, then we will pick one of these two notes randomly, with these probabilities. We will also add a "temperature" parameter that will control how "hot" we want the system to feel. A high temperature will bring the predicted probabilities closer together, reducing the probability of the likely notes and increasing the probability of the unlikely ones.

```
def generate_chorale_hot(model, seed_chords, length, temperature=1):
    arpeggio = preprocess(tf.constant(seed_chords, dtype=tf.int64))
    arpeggio = tf.reshape(arpeggio, [1, -1])
    for chord in range(length):
        for note in range(4):
            next_note_probas = model.predict(arpeggio)[0, -1:]
            rescaled_logits = tf.math.log(next_note_probas) / temperature
            next_note = tf.random.categorical(rescaled_logits, num_samples=1)
            arpeggio = tf.concat([arpeggio, next_note], axis=1)
    arpeggio = tf.where(arpeggio == 0, arpeggio, arpeggio + min_note - 1)
    return tf.reshape(arpeggio, shape=[-1, 4])
```

Let's generate 3 chorales: one cold, one medium, and one hot and saves each chorale to a separate file. You can run these cells over and over again until you generate a masterpiece!

```
new_chorale_cold = generate_chorale_hot(model, seed_chords, 56, temperature=0.8)
play_chords(new_chorale_cold, filepath="./data/bach_cold.wav")
```

```
1/1 [=====] - 0s 43ms/step
1/1 [=====] - 0s 26ms/step
1/1 [=====] - 0s 26ms/step
1/1 [=====] - 0s 28ms/step
1/1 [=====] - 0s 26ms/step
1/1 [=====] - 0s 25ms/step
1/1 [=====] - 0s 23ms/step
1/1 [=====] - 0s 26ms/step
1/1 [=====] - 0s 22ms/step
```

```
1/1 [=====] - 0s 21ms/step
1/1 [=====] - 0s 21ms/step
1/1 [=====] - 0s 23ms/step
1/1 [=====] - 0s 21ms/step
1/1 [=====] - 0s 21ms/step
1/1 [=====] - 0s 22ms/step
1/1 [=====] - 0s 24ms/step
1/1 [=====] - 0s 29ms/step
1/1 [=====] - 0s 26ms/step
1/1 [=====] - 0s 25ms/step
1/1 [=====] - 0s 23ms/step
1/1 [=====] - 0s 25ms/step
1/1 [=====] - 0s 22ms/step
1/1 [=====] - 0s 21ms/step
1/1 [=====] - 0s 21ms/step
1/1 [=====] - 0s 23ms/step
1/1 [=====] - 0s 28ms/step
1/1 [=====] - 0s 29ms/step
1/1 [=====] - 0s 28ms/step
1/1 [=====] - 0s 29ms/step
1/1 [=====] - 0s 28ms/step
1/1 [=====] - 0s 25ms/step
1/1 [=====] - 0s 22ms/step
1/1 [=====] - 0s 21ms/step
1/1 [=====] - 0s 21ms/step
1/1 [=====] - 0s 23ms/step
1/1 [=====] - 0s 21ms/step
1/1 [=====] - 0s 21ms/step
1/1 [=====] - 0s 24ms/step
1/1 [=====] - 0s 22ms/step
1/1 [=====] - 0s 32ms/step
1/1 [=====] - 0s 33ms/step
1/1 [=====] - 0s 30ms/step
1/1 [=====] - 0s 33ms/step
1/1 [=====] - 0s 27ms/step
1/1 [=====] - 0s 27ms/step
1/1 [=====] - 0s 25ms/step
1/1 [=====] - 0s 22ms/step
```

```
1/1 [=====] - 0s 22ms/step
1/1 [=====] - 0s 23ms/step
1/1 [=====] - 0s 22ms/step
1/1 [=====] - 0s 23ms/step
1/1 [=====] - 0s 22ms/step
1/1 [=====] - 0s 24ms/step
1/1 [=====] - 0s 23ms/step
1/1 [=====] - 0s 24ms/step
1/1 [=====] - 0s 23ms/step
1/1 [=====] - 0s 24ms/step
1/1 [=====] - 0s 25ms/step
1/1 [=====] - 0s 23ms/step
1/1 [=====] - 0s 24ms/step
1/1 [=====] - 0s 24ms/step
1/1 [=====] - 0s 25ms/step
1/1 [=====] - 0s 23ms/step
1/1 [=====] - 0s 24ms/step
1/1 [=====] - 0s 28ms/step
1/1 [=====] - 0s 24ms/step
1/1 [=====] - 0s 24ms/step
1/1 [=====] - 0s 24ms/step
1/1 [=====] - 0s 23ms/step
1/1 [=====] - 0s 24ms/step
1/1 [=====] - 0s 25ms/step
1/1 [=====] - 0s 24ms/step
1/1 [=====] - 0s 24ms/step
1/1 [=====] - 0s 24ms/step
1/1 [=====] - 0s 26ms/step
1/1 [=====] - 0s 25ms/step
1/1 [=====] - 0s 28ms/step
1/1 [=====] - 0s 25ms/step
1/1 [=====] - 0s 26ms/step
1/1 [=====] - 0s 24ms/step
1/1 [=====] - 0s 24ms/step
1/1 [=====] - 0s 39ms/step
1/1 [=====] - 0s 25ms/step
1/1 [=====] - 0s 26ms/step
```

```
1/1 [=====] - 0s 26ms/step
1/1 [=====] - 0s 27ms/step
1/1 [=====] - 0s 25ms/step
1/1 [=====] - 0s 25ms/step
1/1 [=====] - 0s 27ms/step
1/1 [=====] - 0s 26ms/step
1/1 [=====] - 0s 26ms/step
1/1 [=====] - 0s 25ms/step
1/1 [=====] - 0s 27ms/step
1/1 [=====] - 0s 25ms/step
1/1 [=====] - 0s 26ms/step
1/1 [=====] - 0s 26ms/step
1/1 [=====] - 0s 25ms/step
1/1 [=====] - 0s 27ms/step
1/1 [=====] - 0s 26ms/step
1/1 [=====] - 0s 27ms/step
1/1 [=====] - 0s 26ms/step
1/1 [=====] - 0s 25ms/step
1/1 [=====] - 0s 27ms/step
1/1 [=====] - 0s 25ms/step
1/1 [=====] - 0s 26ms/step
1/1 [=====] - 0s 28ms/step
1/1 [=====] - 0s 26ms/step
1/1 [=====] - 0s 27ms/step
1/1 [=====] - 0s 30ms/step
1/1 [=====] - 0s 27ms/step
1/1 [=====] - 0s 25ms/step
1/1 [=====] - 0s 26ms/step
1/1 [=====] - 0s 27ms/step
1/1 [=====] - 0s 29ms/step
1/1 [=====] - 0s 26ms/step
1/1 [=====] - 0s 27ms/step
1/1 [=====] - 0s 27ms/step
1/1 [=====] - 0s 28ms/step
1/1 [=====] - 0s 27ms/step
1/1 [=====] - 0s 28ms/step
1/1 [=====] - 0s 26ms/step
1/1 [=====] - 0s 26ms/step
1/1 [=====] - 0s 27ms/step
1/1 [=====] - 0s 26ms/step
1/1 [=====] - 0s 28ms/step
```

```
1/1 [=====] - 0s 32ms/step
1/1 [=====] - 0s 28ms/step
1/1 [=====] - 0s 30ms/step
1/1 [=====] - 0s 30ms/step
1/1 [=====] - 0s 26ms/step
1/1 [=====] - 0s 27ms/step
1/1 [=====] - 0s 29ms/step
1/1 [=====] - 0s 27ms/step
1/1 [=====] - 0s 27ms/step
1/1 [=====] - 0s 27ms/step
1/1 [=====] - 0s 28ms/step
1/1 [=====] - 0s 29ms/step
1/1 [=====] - 0s 29ms/step
1/1 [=====] - 0s 27ms/step
1/1 [=====] - 0s 27ms/step
1/1 [=====] - 0s 32ms/step
1/1 [=====] - 0s 28ms/step
1/1 [=====] - 0s 29ms/step
1/1 [=====] - 0s 28ms/step
1/1 [=====] - 0s 28ms/step
1/1 [=====] - 0s 31ms/step
1/1 [=====] - 0s 30ms/step
1/1 [=====] - 0s 36ms/step
1/1 [=====] - 0s 33ms/step
1/1 [=====] - 0s 32ms/step
1/1 [=====] - 0s 33ms/step
1/1 [=====] - 0s 35ms/step
1/1 [=====] - 0s 29ms/step
1/1 [=====] - 0s 31ms/step
1/1 [=====] - 0s 31ms/step
1/1 [=====] - 0s 29ms/step
1/1 [=====] - 0s 29ms/step
1/1 [=====] - 0s 30ms/step
1/1 [=====] - 0s 31ms/step
1/1 [=====] - 0s 31ms/step
1/1 [=====] - 0s 32ms/step
1/1 [=====] - 0s 30ms/step
1/1 [=====] - 0s 54ms/step
```

```
1/1 [=====] - 0s 31ms/step
1/1 [=====] - 0s 31ms/step
1/1 [=====] - 0s 29ms/step
1/1 [=====] - 0s 31ms/step
1/1 [=====] - 0s 32ms/step
1/1 [=====] - 0s 34ms/step
1/1 [=====] - 0s 34ms/step
1/1 [=====] - 0s 33ms/step
1/1 [=====] - 0s 34ms/step
1/1 [=====] - 0s 33ms/step
1/1 [=====] - 0s 34ms/step
1/1 [=====] - 0s 32ms/step
1/1 [=====] - 0s 33ms/step
1/1 [=====] - 0s 32ms/step
1/1 [=====] - 0s 30ms/step
1/1 [=====] - 0s 30ms/step
1/1 [=====] - 0s 35ms/step
1/1 [=====] - 0s 31ms/step
1/1 [=====] - 0s 32ms/step
1/1 [=====] - 0s 29ms/step
1/1 [=====] - 0s 30ms/step
1/1 [=====] - 0s 31ms/step
1/1 [=====] - 0s 32ms/step
1/1 [=====] - 0s 31ms/step
1/1 [=====] - 0s 34ms/step
1/1 [=====] - 0s 31ms/step
1/1 [=====] - 0s 31ms/step
1/1 [=====] - 0s 31ms/step
1/1 [=====] - 0s 34ms/step
1/1 [=====] - 0s 33ms/step
1/1 [=====] - 0s 36ms/step
1/1 [=====] - 0s 33ms/step
1/1 [=====] - 0s 31ms/step
1/1 [=====] - 0s 32ms/step
1/1 [=====] - 0s 32ms/step
1/1 [=====] - 0s 35ms/step
1/1 [=====] - 0s 30ms/step
1/1 [=====] - 0s 33ms/step
1/1 [=====] - 0s 33ms/step
```

```
1/1 [=====] - 0s 31ms/step
1/1 [=====] - 0s 35ms/step
1/1 [=====] - 0s 32ms/step
1/1 [=====] - 0s 34ms/step
1/1 [=====] - 0s 32ms/step
1/1 [=====] - 0s 34ms/step
1/1 [=====] - 0s 34ms/step
1/1 [=====] - 0s 31ms/step
1/1 [=====] - 0s 34ms/step
1/1 [=====] - 0s 36ms/step
1/1 [=====] - 0s 33ms/step
1/1 [=====] - 0s 34ms/step
1/1 [=====] - 0s 33ms/step
1/1 [=====] - 0s 33ms/step
1/1 [=====] - 0s 34ms/step
1/1 [=====] - 0s 34ms/step
1/1 [=====] - 0s 35ms/step
1/1 [=====] - 0s 36ms/step
1/1 [=====] - 0s 34ms/step
1/1 [=====] - 0s 36ms/step
1/1 [=====] - 0s 33ms/step
1/1 [=====] - 0s 32ms/step
1/1 [=====] - 0s 32ms/step
1/1 [=====] - 0s 41ms/step
1/1 [=====] - 0s 34ms/step
```

<IPython.lib.display.Audio object>

```
new_chorale_medium = generate_chorale_hot(model, seed_chords, 56, temperature=1.0)
play_chords(new_chorale_medium, filepath=".//data/bach_medium.wav")
```

```
1/1 [=====] - 0s 22ms/step
1/1 [=====] - 0s 22ms/step
1/1 [=====] - 0s 23ms/step
1/1 [=====] - 0s 24ms/step
1/1 [=====] - 0s 21ms/step
1/1 [=====] - 0s 21ms/step
1/1 [=====] - 0s 21ms/step
```

```
1/1 [=====] - 0s 21ms/step
1/1 [=====] - 0s 21ms/step
1/1 [=====] - 0s 22ms/step
1/1 [=====] - 0s 20ms/step
1/1 [=====] - 0s 21ms/step
1/1 [=====] - 0s 21ms/step
1/1 [=====] - 0s 22ms/step
1/1 [=====] - 0s 22ms/step
1/1 [=====] - 0s 24ms/step
1/1 [=====] - 0s 21ms/step
1/1 [=====] - 0s 22ms/step
1/1 [=====] - 0s 21ms/step
1/1 [=====] - 0s 23ms/step
1/1 [=====] - 0s 22ms/step
1/1 [=====] - 0s 22ms/step
1/1 [=====] - 0s 21ms/step
1/1 [=====] - 0s 22ms/step
1/1 [=====] - 0s 22ms/step
1/1 [=====] - 0s 23ms/step
1/1 [=====] - 0s 21ms/step
1/1 [=====] - 0s 23ms/step
1/1 [=====] - 0s 23ms/step
1/1 [=====] - 0s 27ms/step
1/1 [=====] - 0s 28ms/step
1/1 [=====] - 0s 28ms/step
1/1 [=====] - 0s 28ms/step
1/1 [=====] - 0s 30ms/step
1/1 [=====] - 0s 25ms/step
1/1 [=====] - 0s 22ms/step
1/1 [=====] - 0s 23ms/step
1/1 [=====] - 0s 22ms/step
1/1 [=====] - 0s 23ms/step
1/1 [=====] - 0s 24ms/step
```

```
1/1 [=====] - 0s 22ms/step
1/1 [=====] - 0s 24ms/step
1/1 [=====] - 0s 23ms/step
1/1 [=====] - 0s 23ms/step
1/1 [=====] - 0s 23ms/step
1/1 [=====] - 0s 24ms/step
1/1 [=====] - 0s 23ms/step
1/1 [=====] - 0s 23ms/step
1/1 [=====] - 0s 24ms/step
1/1 [=====] - 0s 23ms/step
1/1 [=====] - 0s 23ms/step
1/1 [=====] - 0s 25ms/step
1/1 [=====] - 0s 52ms/step
1/1 [=====] - 0s 26ms/step
1/1 [=====] - 0s 23ms/step
1/1 [=====] - 0s 24ms/step
1/1 [=====] - 0s 27ms/step
1/1 [=====] - 0s 24ms/step
1/1 [=====] - 0s 23ms/step
1/1 [=====] - 0s 23ms/step
1/1 [=====] - 0s 23ms/step
1/1 [=====] - 0s 25ms/step
1/1 [=====] - 0s 23ms/step
1/1 [=====] - 0s 24ms/step
1/1 [=====] - 0s 24ms/step
1/1 [=====] - 0s 25ms/step
1/1 [=====] - 0s 24ms/step
1/1 [=====] - 0s 25ms/step
1/1 [=====] - 0s 24ms/step
1/1 [=====] - 0s 26ms/step
1/1 [=====] - 0s 29ms/step
1/1 [=====] - 0s 25ms/step
1/1 [=====] - 0s 25ms/step
1/1 [=====] - 0s 25ms/step
1/1 [=====] - 0s 26ms/step
1/1 [=====] - 0s 26ms/step
1/1 [=====] - 0s 30ms/step
```

1/1 [=====] - 0s 32ms/step
1/1 [=====] - 0s 28ms/step
1/1 [=====] - 0s 25ms/step
1/1 [=====] - 0s 25ms/step
1/1 [=====] - 0s 26ms/step
1/1 [=====] - 0s 24ms/step
1/1 [=====] - 0s 27ms/step
1/1 [=====] - 0s 26ms/step
1/1 [=====] - 0s 26ms/step
1/1 [=====] - 0s 27ms/step
1/1 [=====] - 0s 26ms/step
1/1 [=====] - 0s 27ms/step
1/1 [=====] - 0s 25ms/step
1/1 [=====] - 0s 27ms/step
1/1 [=====] - 0s 26ms/step
1/1 [=====] - 0s 28ms/step
1/1 [=====] - 0s 26ms/step
1/1 [=====] - 0s 27ms/step
1/1 [=====] - 0s 26ms/step
1/1 [=====] - 0s 27ms/step
1/1 [=====] - 0s 26ms/step
1/1 [=====] - 0s 29ms/step
1/1 [=====] - 0s 26ms/step
1/1 [=====] - 0s 26ms/step
1/1 [=====] - 0s 28ms/step
1/1 [=====] - 0s 26ms/step
1/1 [=====] - 0s 27ms/step
1/1 [=====] - 0s 26ms/step
1/1 [=====] - 0s 27ms/step
1/1 [=====] - 0s 28ms/step
1/1 [=====] - 0s 28ms/step
1/1 [=====] - 0s 27ms/step

```
1/1 [=====] - 0s 28ms/step
1/1 [=====] - 0s 27ms/step
1/1 [=====] - 0s 29ms/step
1/1 [=====] - 0s 26ms/step
1/1 [=====] - 0s 27ms/step
1/1 [=====] - 0s 30ms/step
1/1 [=====] - 0s 30ms/step
1/1 [=====] - 0s 34ms/step
1/1 [=====] - 0s 30ms/step
1/1 [=====] - 0s 31ms/step
1/1 [=====] - 0s 29ms/step
1/1 [=====] - 0s 28ms/step
1/1 [=====] - 0s 28ms/step
1/1 [=====] - 0s 29ms/step
1/1 [=====] - 0s 28ms/step
1/1 [=====] - 0s 27ms/step
1/1 [=====] - 0s 28ms/step
1/1 [=====] - 0s 28ms/step
1/1 [=====] - 0s 30ms/step
1/1 [=====] - 0s 27ms/step
1/1 [=====] - 0s 34ms/step
1/1 [=====] - 0s 33ms/step
1/1 [=====] - 0s 35ms/step
1/1 [=====] - 0s 32ms/step
1/1 [=====] - 0s 31ms/step
1/1 [=====] - 0s 30ms/step
1/1 [=====] - 0s 32ms/step
1/1 [=====] - 0s 33ms/step
1/1 [=====] - 0s 28ms/step
1/1 [=====] - 0s 31ms/step
1/1 [=====] - 0s 36ms/step
1/1 [=====] - 0s 28ms/step
1/1 [=====] - 0s 31ms/step
1/1 [=====] - 0s 32ms/step
1/1 [=====] - 0s 33ms/step
1/1 [=====] - 0s 33ms/step
1/1 [=====] - 0s 35ms/step
1/1 [=====] - 0s 34ms/step
```

```
1/1 [=====] - 0s 31ms/step
1/1 [=====] - 0s 33ms/step
1/1 [=====] - 0s 32ms/step
1/1 [=====] - 0s 36ms/step
1/1 [=====] - 0s 32ms/step
1/1 [=====] - 0s 31ms/step
1/1 [=====] - 0s 31ms/step
1/1 [=====] - 0s 35ms/step
1/1 [=====] - 0s 31ms/step
1/1 [=====] - 0s 33ms/step
1/1 [=====] - 0s 33ms/step
1/1 [=====] - 0s 30ms/step
1/1 [=====] - 0s 55ms/step
1/1 [=====] - 0s 32ms/step
1/1 [=====] - 0s 30ms/step
1/1 [=====] - 0s 32ms/step
1/1 [=====] - 0s 31ms/step
1/1 [=====] - 0s 33ms/step
1/1 [=====] - 0s 33ms/step
1/1 [=====] - 0s 33ms/step
1/1 [=====] - 0s 31ms/step
1/1 [=====] - 0s 31ms/step
1/1 [=====] - 0s 32ms/step
1/1 [=====] - 0s 30ms/step
1/1 [=====] - 0s 33ms/step
1/1 [=====] - 0s 32ms/step
1/1 [=====] - 0s 31ms/step
1/1 [=====] - 0s 31ms/step
1/1 [=====] - 0s 32ms/step
```

```
1/1 [=====] - 0s 32ms/step
1/1 [=====] - 0s 34ms/step
1/1 [=====] - 0s 32ms/step
1/1 [=====] - 0s 35ms/step
1/1 [=====] - 0s 36ms/step
1/1 [=====] - 0s 34ms/step
1/1 [=====] - 0s 33ms/step
1/1 [=====] - 0s 36ms/step
1/1 [=====] - 0s 32ms/step
1/1 [=====] - 0s 32ms/step
1/1 [=====] - 0s 32ms/step
1/1 [=====] - 0s 33ms/step
1/1 [=====] - 0s 35ms/step
1/1 [=====] - 0s 32ms/step
1/1 [=====] - 0s 34ms/step
1/1 [=====] - 0s 33ms/step
1/1 [=====] - 0s 35ms/step
1/1 [=====] - 0s 34ms/step
1/1 [=====] - 0s 35ms/step
1/1 [=====] - 0s 33ms/step
1/1 [=====] - 0s 31ms/step
1/1 [=====] - 0s 33ms/step
1/1 [=====] - 0s 36ms/step
1/1 [=====] - 0s 33ms/step
```

<IPython.lib.display.Audio object>

```
new_chorale_hot = generate_chorale_hot(model, seed_chords, 56, temperature=1.5)
play_chords(new_chorale_hot, filepath="./data/bach_hot.wav")
```

```
1/1 [=====] - 0s 26ms/step
1/1 [=====] - 0s 20ms/step
1/1 [=====] - 0s 24ms/step
1/1 [=====] - 0s 27ms/step
1/1 [=====] - 0s 28ms/step
```

```
1/1 [=====] - 0s 26ms/step
1/1 [=====] - 0s 24ms/step
1/1 [=====] - 0s 24ms/step
1/1 [=====] - 0s 25ms/step
1/1 [=====] - 0s 22ms/step
1/1 [=====] - 0s 23ms/step
1/1 [=====] - 0s 21ms/step
1/1 [=====] - 0s 21ms/step
1/1 [=====] - 0s 23ms/step
1/1 [=====] - 0s 21ms/step
1/1 [=====] - 0s 23ms/step
1/1 [=====] - 0s 21ms/step
1/1 [=====] - 0s 22ms/step
1/1 [=====] - 0s 29ms/step
1/1 [=====] - 0s 26ms/step
1/1 [=====] - 0s 25ms/step
1/1 [=====] - 0s 23ms/step
1/1 [=====] - 0s 21ms/step
1/1 [=====] - 0s 21ms/step
1/1 [=====] - 0s 22ms/step
1/1 [=====] - 0s 21ms/step
1/1 [=====] - 0s 22ms/step
1/1 [=====] - 0s 22ms/step
1/1 [=====] - 0s 23ms/step
1/1 [=====] - 0s 22ms/step
1/1 [=====] - 0s 22ms/step
1/1 [=====] - 0s 25ms/step
1/1 [=====] - 0s 27ms/step
1/1 [=====] - 0s 23ms/step
1/1 [=====] - 0s 23ms/step
1/1 [=====] - 0s 22ms/step
1/1 [=====] - 0s 26ms/step
1/1 [=====] - 0s 24ms/step
1/1 [=====] - 0s 23ms/step
```

```
1/1 [=====] - 0s 24ms/step
1/1 [=====] - 0s 24ms/step
1/1 [=====] - 0s 26ms/step
1/1 [=====] - 0s 24ms/step
1/1 [=====] - 0s 24ms/step
1/1 [=====] - 0s 23ms/step
1/1 [=====] - 0s 24ms/step
1/1 [=====] - 0s 23ms/step
1/1 [=====] - 0s 26ms/step
1/1 [=====] - 0s 32ms/step
1/1 [=====] - 0s 30ms/step
1/1 [=====] - 0s 31ms/step
1/1 [=====] - 0s 33ms/step
1/1 [=====] - 0s 31ms/step
1/1 [=====] - 0s 29ms/step
1/1 [=====] - 0s 29ms/step
1/1 [=====] - 0s 31ms/step
1/1 [=====] - 0s 30ms/step
1/1 [=====] - 0s 28ms/step
1/1 [=====] - 0s 24ms/step
1/1 [=====] - 0s 24ms/step
1/1 [=====] - 0s 25ms/step
1/1 [=====] - 0s 24ms/step
1/1 [=====] - 0s 29ms/step
1/1 [=====] - 0s 24ms/step
1/1 [=====] - 0s 27ms/step
1/1 [=====] - 0s 31ms/step
1/1 [=====] - 0s 33ms/step
1/1 [=====] - 0s 34ms/step
1/1 [=====] - 0s 34ms/step
1/1 [=====] - 0s 35ms/step
1/1 [=====] - 0s 59ms/step
1/1 [=====] - 0s 30ms/step
1/1 [=====] - 0s 32ms/step
1/1 [=====] - 0s 30ms/step
1/1 [=====] - 0s 29ms/step
1/1 [=====] - 0s 36ms/step
1/1 [=====] - 0s 34ms/step
```

```
1/1 [=====] - 0s 34ms/step
1/1 [=====] - 0s 35ms/step
1/1 [=====] - 0s 33ms/step
1/1 [=====] - 0s 34ms/step
1/1 [=====] - 0s 33ms/step
1/1 [=====] - 0s 34ms/step
1/1 [=====] - 0s 35ms/step
1/1 [=====] - 0s 35ms/step
1/1 [=====] - 0s 32ms/step
1/1 [=====] - 0s 49ms/step
1/1 [=====] - 0s 33ms/step
1/1 [=====] - 0s 33ms/step
1/1 [=====] - 0s 35ms/step
1/1 [=====] - 0s 31ms/step
1/1 [=====] - 0s 26ms/step
1/1 [=====] - 0s 33ms/step
1/1 [=====] - 0s 27ms/step
1/1 [=====] - 0s 33ms/step
1/1 [=====] - 0s 34ms/step
1/1 [=====] - 0s 34ms/step
1/1 [=====] - 0s 35ms/step
1/1 [=====] - 0s 35ms/step
1/1 [=====] - 0s 33ms/step
1/1 [=====] - 0s 34ms/step
1/1 [=====] - 0s 60ms/step
1/1 [=====] - 0s 36ms/step
1/1 [=====] - 0s 35ms/step
1/1 [=====] - 0s 34ms/step
1/1 [=====] - 0s 37ms/step
1/1 [=====] - 0s 29ms/step
1/1 [=====] - 0s 36ms/step
1/1 [=====] - 0s 36ms/step
1/1 [=====] - 0s 40ms/step
1/1 [=====] - 0s 36ms/step
1/1 [=====] - 0s 39ms/step
1/1 [=====] - 0s 41ms/step
```

```
1/1 [=====] - 0s 39ms/step
1/1 [=====] - 0s 35ms/step
1/1 [=====] - 0s 73ms/step
1/1 [=====] - 0s 34ms/step
1/1 [=====] - 0s 45ms/step
1/1 [=====] - 0s 40ms/step
1/1 [=====] - 0s 35ms/step
1/1 [=====] - 0s 37ms/step
1/1 [=====] - 0s 38ms/step
1/1 [=====] - 0s 39ms/step
1/1 [=====] - 0s 39ms/step
1/1 [=====] - 0s 49ms/step
1/1 [=====] - 0s 44ms/step
1/1 [=====] - 0s 39ms/step
1/1 [=====] - 0s 40ms/step
1/1 [=====] - 0s 38ms/step
1/1 [=====] - 0s 35ms/step
1/1 [=====] - 0s 34ms/step
1/1 [=====] - 0s 35ms/step
1/1 [=====] - 0s 37ms/step
1/1 [=====] - 0s 38ms/step
1/1 [=====] - 0s 44ms/step
1/1 [=====] - 0s 42ms/step
1/1 [=====] - 0s 33ms/step
1/1 [=====] - 0s 31ms/step
1/1 [=====] - 0s 39ms/step
1/1 [=====] - 0s 40ms/step
1/1 [=====] - 0s 47ms/step
1/1 [=====] - 0s 41ms/step
1/1 [=====] - 0s 42ms/step
1/1 [=====] - 0s 87ms/step
1/1 [=====] - 0s 53ms/step
1/1 [=====] - 0s 42ms/step
1/1 [=====] - 0s 41ms/step
1/1 [=====] - 0s 39ms/step
1/1 [=====] - 0s 42ms/step
1/1 [=====] - 0s 38ms/step
1/1 [=====] - 0s 40ms/step
```



```
1/1 [=====] - 0s 44ms/step
1/1 [=====] - 0s 46ms/step
1/1 [=====] - 0s 50ms/step
1/1 [=====] - 0s 44ms/step
1/1 [=====] - 0s 33ms/step
1/1 [=====] - 0s 45ms/step
1/1 [=====] - 0s 45ms/step
1/1 [=====] - 0s 43ms/step
1/1 [=====] - 0s 41ms/step
1/1 [=====] - 0s 36ms/step
1/1 [=====] - 0s 40ms/step
1/1 [=====] - 0s 42ms/step
1/1 [=====] - 0s 47ms/step
1/1 [=====] - 0s 41ms/step
1/1 [=====] - 0s 34ms/step
1/1 [=====] - 0s 44ms/step
1/1 [=====] - 0s 45ms/step
1/1 [=====] - 0s 87ms/step
1/1 [=====] - 0s 41ms/step
1/1 [=====] - 0s 44ms/step
1/1 [=====] - 0s 48ms/step
1/1 [=====] - 0s 46ms/step
1/1 [=====] - 0s 46ms/step
1/1 [=====] - 0s 44ms/step
1/1 [=====] - 0s 47ms/step
1/1 [=====] - 0s 48ms/step
1/1 [=====] - 0s 49ms/step
1/1 [=====] - 0s 45ms/step
1/1 [=====] - 0s 54ms/step
```

<IPython.lib.display.Audio object>